

# DNA Sequencing, Mapping and Assembly

---

DAVID L. TABB, PH.D.


# Overview

---

- Sequencing chemistry
- Base-calling errors: the Phred algorithm
- Mapping versus Assembly strategies
- Mapping via Burrows-Wheeler algorithm
- Assembly via  $k$ -mer graphs
- FASTQ, SAM/BAM, FASTA


# DNA backbone structure

---



# Sanger sequencing

- Given a template, generate complementary sequence.
- After dideoxynucleotide is incorporated, no more extension is possible.
- Fragment ladder is separated through electrophoresis.




# Massively parallel sequencing

- Detects nucleotide incorporation during DNA strand synthesis for millions of templates
- Base detection and strand synthesis typically differentiate competing technologies.

How it works:

- Separate millions of single-stranded DNA templates
- Fix the templates' location on a substrate & amplify
- Detect the incorporation of each base in each location



# Sequencing jargon

---

- A “template” is a piece of DNA to be sequenced, often ends of an “insert.”
- A “read” is a sequence corresponding to a single template, output by sequencer.
- “Shotgun sequencing” generates reads at random locations in the target DNA.
- “Fold coverage” divides the sum of read lengths by the target DNA length.

I prefer “massively parallel sequencing” to “next-gen sequencing.”

# Old sequencing versus new

---

- Sanger sequencing typically produces longer reads (600 bp versus 150 bp).
- Sanger sequencing produces more accurate base calls for individual reads, but massively parallel sequencing overlaps many reads for each position.
- Massively parallel sequencing produces sequence from an incomparably larger number of templates in each experiment (millions rather than tens).

# Electropherogram output



Model 377  
Version 3.0  
SemiAdaptive  
Version 3.0

Uu3 HHH-4348 x  
11055-18  
Uu3 HHH-4348 x  
Lane 6

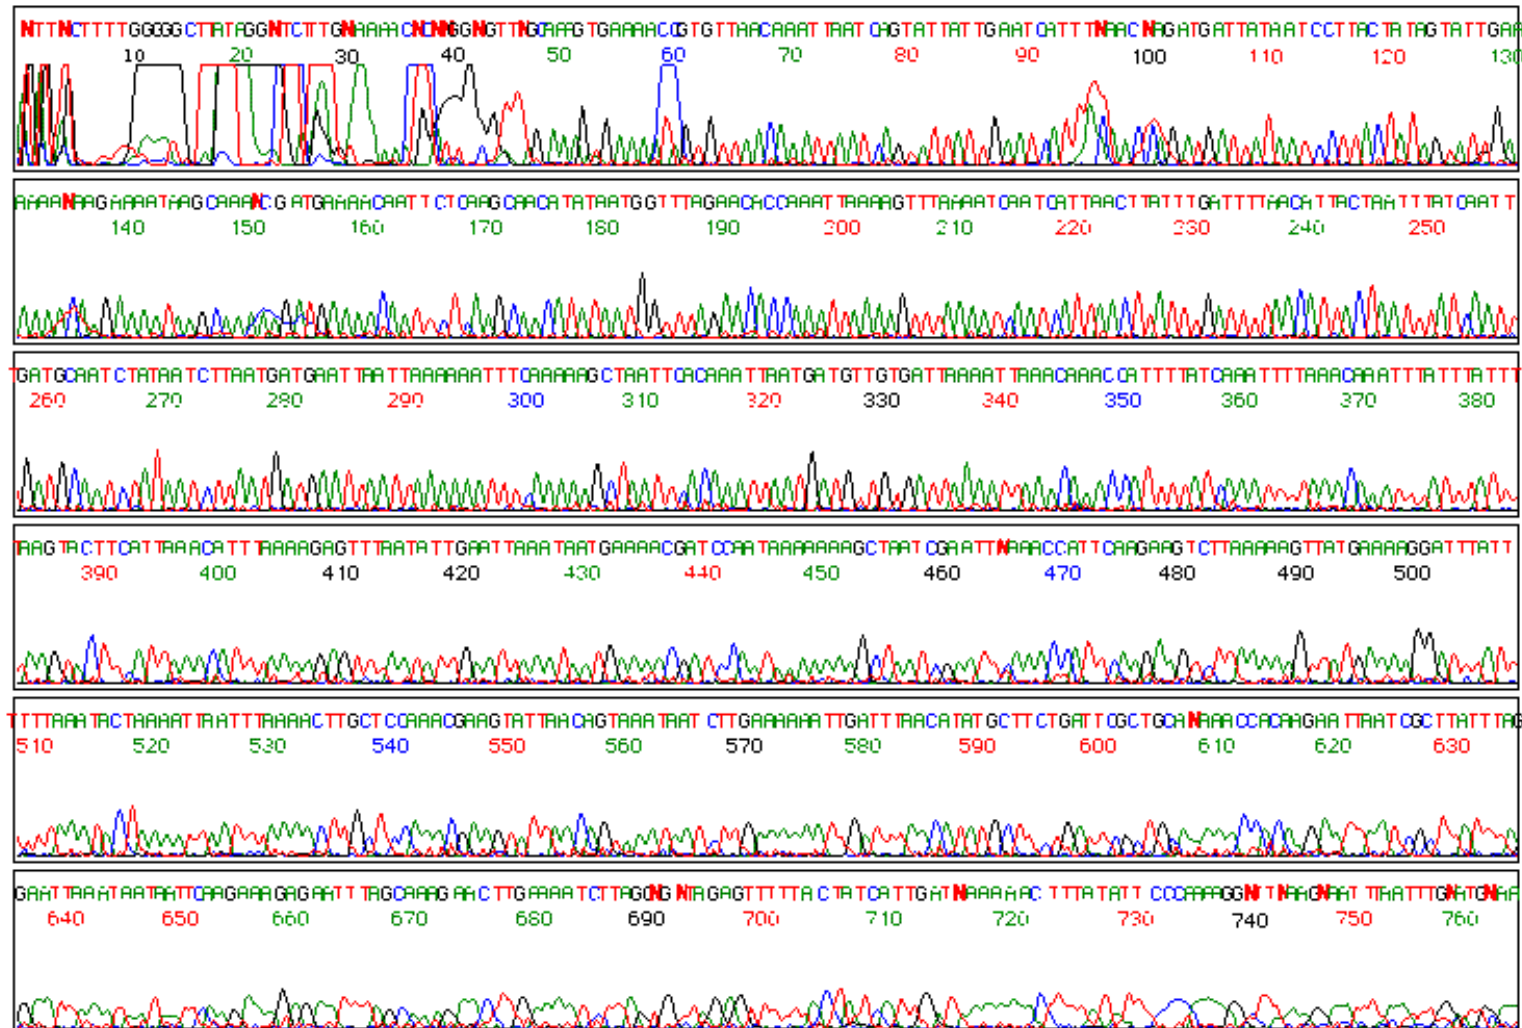
Signal G:56 A:68 T:41 C:60  
D TBigDyes{E Set-AnyPrimer}  
Zebra dRhod Matrix  
Points 1105 to 14760 Base 1:1105

Page 1 of 2

Tue, Oct 7, 1997 5:10 PM

Wed, Jul 16, 1997 4:11 PM

Spacing: 13.81{13.81}





# Phred: estimate basecalling errors


---

- Infer sequence from electropherogram
- Associate each basecall with a probability of error for this letter.
- A good basecall has these properties:
  - Peak matches the “beat” of its neighbors.
  - Only one trace is concave down at this call.
  - Highly-probable errors are far away.

B Ewing et al. *Genome Res.* (1998) 8: 175-185.


B Ewing et al. *Genome Res.* (1998) 8: 186-194.

# Fourier Transforms shift signal to sums of frequencies



- FTs decompose signals in time to frequencies.
- FTs are ubiquitous for recognizing frequencies.
- MP3 encoding uses FT to compress music.
- Sequencing requires that we recognize frequencies of base calls.

# Aligning observed and predicted peaks



# Negative log scores help differentiate rare events

---


- $P$  = probability of base call error, ranging from 0 to 1. Proximity to zero matters a lot.
- $Q = -10 \log_{10} P$
- If  $P=.01$  (1E-2) (1%)  $Q=20$ .
- If  $P=.001$  (1E-3) (0.1%)  $Q=30$ .
- If  $P=.0001$  (1E-4) (0.01%)  $Q=40$ .
- $Q$  is the “Phred score.”

# FASTQ: the output from DNA sequencers

---

```
@NB501496:55:HMMT2AFX:1:11101:13769:2030 1:N:0:CGCTCATT+AGGCTATA  
CCCGCACTTCACATACCGAAGCCGCCTGTGCCGCTCCTGACCGCCTAATCCCGGAGGGGGGGTGAGTGTGTGTT  
+  
AAAAAEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEEA/
```

- Line1: '@' followed by a record identifier
- Line2: Sequencing read base calls
- Line3: '+' means next line is quality scores
- Line4: Phred scores, from '!' to '~'; A=32
- A FASTQ typically contains millions of reads and requires compression (.fastq.gz or .fastq.bz2)



# Intermission

---

# Two chief uses for reads: *Mapping and Assembly*

---

## MAPPING

- Determines where each read matches to the annotation
- Basis for exome / WGS variant calling

## ASSEMBLY

- Infers large “contig” regions from overlapping reads
- Necessary for non-model organisms and for unmapped reads



# Why emphasize *mapping*: aligning short reads to reference?

---

- Assembling reads *de novo* is time-consuming and relatively error-prone.
- Recognition of sequence variants is easier when annotation provides typical sequence.
- Short read massively parallel sequencers (e.g. Illumina) produce >100 Gbp per day.
- QC, trimming and alignment are assumed for many downstream tools.

# Burrows-Wheeler Transform prepares a genome index.

- A suffix array shows all truncations from 5' of sequence.
- The array of suffixes is sorted.
- BW transform stores the letter preceding truncated sequence and where it appears in original.

Sorted suffixes

13	6	8	10	1	4	12	5	7	9	0	3	11	2
\$	A	A	A	A	C	G	G	G	G	G	G	T	T
A	G	G	T	T	G	\$	A	A	A	A	C	G	G
G	A	A	G	G	A		G	G	T	T	G	\$	C
A	G	T	\$	C	G		A	A	G	G	A		G
T	\$	\$		G	A		A	T	\$	C	A		A
\$				A	T		T	\$		G	G		G
				A	G		\$			A	A		A
				A	A					A	T		T
				T	G					A	G		G
				\$	\$					T	\$		\$

Suffix array of GATGCGAGAGATG



<http://blog.thegrandlocus.com/2016/07/a-tutorial-on-burrows-wheeler-indexing-methods>



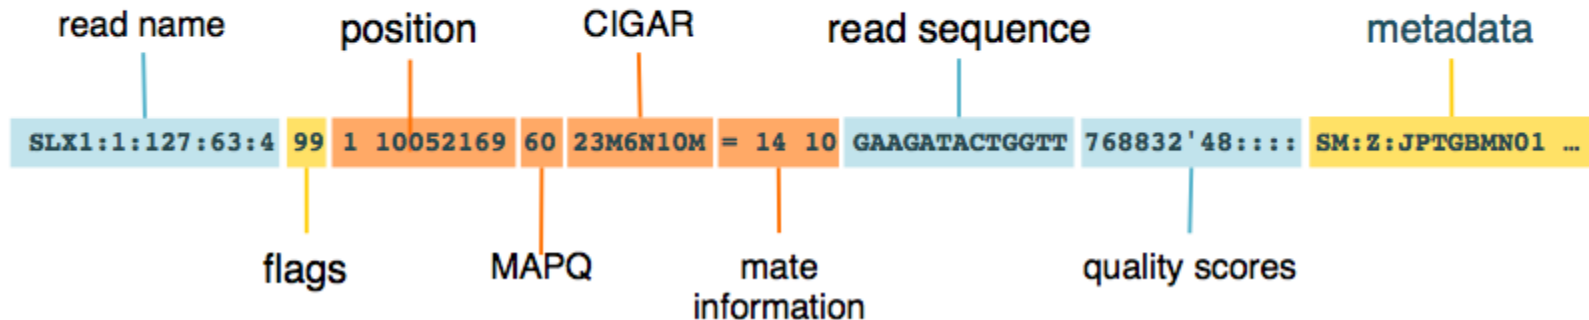
$i$   $S(i)$   $B[i]$   
 $\downarrow$   $\downarrow$   
 $(6, 3, 0, 5, 2, 4, 1)$

If mismatches are allowed, traversal takes longer!

# BAM files: Sequence Alignment Maps, in binary

**HEADER** containing metadata (sequence dictionary, read group definitions etc)

**RECORDS** containing structured read information (1 line per read record)



<https://gatk.broadinstitute.org/hc/en-us/articles/360035890791-SAM-or-BAM-or-CRAM-Mapped-sequence-data-formats>

MAPQ is mapping quality; CIGAR reflects diffs versus annotation. Mate Info relates paired-end sequences. Basecall PHRED scores are given in single-character format (ASCII-33 through 126).

# Sequence Assembly

---

# Why do we need *assembly*?

---

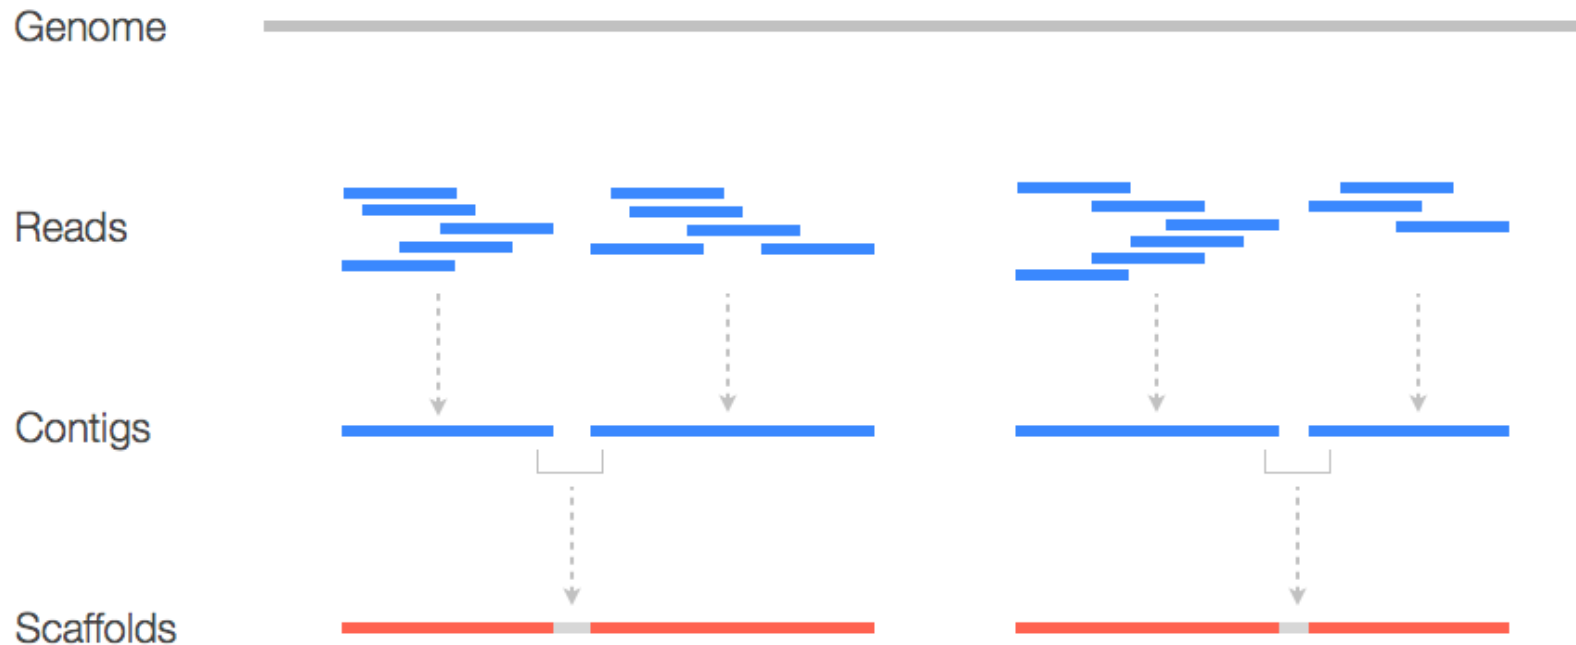
- We have millions of short reads, but we need scaffolds, ideally one per chromosome.
- Even if we map our reads to an annotation, many *unmapped reads* may not align to our selected reference annotation.
- We may work in a non-model organism that lacks a reference genome annotation!

# Assembly definitions

---

- A ***contig*** is a contiguous length of genomic sequence in which the order of bases is known to a high confidence level.
- ***Scaffolds*** are composed of contigs and gaps.
- Within a scaffold, the ordering and orientation of contigs with respect to each other has been established.

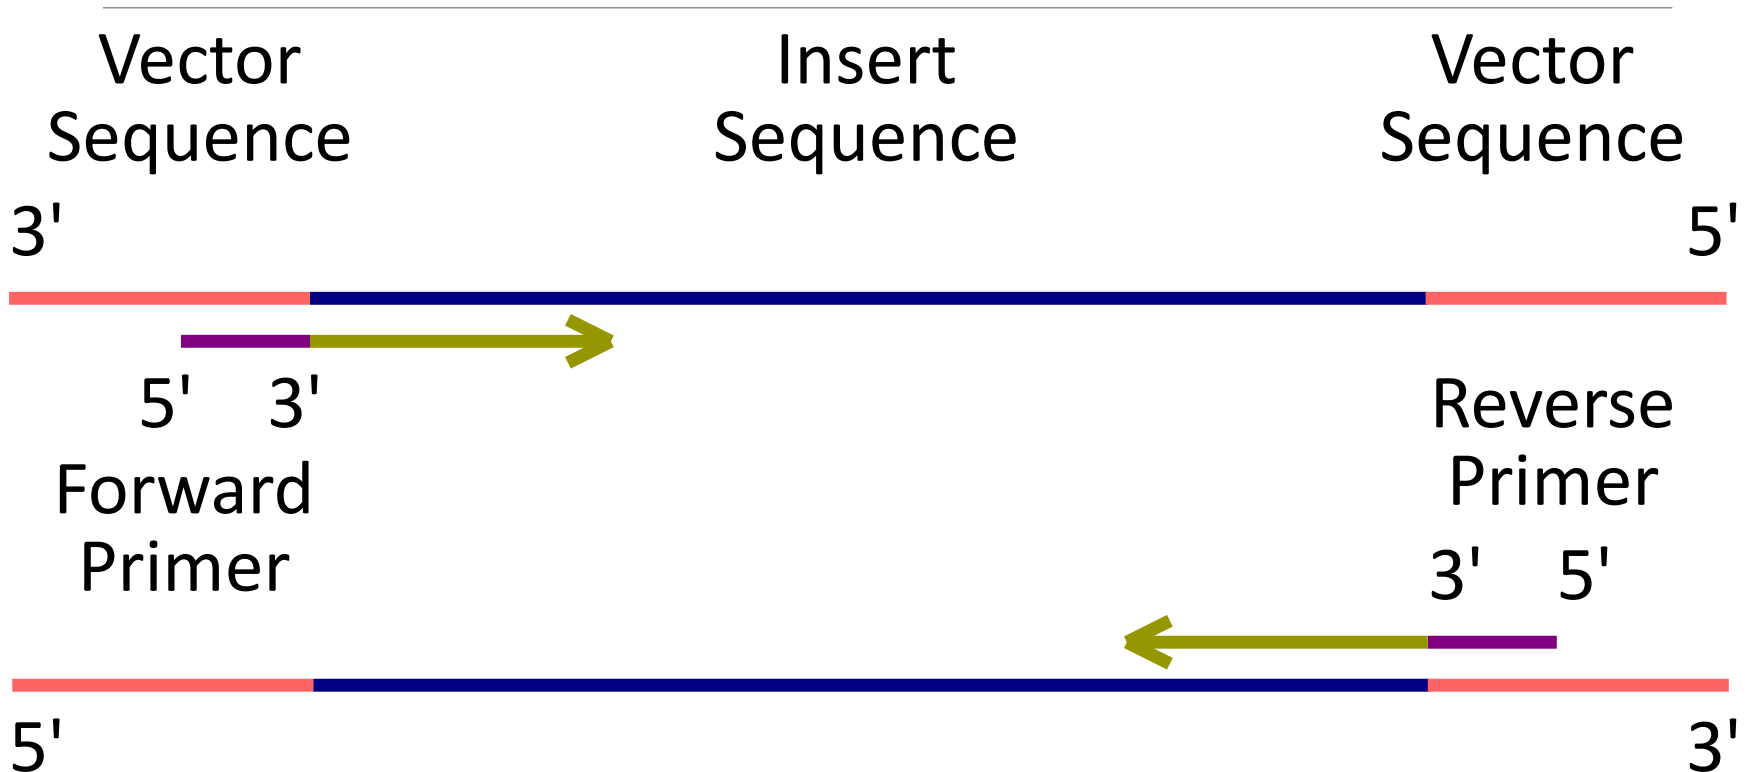
# Assembly builds long sequences from many short ones



If reads are independent, how do we judge which contigs are near each other?

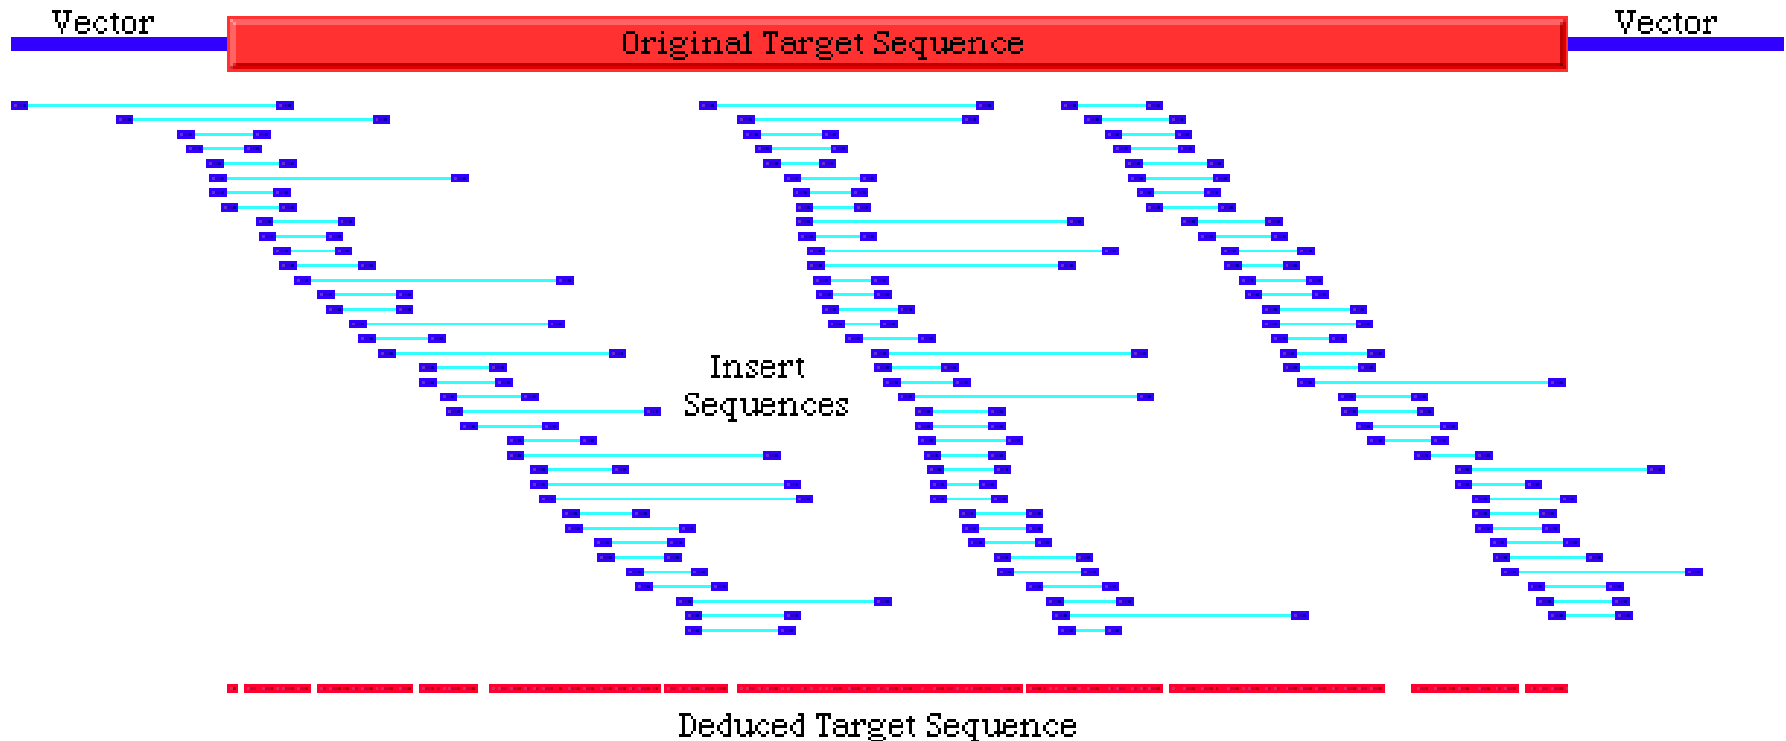


# Paired End Sequencing



Now the reads come in pairs, since each comes from the opposite end of a single piece of DNA.

# Shotgun sequencing of genomes

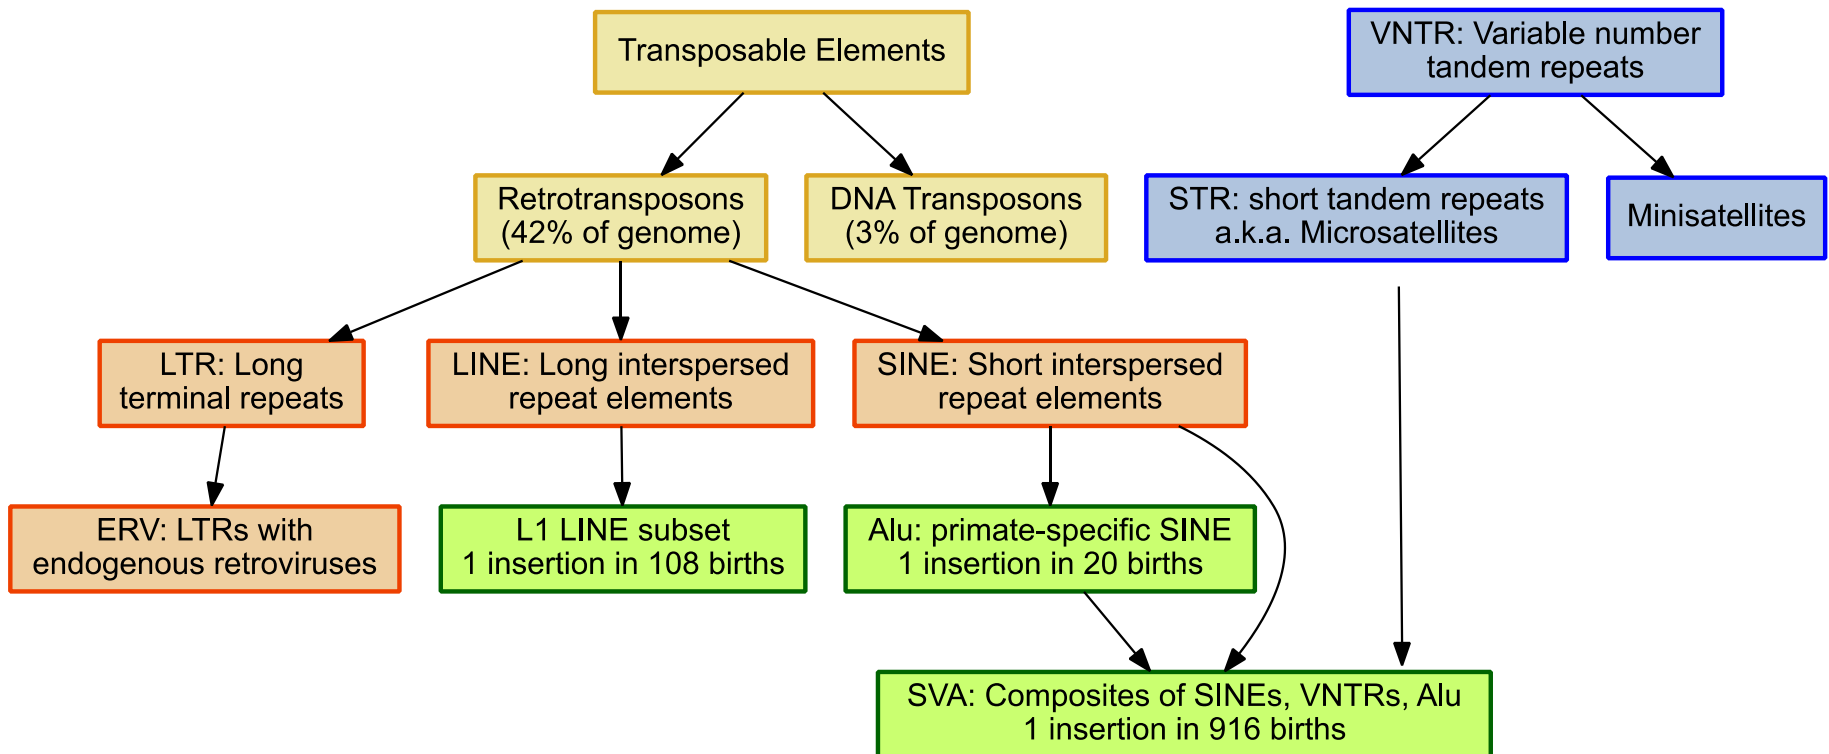


Given millions of reads of 100 nucleotides, assemble contigs of overlapping sequences. Determine which contigs are neighbors.

*Figure courtesy of Jared Roach*

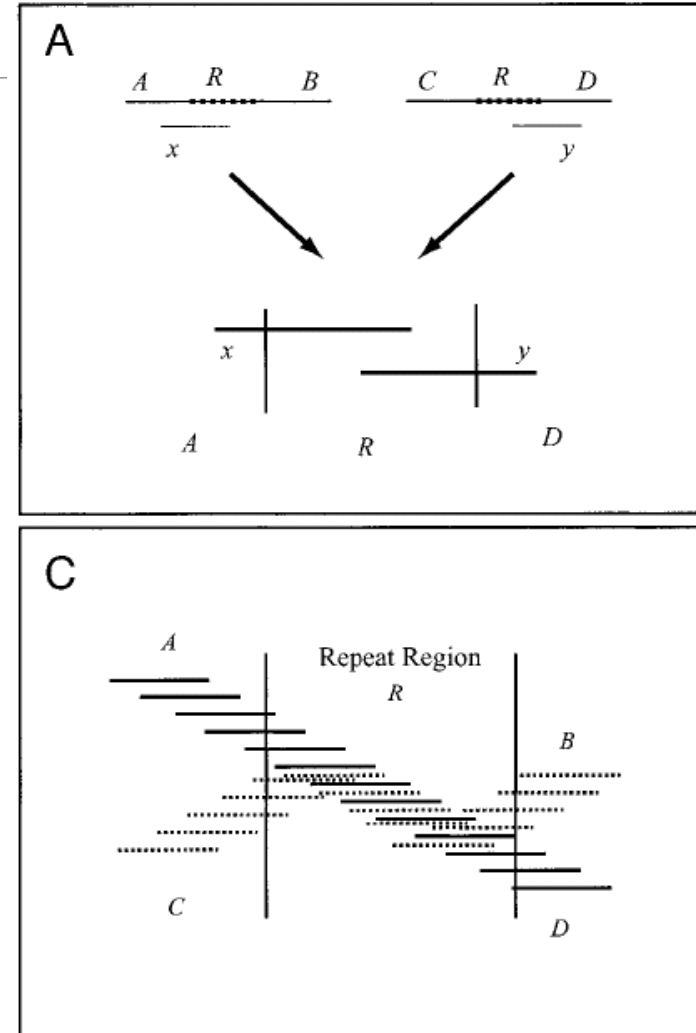
South Africa operates >14 sequencers in 9 facilities.

# Noncoding DNA



# ARACHNE assembler: The danger of assembling repeats


- The same repeat appears between genes A and B and between C and D.
- If these two repeats are treated as one sequence, neither upstream nor downstream genes will assemble correctly.



# Initial assemblers: Overlap-Layout-Consensus

Each read (a) is a vertex in overlap graph (b). Edges represent overlap.

Old assemblers sought *Hamiltonian Path*, visiting every vertex once. This is *NP-complete*.



# Making a “ $k$ -mers” catalog

---

- Create a sorted list of all  $k$  bp sequences along with their positions within reads.
- Exclude high-count  $k$ -mers as repeats.
- Overlapping reads will share  $k$ -mers unless:
  - Overlap is less than  $k$  in length
  - Base calling errors obfuscate overlap

# k-mer traversal yields contig

---

TGGTTTTGATTATTTGCTGGTTGCC  
GGTTTTGATTATTTGCTGGTTGCCA ← k-mers of 25bp  
GTTTTGATTATTTGCTGGTTGCCAA  
TTTTGATTATTTGCTGGTTGCCAAA  
TTTGATTATTTGCTGGTTGCCAAAC  
TTGATTATTTGCTGGTTGCCAAACA  
TGATTATTTGCTGGTTGCCAAACAT  
GATTATTTGCTGGTTGCCAAACATC  
TGGTTTTGATTATTTGCTGGTTGCCAAACATC ← Contig of 32bp

# *Eulerian path* through de Bruijn graph visits each edge once


sequence

**ATGGAAGTCGCGGAATC**

7mers

ATGGAAG  
TGGAAGT  
GGAAGTC  
GAAGTCG  
AAGTCGC  
AGTCGCG  
GTCGCGG  
TCGCGGA  
CGCGGAA  
GCGGAAT  
CGGAATC

de Bruijn graph






# Paired reads enable contig linking


A

If two separate contigs contain each end of a particular insert, those contigs are near each other.




B

A set of neighboring contigs is a **scaffold**.



# N50: how “chopped up” is my assembly?

Contig size, largest to smallest



- Sort contigs from big to small. Stack them.
- N50 is size of contig at half of length sum.

# FASTA “Database” Format

---

```
>ENSP00000396333.1 pep chromosome:GRCh38:1:154963677:154966490:-1
gene:ENSG00000160691.18 transcript:ENST00000444664.5
gene_symbol:SHC1 description:SHC adaptor protein 1 [Source:HGNC
Symbol;Acc:HGNC:10840]
XDEEEEEPPDHQYYNDFPGKEPPLGGVDMRLREGAAPGAARPTAPNAQTPSHLGATLPV
GQPVGGDPEVRKQMPPPPPCPGRELFDDPSYVNVQNLDKARQAVGGAGPPNPAINGSAPR
DLFDMKPFEDALRVPPPPQSVSMAEQLRGEPWFHGKLSRREA EALLQLNGDFLVRESTTT
PGQYVLTGLQSGQPKHLLLVDP EGV RWGFAMLPKLF LNSRAQVIRLPRPPRVLGLQARTT
MPSLHIF FCTVYTLLRHANFLQVKKG VYSSQLHSFRADVAF AF SHFTDLSIPTTVSF
```

- Line 1: ‘>’ + accession + whitespace + description
- Following lines: sequence

# Closing thoughts

---

- Phred helps us to evaluate each basecall.
- Mapping aligns reads to existing annotation; assembly builds long sequences from reads.
- Mapping speeds increase dramatically when a BWT-based index is available.
- *De novo* assembly starts with a  $k$ -mer catalog and seeks a Eulerian path through it.