

Longest Paths in DAGs


Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Mathematics for Data Science 1
Week 11


Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Directed Acyclic Graphs

- $G = (V, E)$, a directed graph without directed cycles
- **Topological sorting**
 - Enumerate $V = \{0, 1, \dots, n-1\}$ such that for any $(i, j) \in E$, i appears before j
 - Feasible schedule
- Imagine the DAG represents prerequisites between courses
- Each course takes a semester
- Minimum number of semesters to complete the programme?




Longest Path

- Find the longest path in a DAG




Longest Path

- Find the longest path in a DAG
- If $\text{indegree}(i) = 0$,
longest-path-to(i) = 0




Longest Path

- Find the longest path in a DAG
- If $\text{indegree}(i) = 0$,
 $\text{longest-path-to}(i) = 0$
- If $\text{indegree}(i) > 0$, longest path to i is
1 more than longest path to its
incoming neighbours
 $\text{longest-path-to}(i) =$
 $1 + \max\{\text{longest-path-to}(j) \mid (j, i) \in E\}$




Longest Path

- $\text{longest-path-to}(i) = 1 + \max\{\text{longest-path-to}(j) \mid (j, i) \in E\}$




Longest Path

- $\text{longest-path-to}(i) = 1 + \max\{\text{longest-path-to}(j) \mid (j, i) \in E\}$
- To compute $\text{longest-path-to}(i)$, need $\text{longest-path-to}(k)$, for each incoming neighbour k




Longest Path

- $\text{longest-path-to}(i) = 1 + \max\{\text{longest-path-to}(j) \mid (j, i) \in E\}$
- To compute $\text{longest-path-to}(i)$, need $\text{longest-path-to}(k)$, for each incoming neighbour k
- If graph is topologically sorted, k is listed before i




Longest Path

- $\text{longest-path-to}(i) = 1 + \max\{\text{longest-path-to}(j) \mid (j, i) \in E\}$
- To compute $\text{longest-path-to}(i)$, need $\text{longest-path-to}(k)$, for each incoming neighbour k
- If graph is topologically sorted, k is listed before i
- Hence compute $\text{longest-path-to}()$ in topological order




Longest Path

- Let i_0, i_1, \dots, i_{n-1} be a topological ordering of V




Longest Path

- Let i_0, i_1, \dots, i_{n-1} be a topological ordering of V
- All neighbours of i_k appear before it in this list




Longest Path

- Let i_0, i_1, \dots, i_{n-1} be a topological ordering of V
- All neighbours of i_k appear before it in this list
- From left to right, compute longest-path-to(i_k) as
 $1 + \max\{\text{longest-path-to}(i_j) \mid (i_j, i_k) \in E\}$




Longest Path

- Let i_0, i_1, \dots, i_{n-1} be a topological ordering of V
- All neighbours of i_k appear before it in this list
- From left to right, compute $\text{longest-path-to}(i_k)$ as $1 + \max\{\text{longest-path-to}(i_j) \mid (i_j, i_k) \in E\}$
- Overlap this computation with topological sorting




Longest path algorithm

- Compute **indegree** of each vertex




Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest-path-to** to 0 for all vertices




Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG



Longest path algorithm


- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path



Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**




Topological order	1	0
Longest path to	0	0

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, Longest path




Topological order	1	0
Longest path to	0	0

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, Longest path




Topological order	1	0	3
Longest path to	0	0	1

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, Longest path




Topological order	1	0	3
Longest path to	0	0	1

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed


Indegree, **Longest path**



Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**




Topological order	1	0	3	2
Longest path to	0	0	1	1

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**




Topological order	1	0	3	2	5
Longest path to	0	0	1	1	2

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**




Topological order	1	0	3	2	5
Longest path to	0	0	1	1	2

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**




Topological order	1	0	3	2	5	6
Longest path to	0	0	1	1	2	3

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**



Topological order	1	0	3	2	5	6
Longest path to	0	0	1	1	2	3

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**

1
7
4

Topological order	1	0	3	2	5	6	4
Longest path to	0	0	1	1	2	3	1

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**

0
7
4

Topological order	1	0	3	2	5	6	4
Longest path to	0	0	1	1	2	3	1

Longest path algorithm


- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed

Indegree, **Longest path**

Topological order	1	0	3	2	5	6	4	7
Longest path to	0	0	1	1	2	3	1	4

Longest path algorithm

- Compute **indegree** of each vertex
 - Scan each column of the adjacency matrix
- Initialize **longest path to** to 0 for all vertices
- List a vertex with indegree 0 and remove it from the DAG
- Update indegrees, longest path
- Repeat till all vertices are listed



Topological order	1	0	3	2	5	6	4	7
Longest path to	0	0	1	1	2	3	1	4

Summary

- Directed acyclic graphs are a natural way to represent dependencies
- Topological sort gives a feasible schedule that represents dependencies
- In parallel with topological sort, we can compute the longest path

Summary

- Directed acyclic graphs are a natural way to represent dependencies
- Topological sort gives a feasible schedule that represents dependencies
- In parallel with topological sort, we can compute the longest path
- Notion of longest path makes sense even for graphs with cycles
 - No repeated vertices in a path, so path has at most $n - 1$ edges

Summary

- Directed acyclic graphs are a natural way to represent dependencies
- Topological sort gives a feasible schedule that represents dependencies
- In parallel with topological sort, we can compute the longest path
- Notion of longest path makes sense even for graphs with cycles
 - No repeated vertices in a path, so path has at most $n - 1$ edges
- However, computing longest paths in arbitrary graphs is much harder than for DAGs
 - No better strategy known than exhaustively enumerating paths