

3D reconstruction from photographs

(Photogrammetry)

Paul Bourke

<http://paulbourke.net/EPFL2019/>



Outcomes

- A overview of the technology
- When it is appropriate to use it
- Expectations of the outputs
- Workflow if you were to use the technique

Why?

- Creating richer objects (compared to photographs) for recordings in archaeology and heritage.
- Create geometric models suitable for analysis, eg: in geology or geoscience.
- Create models for virtual environments.
- Create models for augmented reality.
- Assets for educational and public outreach.

History

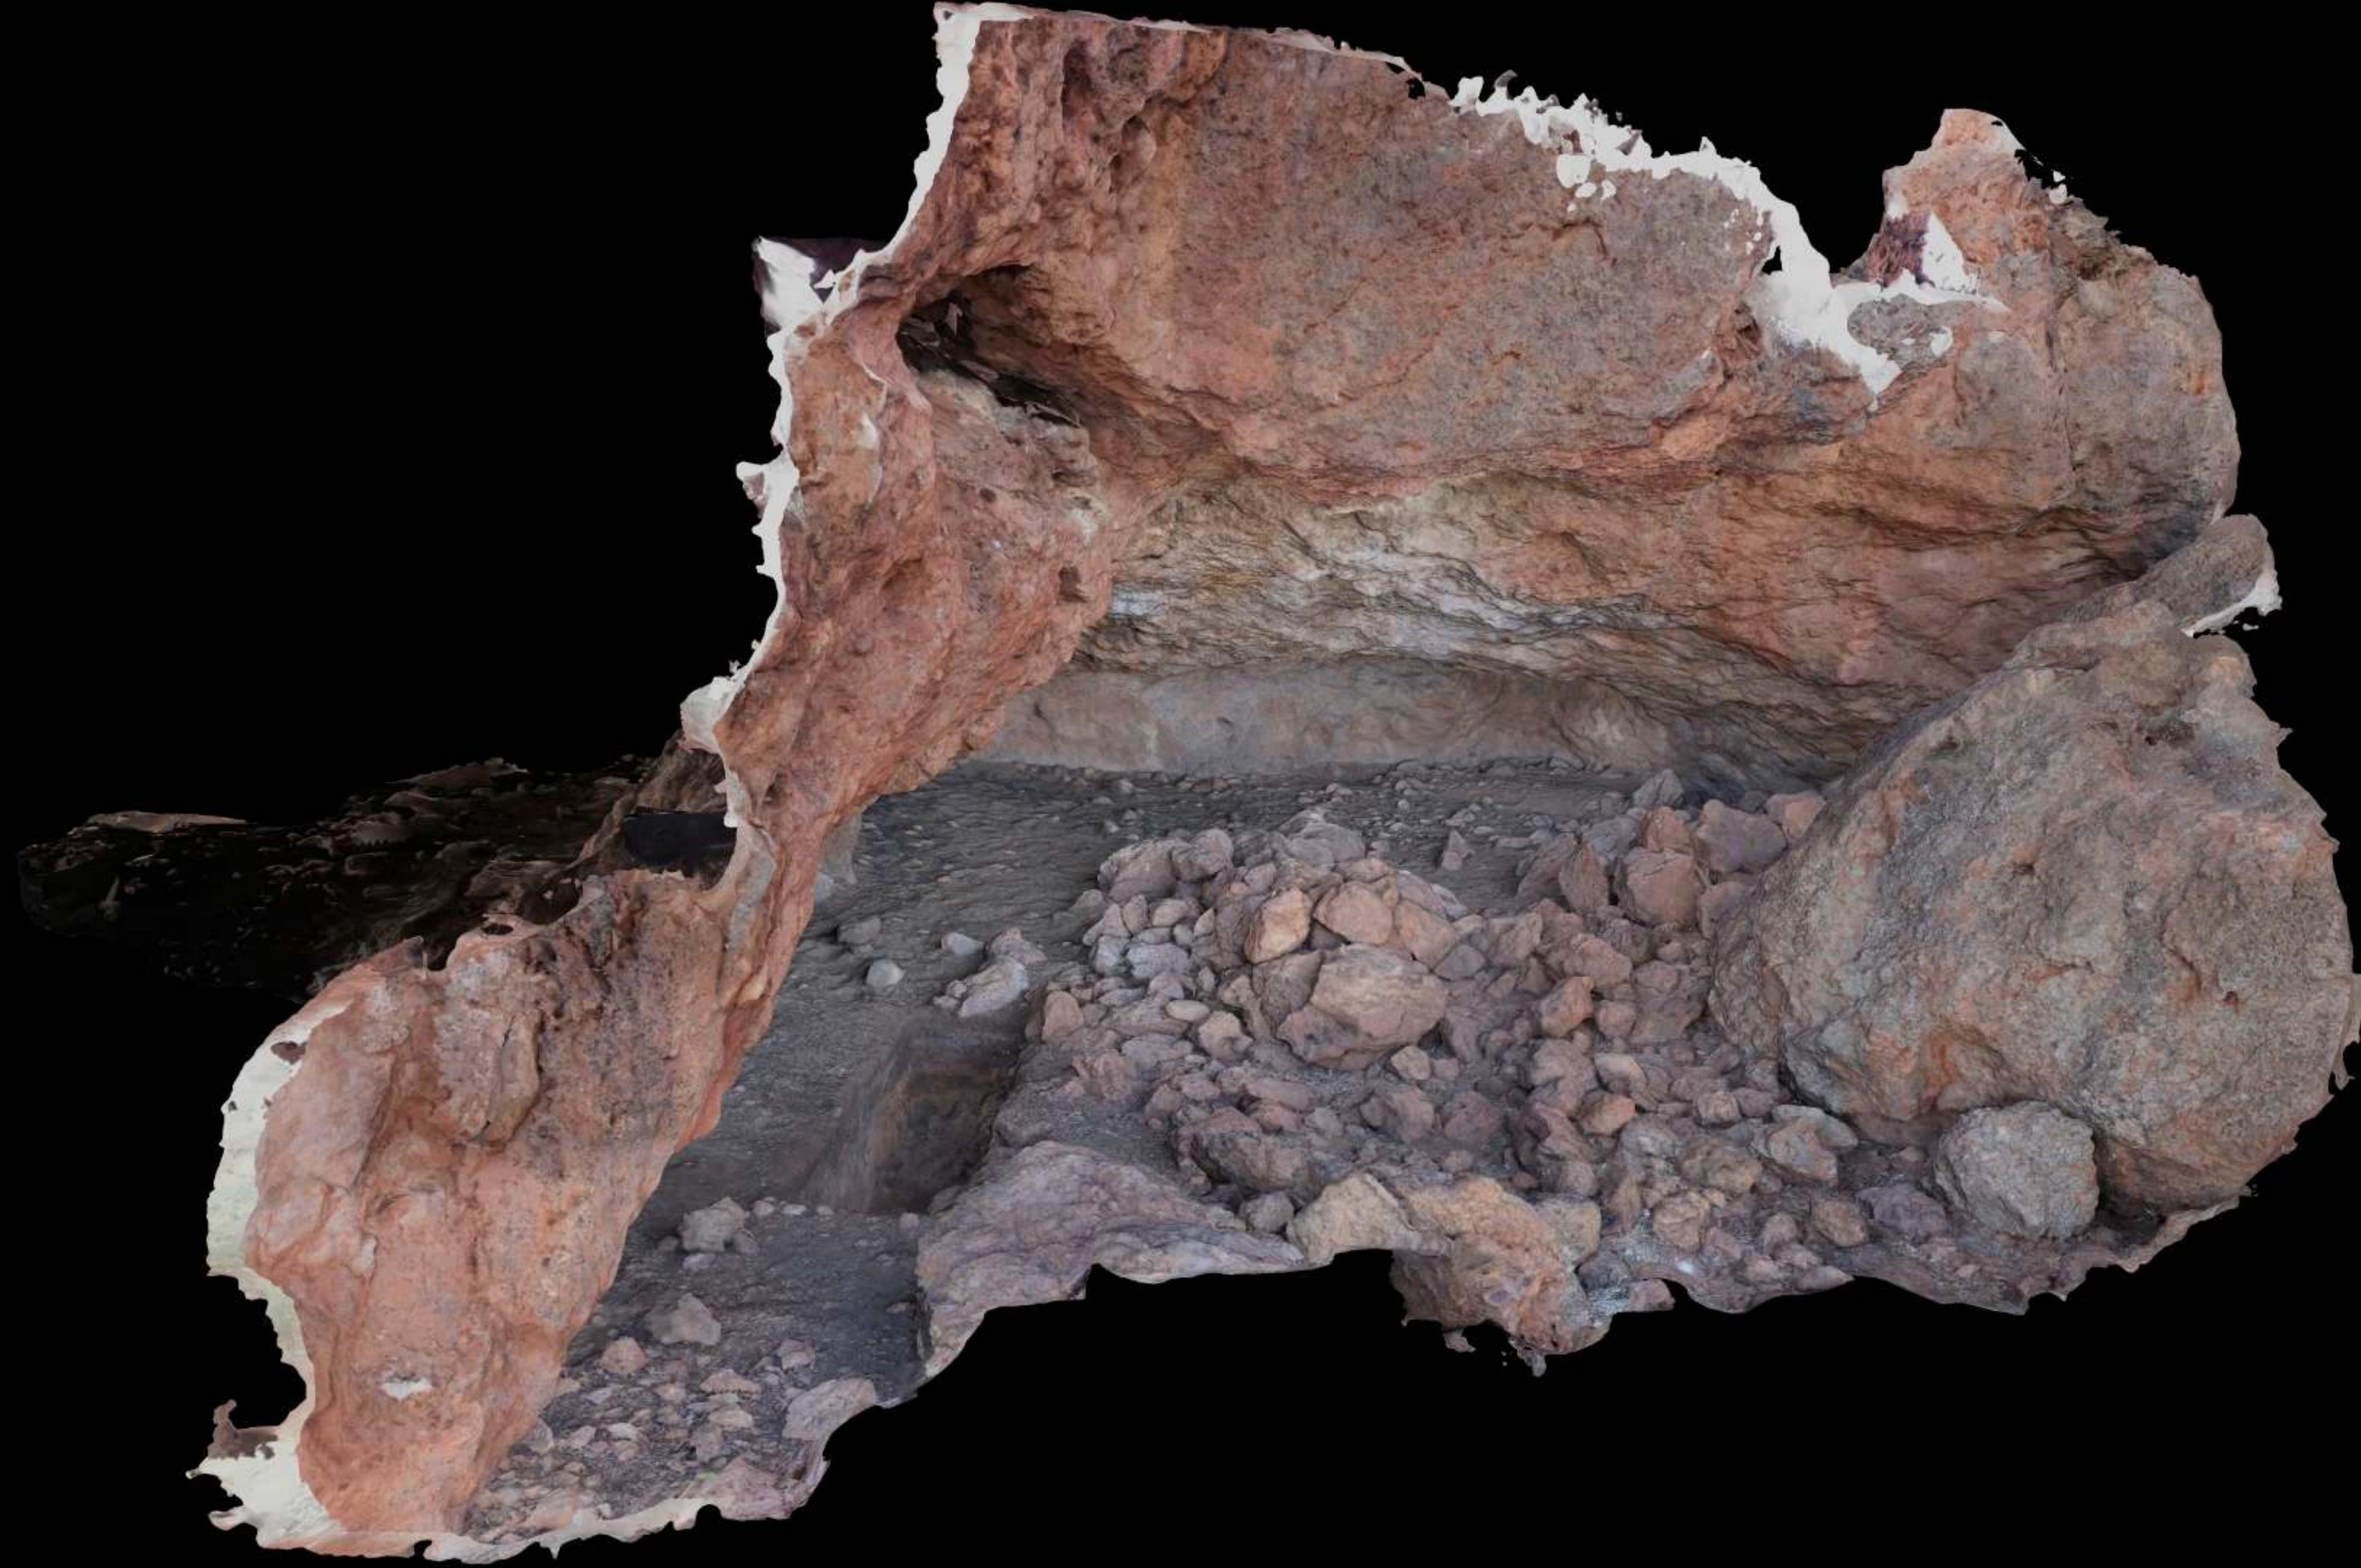
- Photogrammetry is the general term given to deriving geometric information from a series of photographs.
- Initially largely used for aerial surveys, deriving landscape models. Originally only used a stereoscopic pair, that is, just two photographs.
- More recently it has grown through the domain of machine vision, for example: deriving a 3D model of a robots environment.
- Big step forward was the development of SfM algorithms: structure from motion. This generally solves the camera parameters and generation of a 3D point cloud.
- Most common implementation is called Bundler: “bundle adjustment algorithm allows the reconstruction of the 3D geometry of the scene by optimizing the 3D location of key points, the location/orientation of the camera, and its intrinsic parameters”.

Order of business

- Application examples
- When
- Hardware
- Photographic technique
- Software workflow
- Other topics
- Post production
- Worked example
- Your turn

Application examples

- Heritage
- Archaeology
- Forensics
- Medical training
- Replication, 3D printing





Weld range rock shelter



Mowunjum rock art



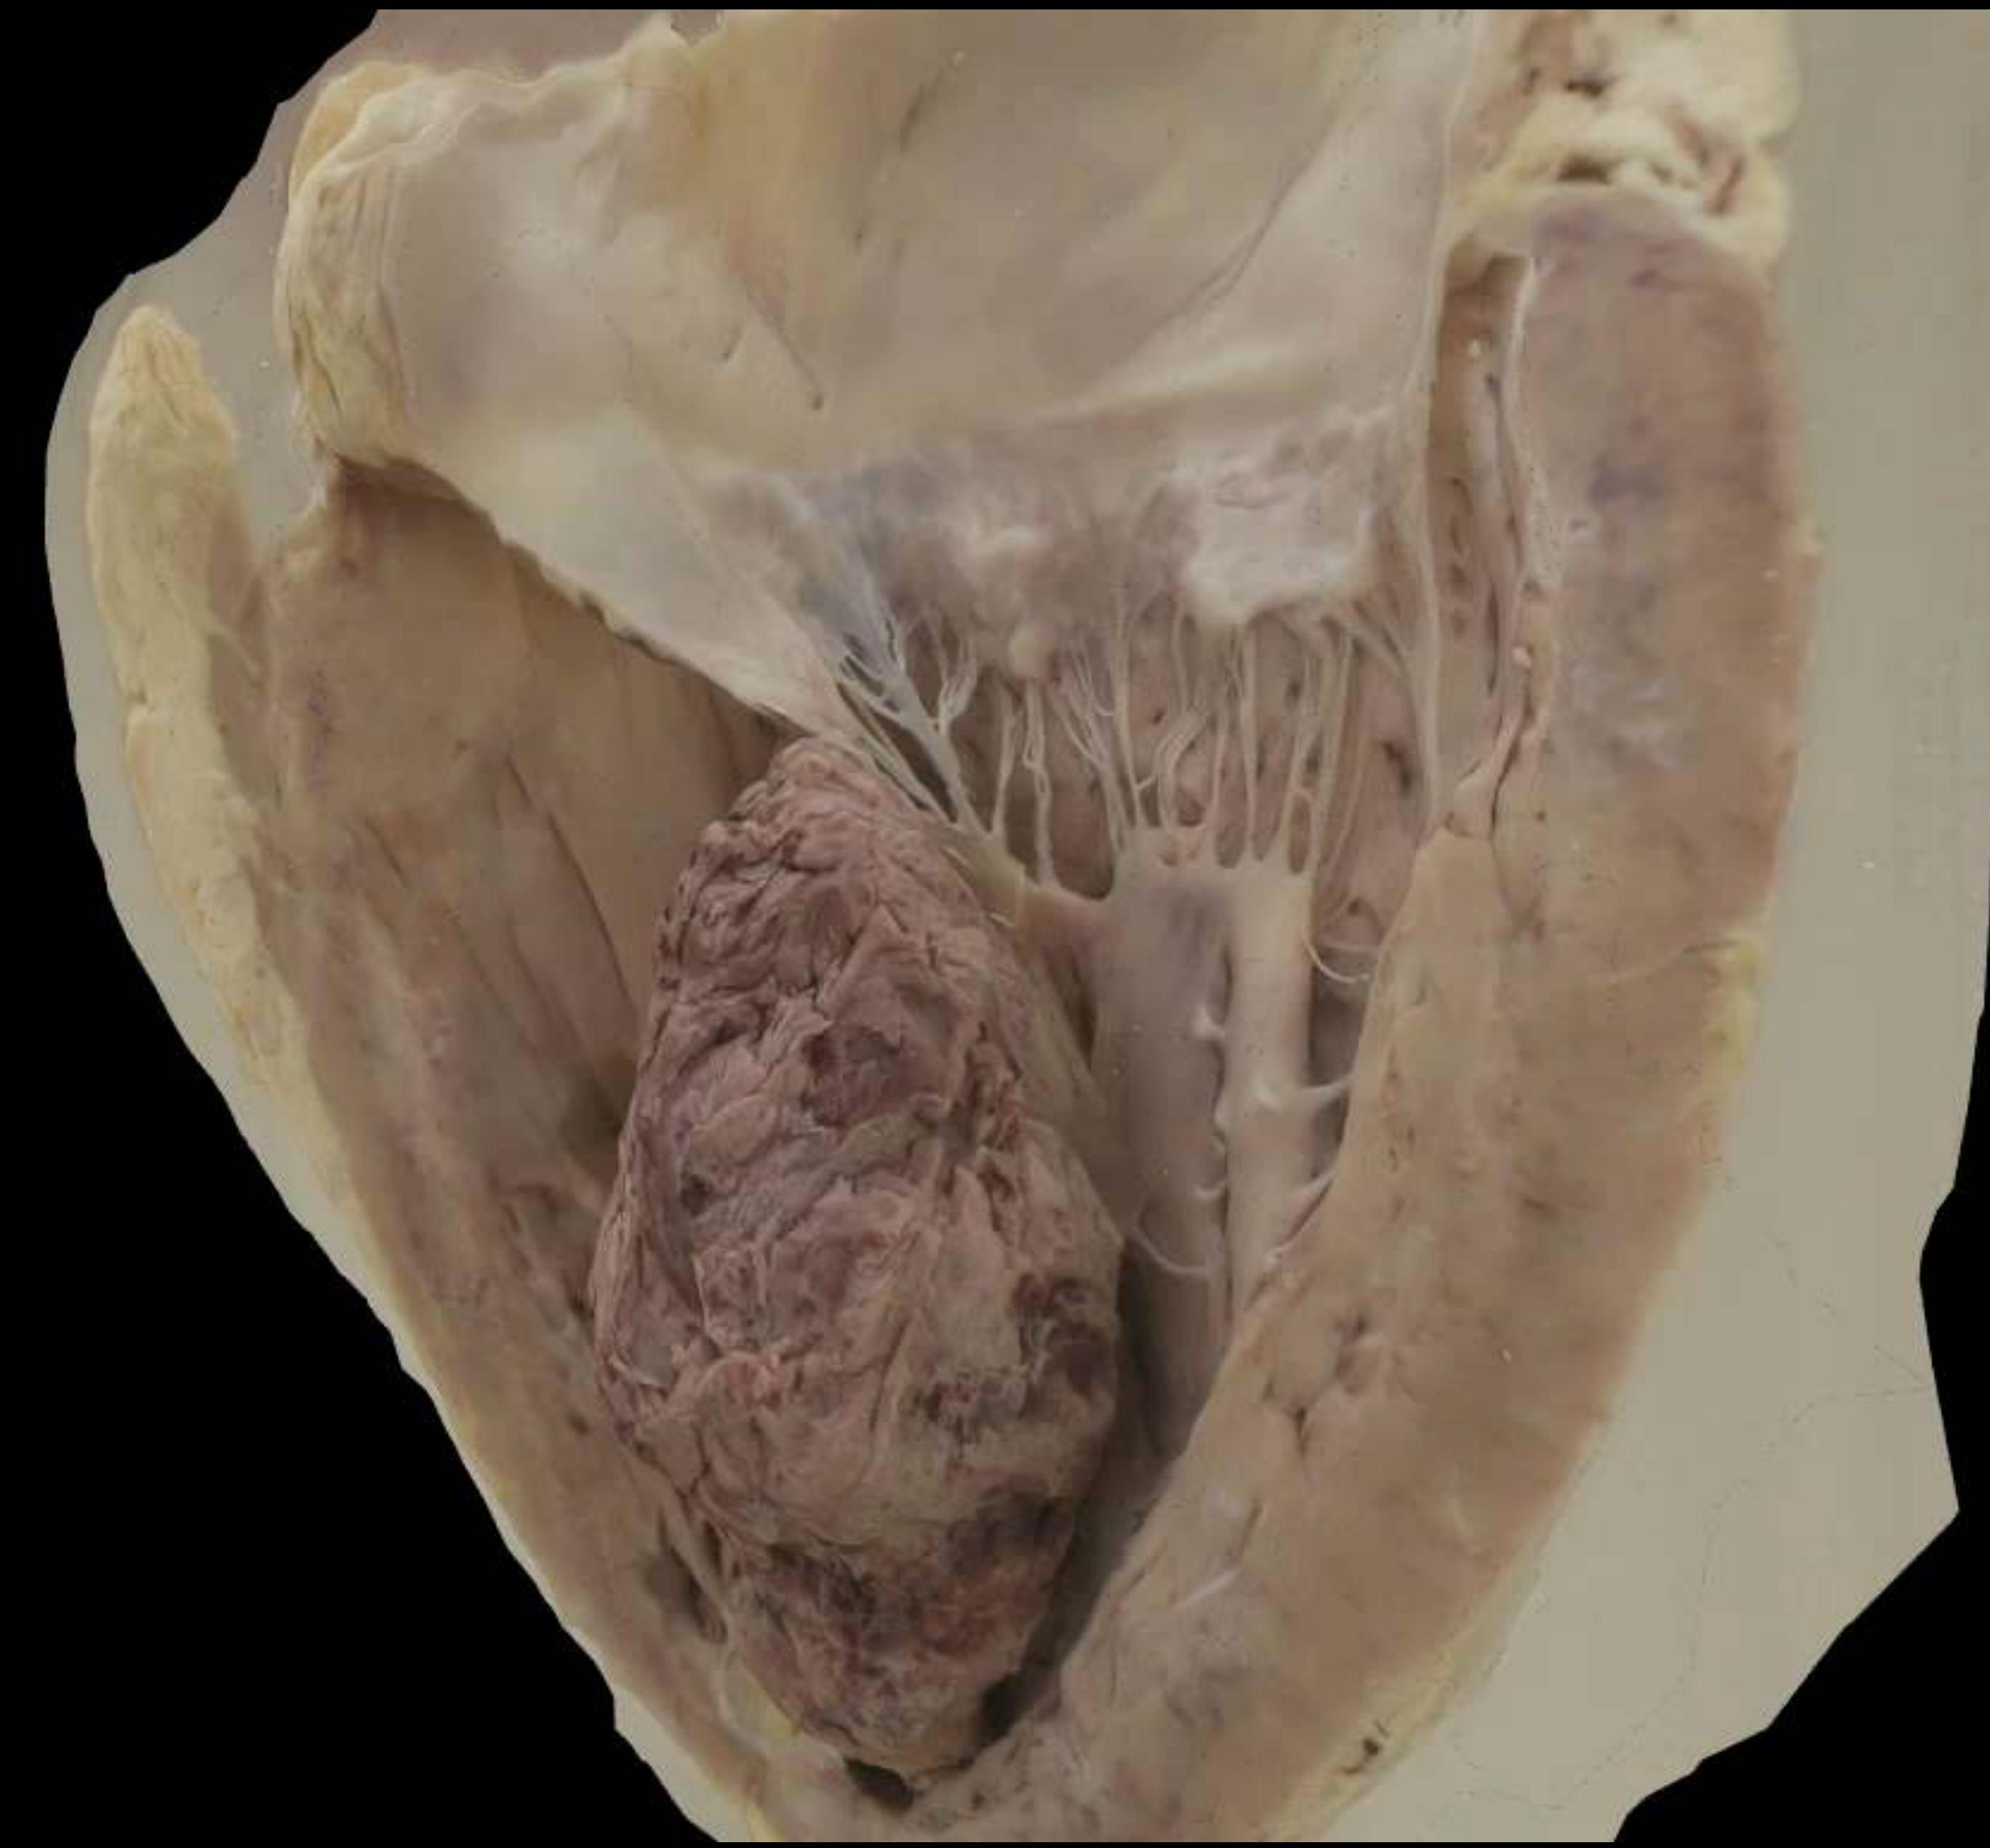
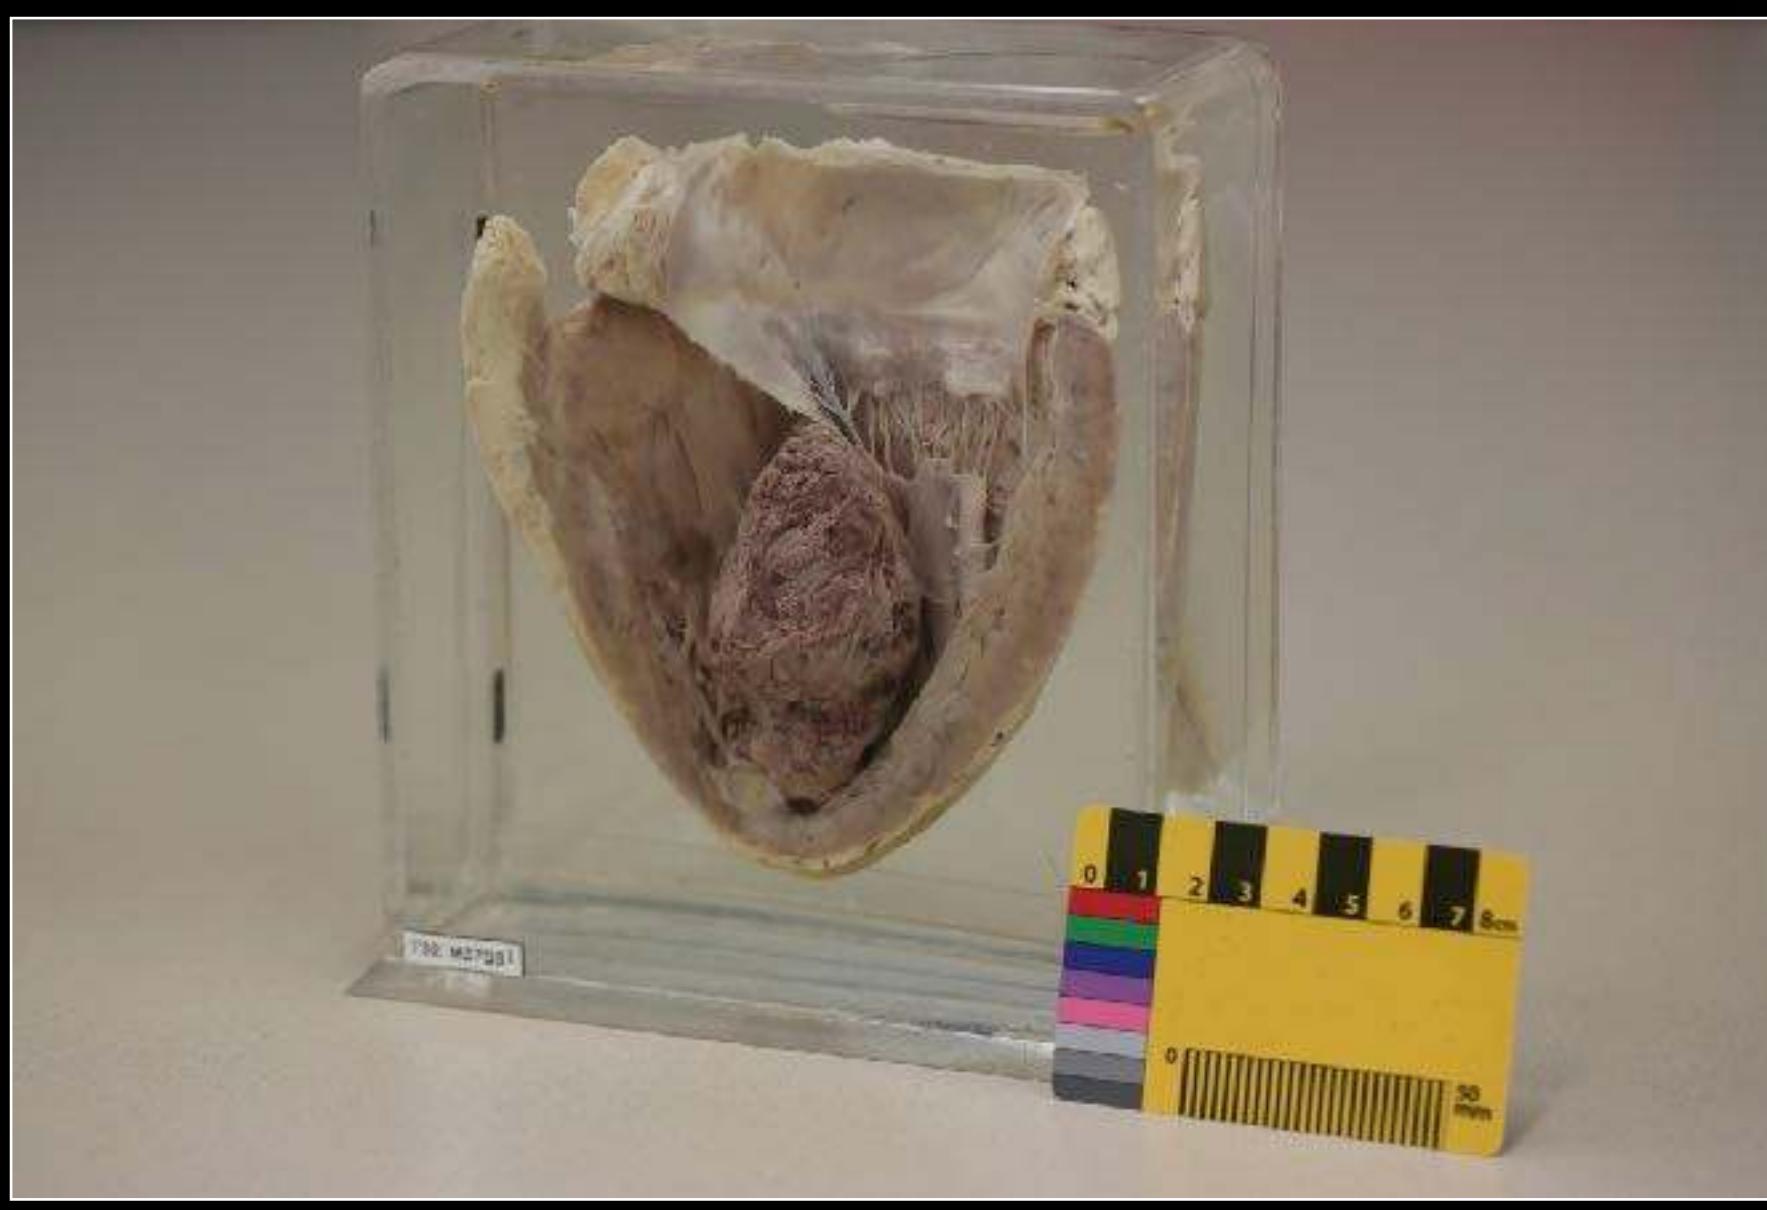




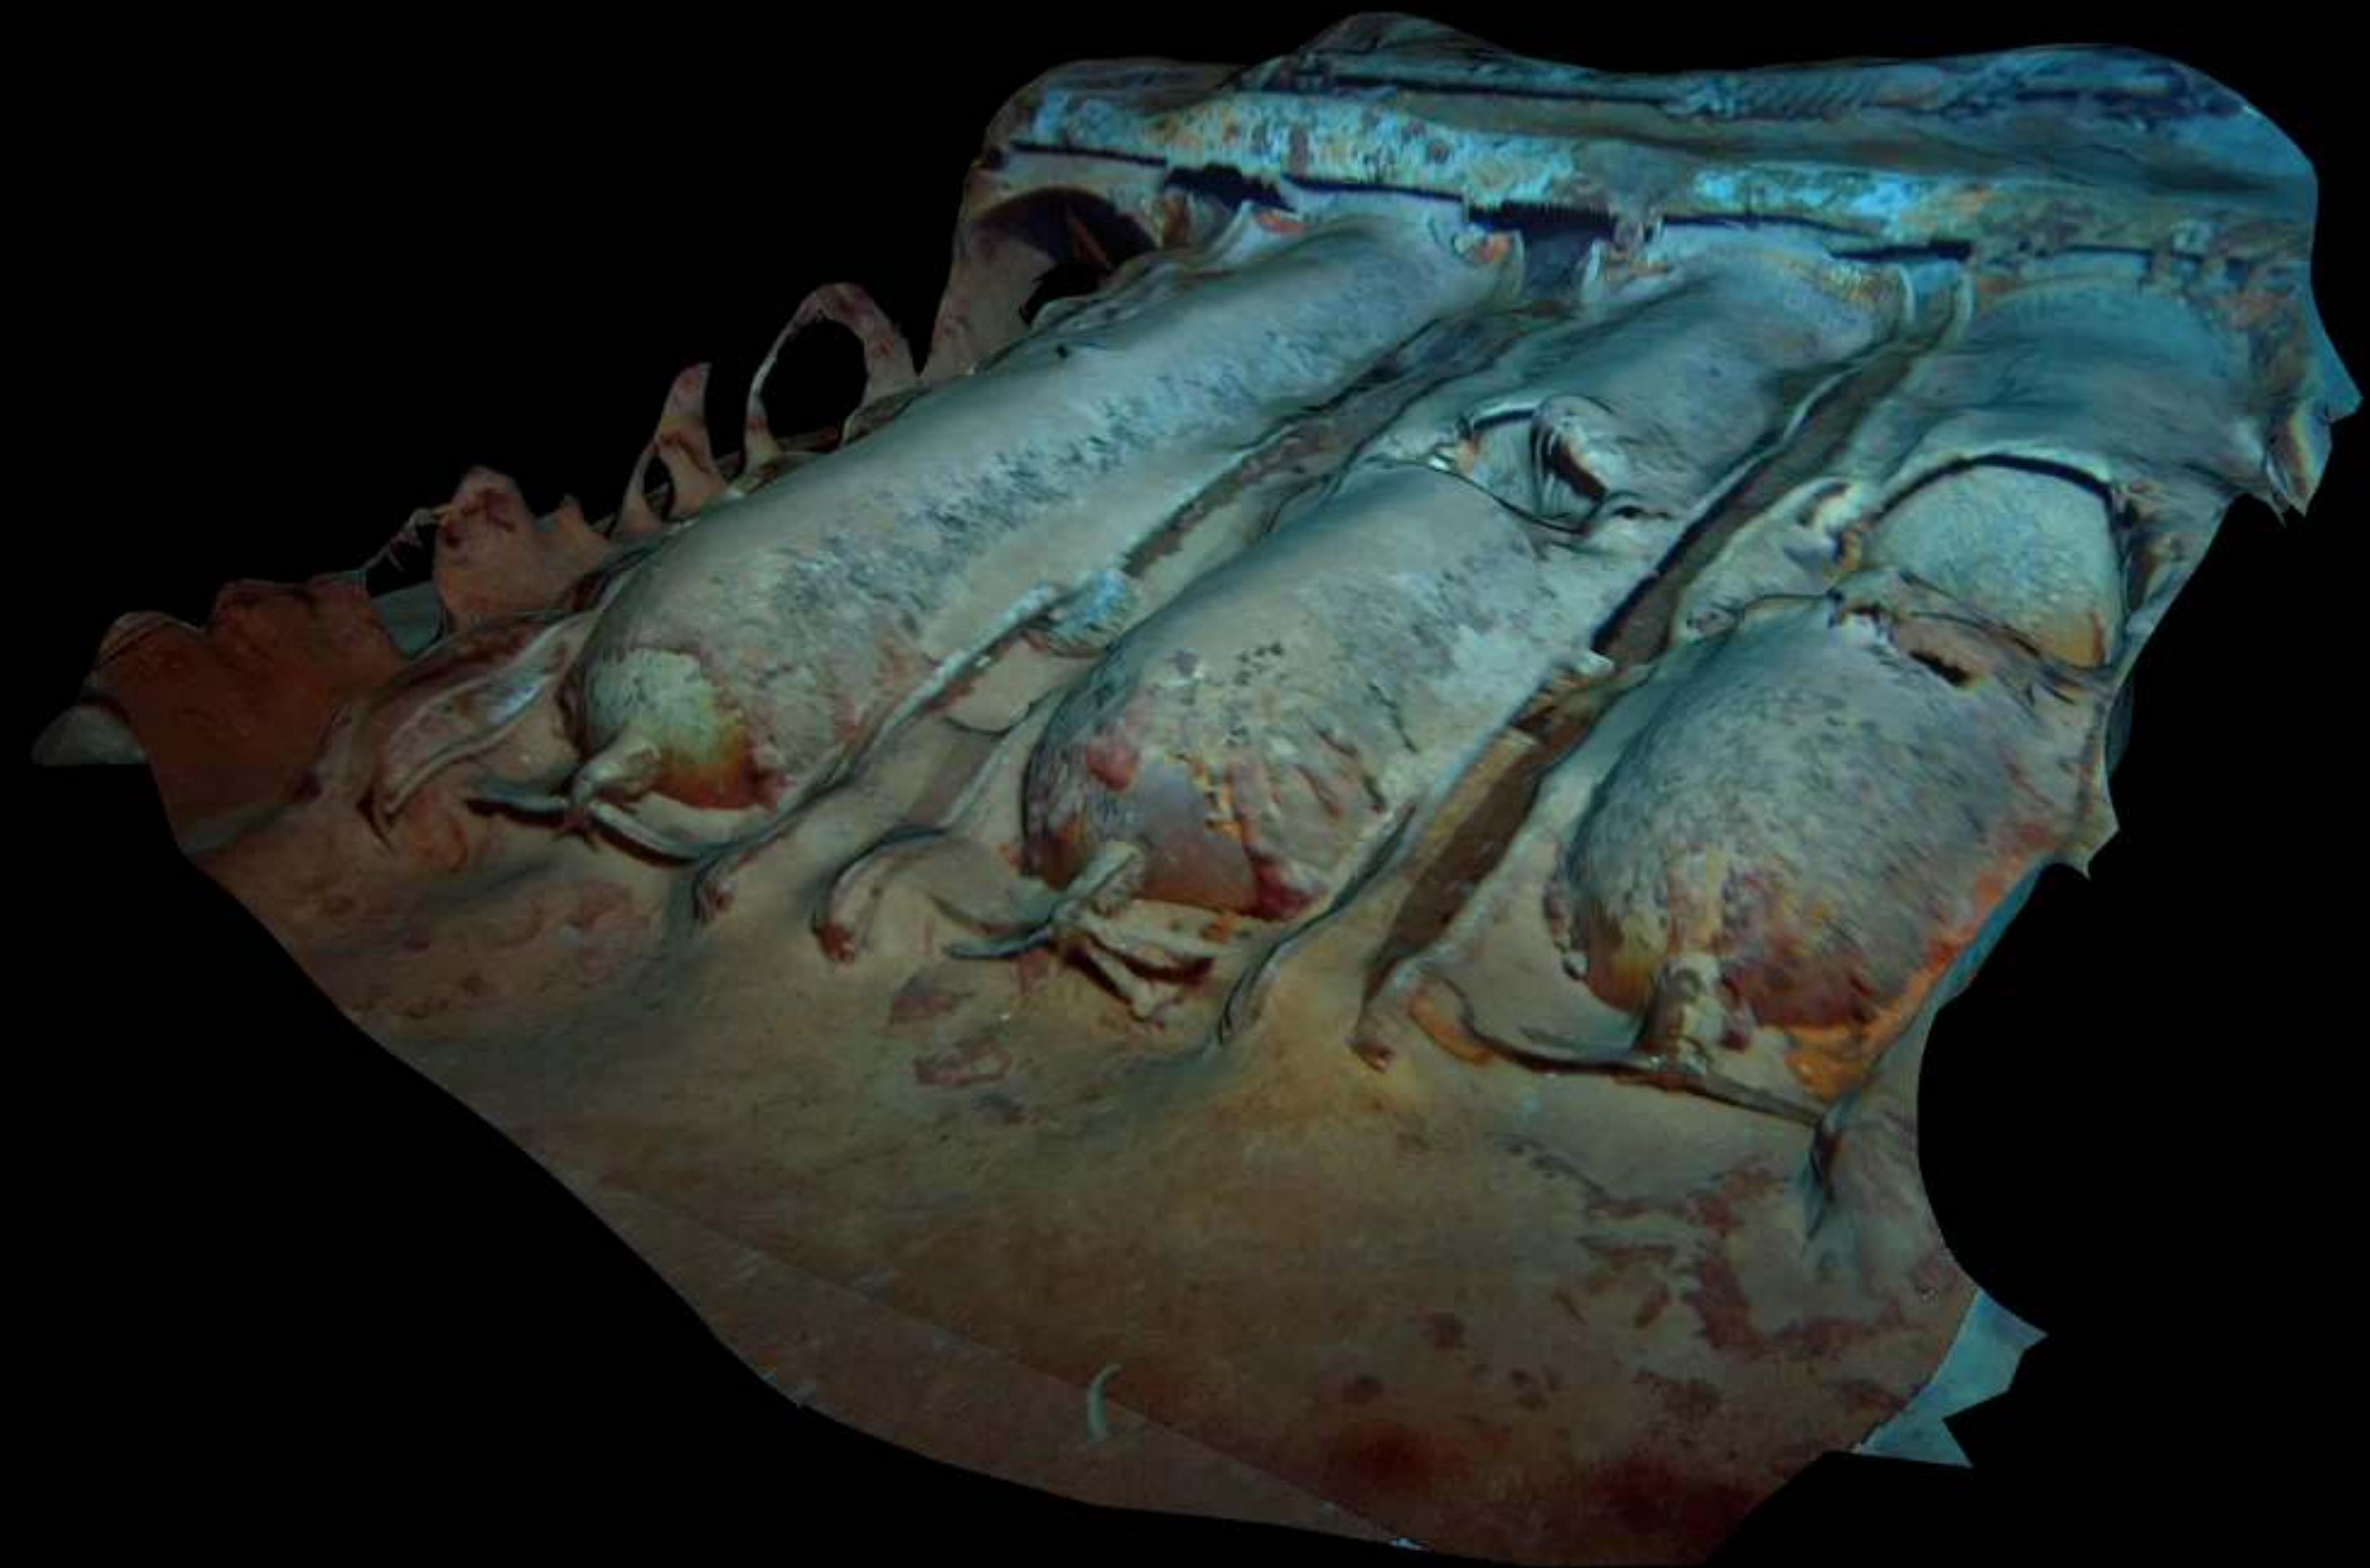
Beacon Island, Batavia shipwreck site



Dragon gardens in Hong Kong



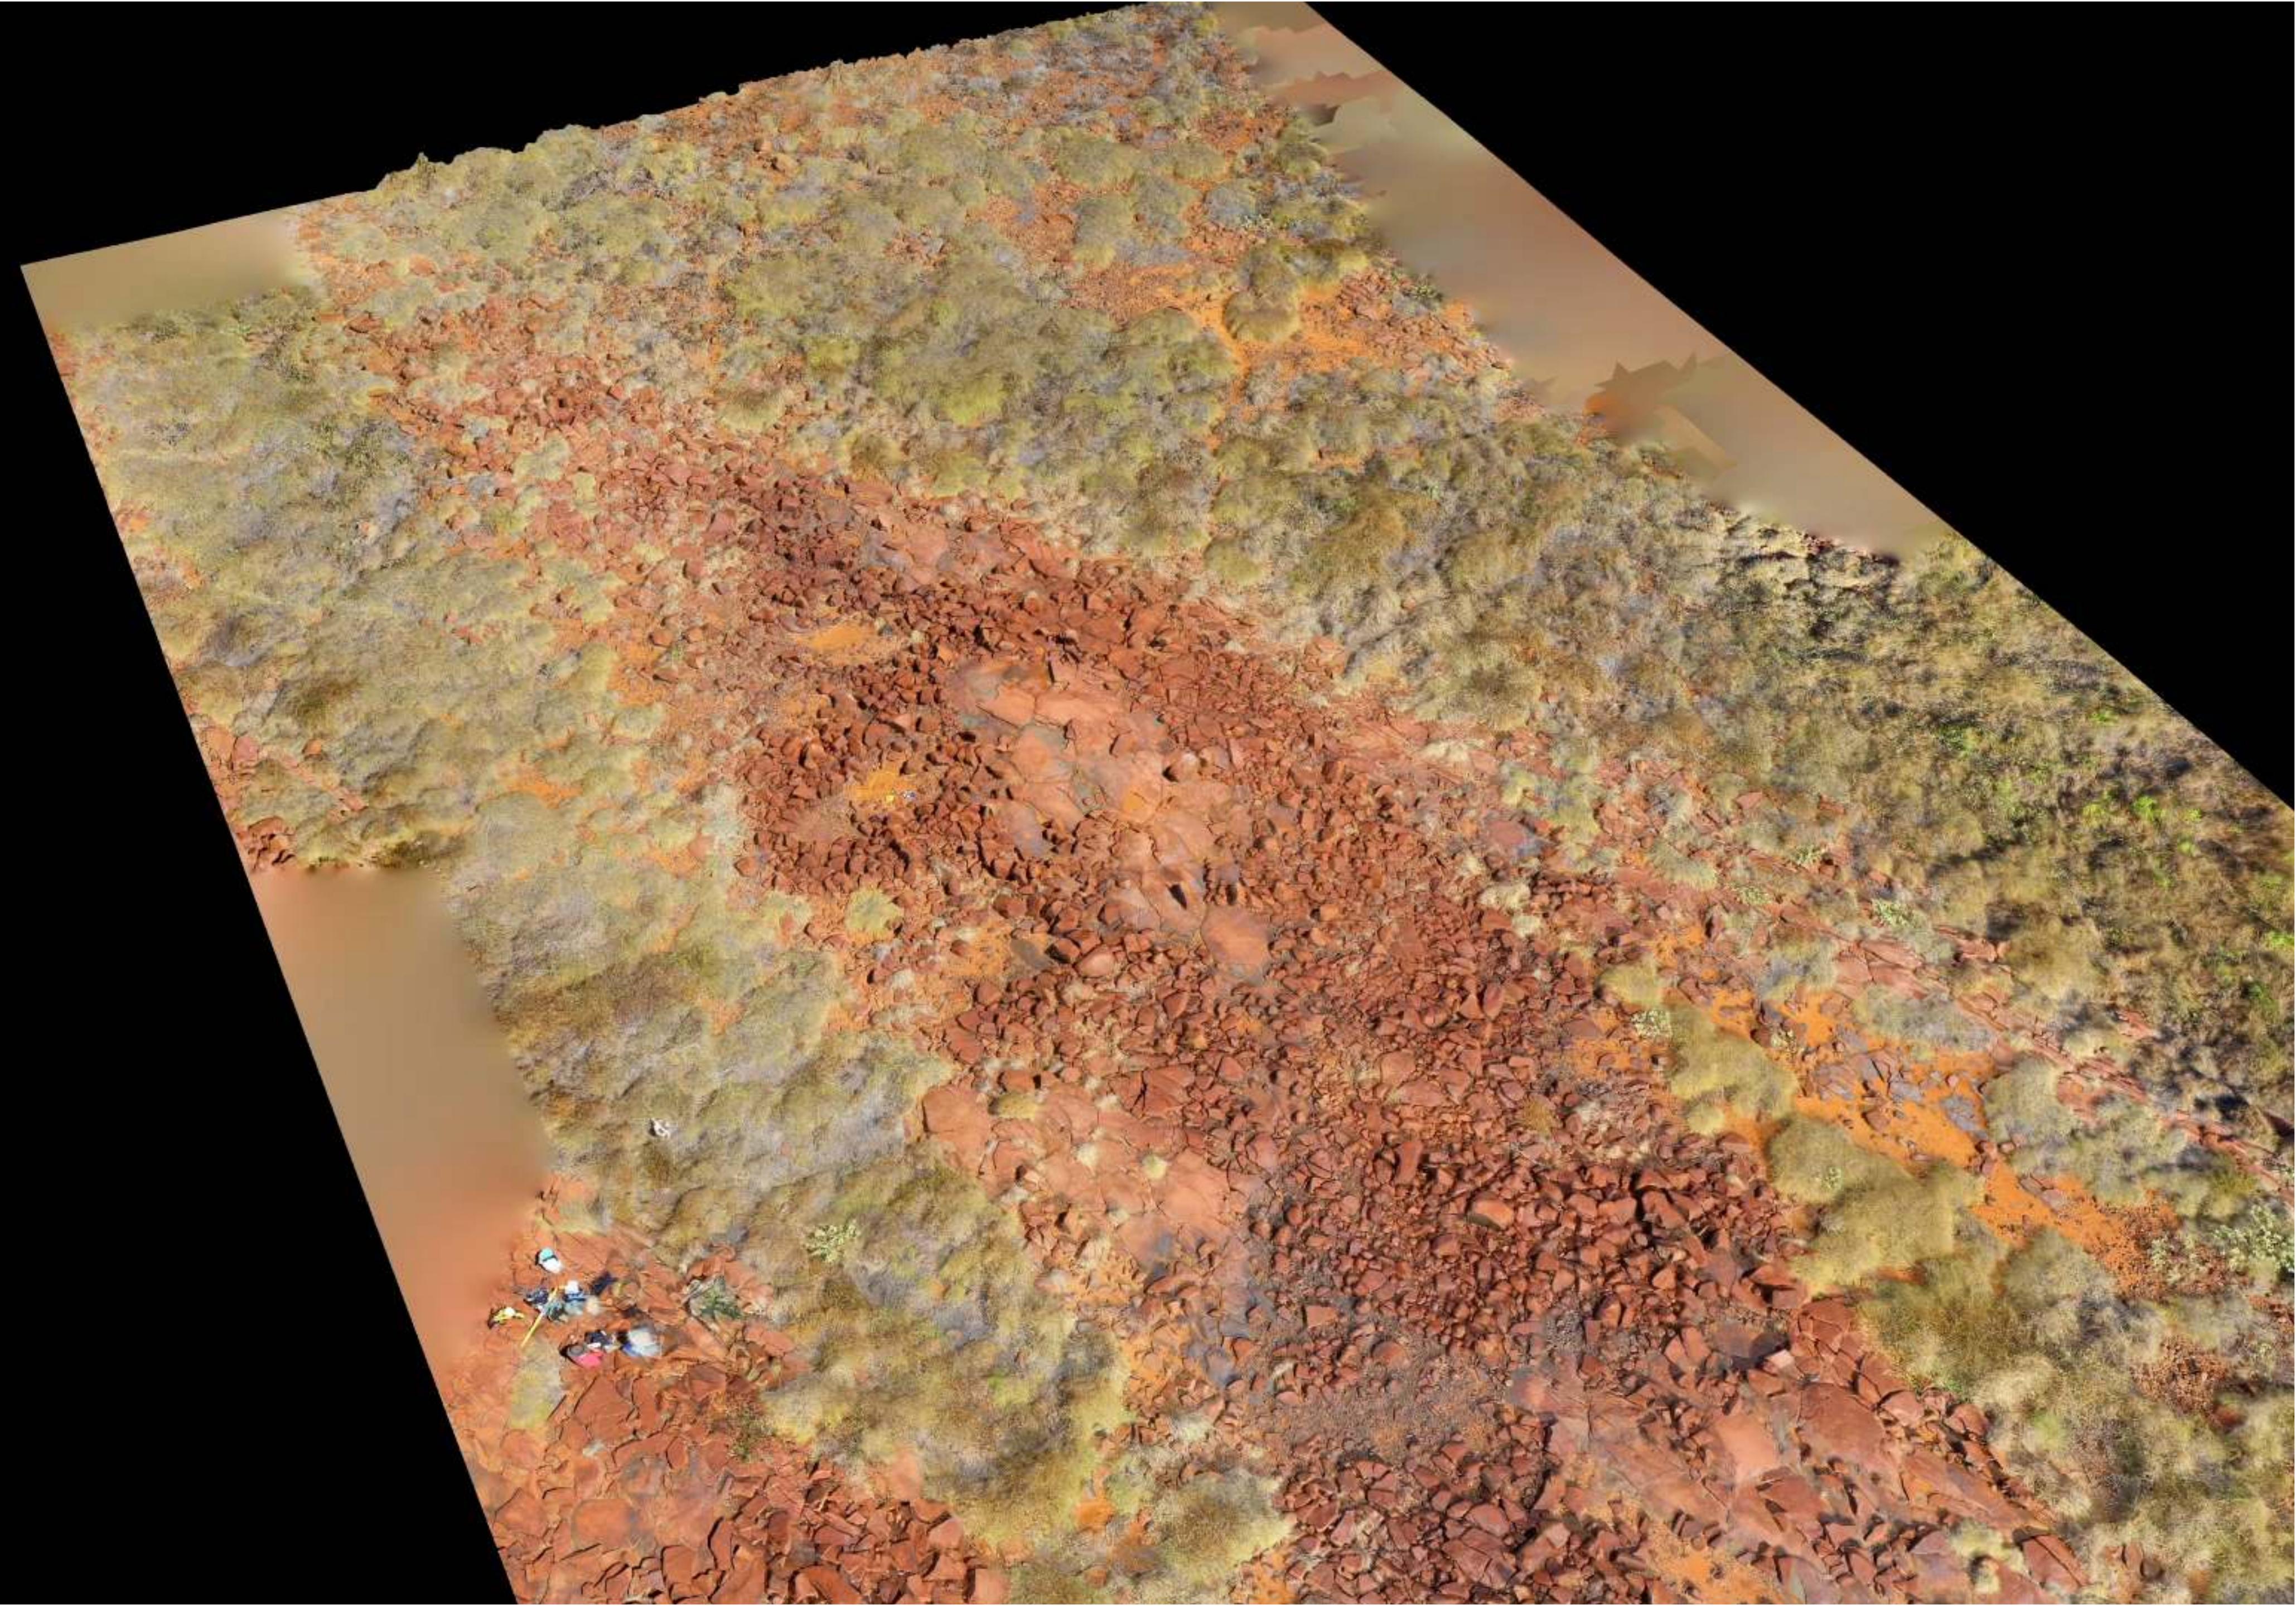




HMAS Sydney







When

- Best suited to organic shapes that are otherwise hard to model.
- Provides an unbiased model compared to human modeller.
- Requires feature point extraction (see later), so well suited to texture rich and noisy surfaces.
- Some models the technique is less suited to
 - Plain flat surfaces
 - Highly reflective or mirror surfaces
- Not suited to extremely thin objects
- Not suited to objects that are moving, unless using a multiple camera rig.

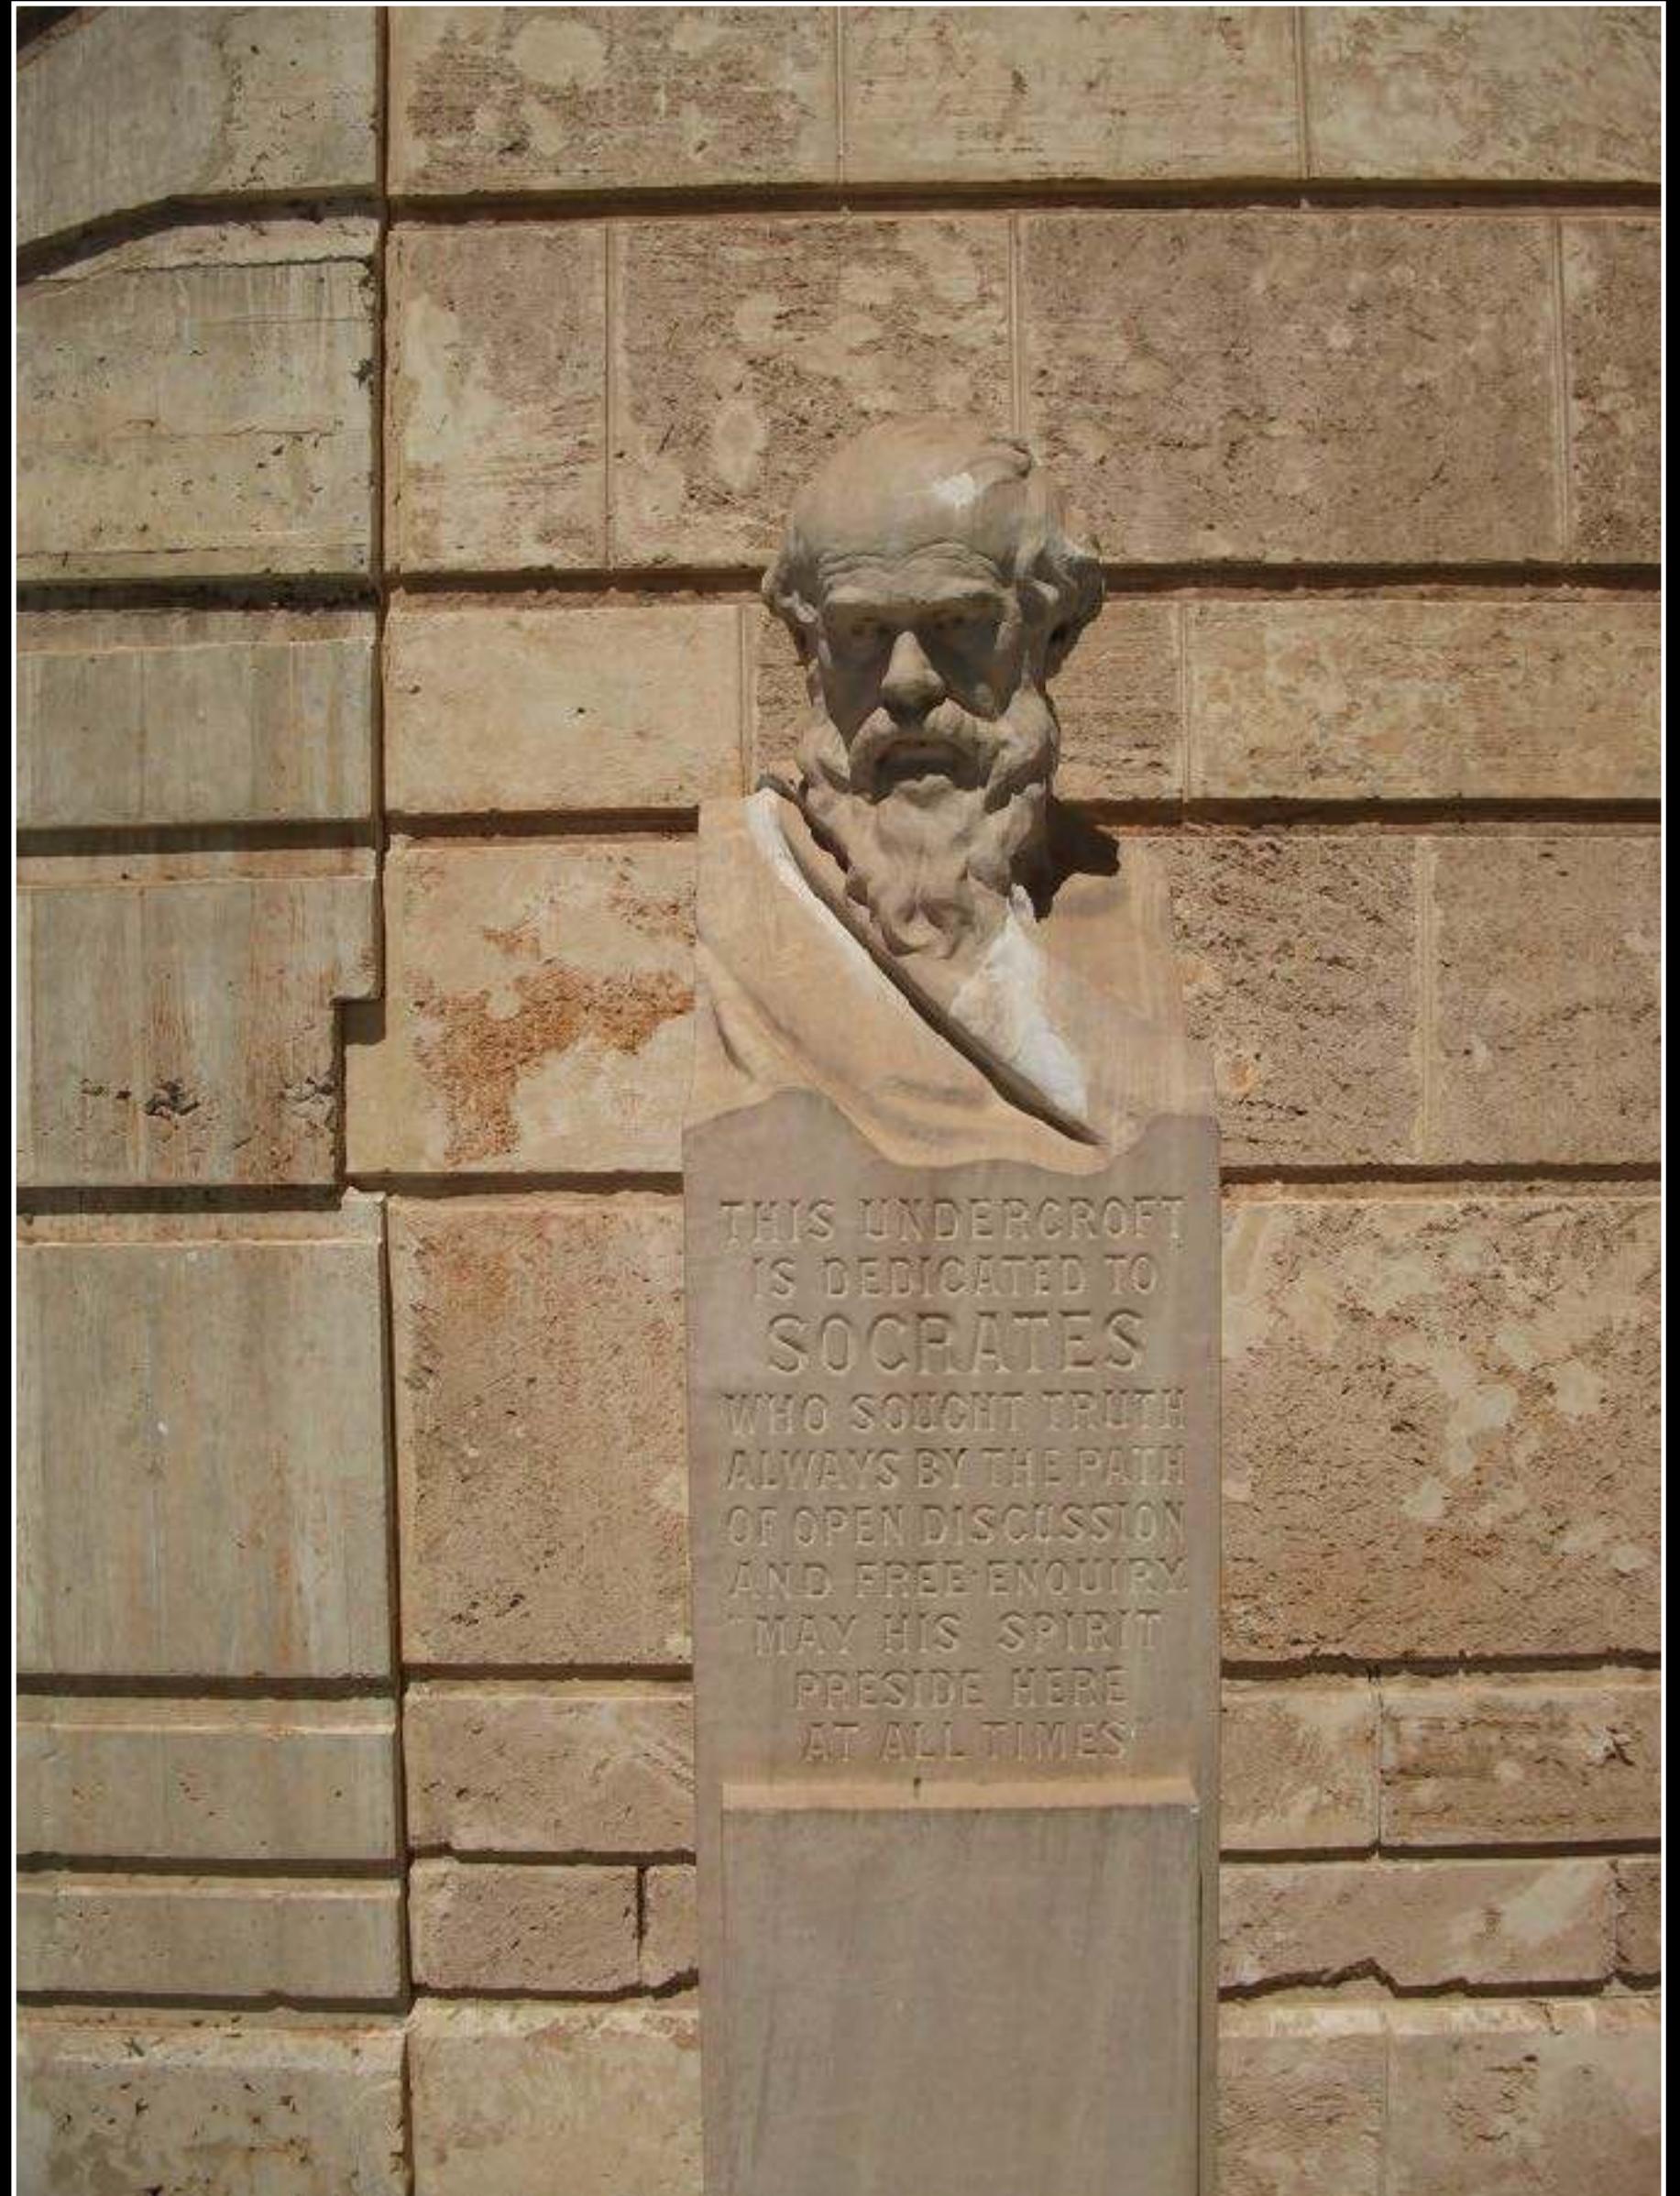
Yes



No



Yes



No



Yes



No





Beacon Island



The University of Western Australia

A small change in camera position results in a large difference in visible objects.



Hardware

- Camera.
- Prime lens, fixed focal point.

Generally better results for better quality camera and optics.

- Tripod for long exposures in low light.
- Possibly a turntable for small objects.

Other methods

- Laser scanning, return of flight. (LIDAR - light detection and ranging)
- Depth cameras, also typically return of flight but different wavelengths.
eg: Primesense (eg: Kinect)
- Structured light.
eg: Artec Scanner
- Light field cameras.

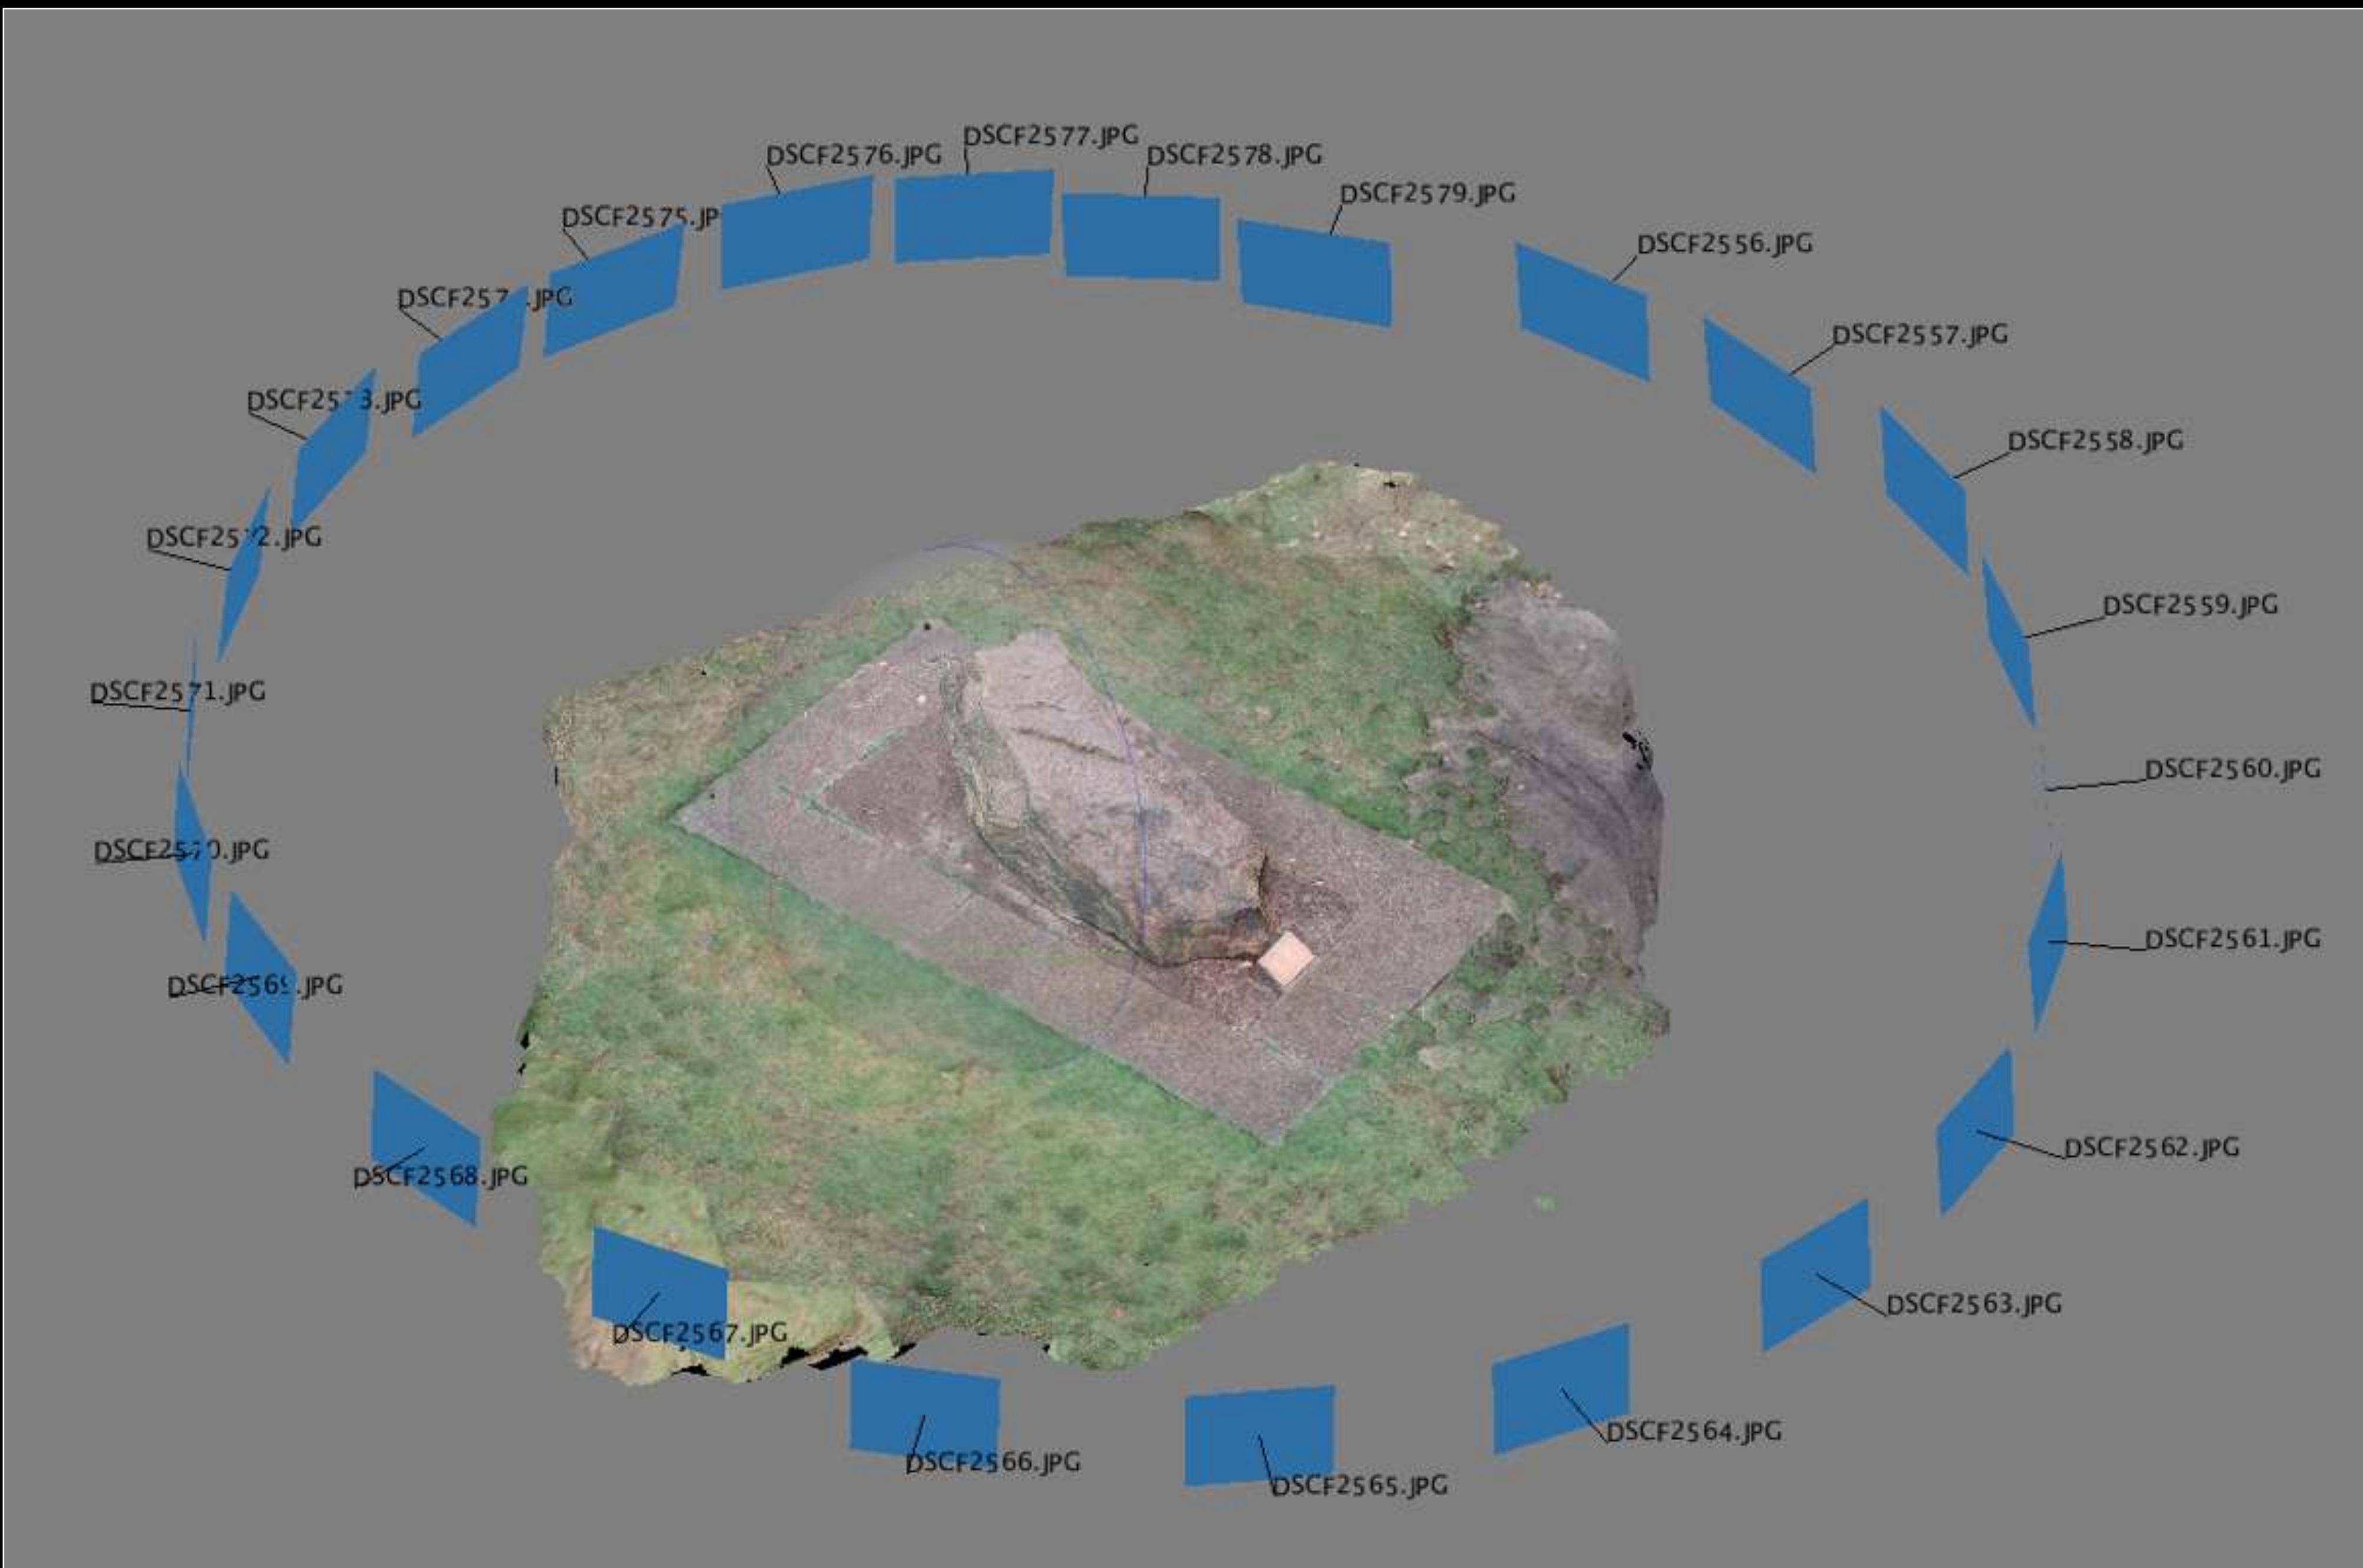
There is no winning technology. Each has limitations and applications where they may be better than other techniques.



Photographic technique

- Sharp in focus photographs with good depth of focus.
 - High resolution photographs.
 - Use manual settings for consistent textures: white point, ISO, exposure time, aperture.
 - Photograph all visible parts of the model (One cannot reconstruct what is not photographed!) + even spread of photographs.
 - Every camera position different.
Opposite to panoramic photography where one tries to photograph from the exact same position.
- Overall aim: any point on the model should be visible from at least 3 photographs, preferably more.
- Different strategies depending on the type of object.
 - Camera orientation does not matter.





DSCF2571.JPG

DSCF2570.JPG

DSCF2569.JPG

DSCF2568.JPG

DSCF2567.JPG

DSCF2566.JPG

DSCF2565.JPG

DSCF2564.JPG

DSCF2563.JPG

DSCF2562.JPG

DSCF2560.JPG

DSCF2561.JPG

DSCF2572.JPG

DSCF2573.JPG

DSCF2575.JPG

DSCF2577.JPG

DSCF2576.JPG

DSCF2578.JPG

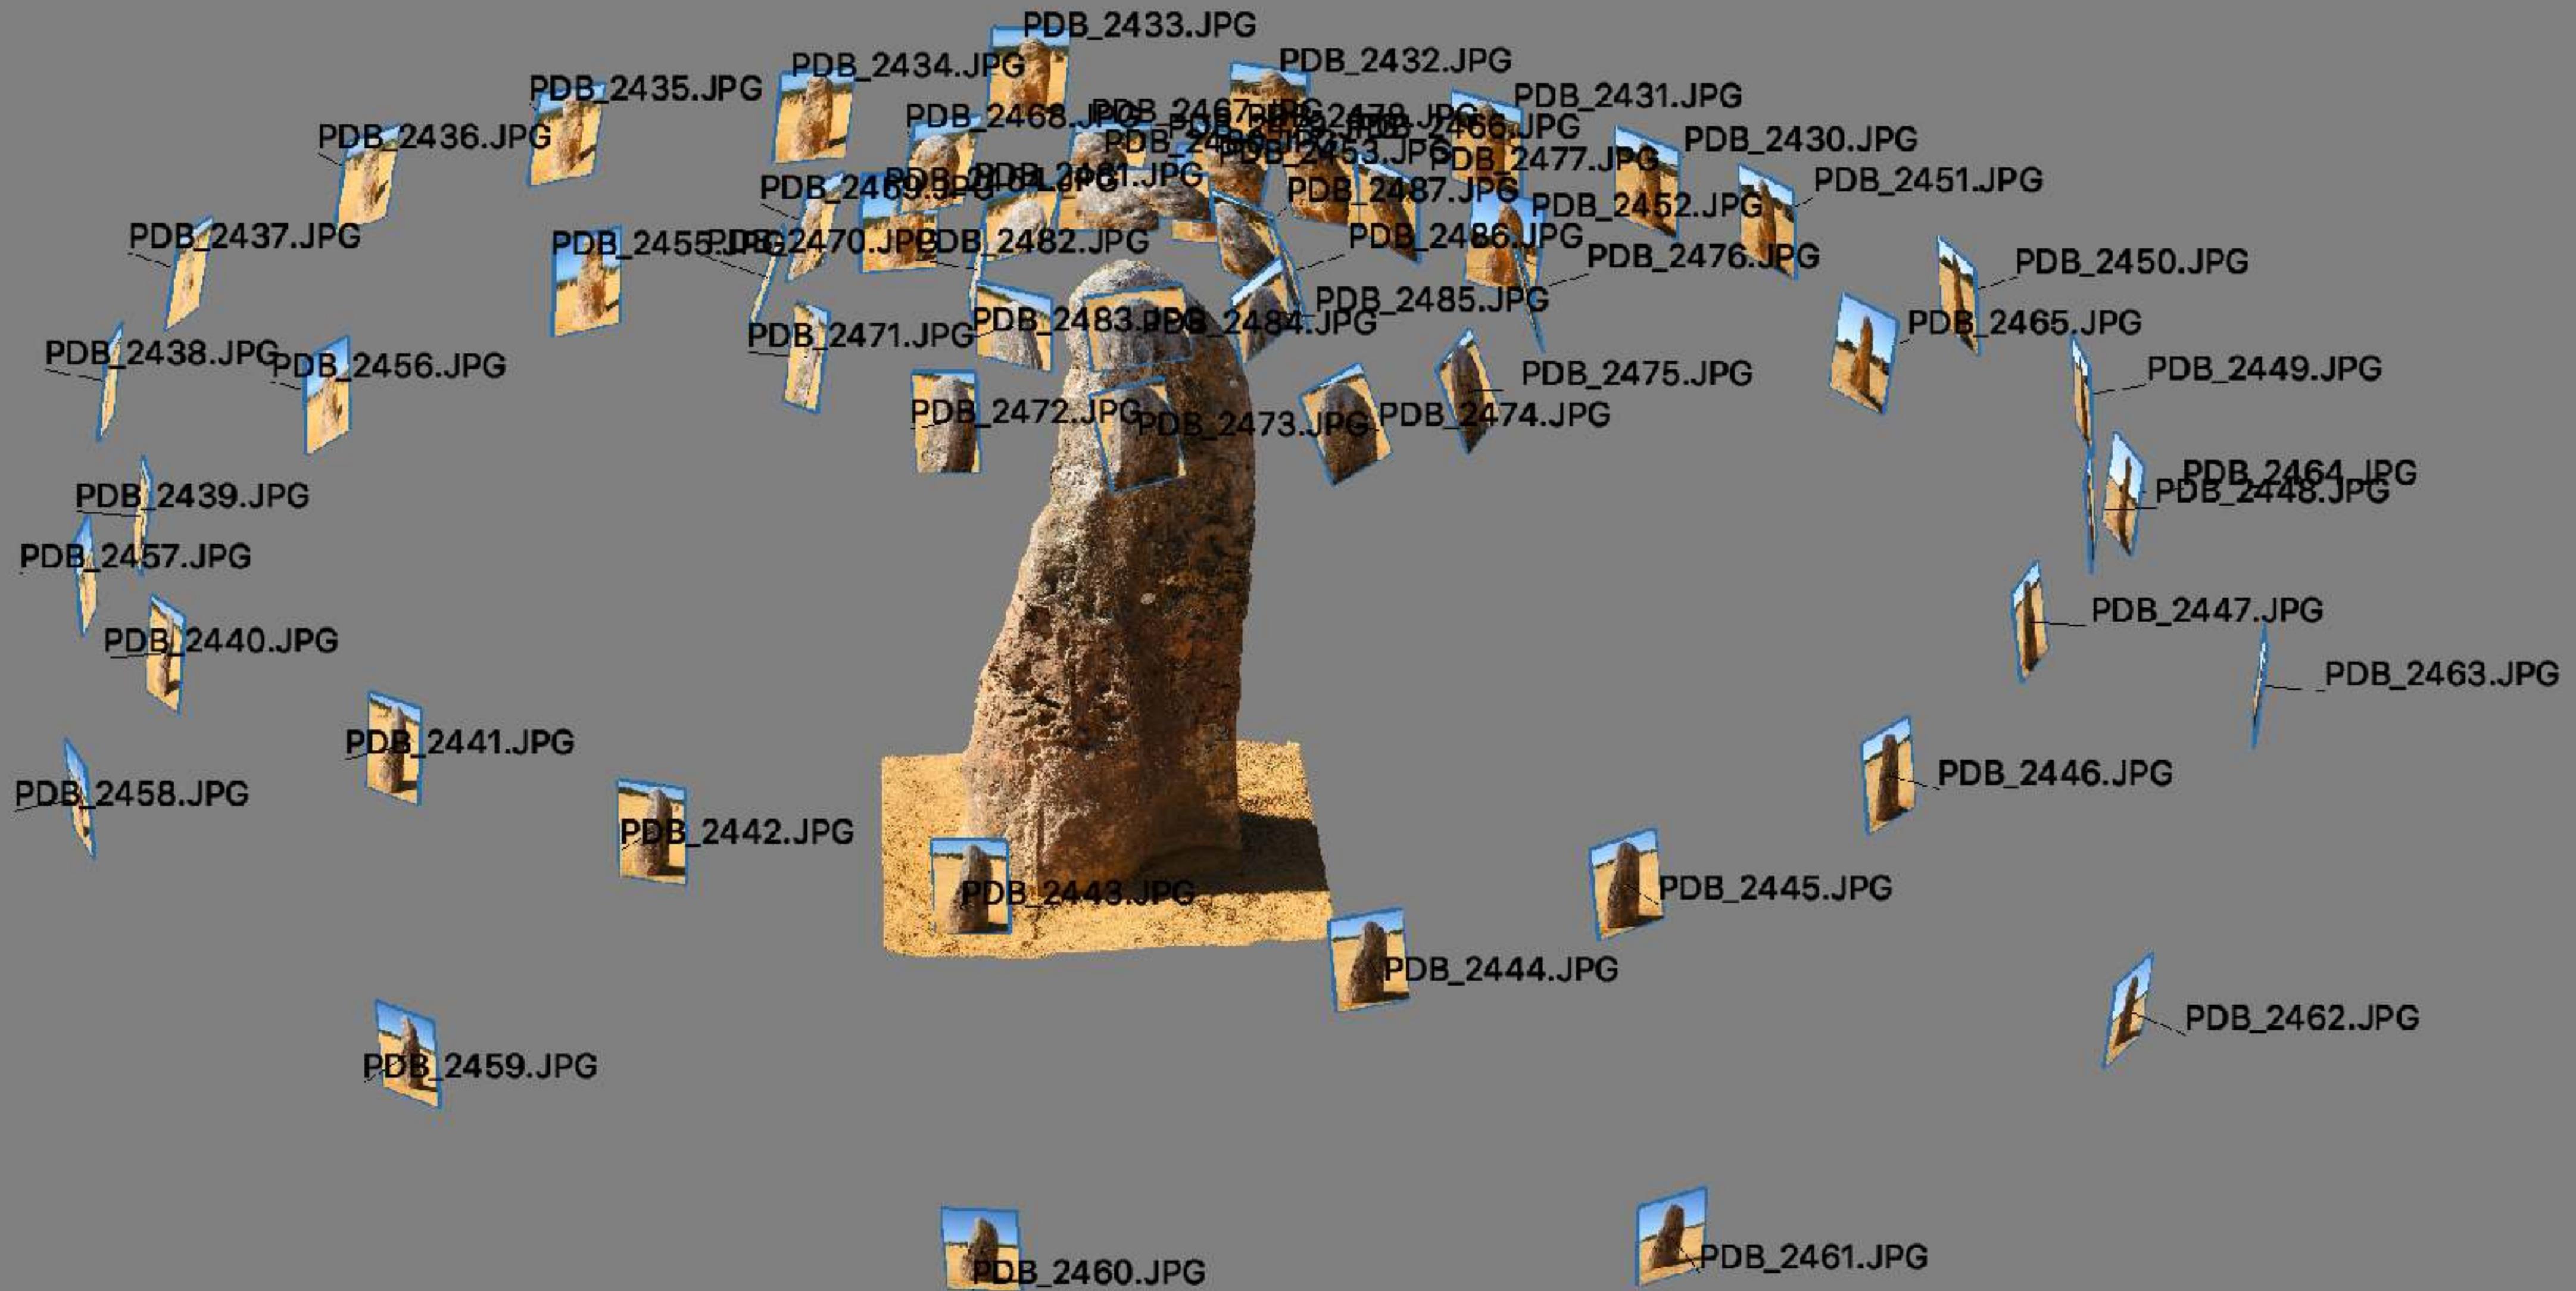
DSCF2579.JPG

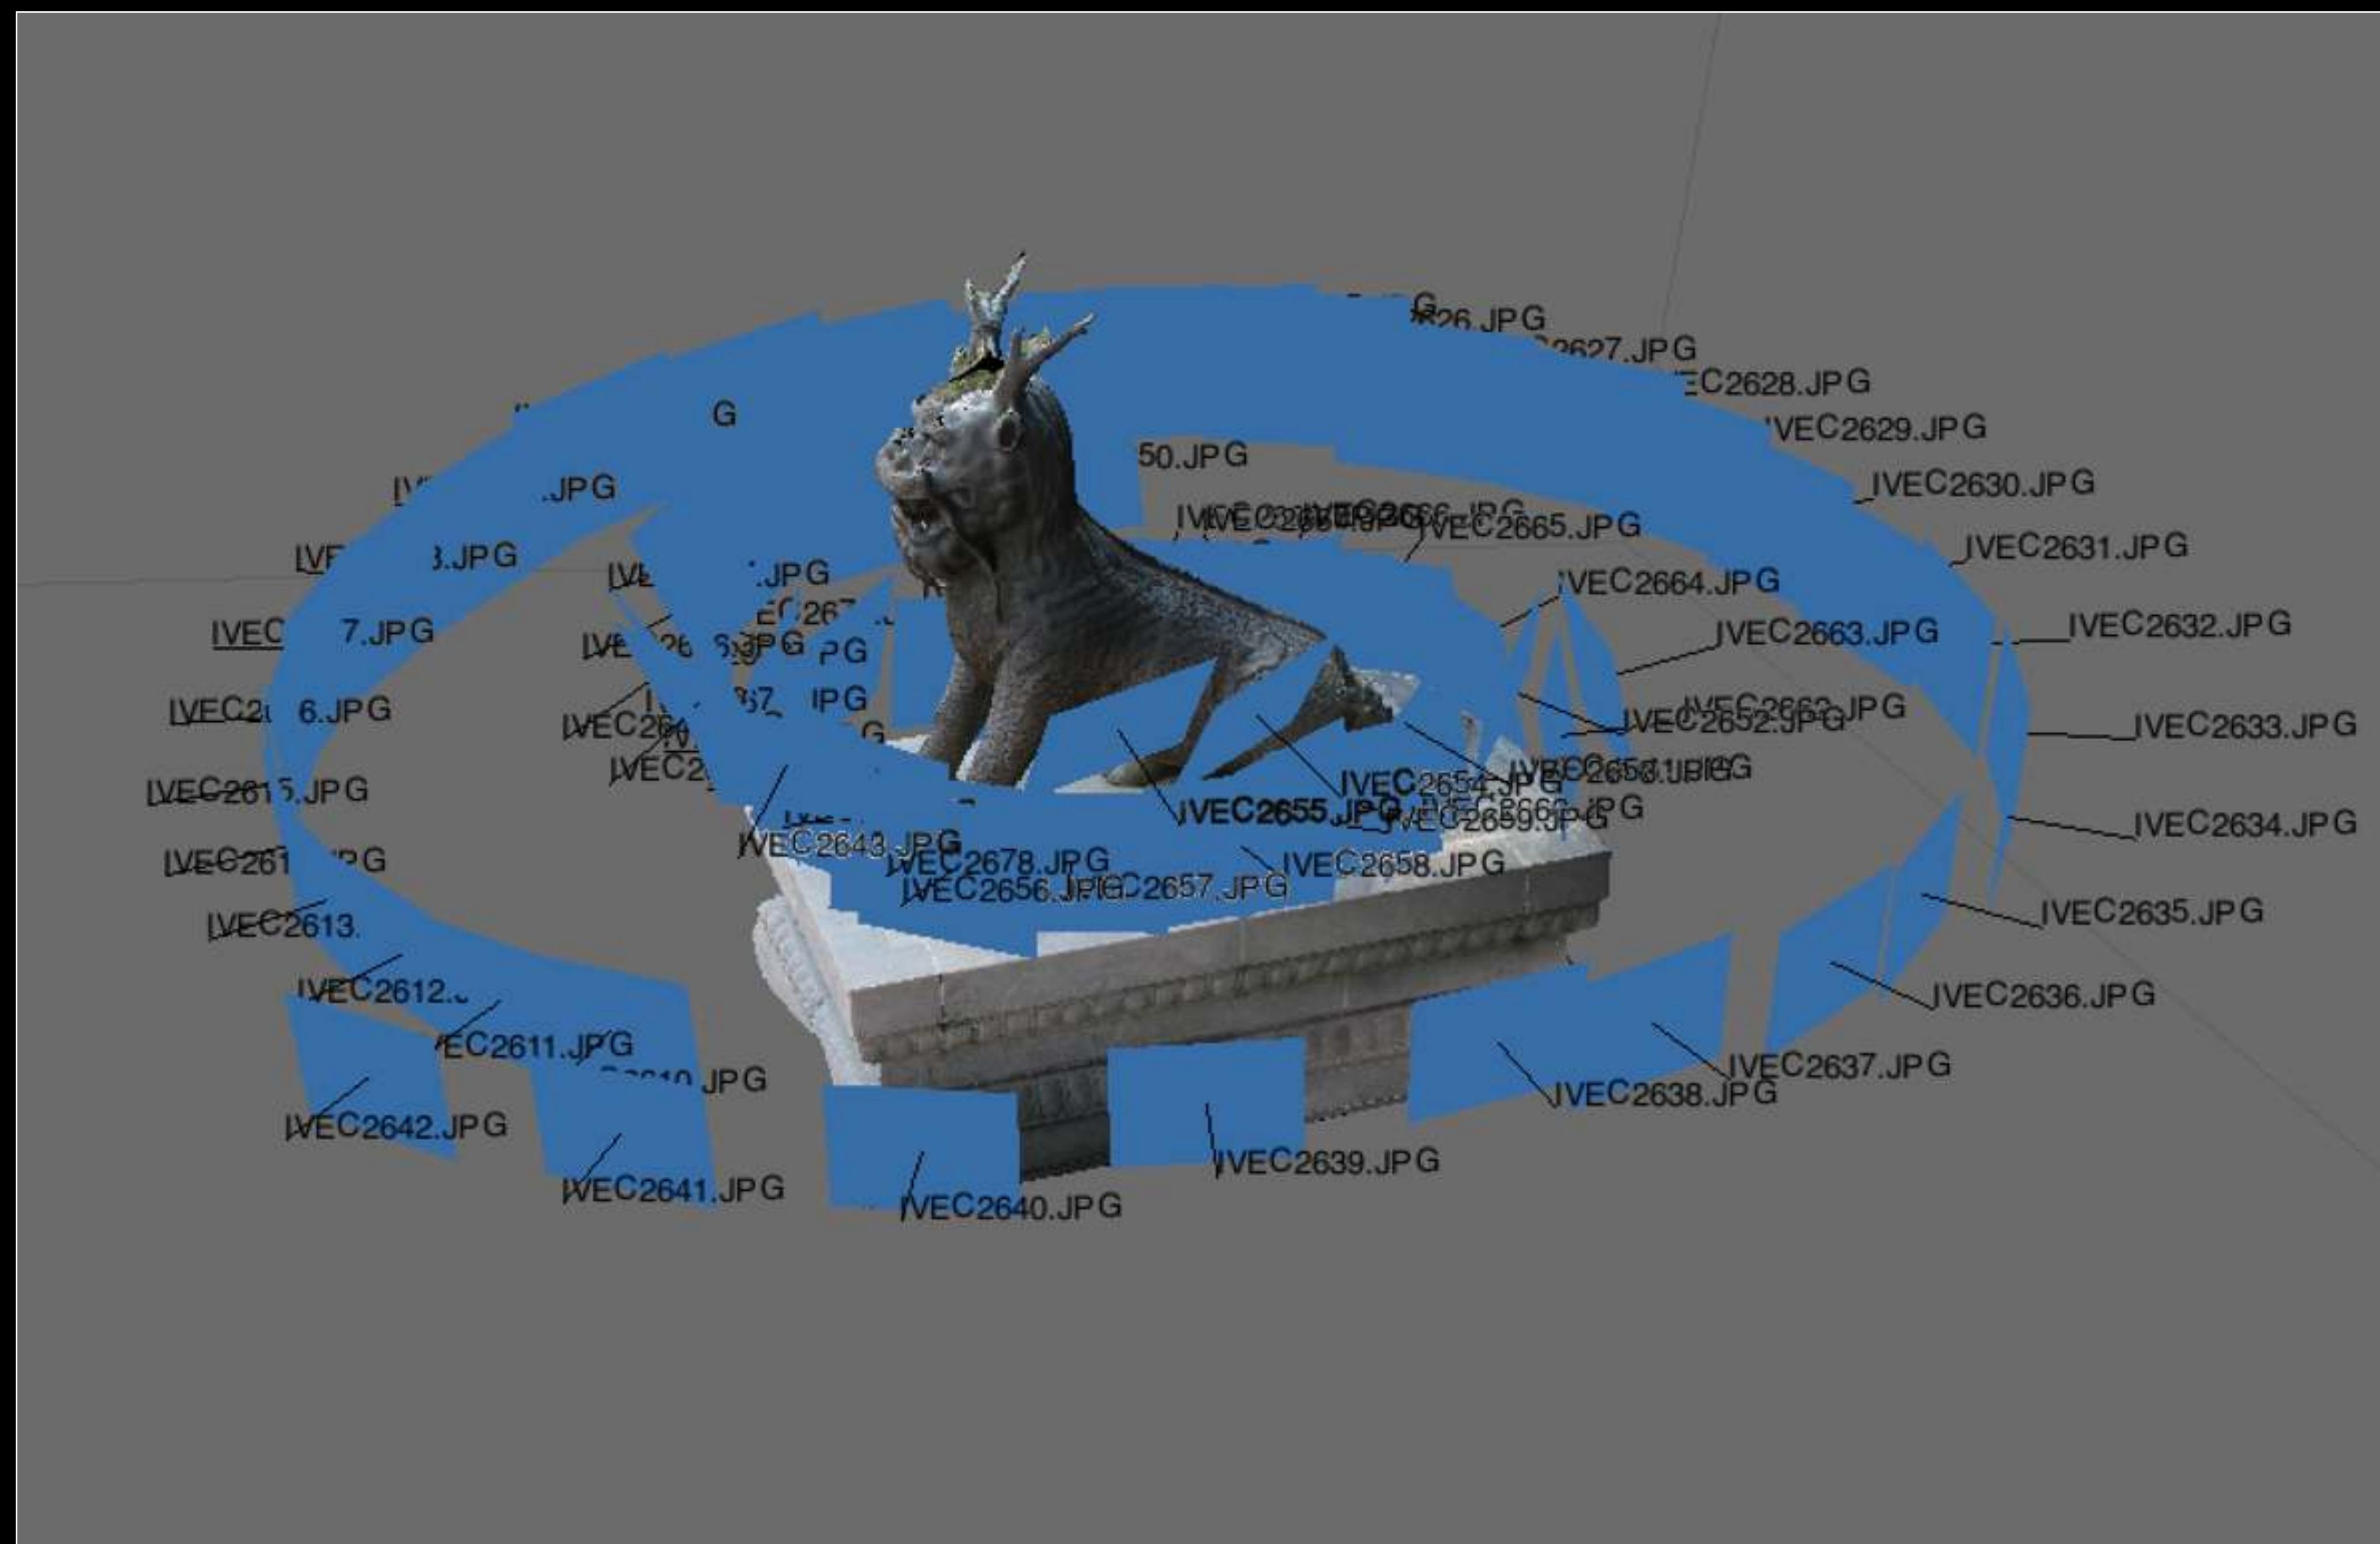
DSCF2556.JPG

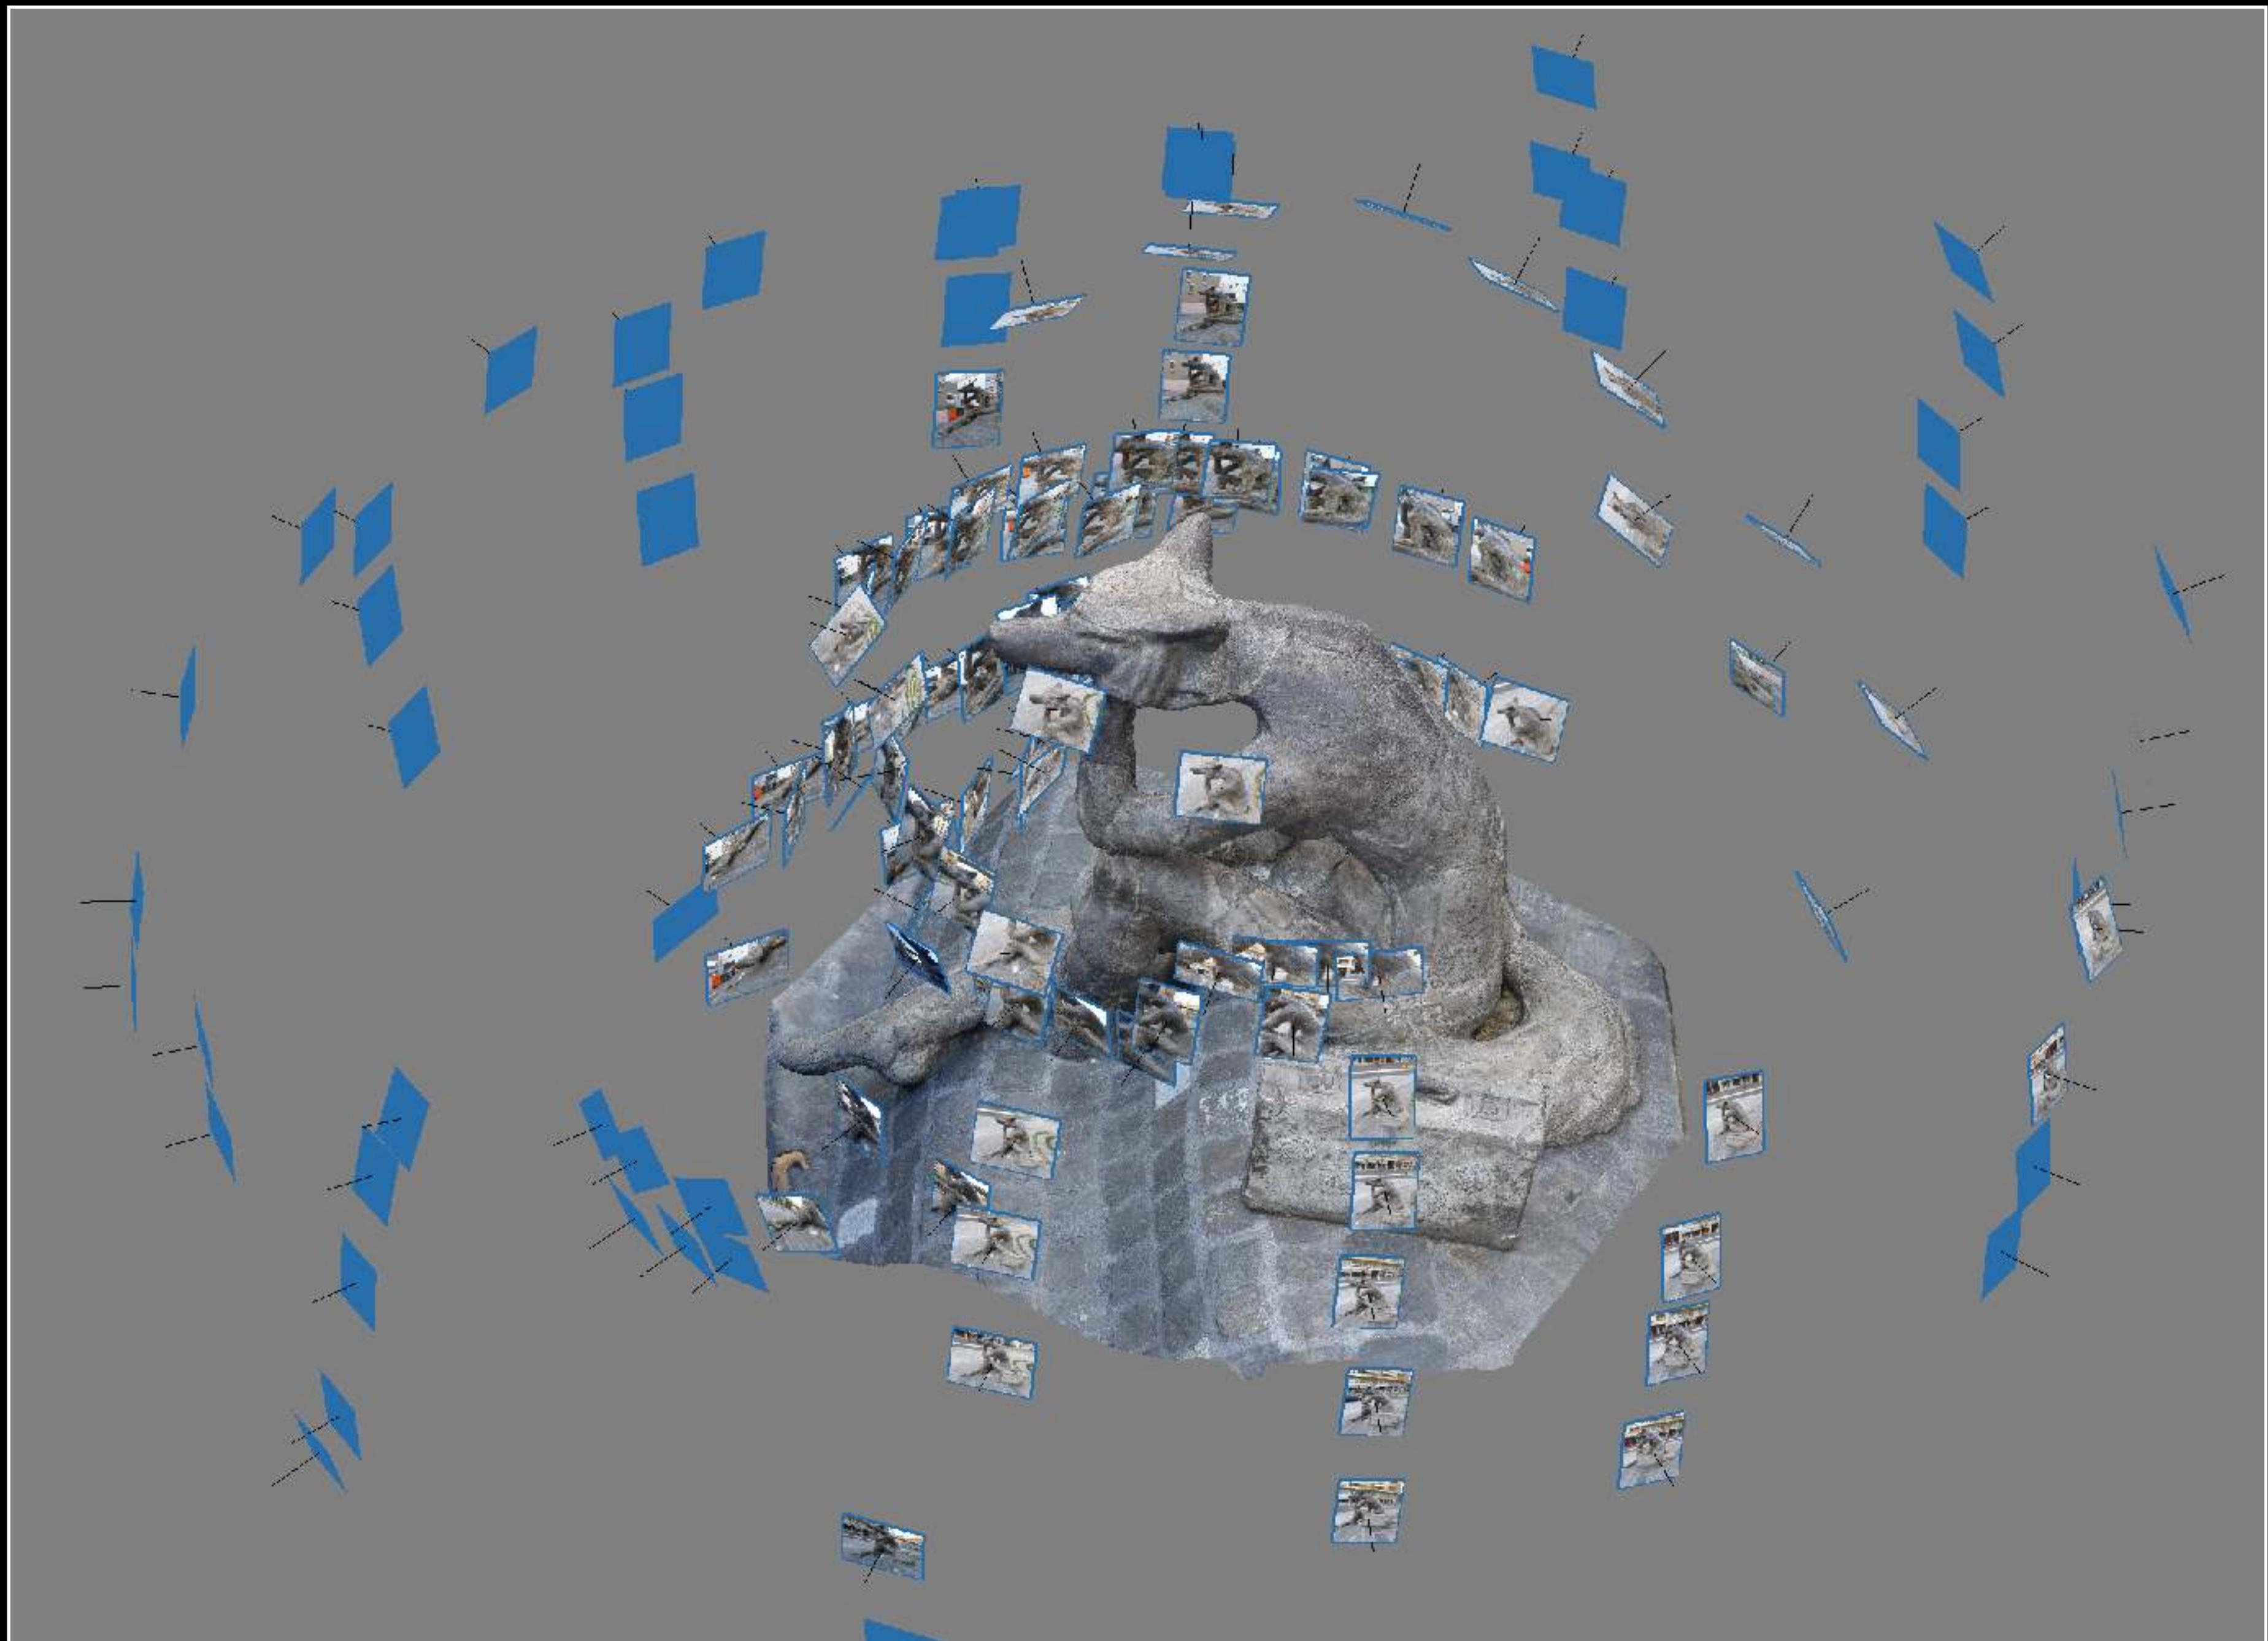
DSCF2557.JPG

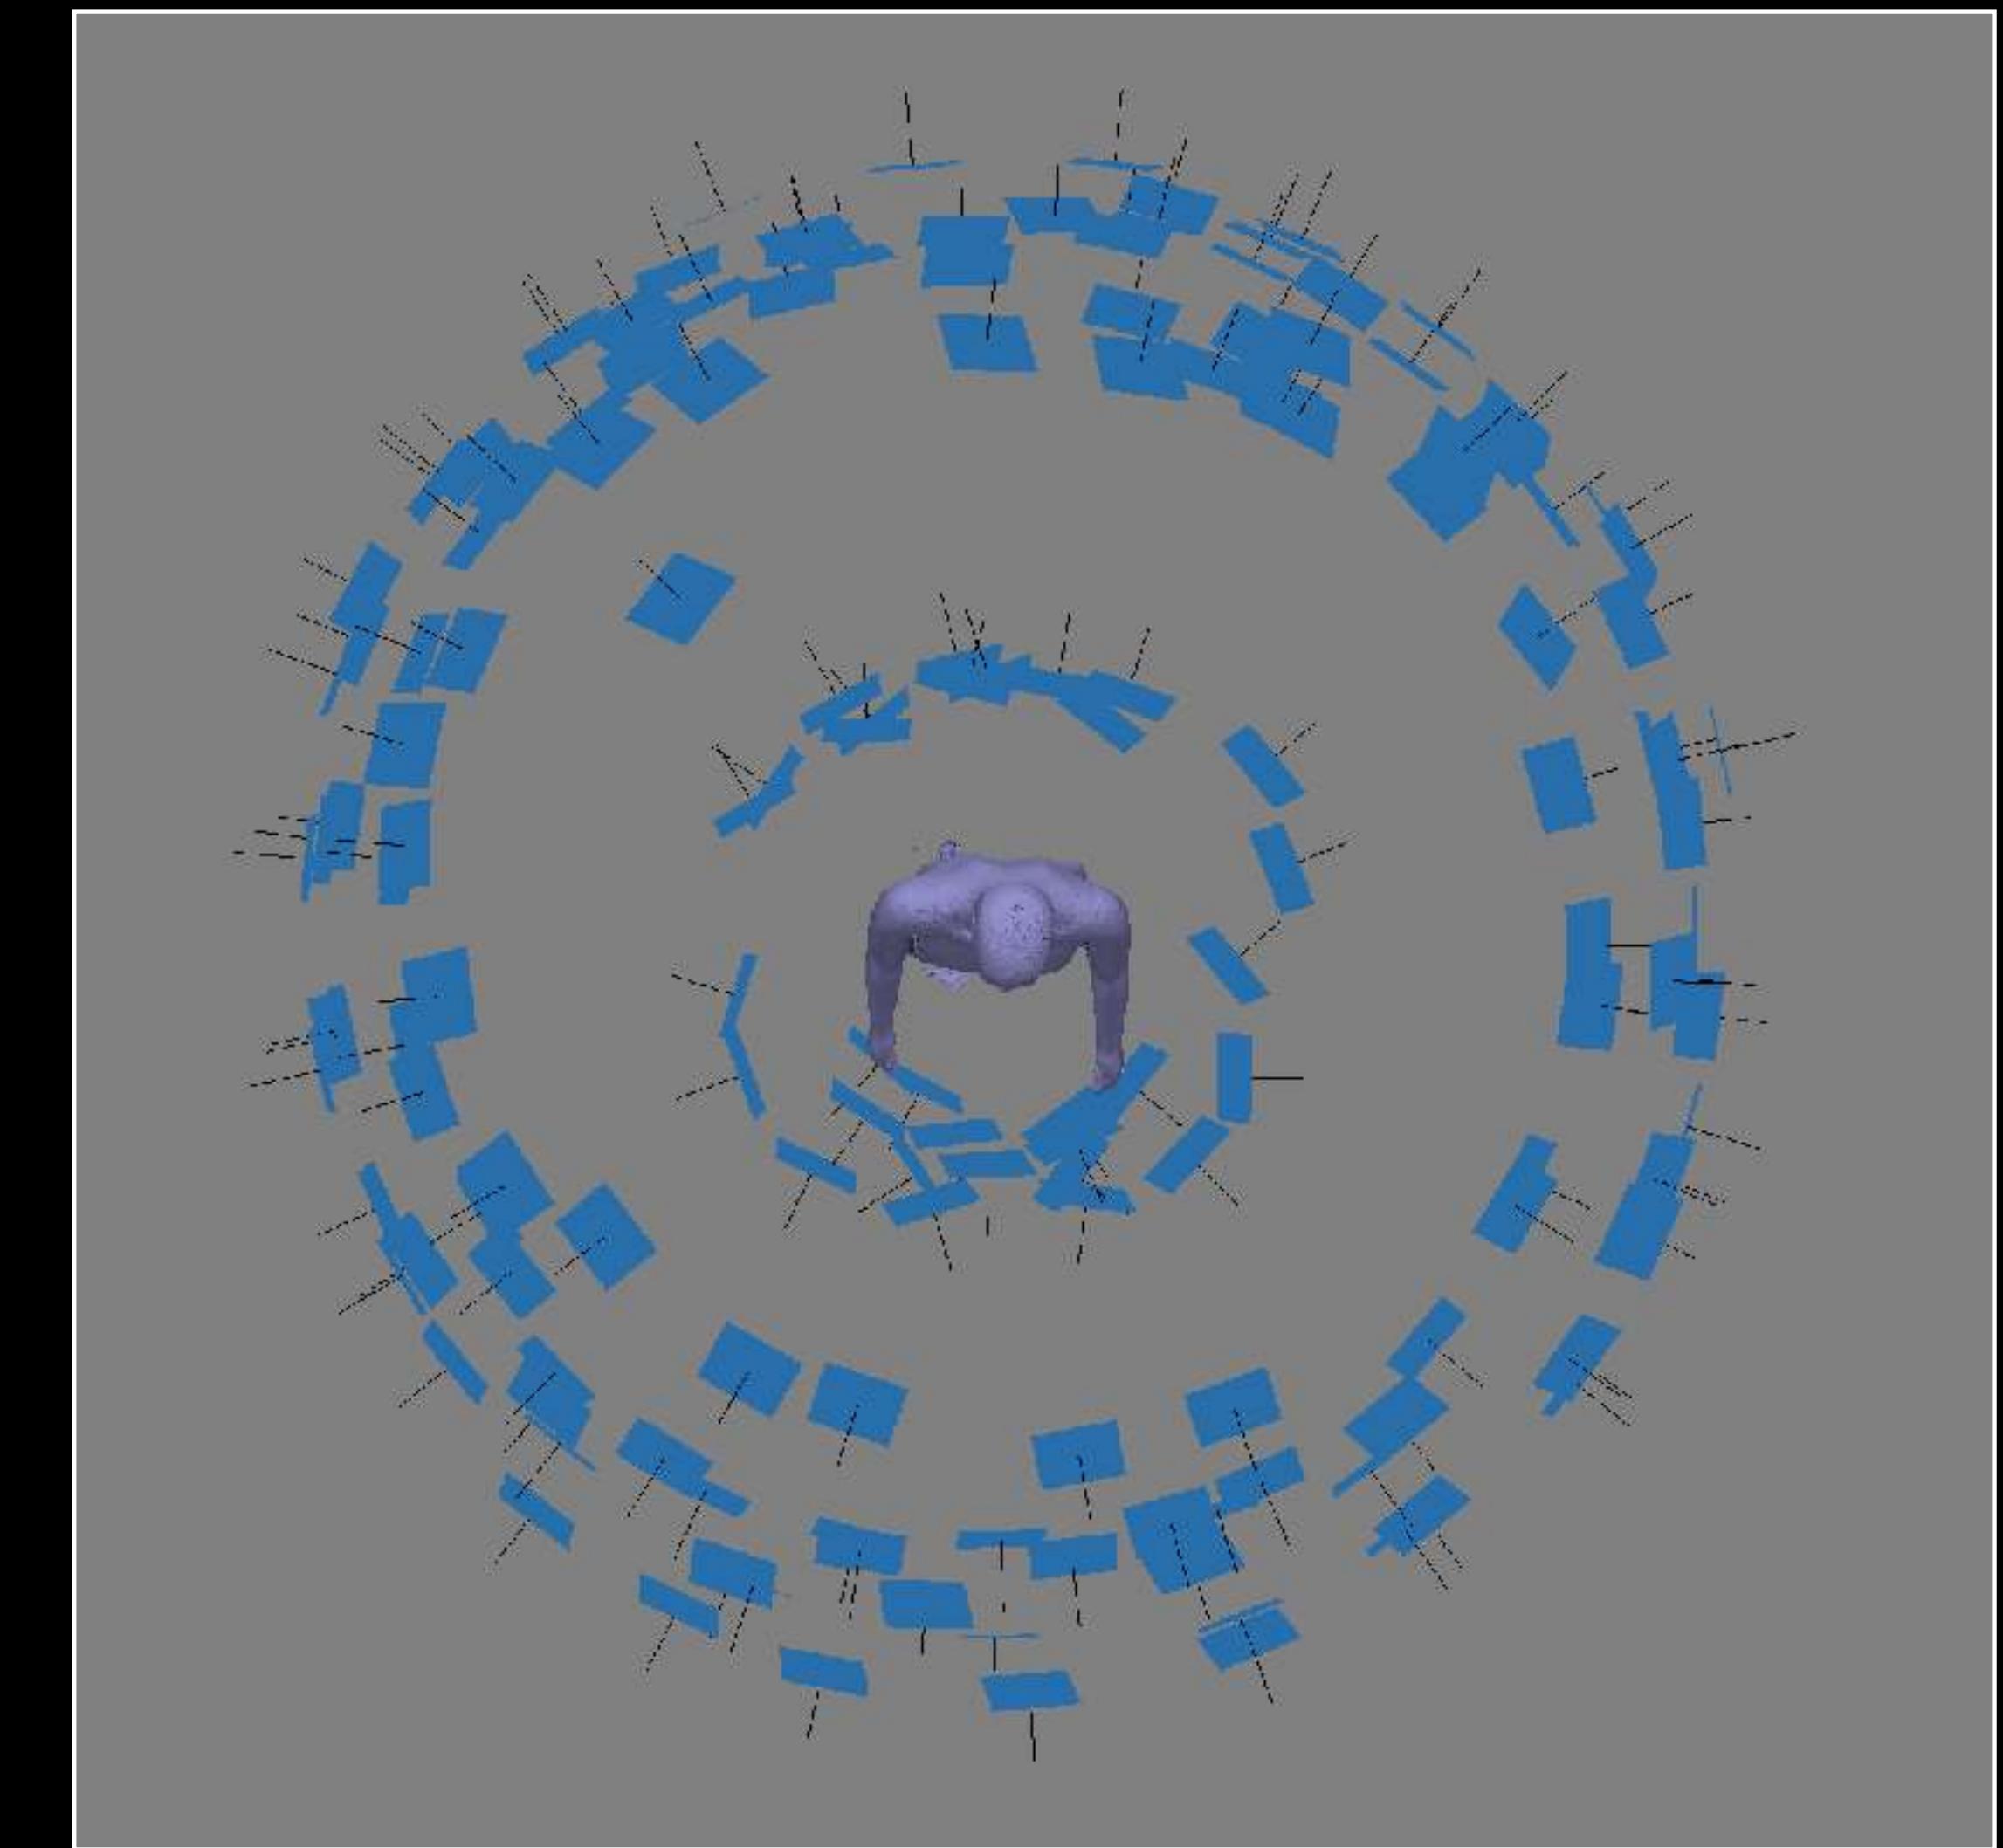
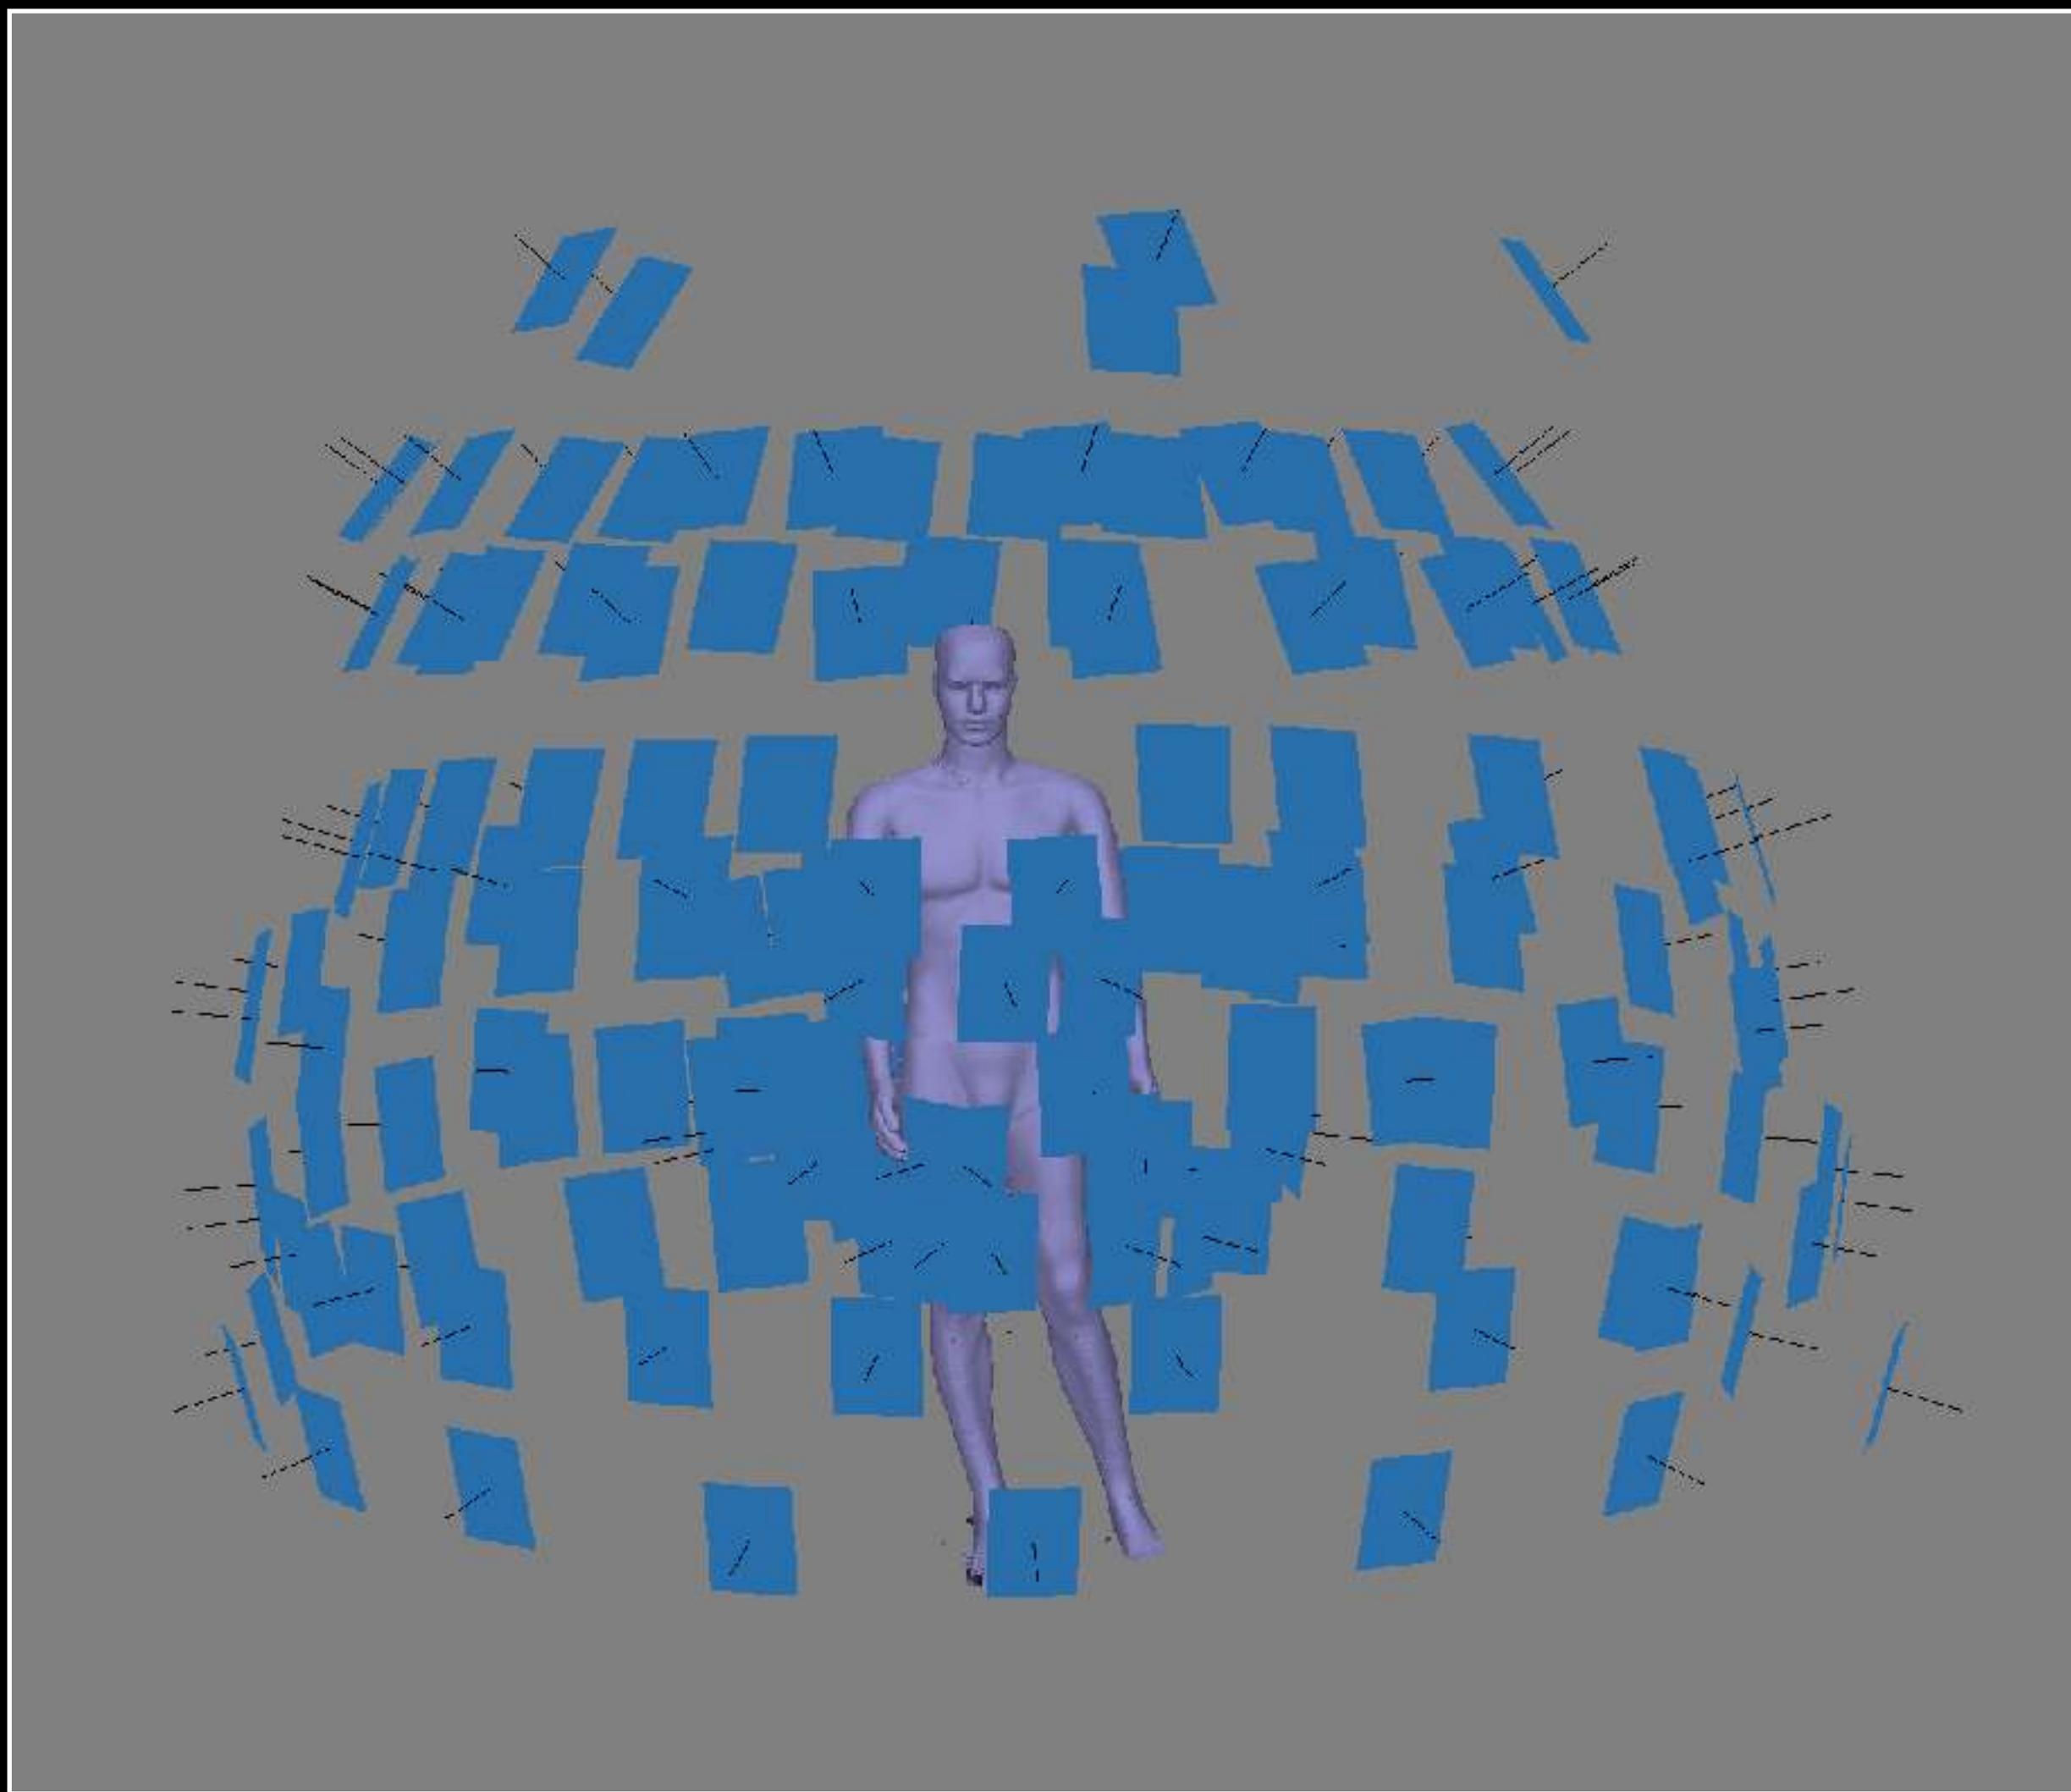
DSCF2558.JPG

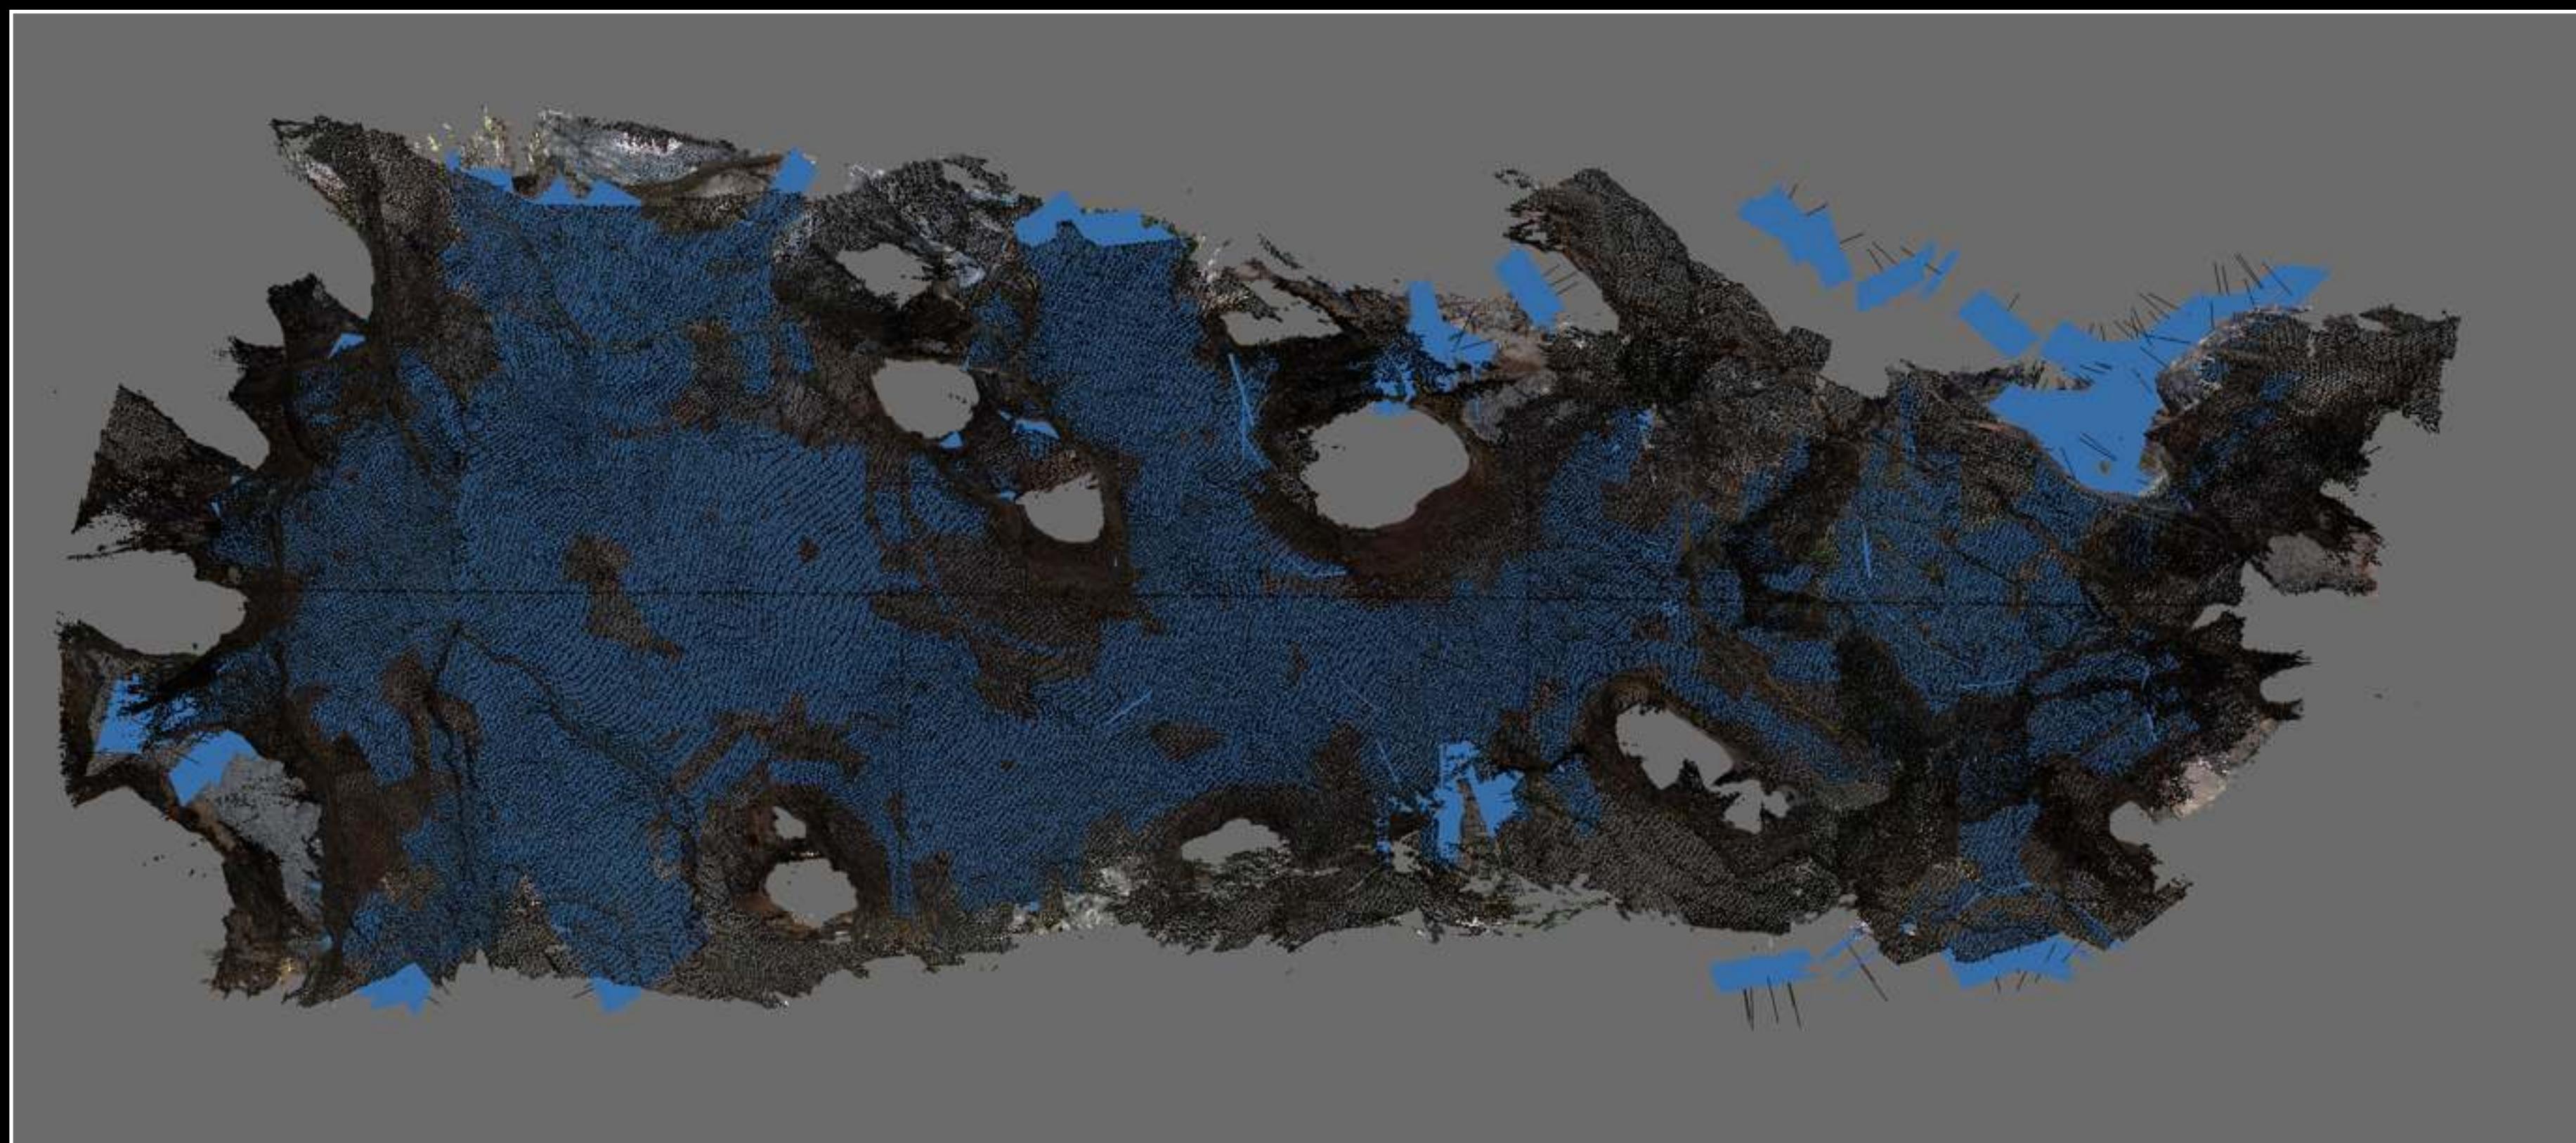
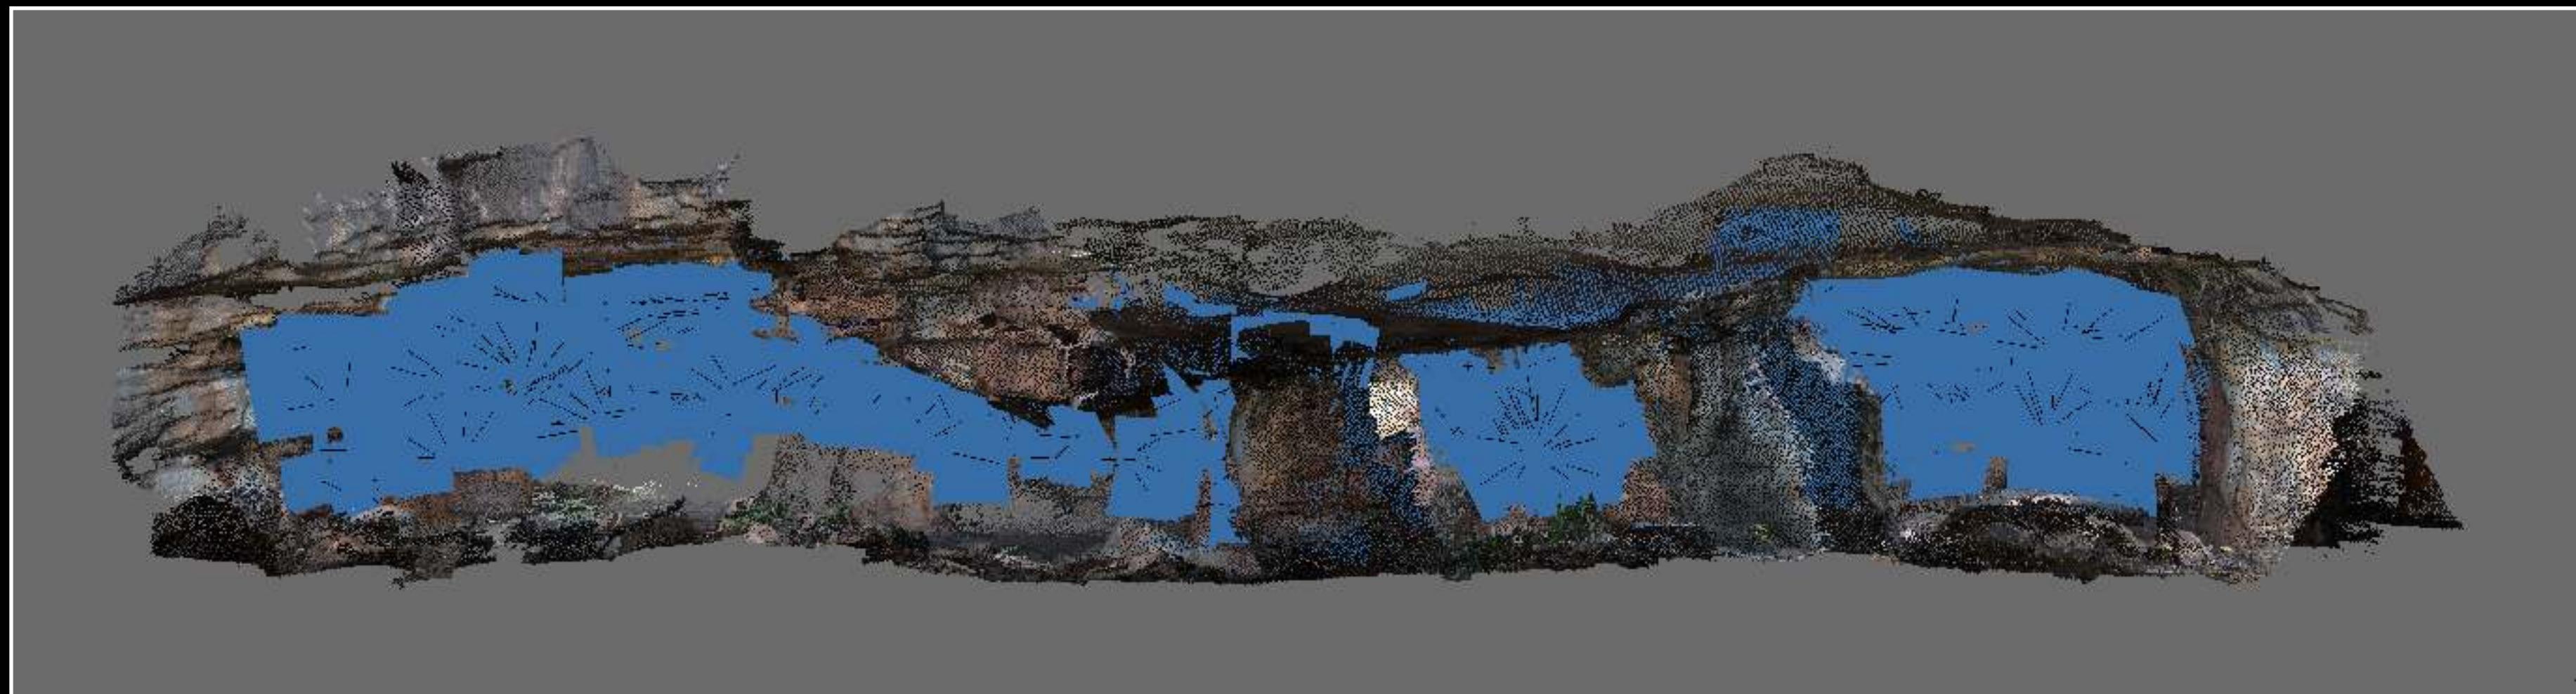
DSCF2559.JPG











Software workflow

- A number of products in the market
 - PhotoModeller Scanner
 - Remake
 - Apero
 - Colmap (free)
 - Alice Meshroom (free)
 - Opensource pipelines
 - MetaShape, was Photoscan
-
- Distinguishing features
 - Degree of automation and human guidedness
 - Implement whole pipeline from photographs to textured mesh
 - Degree of algorithm control
 - Number of photographs that can be handled
 - Computer or cloud based
- We will focus here on perhaps the world leader, MetaShape (Was PhotoScan) from Agisoft.

Typical pipeline

- Feature point detection.
- Bundler for sparse point cloud and camera position and poses.

Possible human intervention to filter sparse points

- Dense point cloud given the extra information available from camera positions.

Possible human intervention to clean dense cloud

- Mesh the dense point cloud.

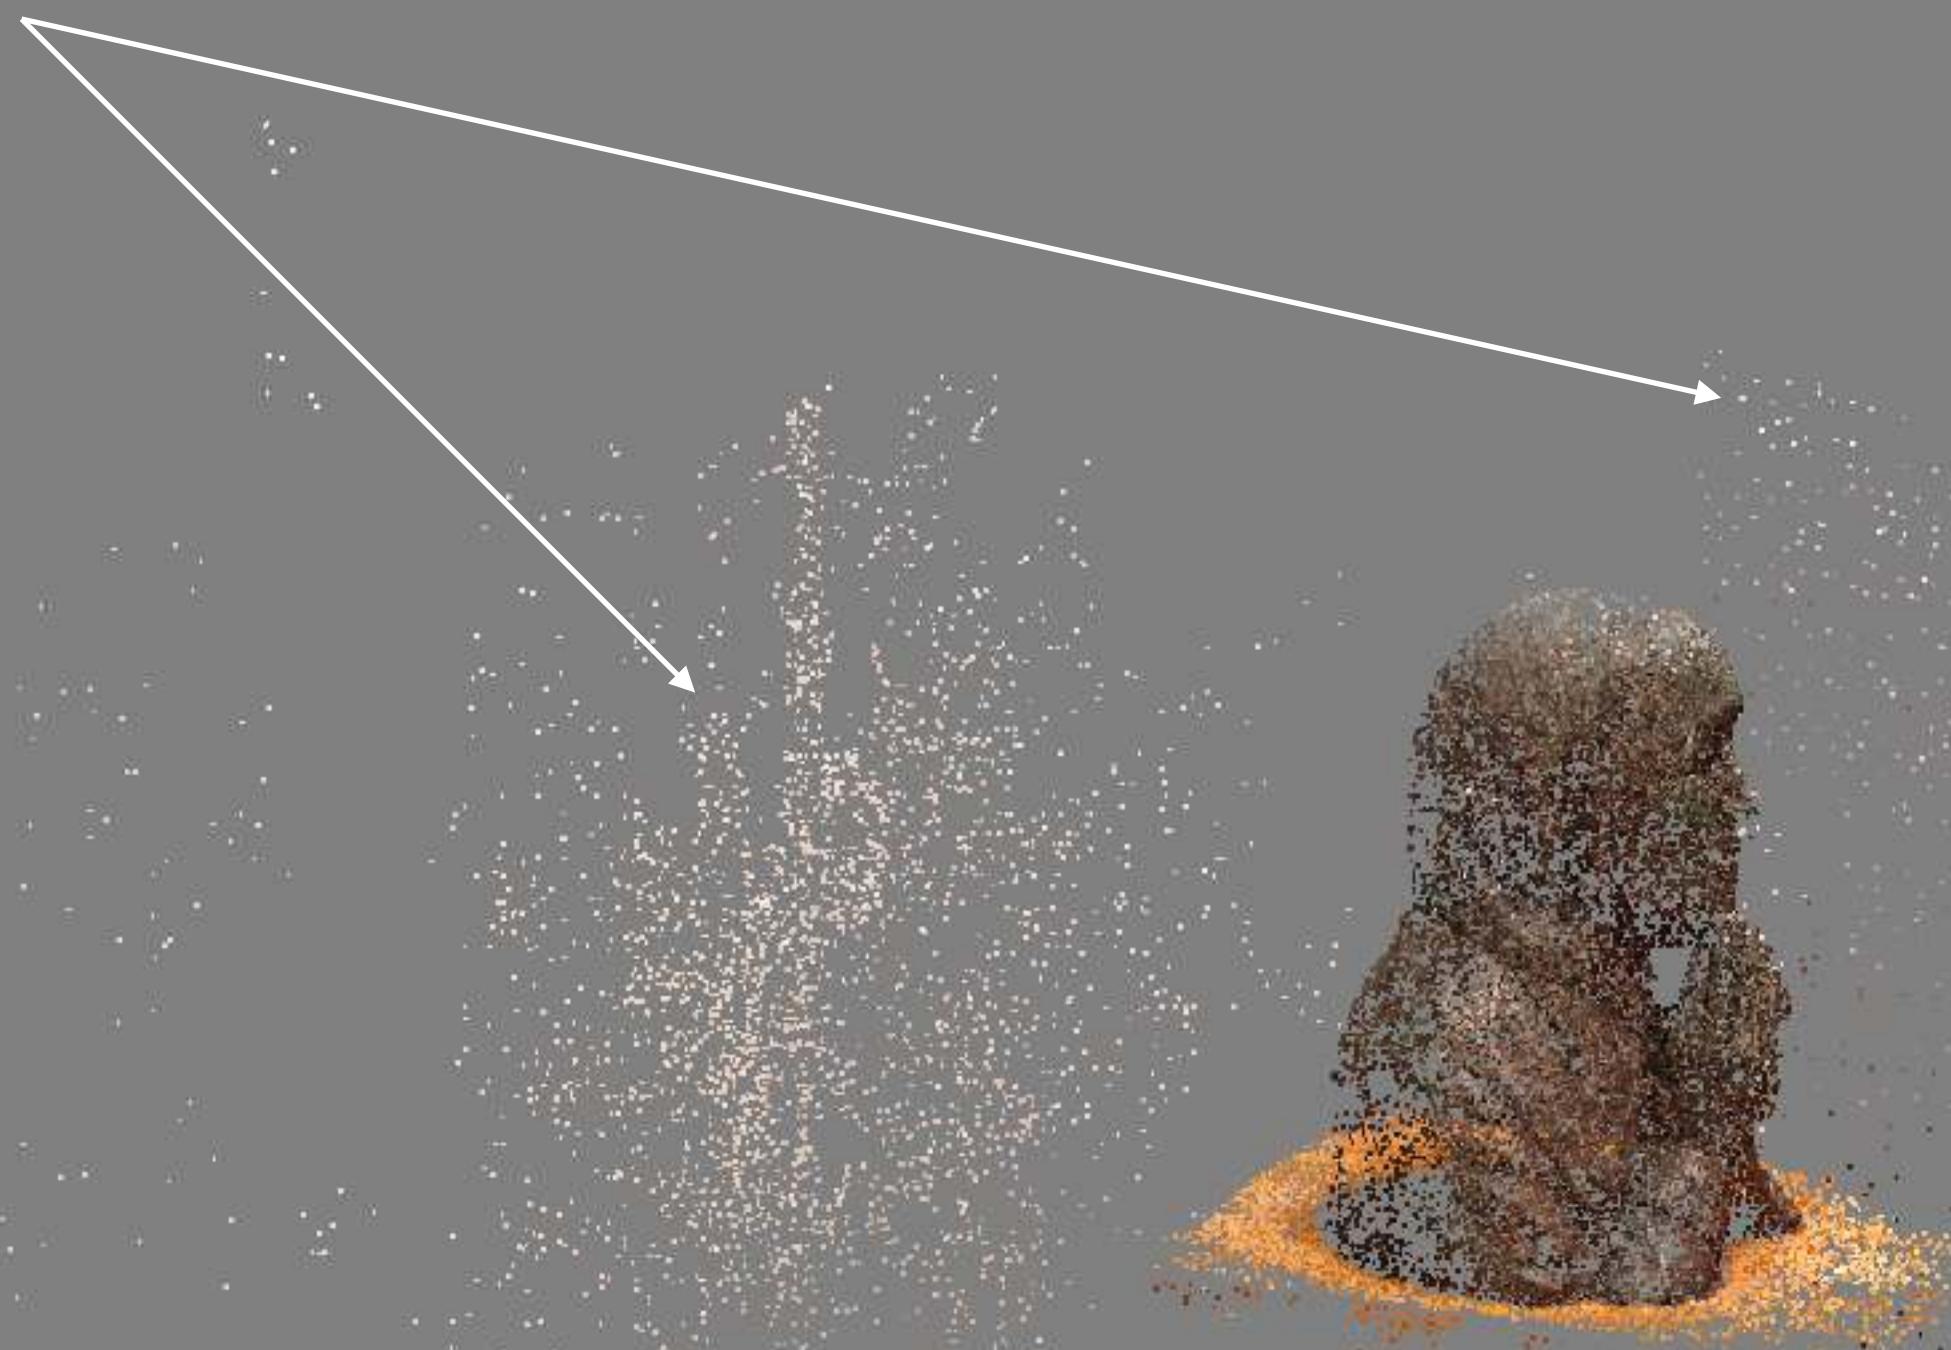
Possible human intervention to edit mesh

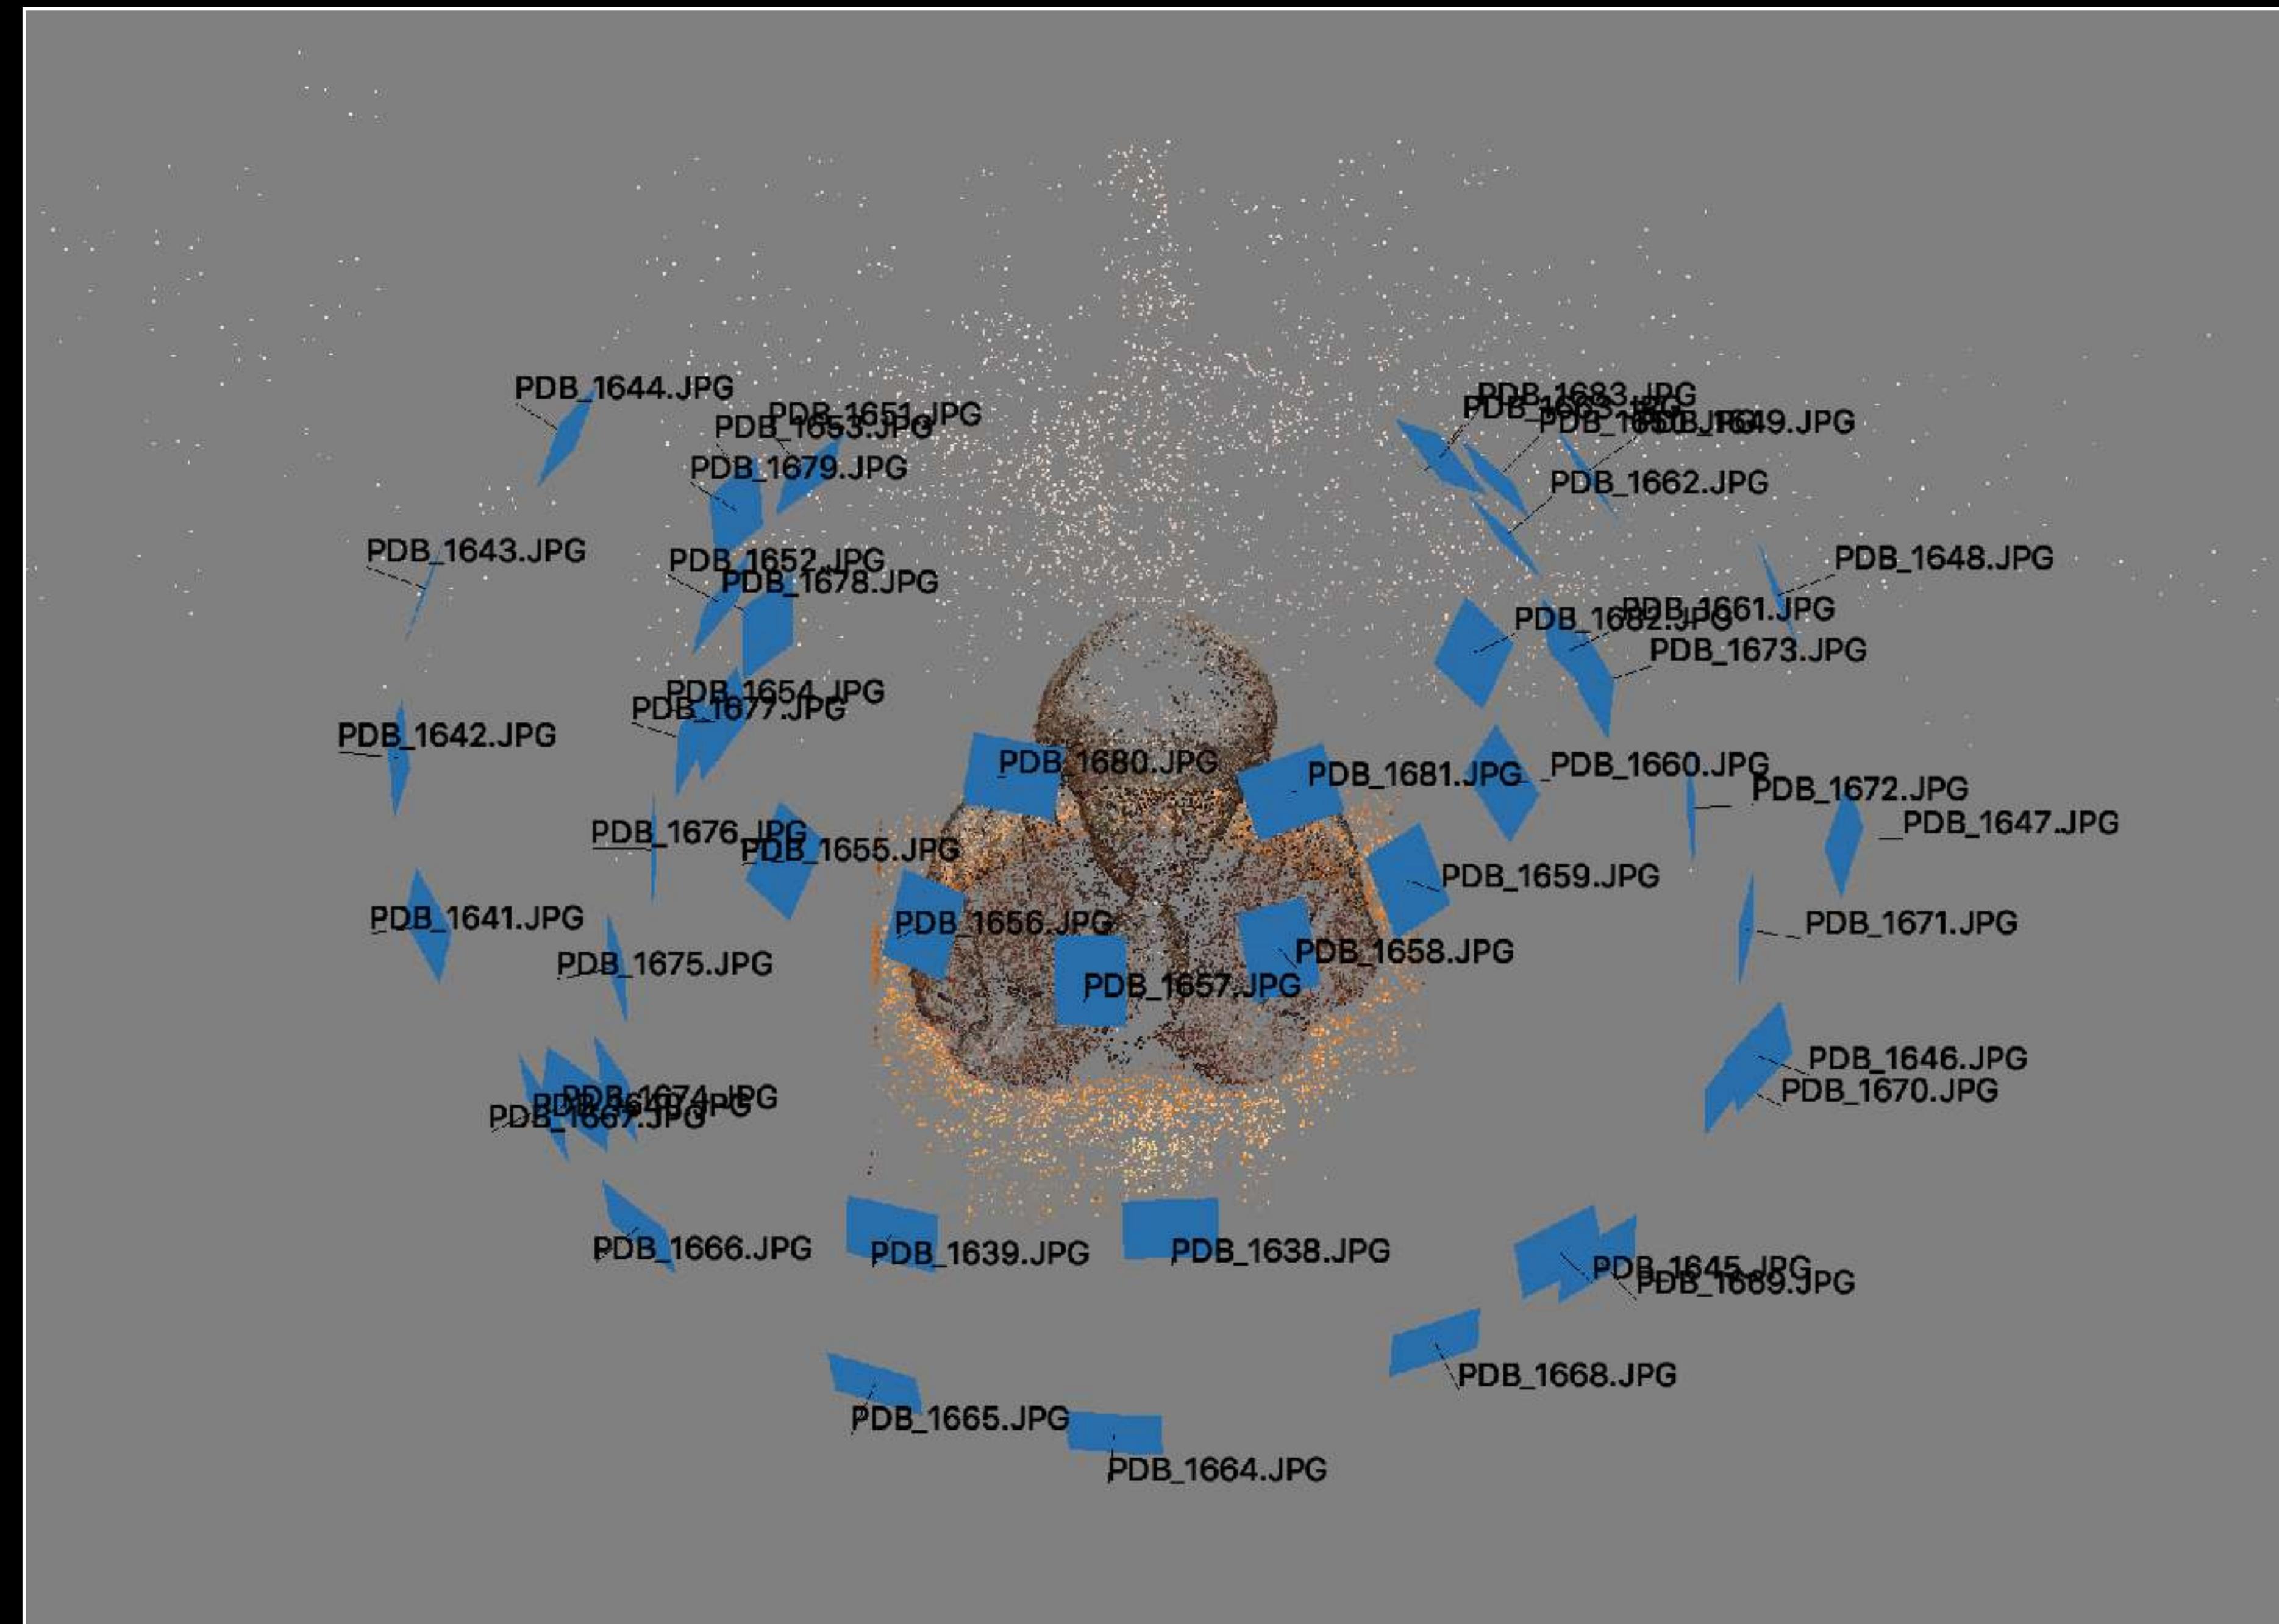
- Preproject photographs from the camera positions onto the mesh.

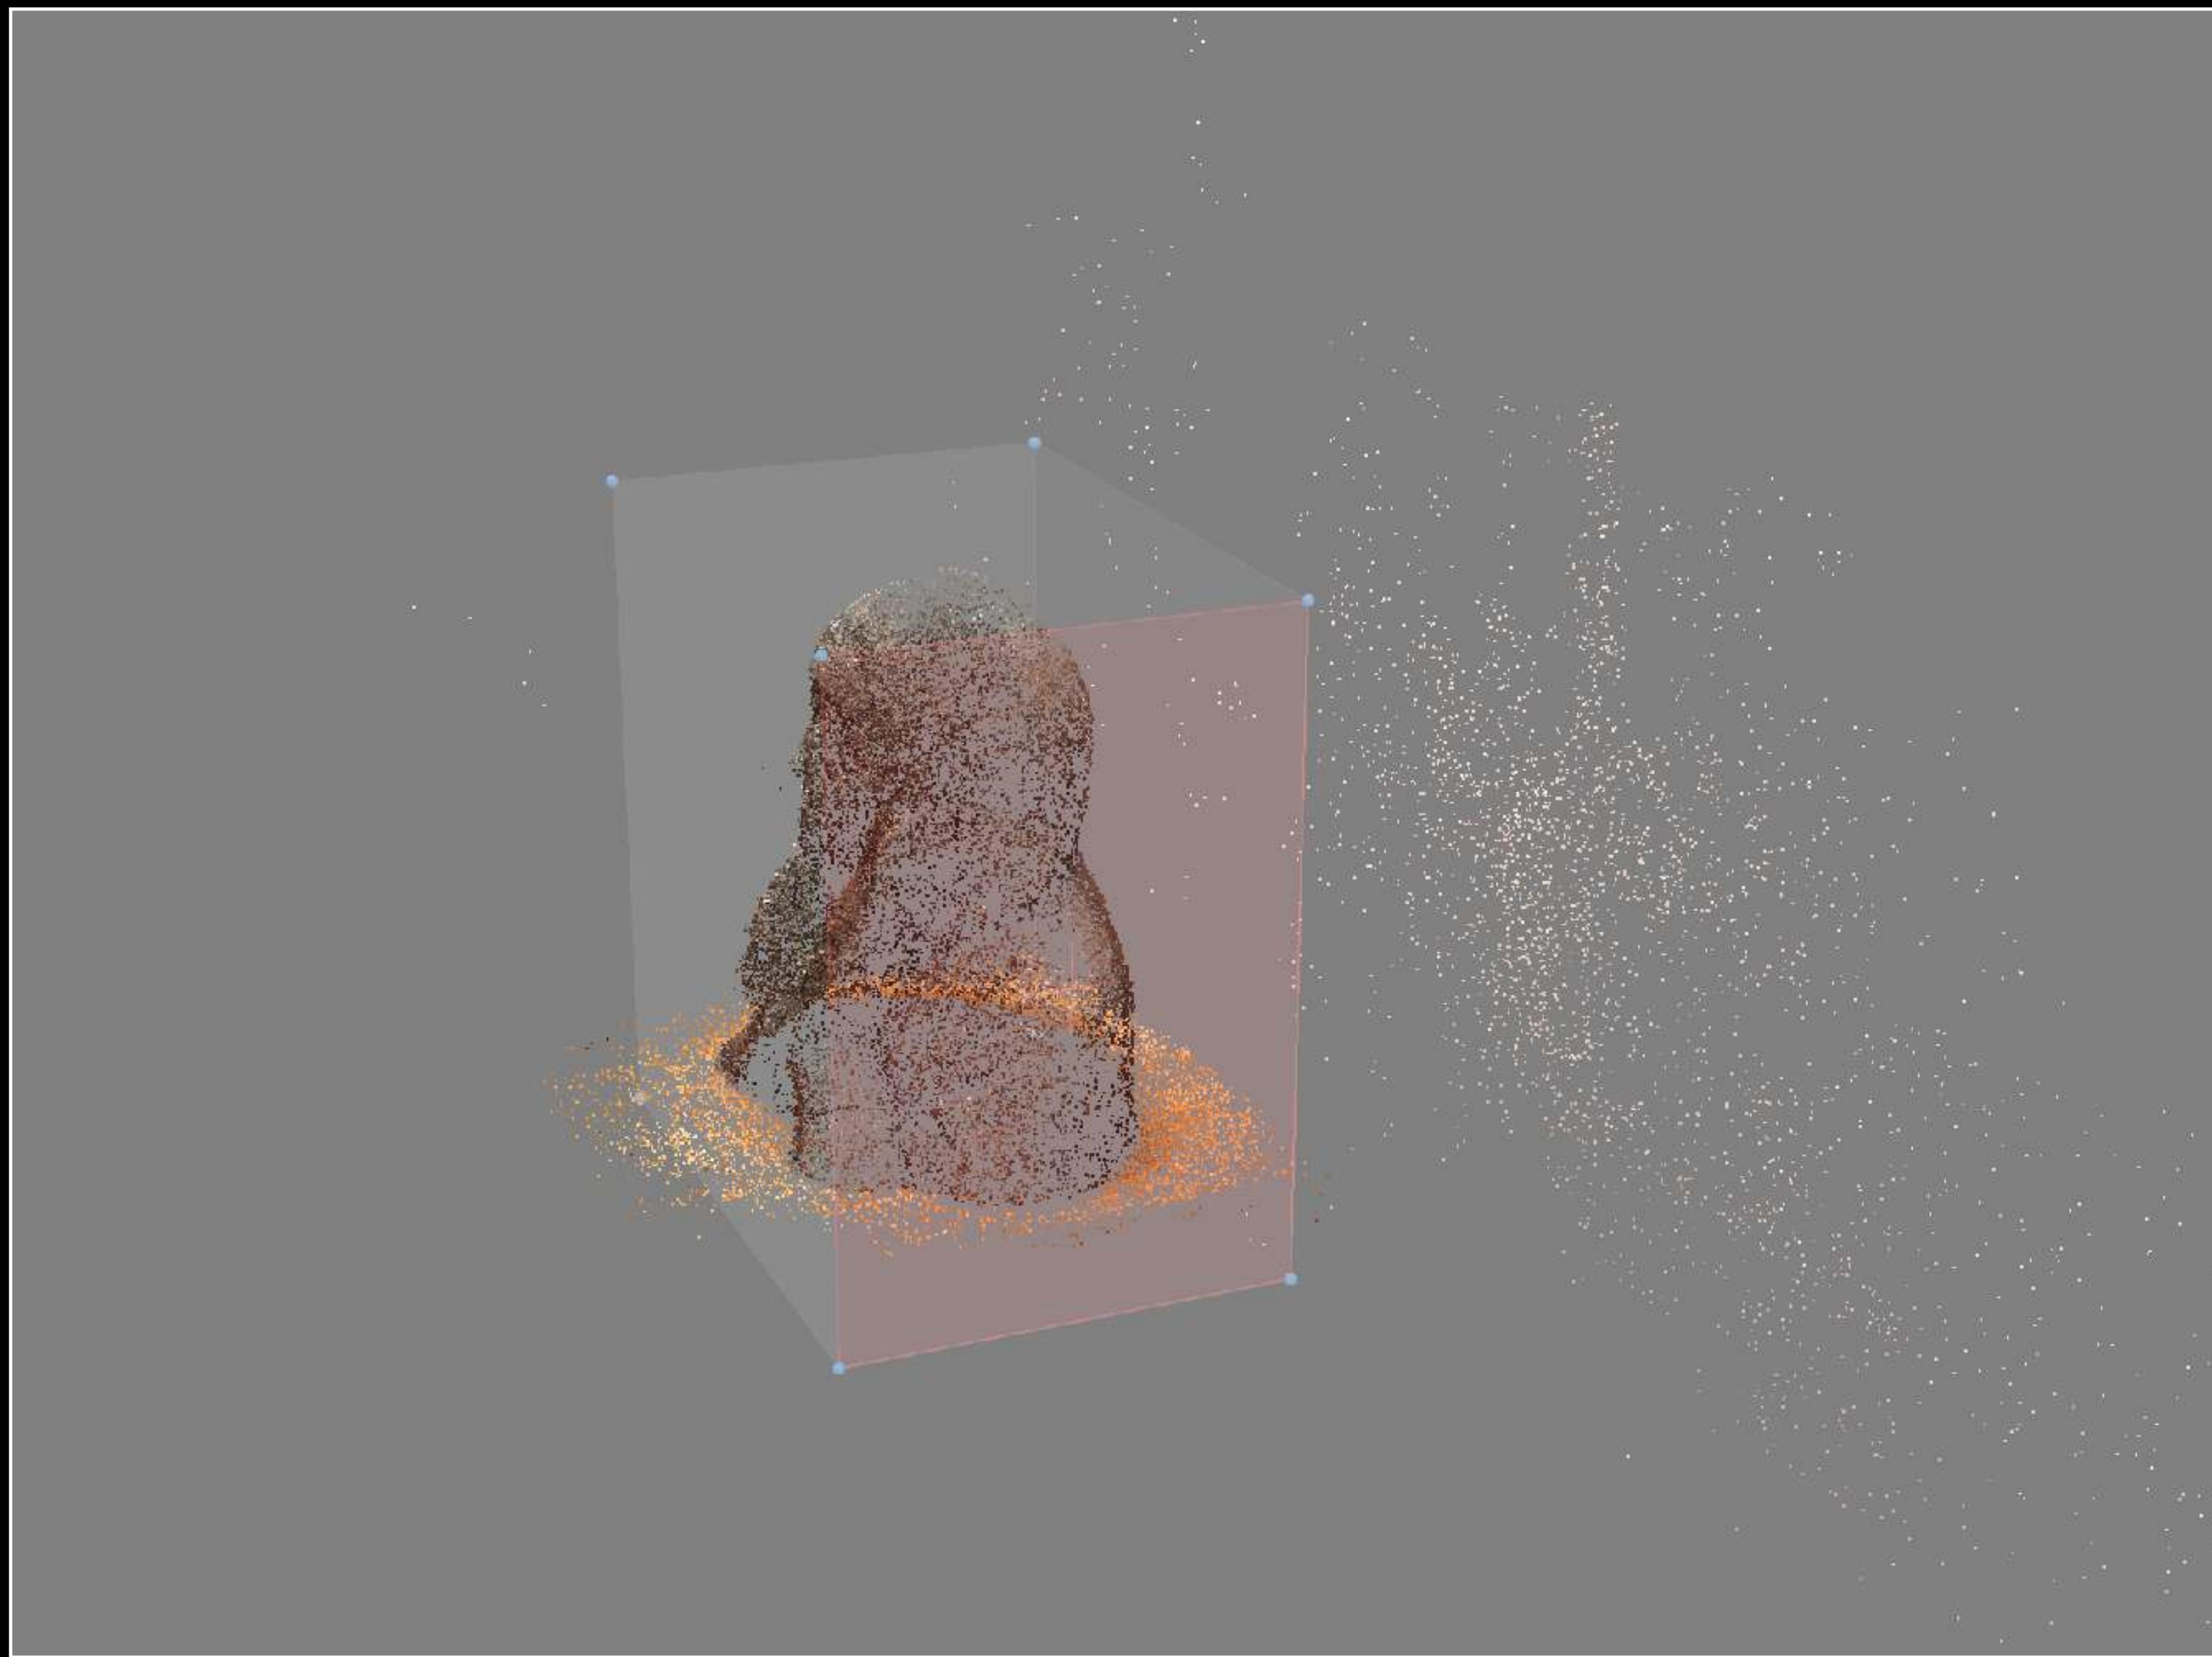
Possible human intervention to adjust texture image file



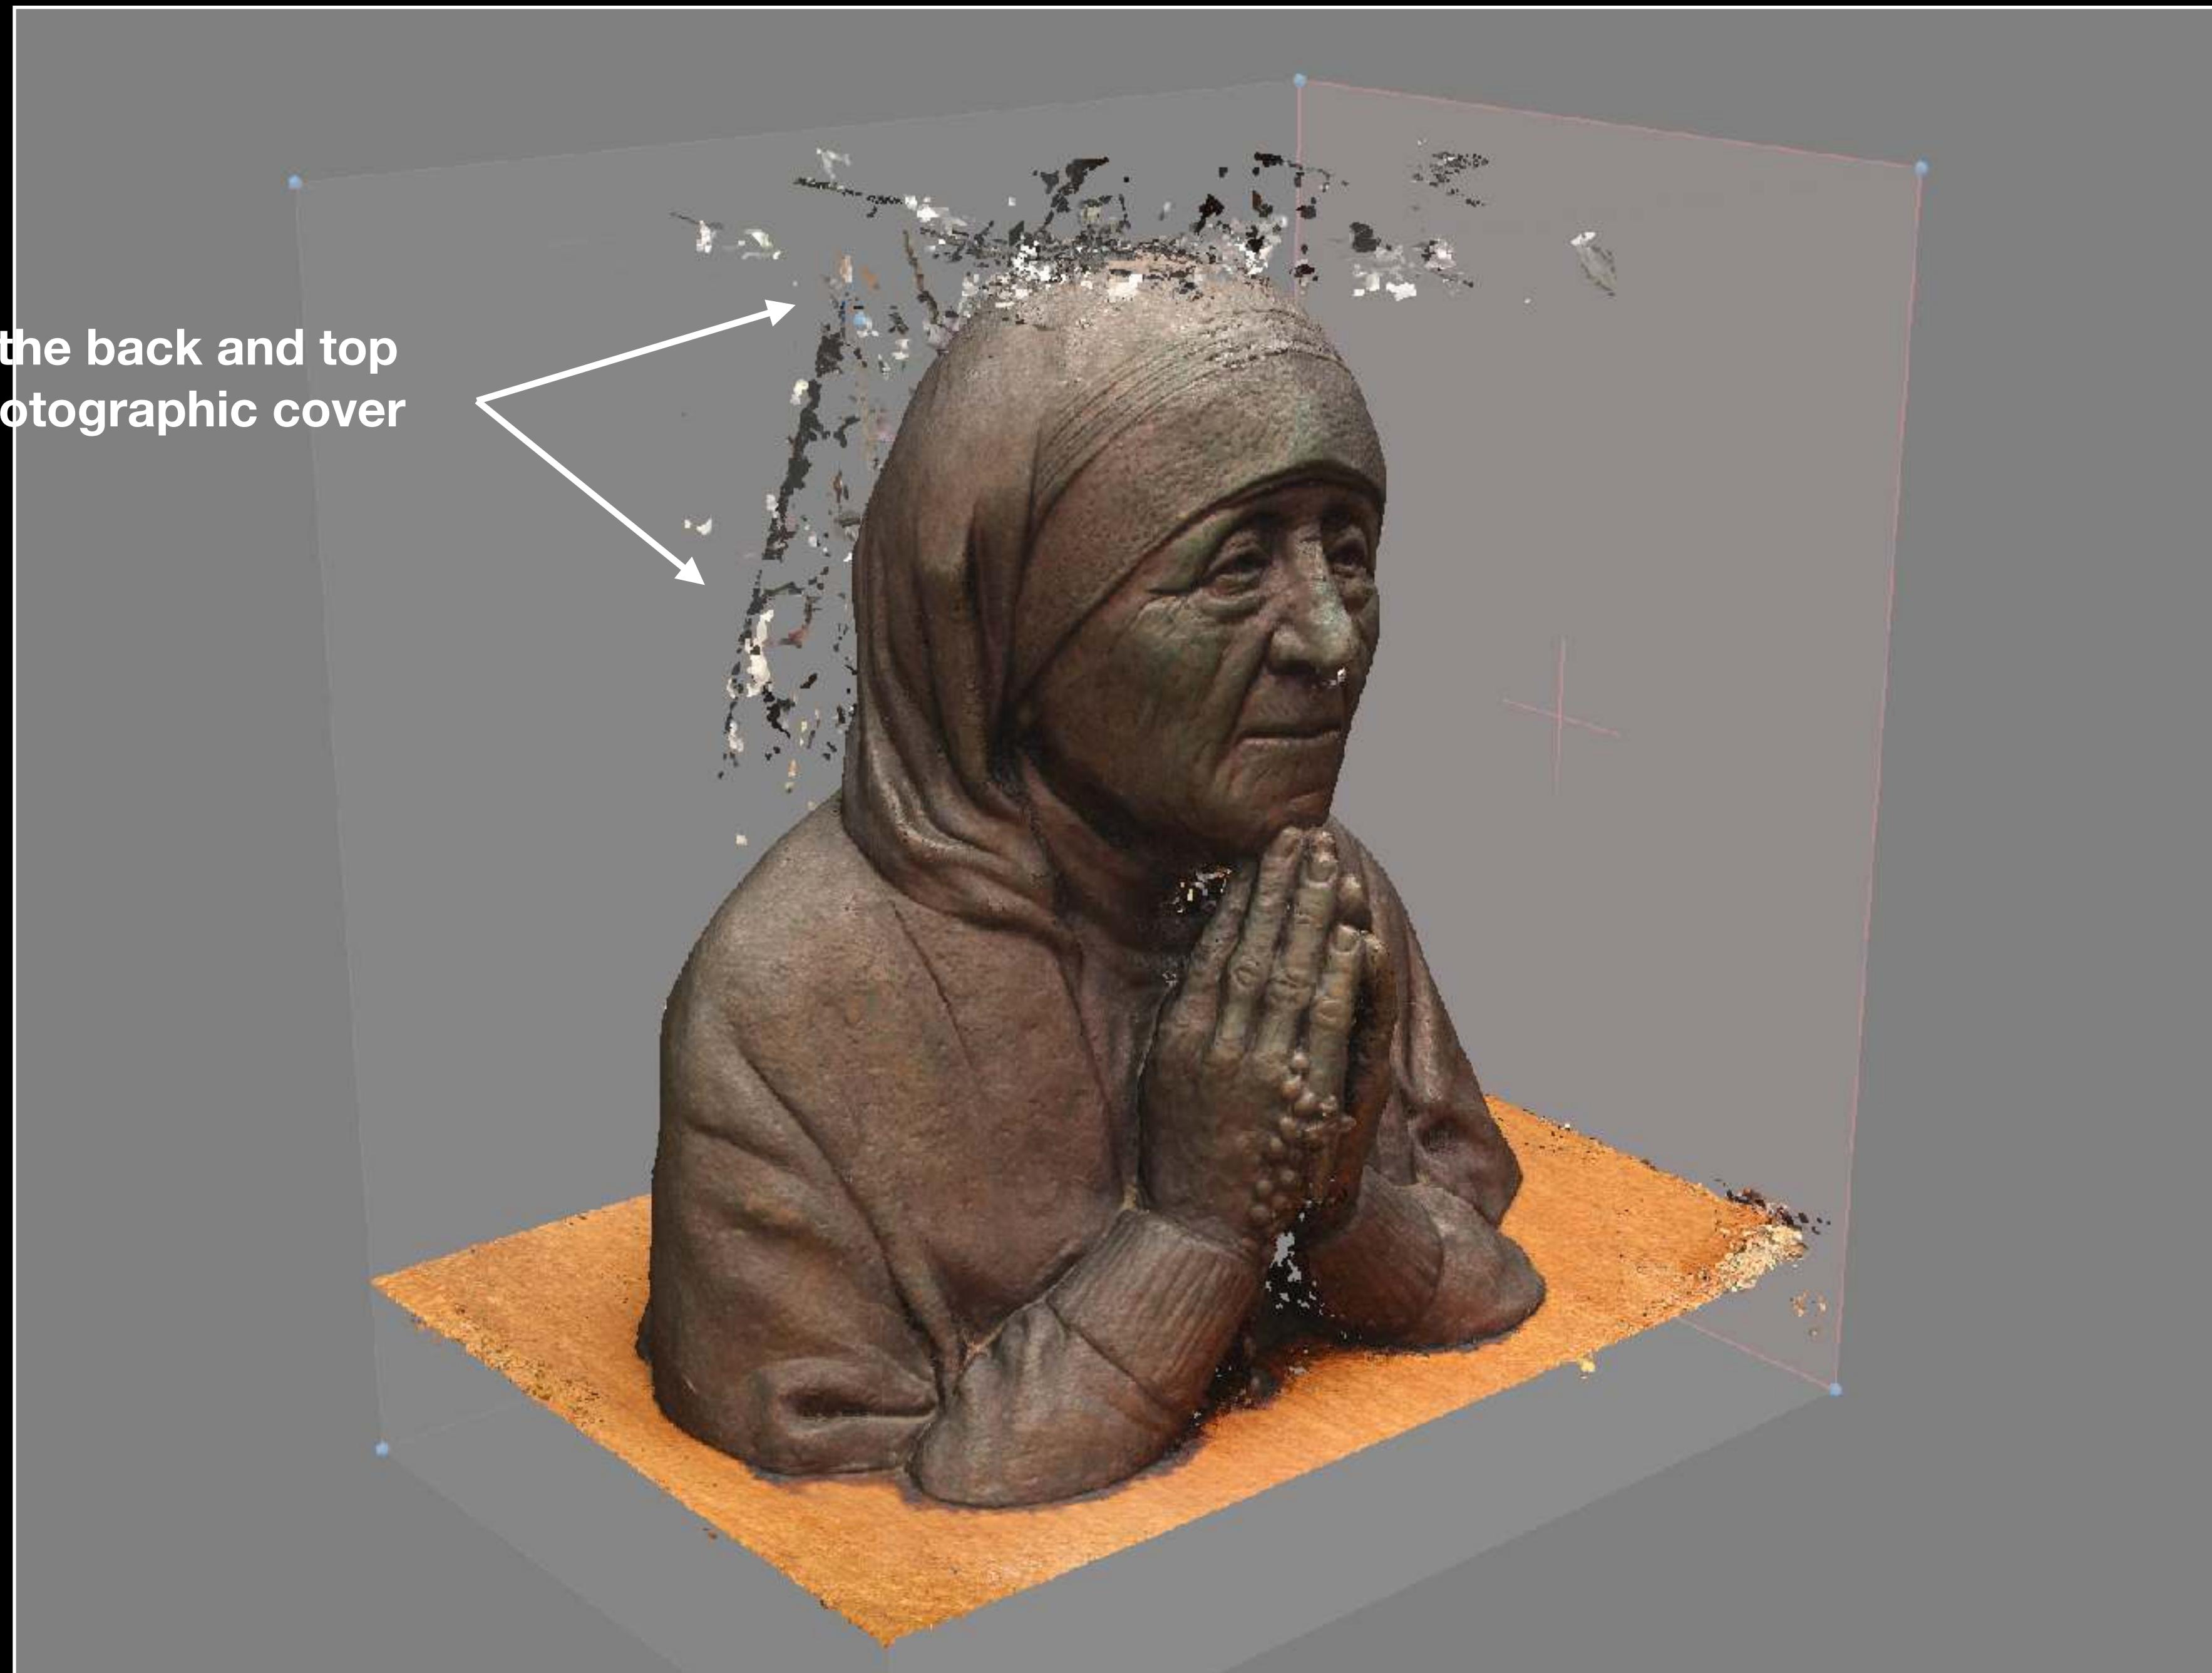
Rear and side walls



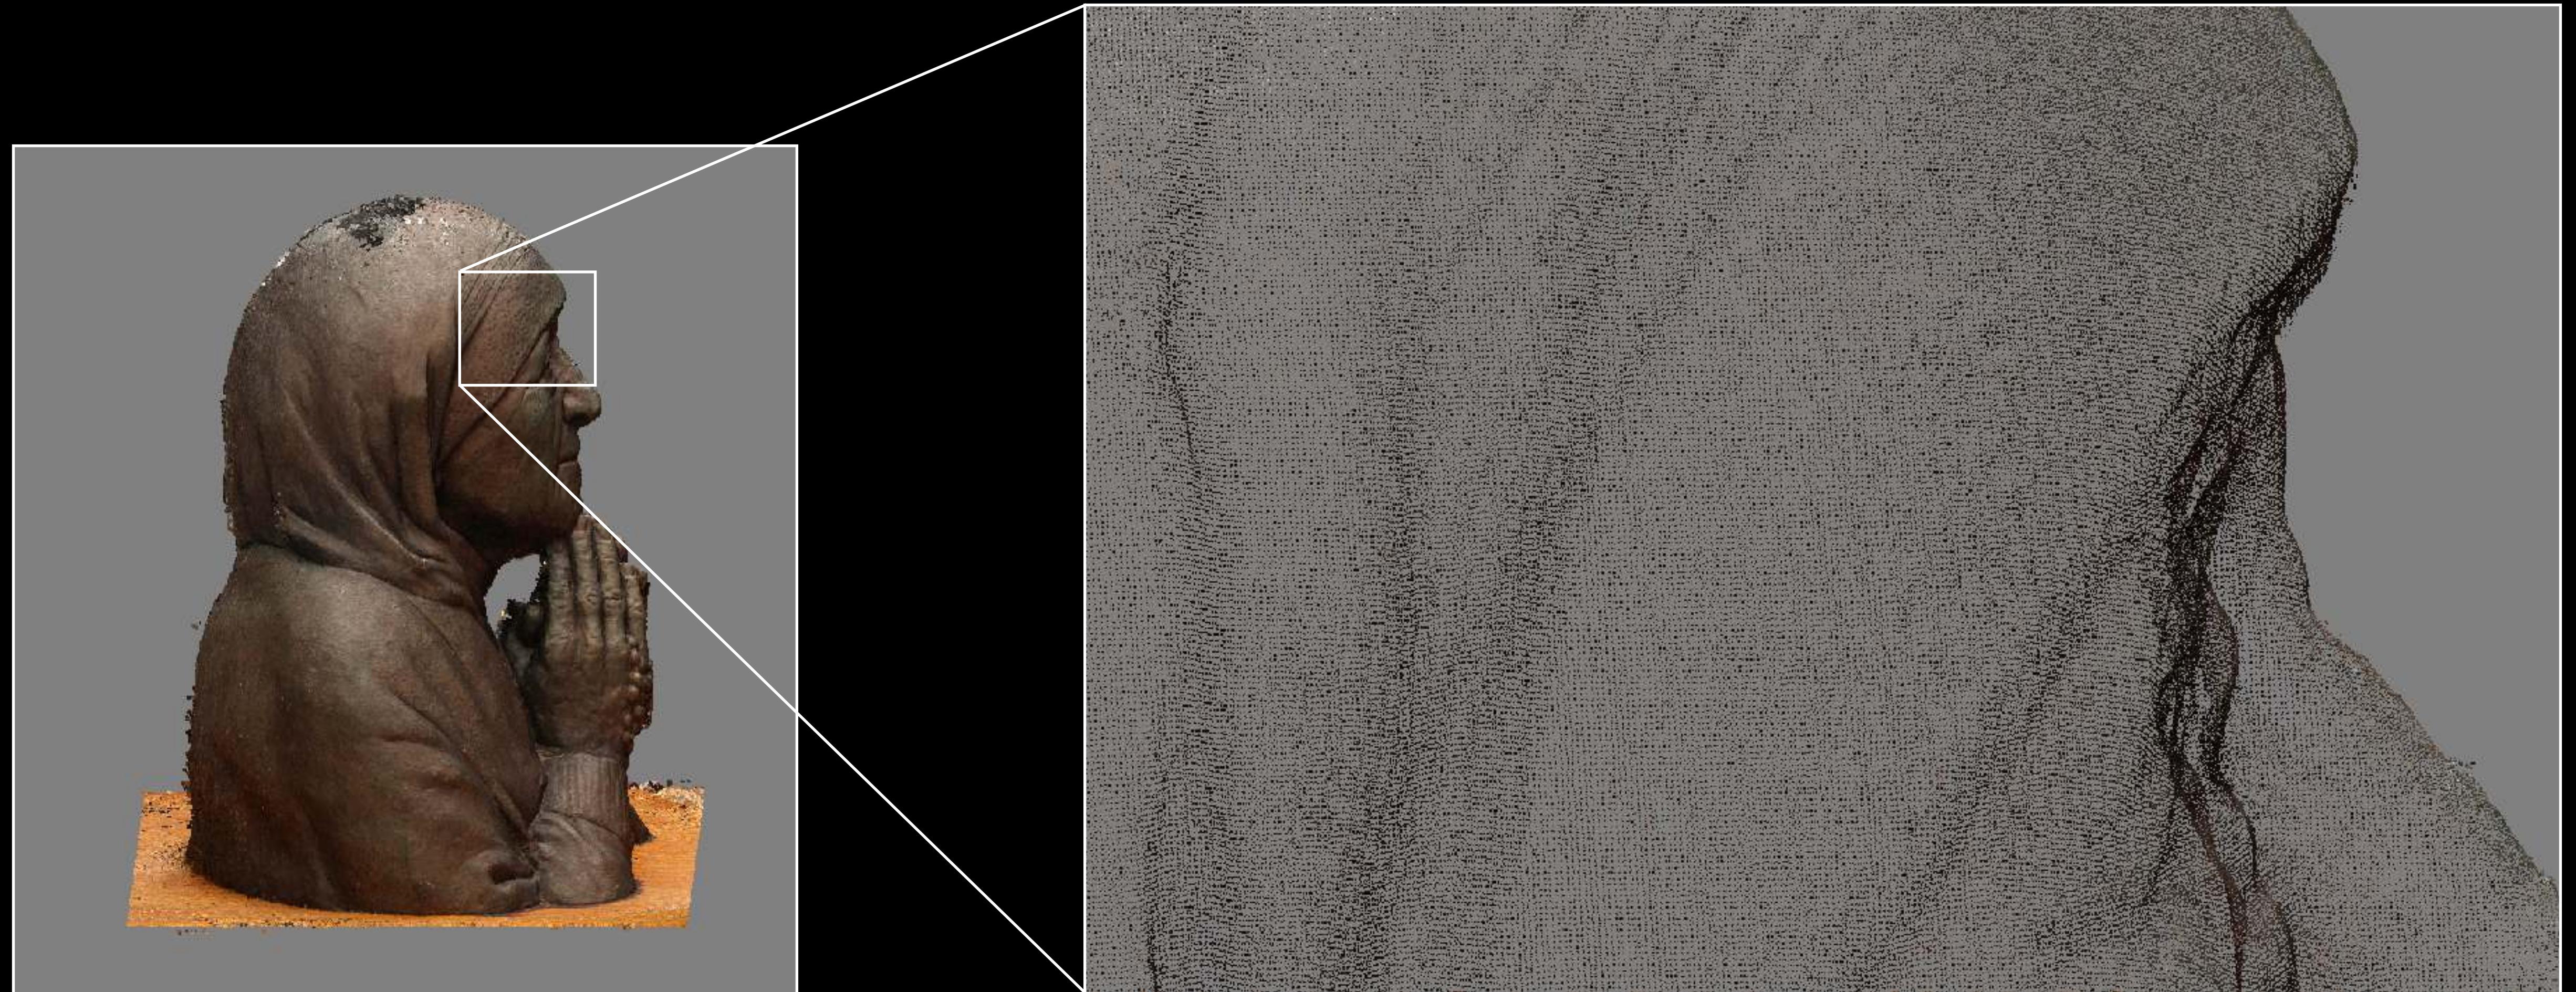




**Stray points at the back and top
due to lower photographic cover**

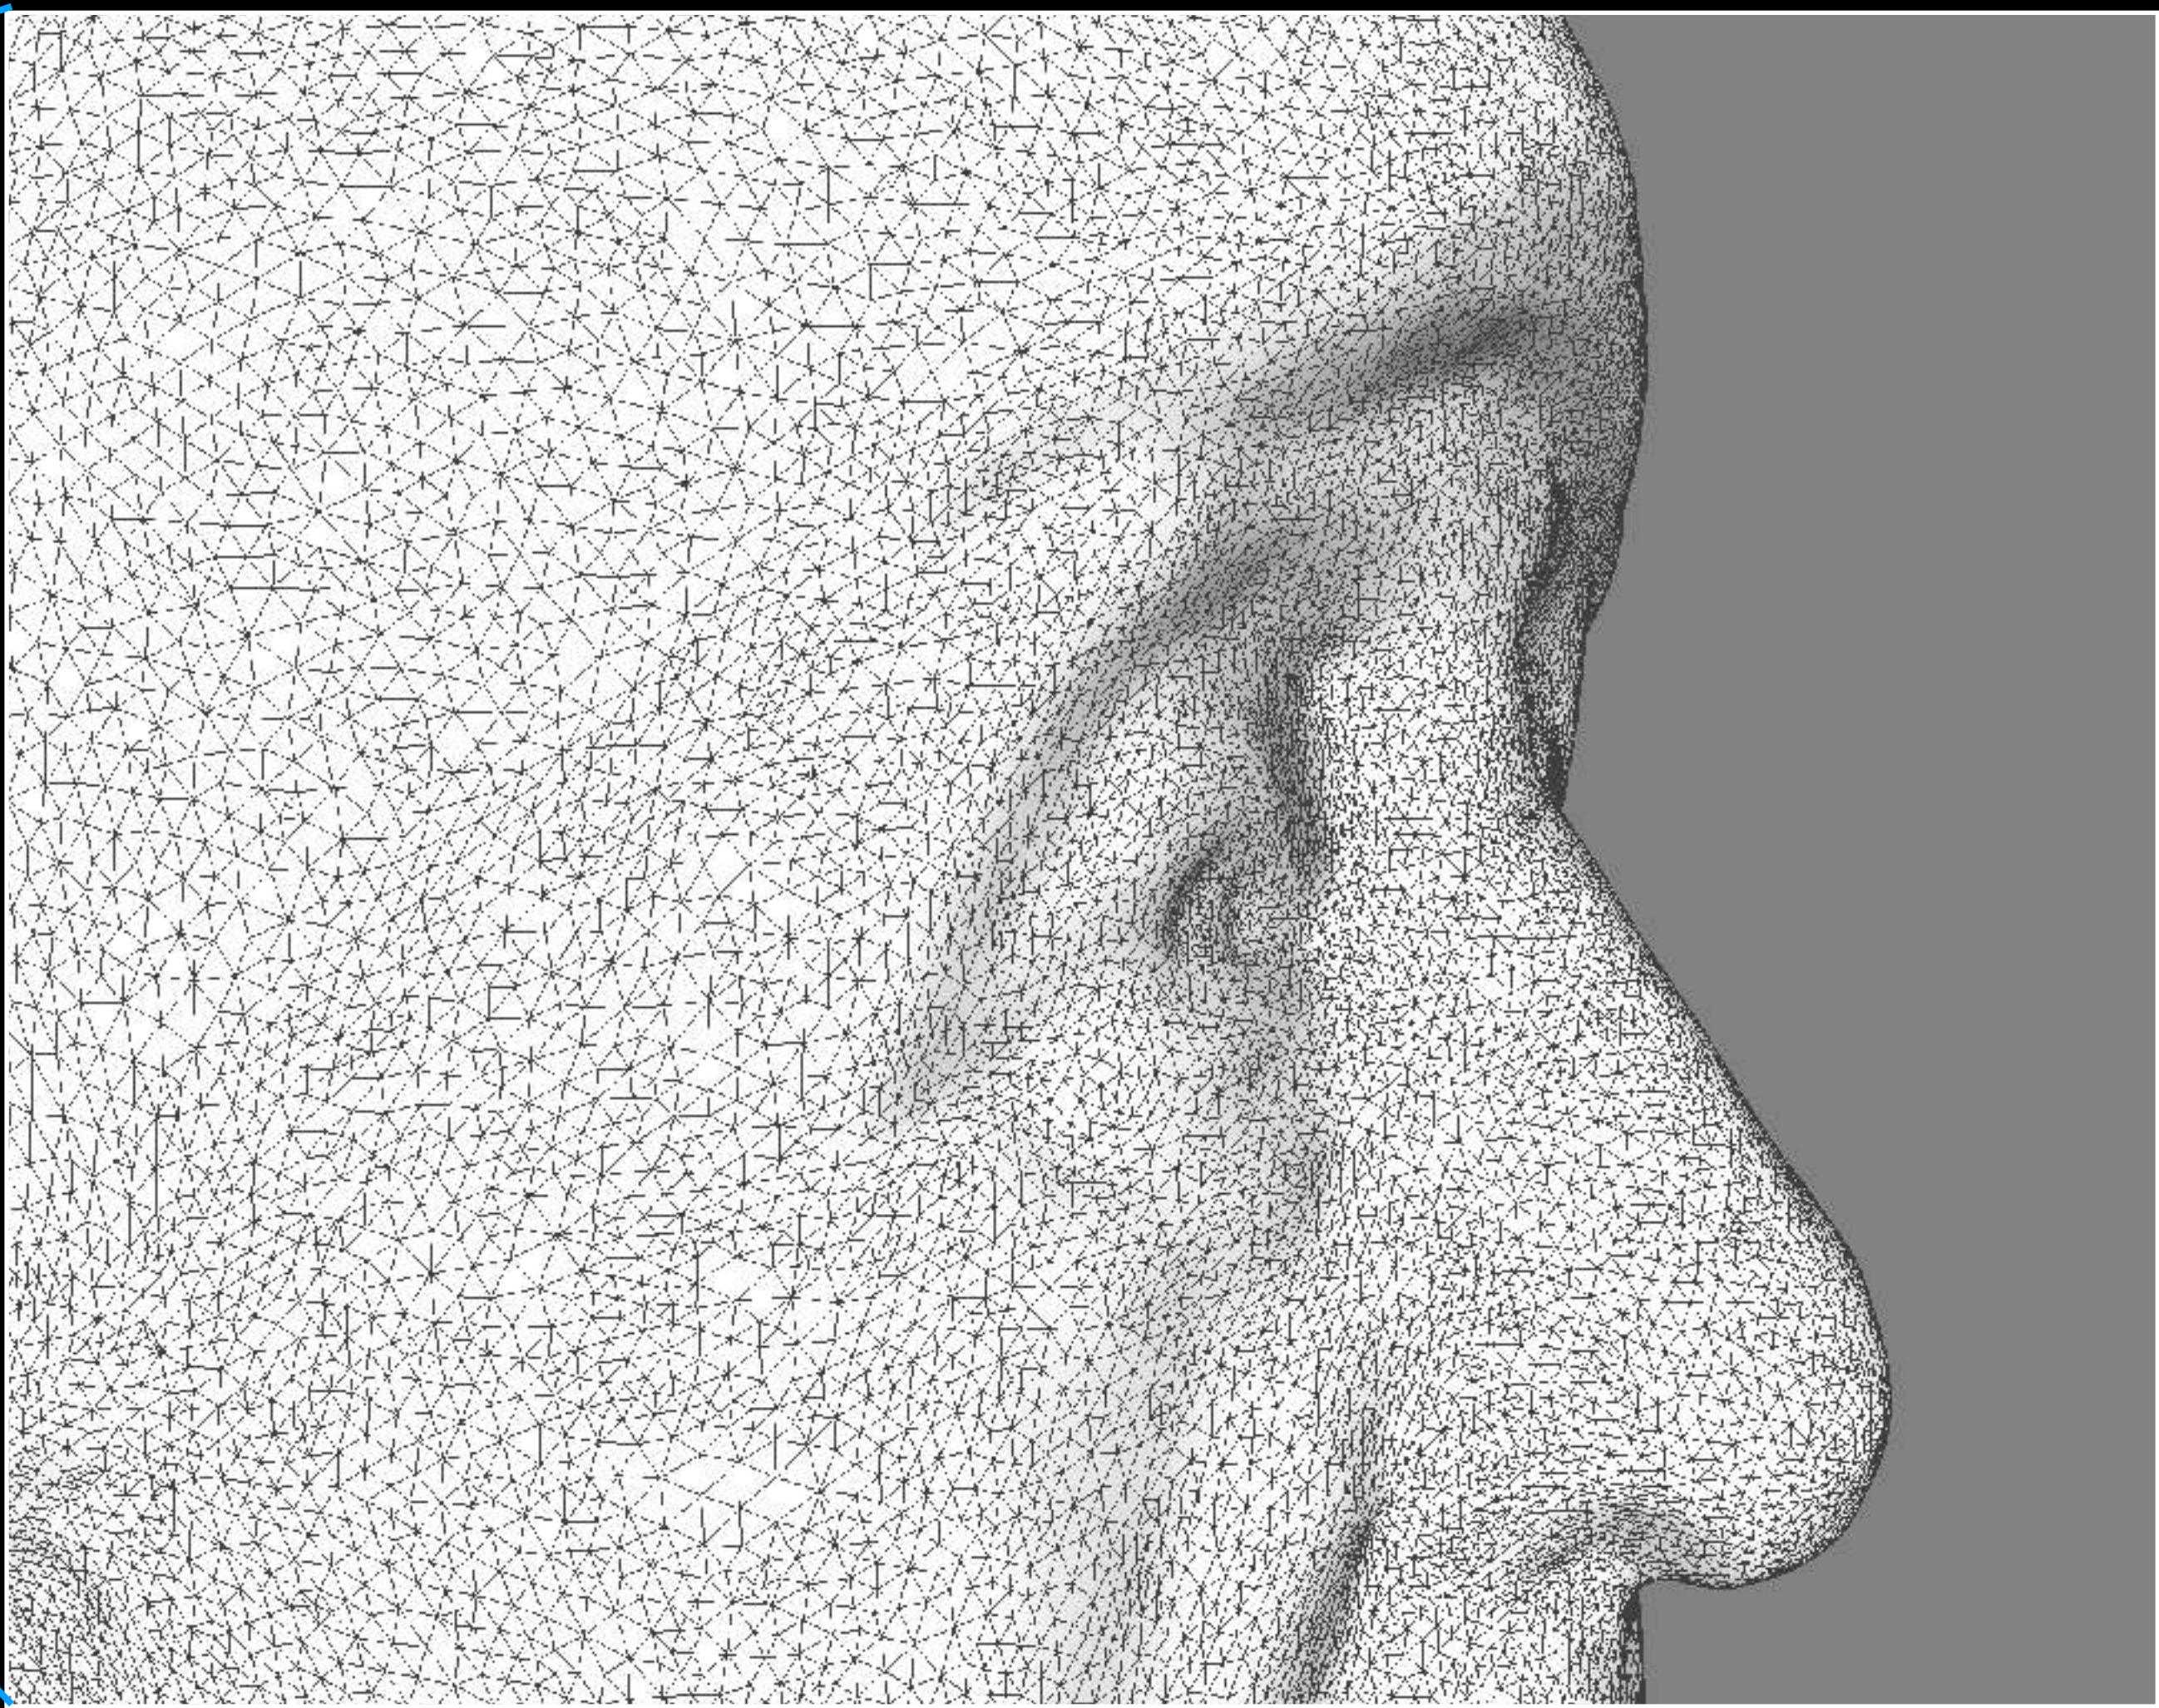
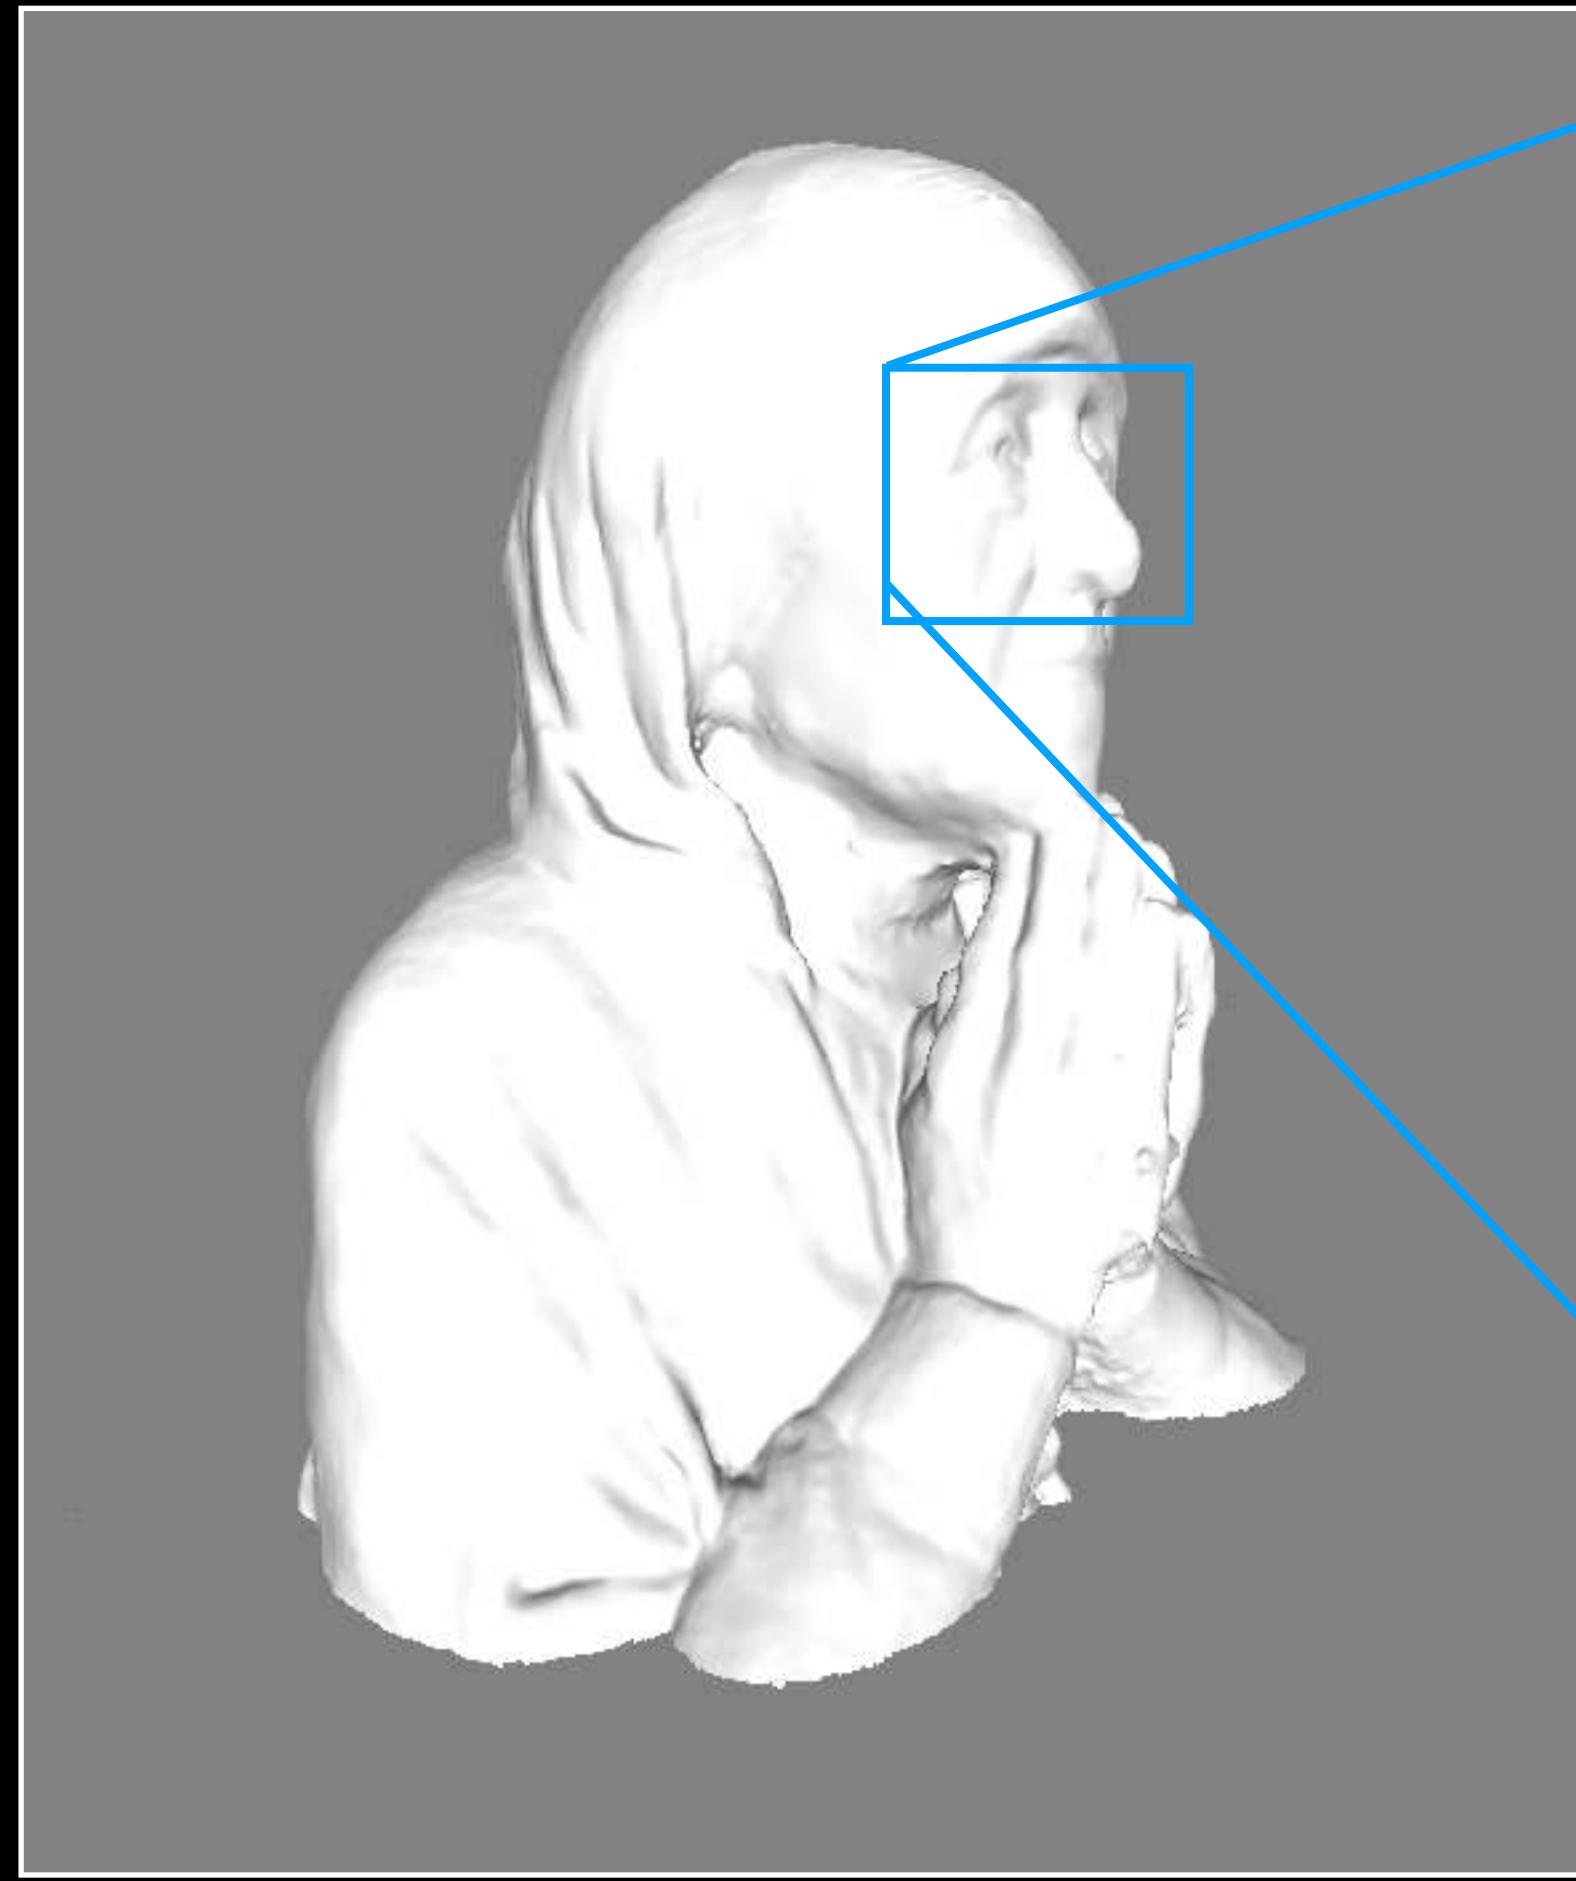






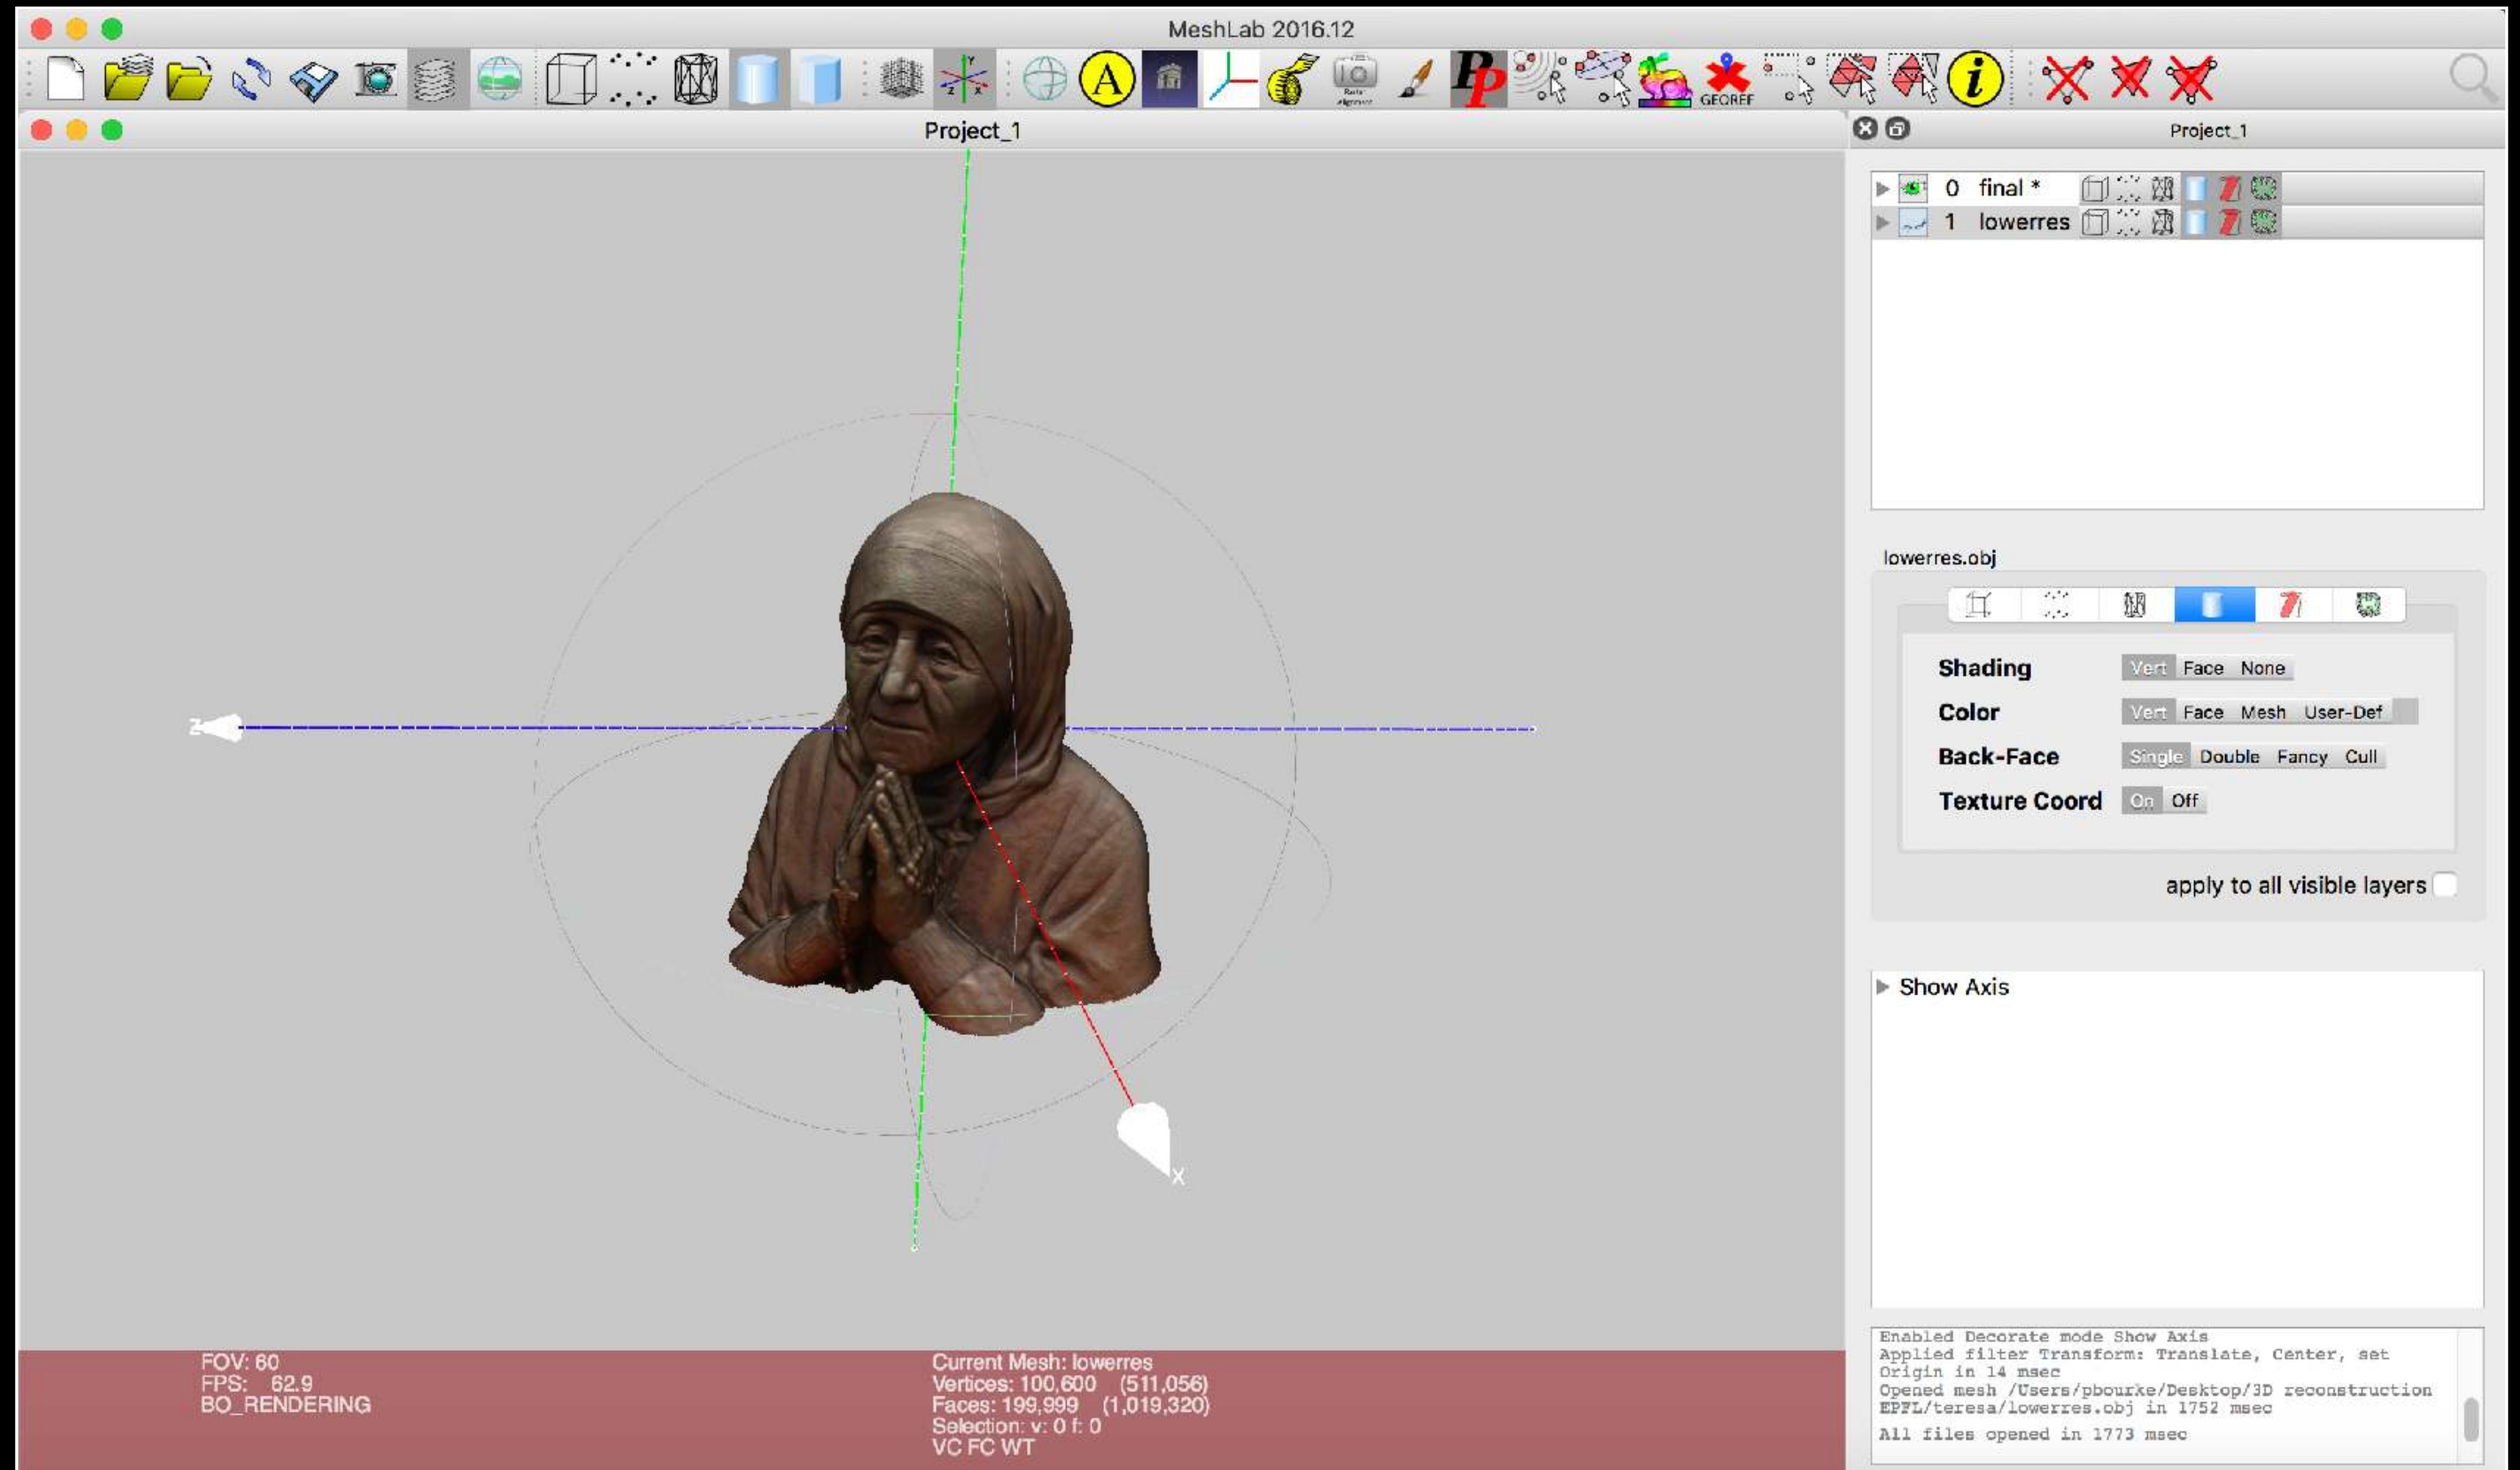
3,500,000 points

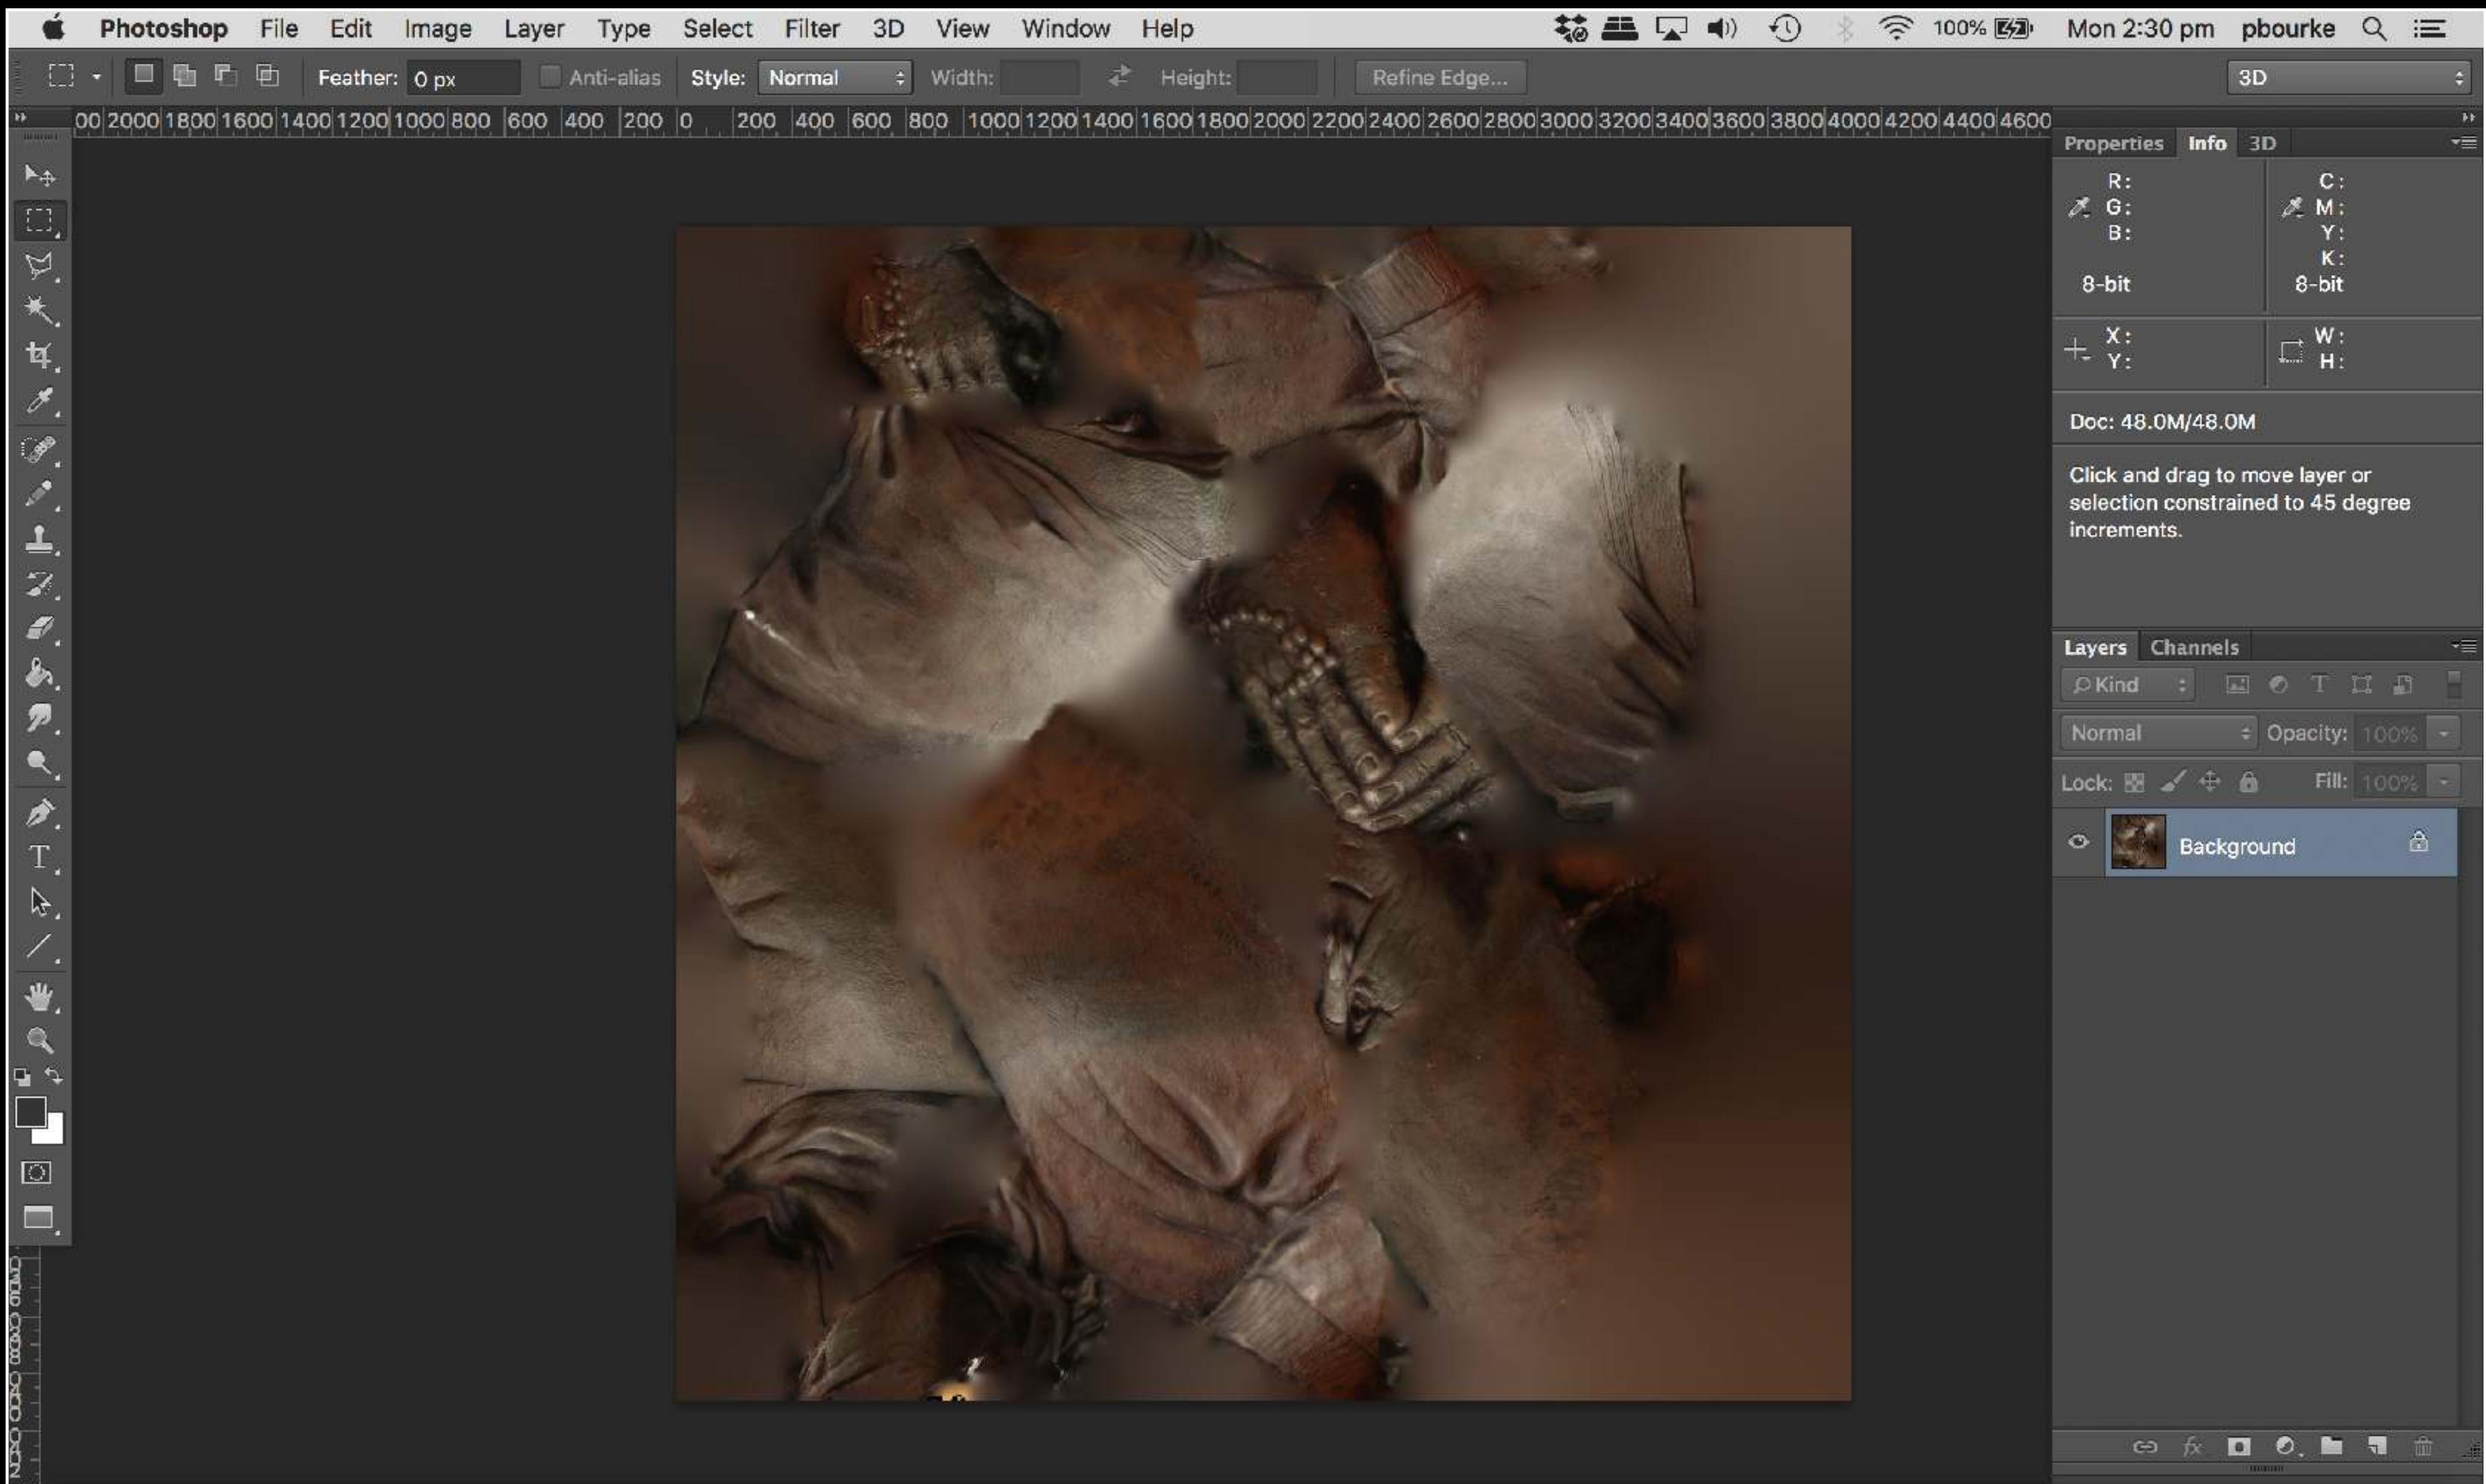














Mother Teresa

3D Model



pbourke PRO

177

★ 7

- Secure payment
- Support from sellers
- Access to future versions

MANAGE THIS MODEL ▾

Status Published

Privacy Public [edit](#)

Download

No

Free

Paid

Price 

License

Standard

[Understand licenses](#)

Additional file

Textures, other
3D formats, etc.[Upload a file](#)**Accepted formats:** .zip, .rar, .7z**Max size:** 500MB

PROPERTIES

3D SETTINGS

MORE ▾

Currently editing **Mother Teresa**[Help Center](#)[EXIT](#)[SAVE SETTINGS](#)[SAVE VIEW](#)

▼ WORLD SCALE

Scale

2.8128

[USE 1 UNIT = 1 METER](#)[RESET SCALE](#)**Tips for world scale**Character is 1.75m tall. A floor tile is 1m².

A smaller world makes your model look larger



A large world makes your model look smaller

▼ FLOOR

 Display floor for teleportation

Floor level

-0.06m

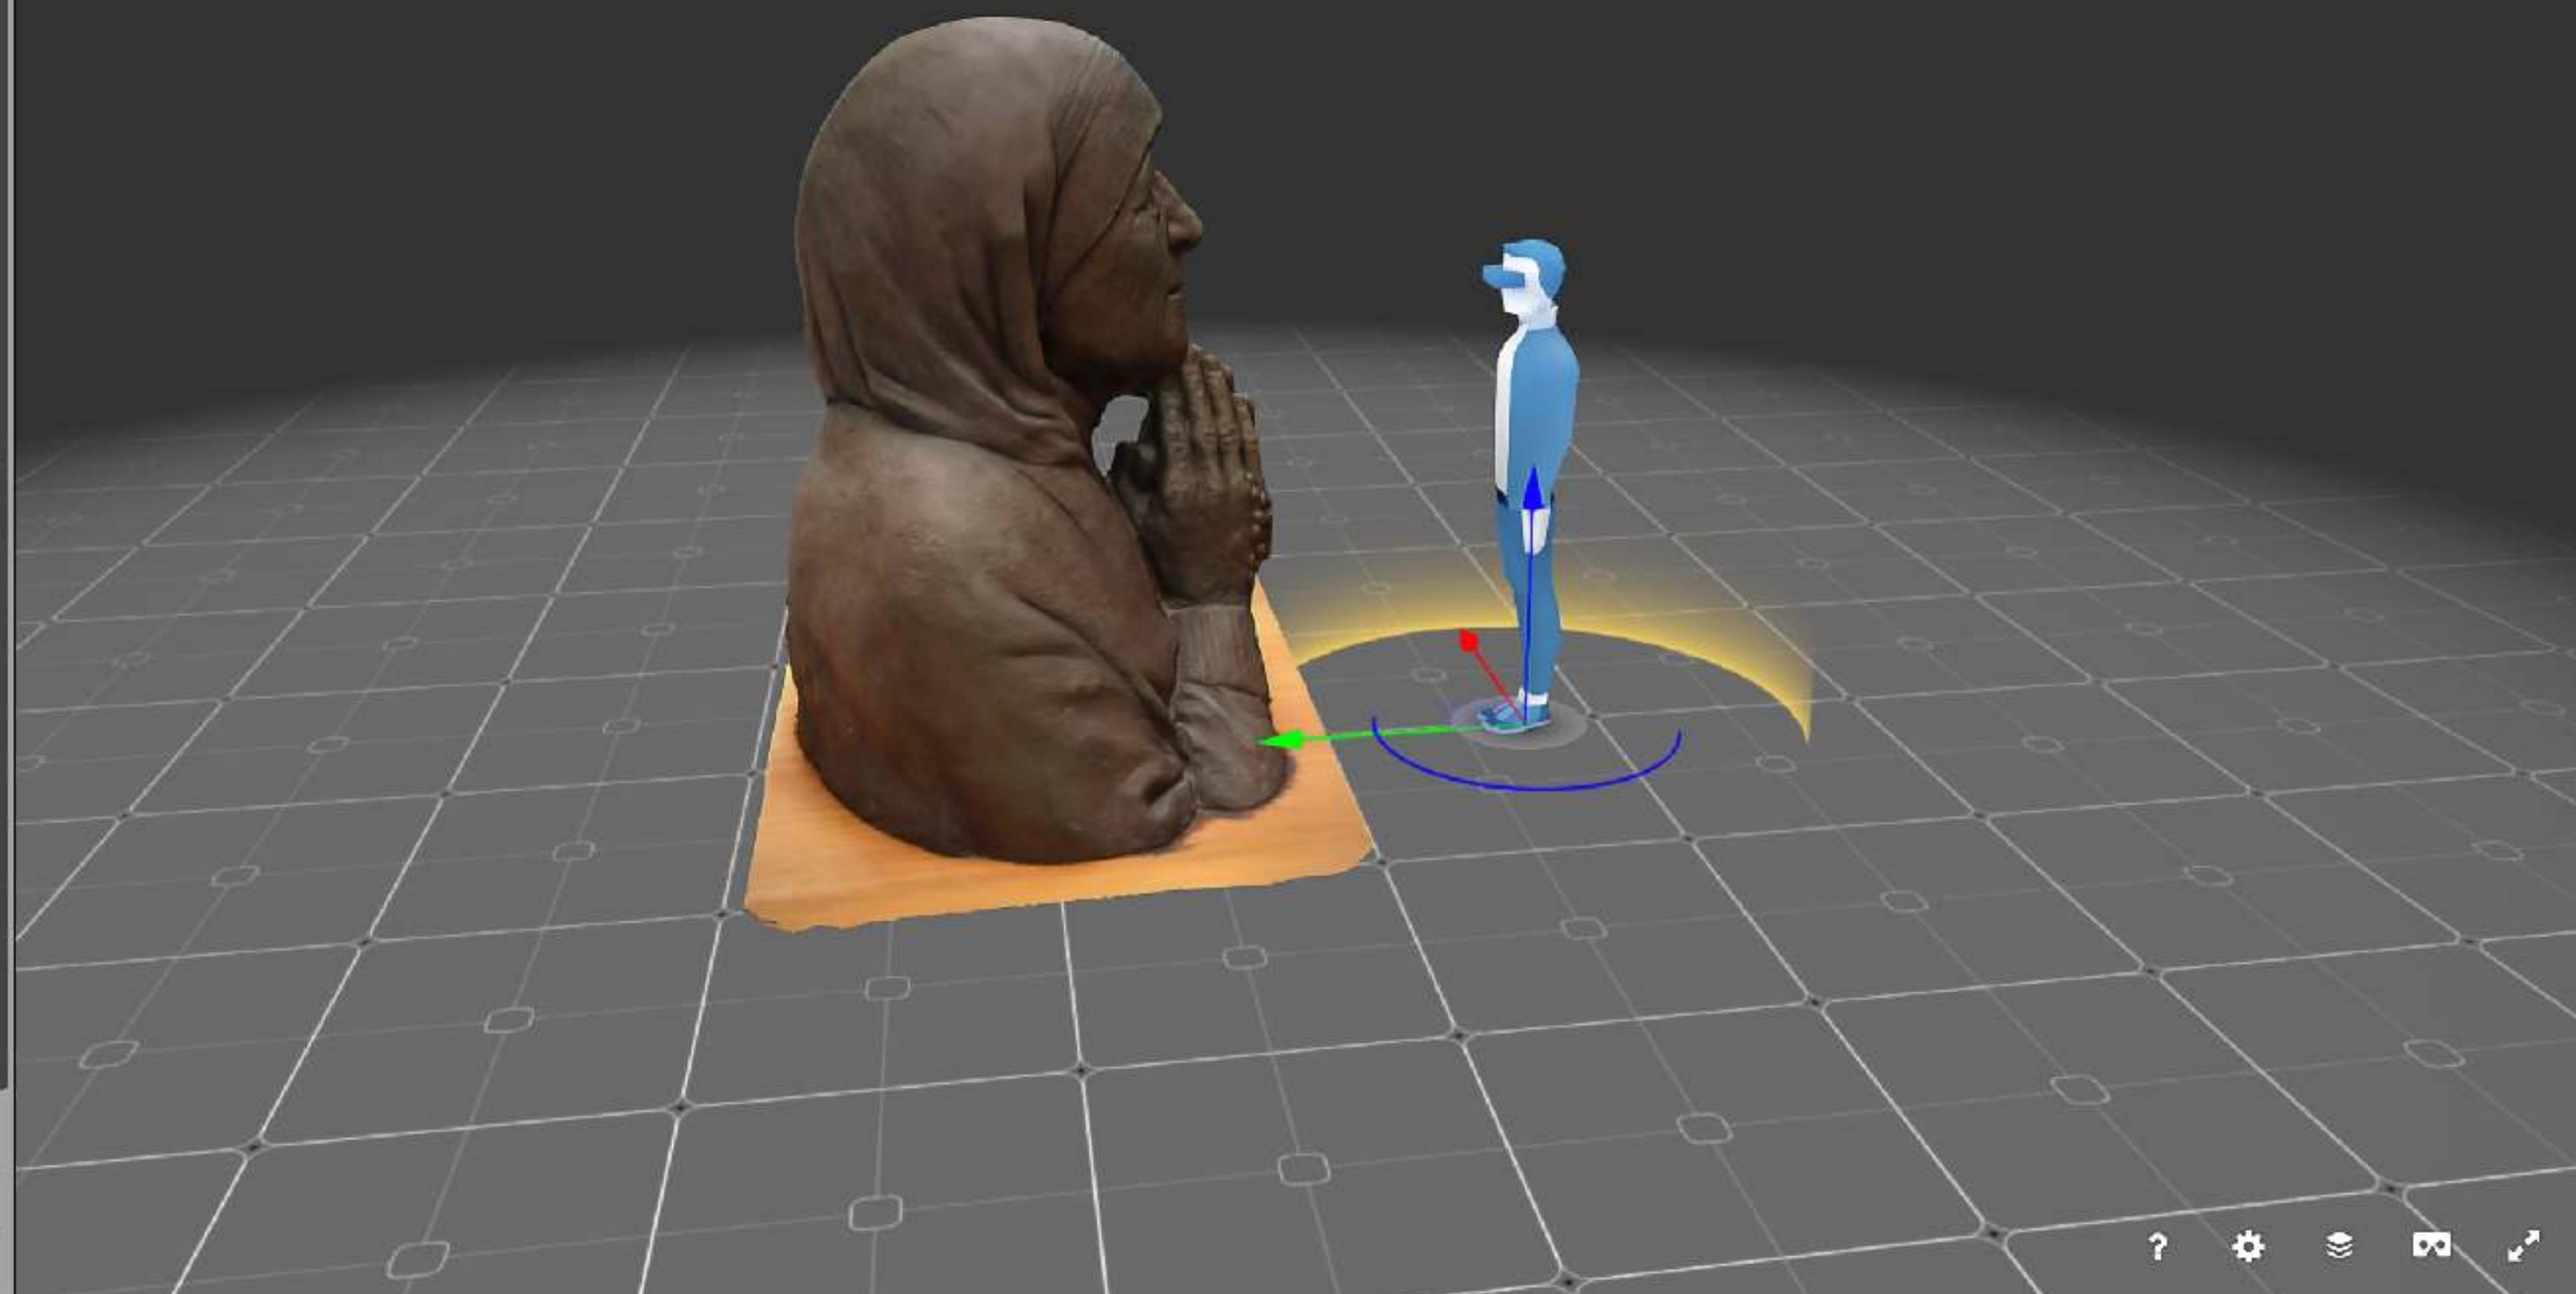
Tips for floor level

Position the floor right below large models



Position the floor further down for smaller objects

▼ INITIAL VIEWING POSITION





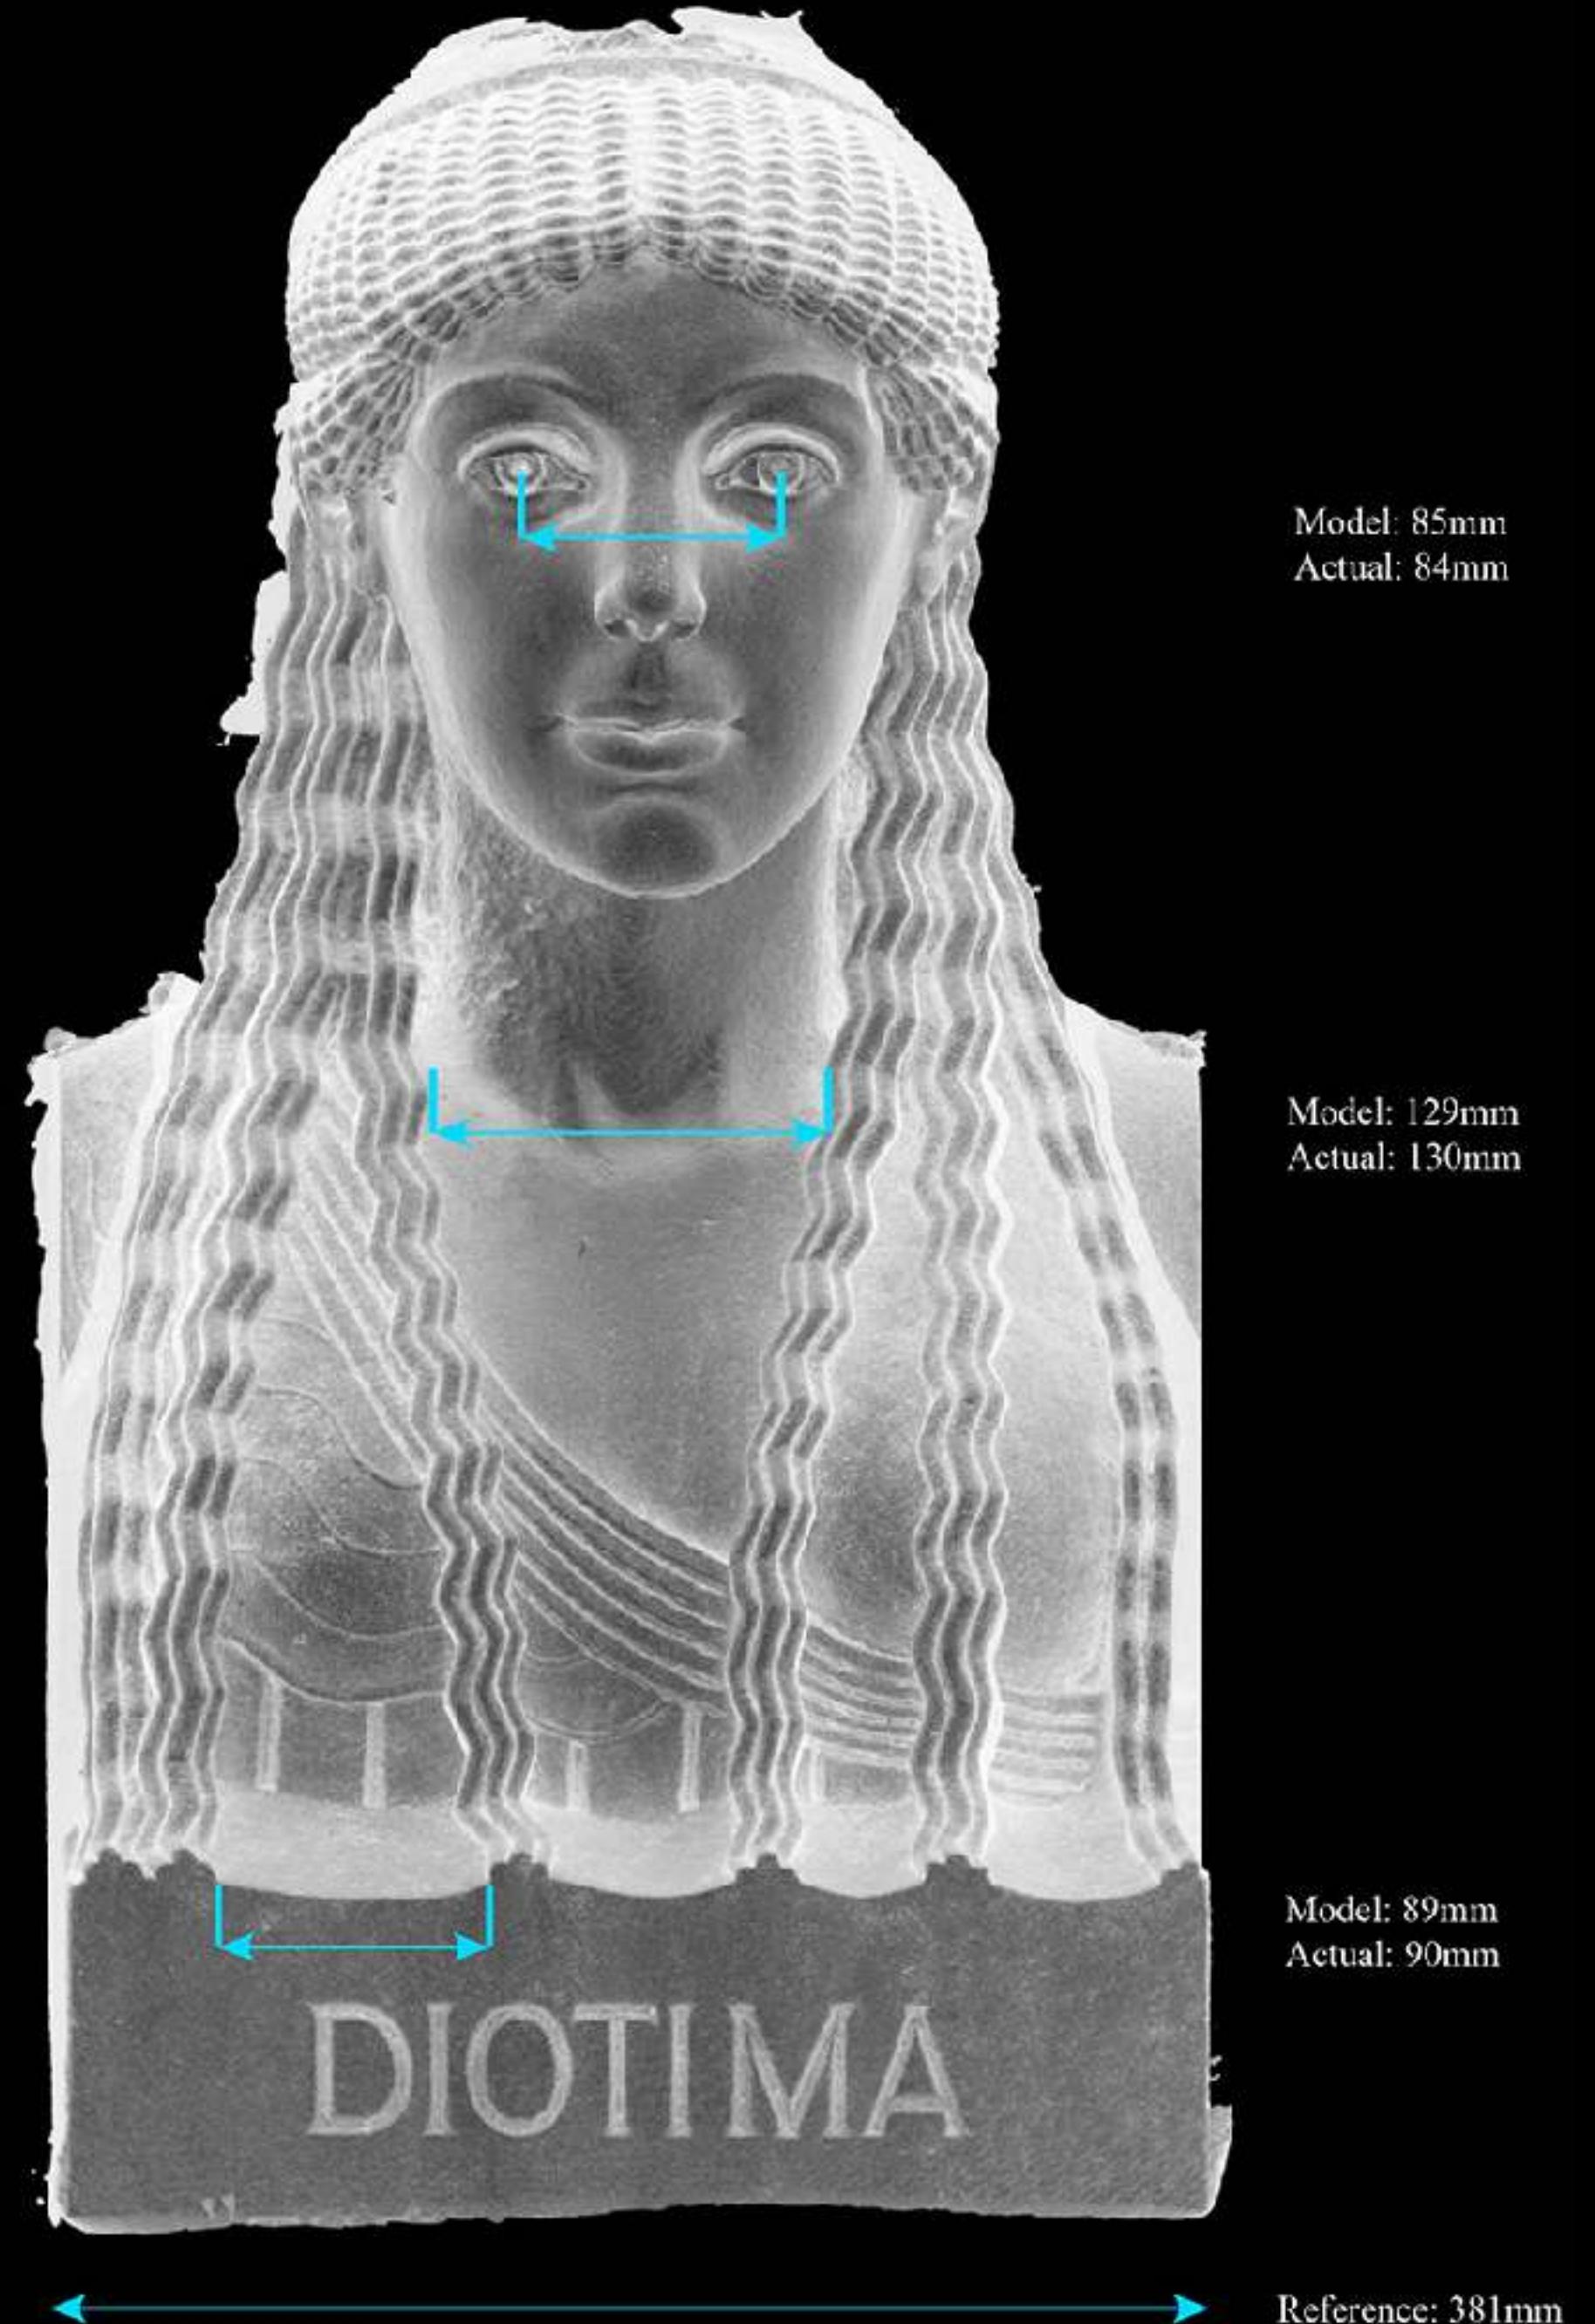
Other topics

- Resolution
- Masking
- Scale bars, colour correction
- Lens calibration
- Merging multiple scans
- Turntables, moving camera vs moving object
- Texture editing
- Data formats
- Limitations

Resolution, Accuracy

- The first question many people ask is how accurate is it?
- Not easy to quantify
 1. Not all parts of a model are equally accurate.
 2. Not always possible to measure ground truth to compare to?
 3. Accuracy can depend on material characteristics of the model itself.
 4. How can you be sure best practice and best technology have been used?
- We have tested three methods to determine accuracy
 1. Reconstruct and compare key measures with known object.
 2. Perform ensemble reconstructions from large image sets, compare variation.
 3. Compare with other scanning techniques: laser scanning, CT, structured light.
 4. Visual comparison of zoomed in photographs of real and reconstructed.

- No absolute scale but use one length as reference.
- Model from 60 images.
- Subsequent measurements accurate to maximum 2mm, average 1mm.





Original photograph



Reconstructed model



Shaded to
emphasise
surface variation



Original photograph

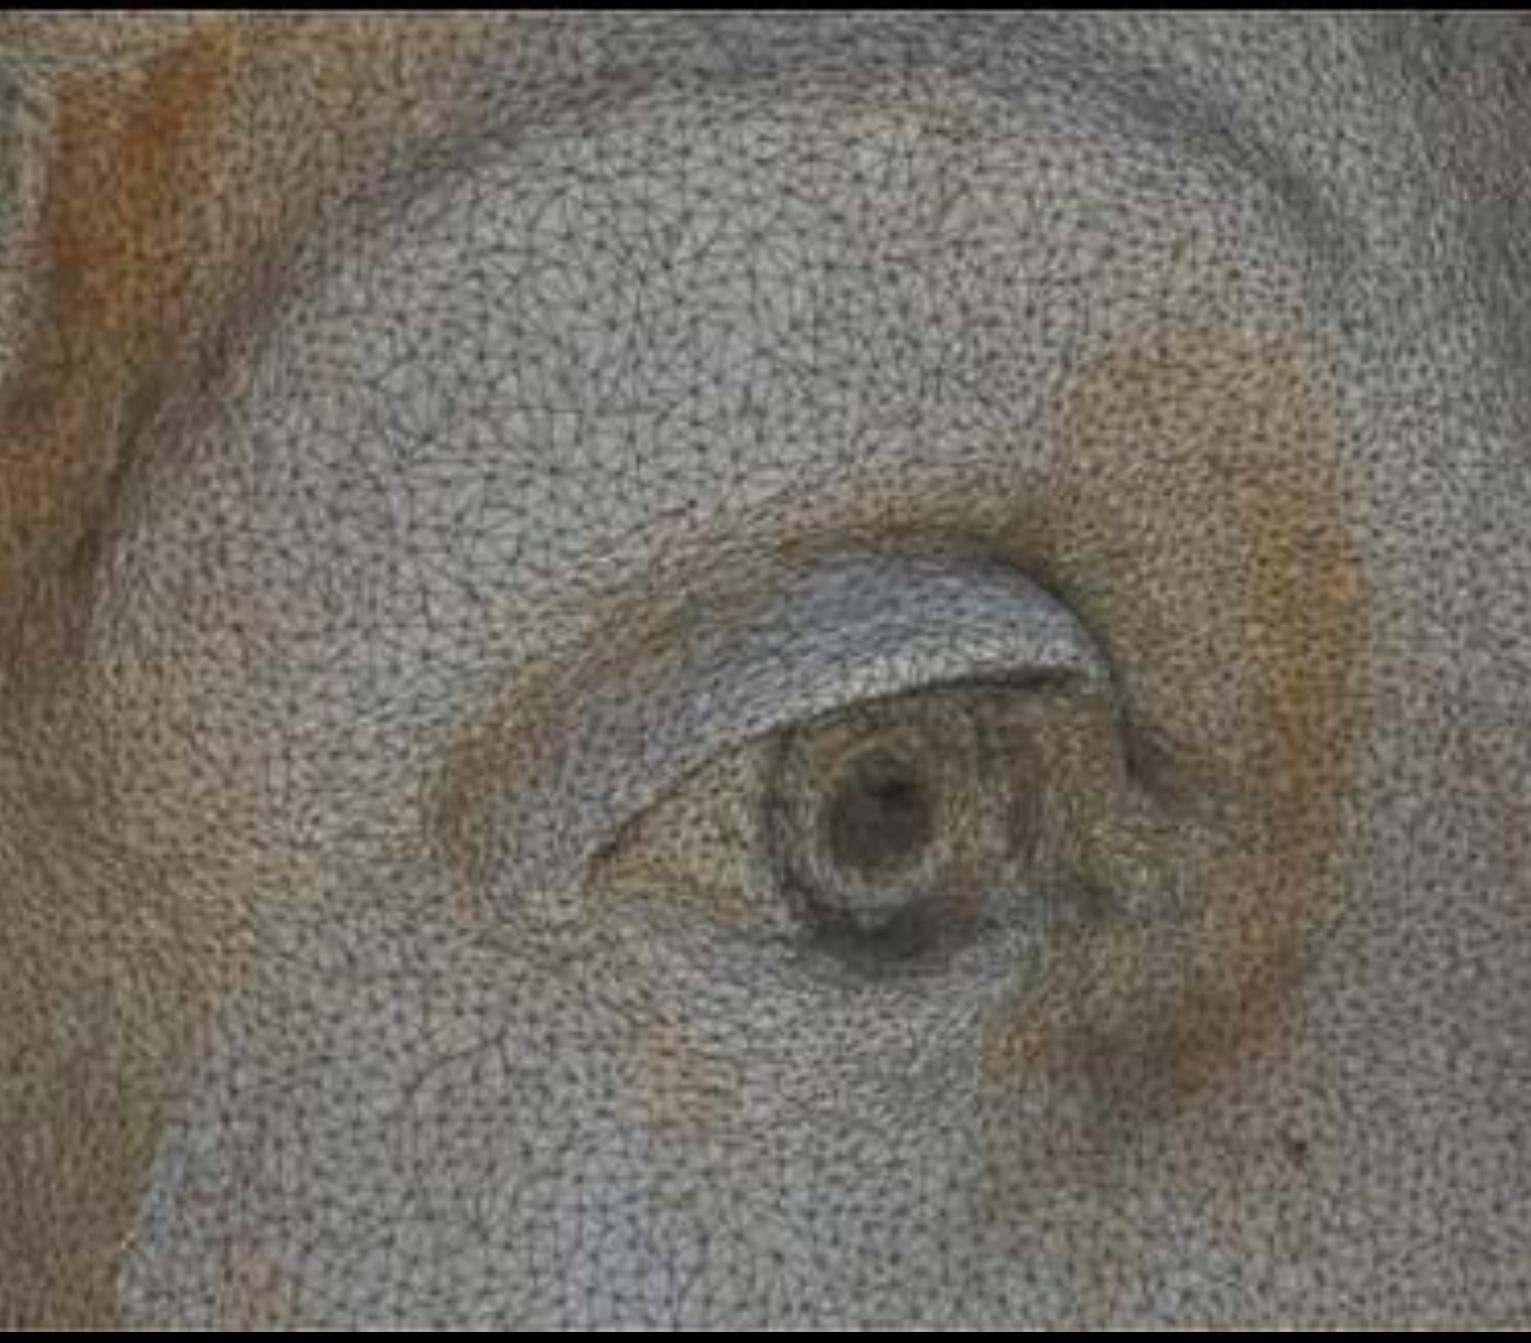


Reconstructed model



Shaded to
emphasise
surface variation

Original



Mesh and texture resolution

- What geometric (mesh) resolution is required to capture the (important) features.
- Texture resolution is similar to the more familiar image resolution, what texture image size is required to capture the visual detail at the intended zoom level.
- Requirements vary depending on the end application.
 - Realtime environments prefer low geometric complexity and high texture detail
 - Analysis generally requires high geometric detail
 - For a digital record one typically wants high geometric and texture detail

	Geometric resolution	Texture resolution
Gaming	Low	High
Analysis	High	Don't care
Education	Medium	High
Archive/heritage	High	High
Online	Low/Average	Low/average



1,000,000 triangles



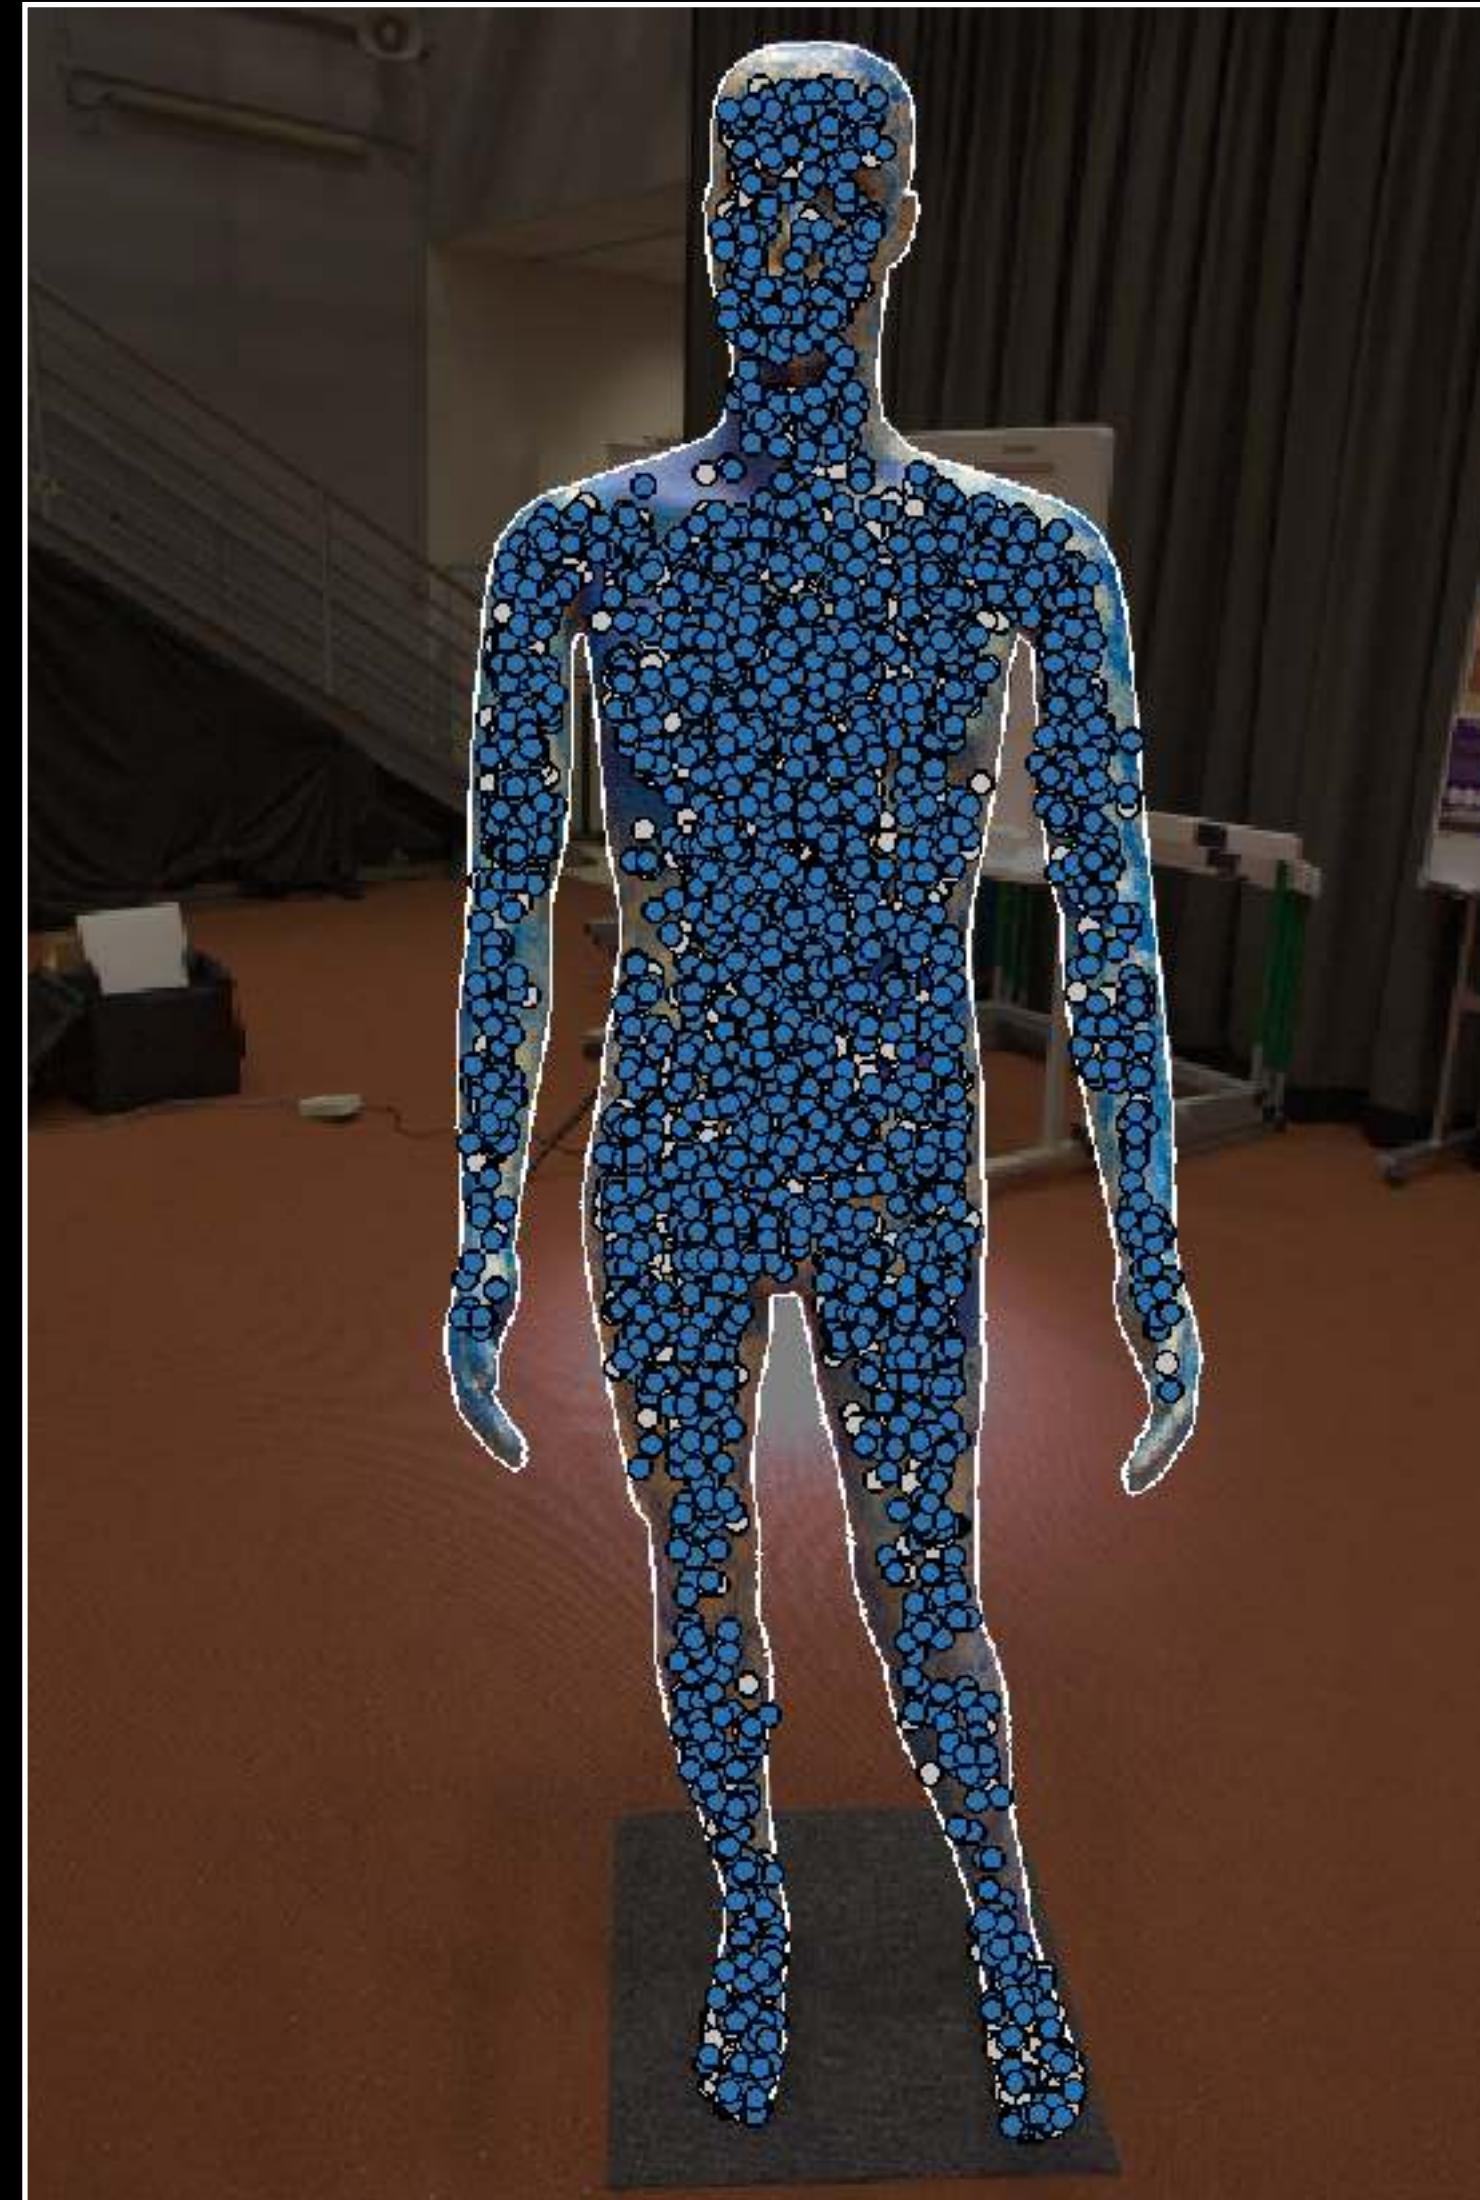
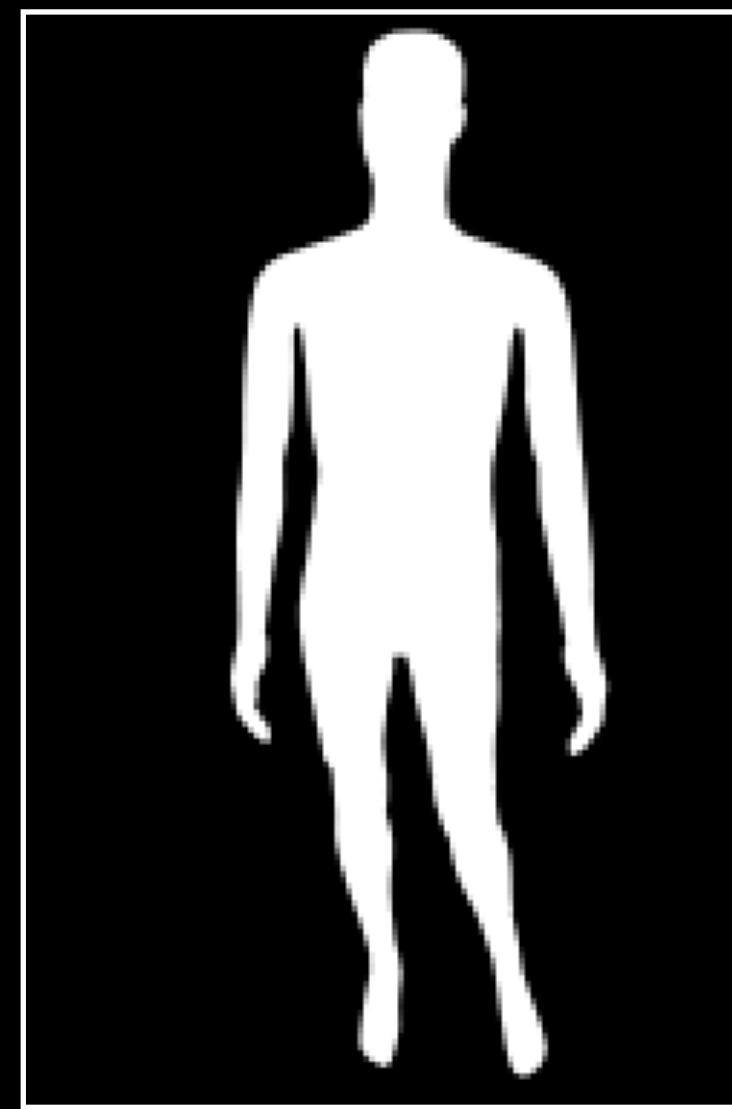
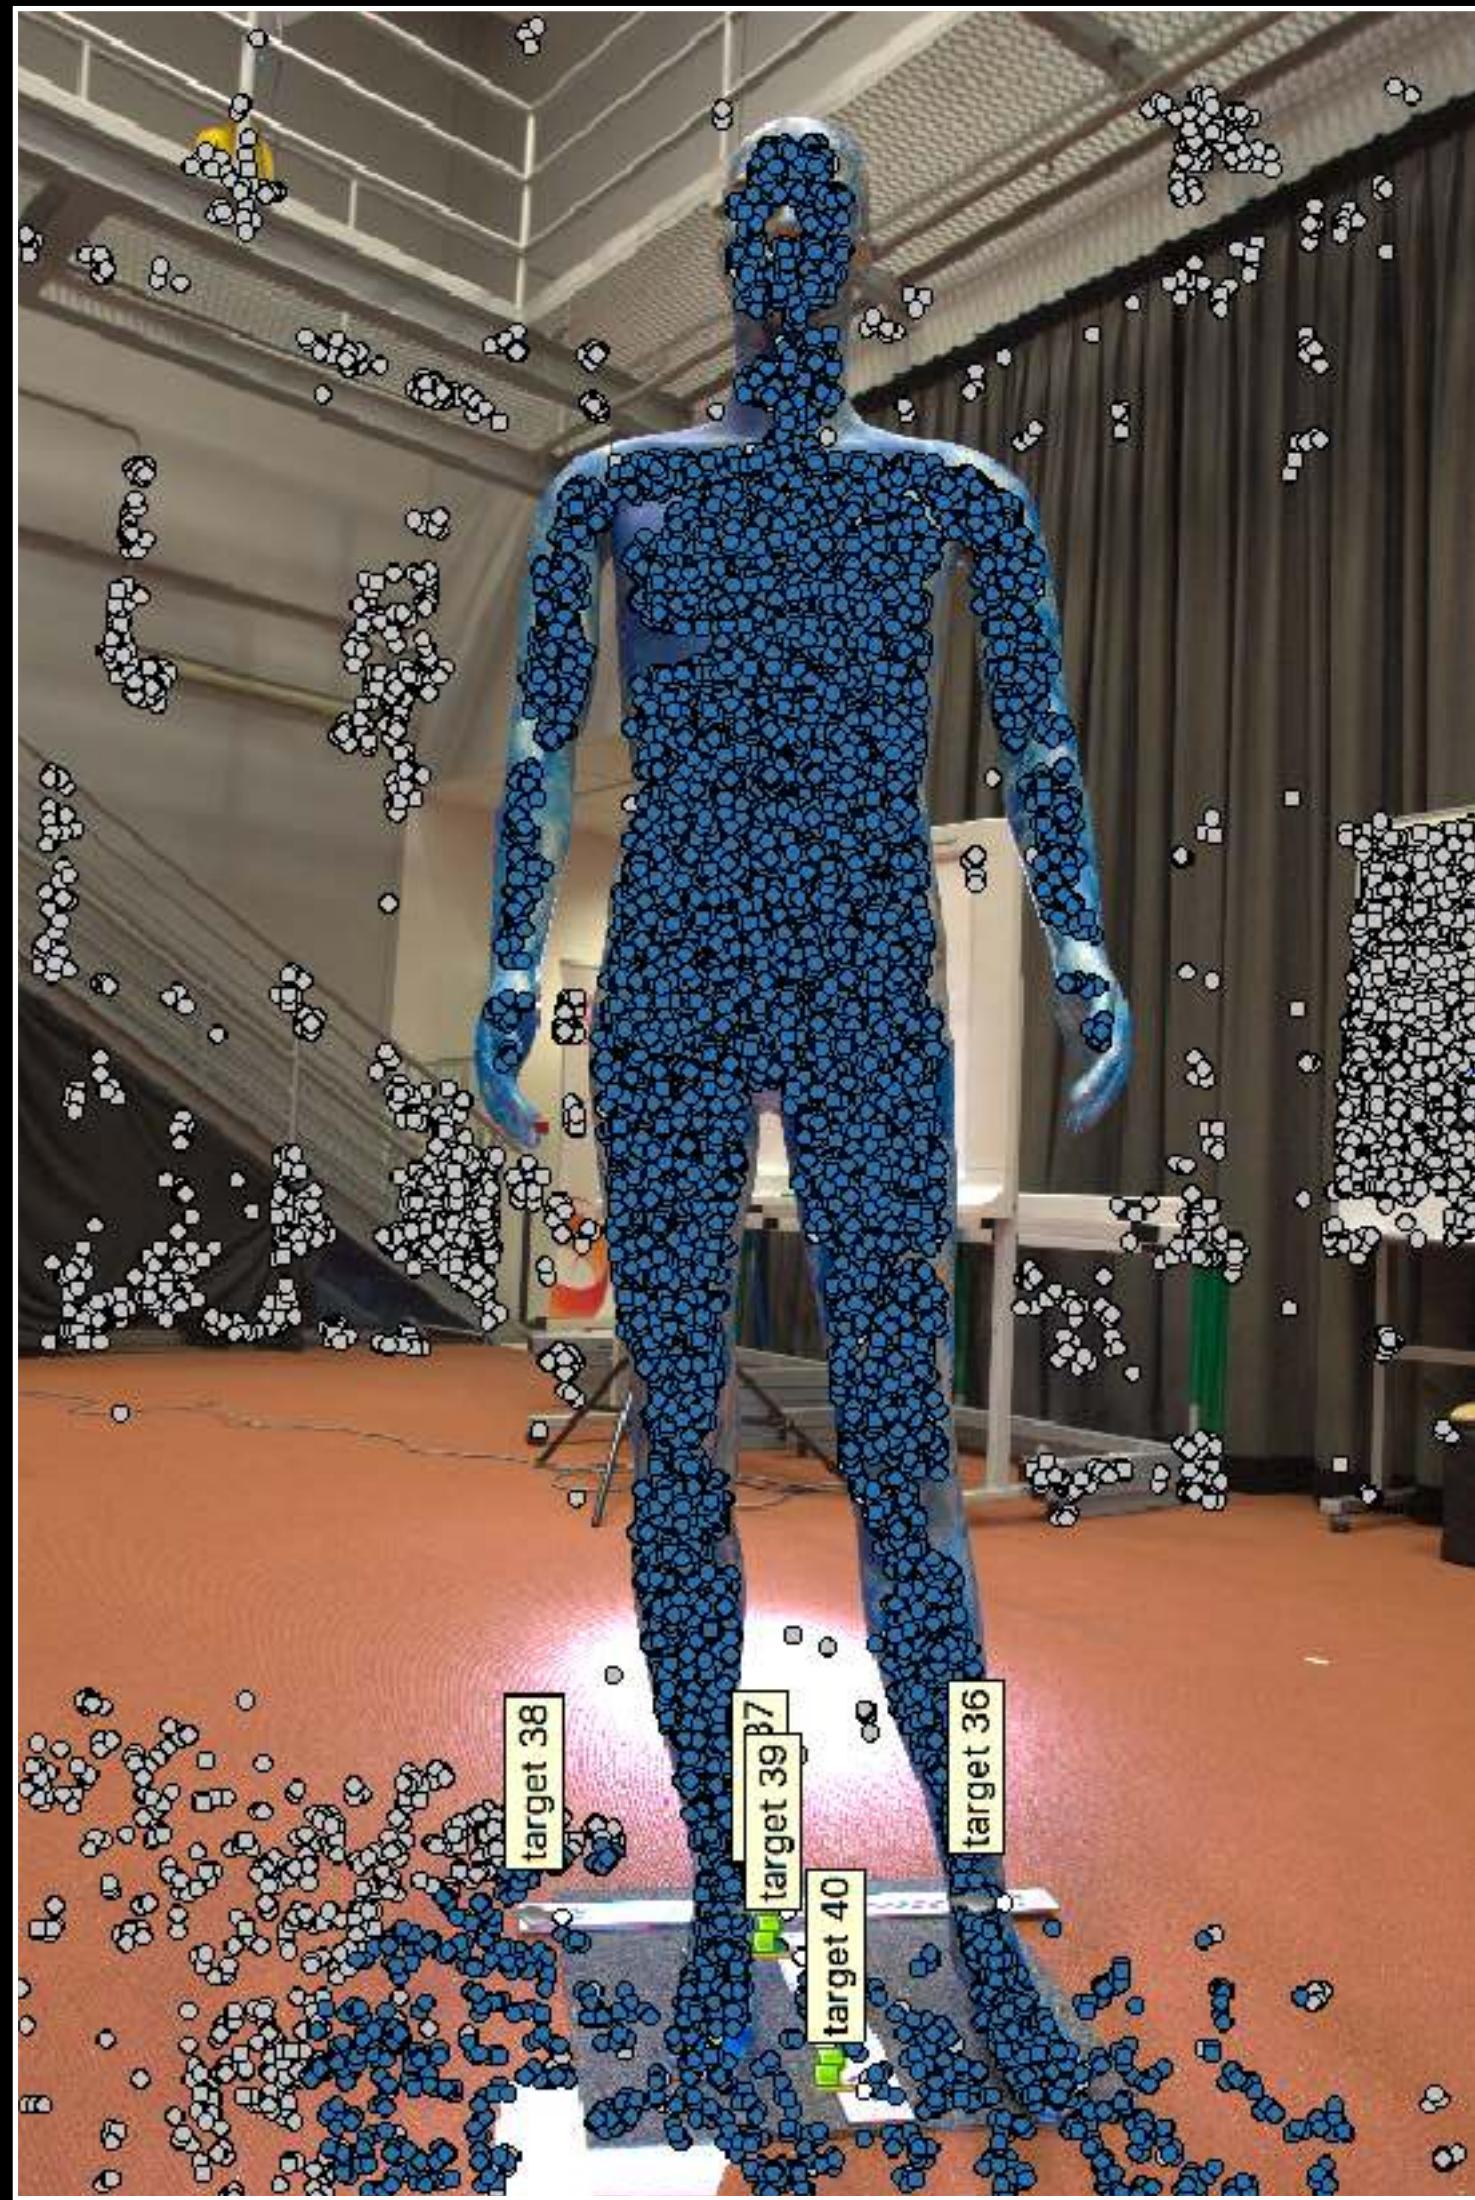
100,000 triangles



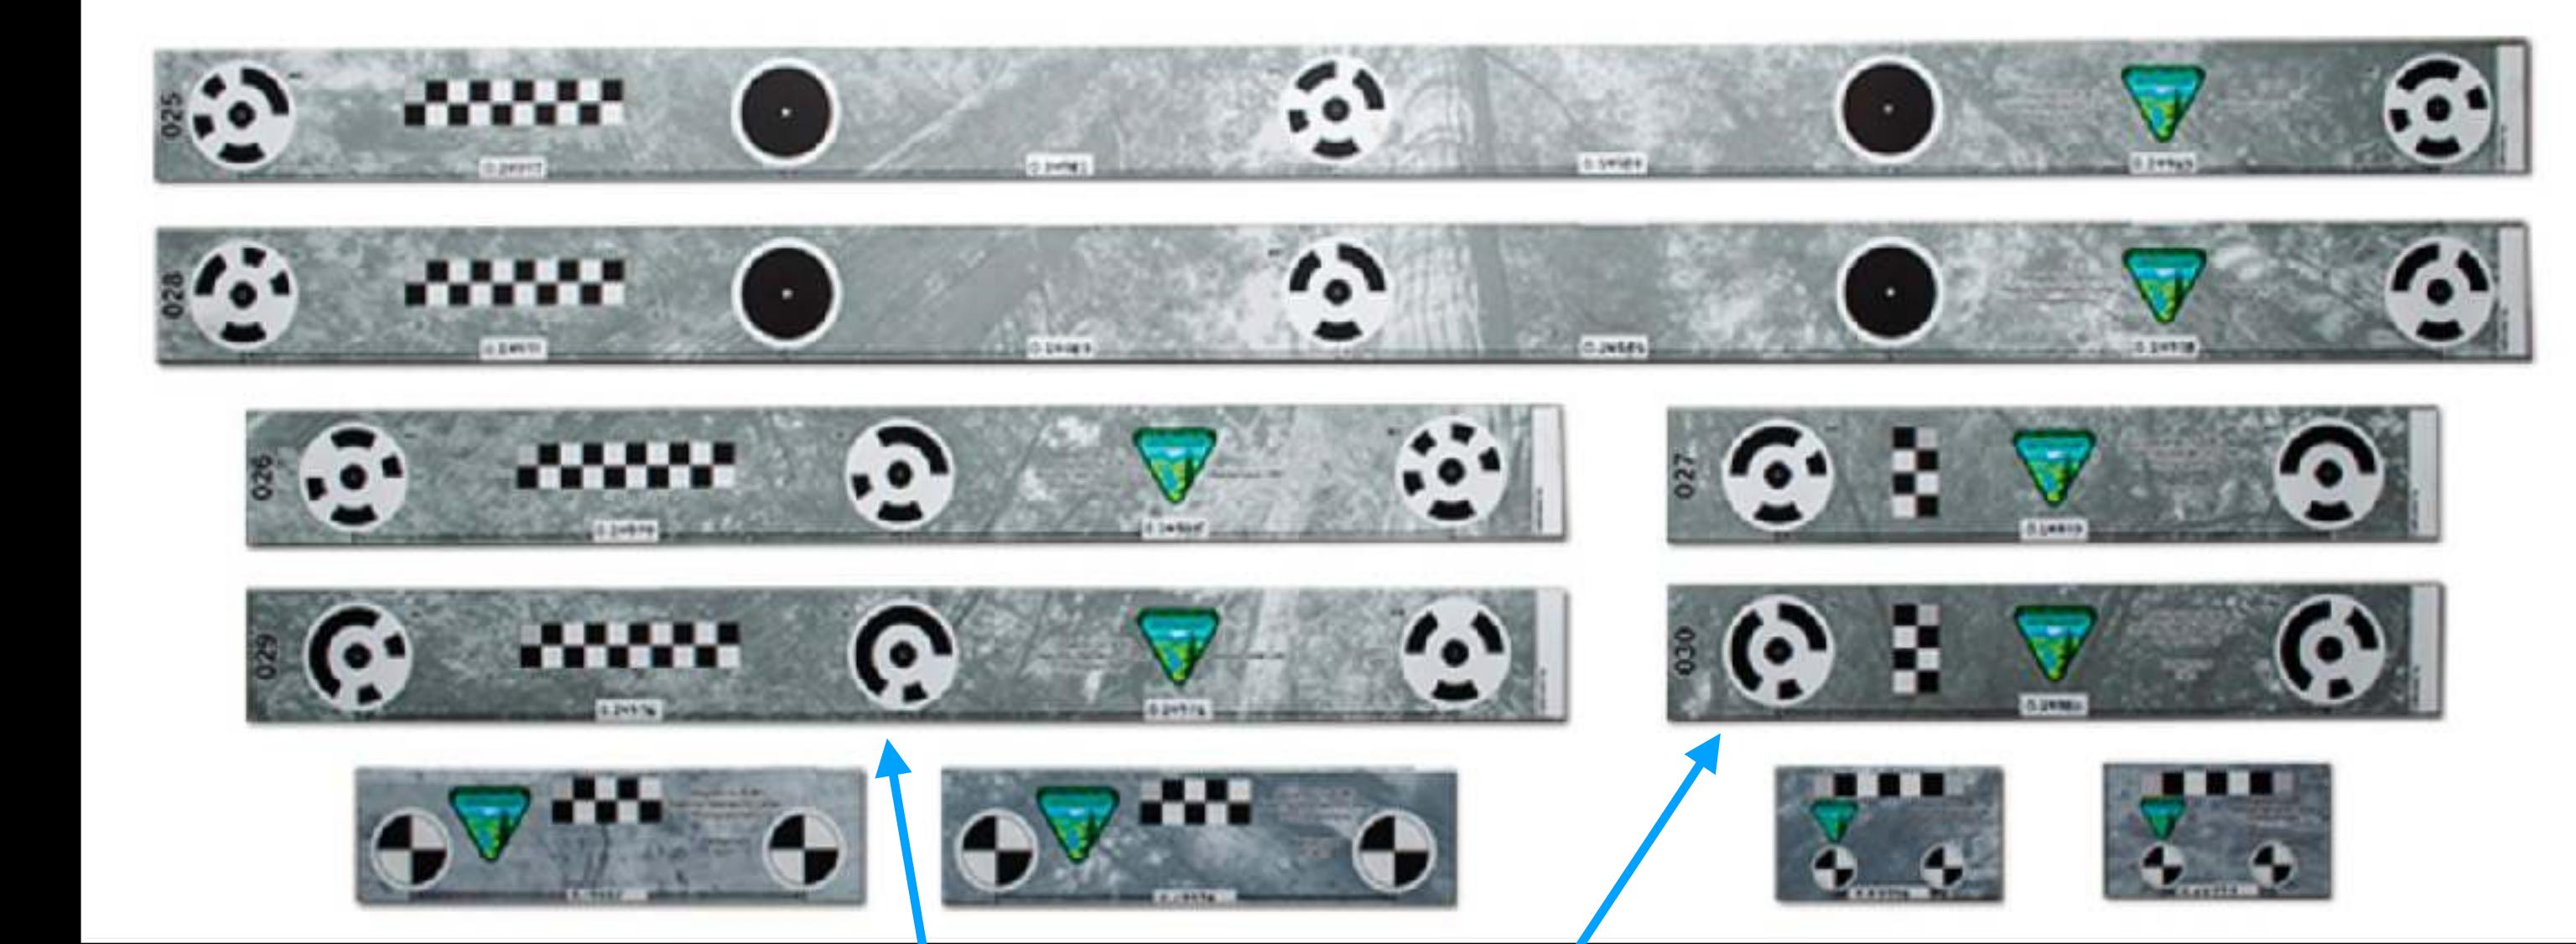
10,000 triangles



Masking

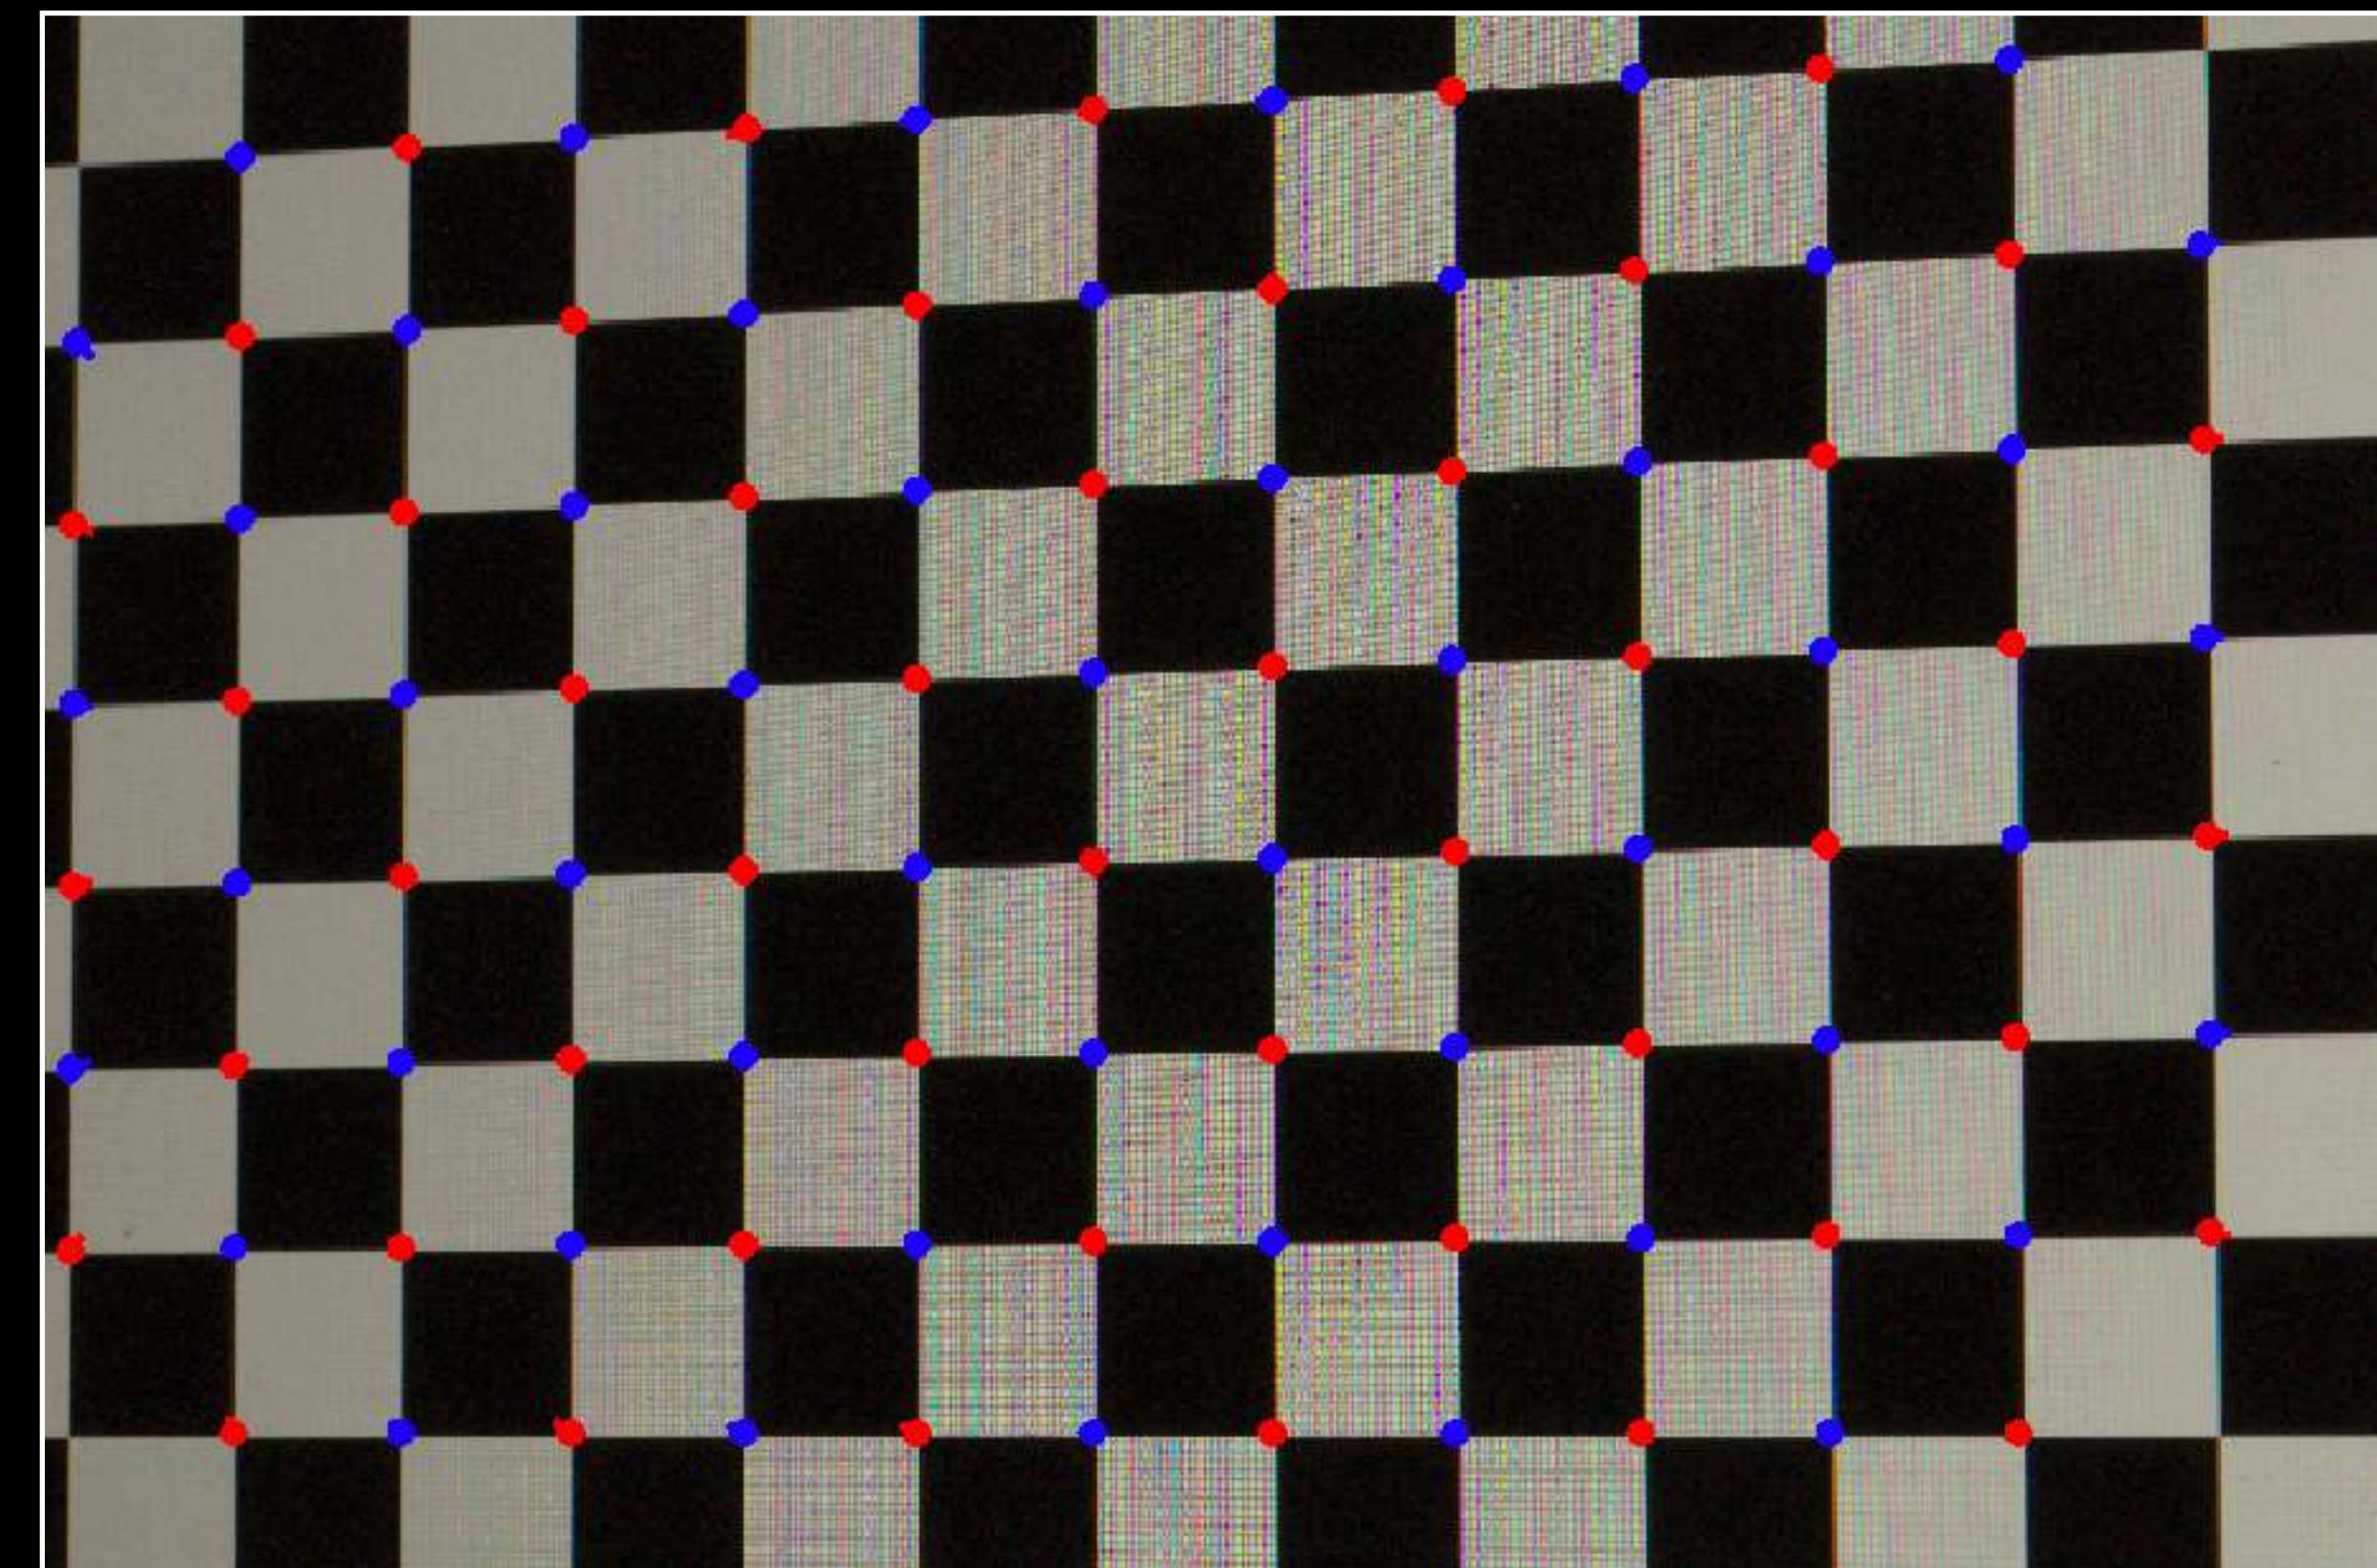


Scale bars



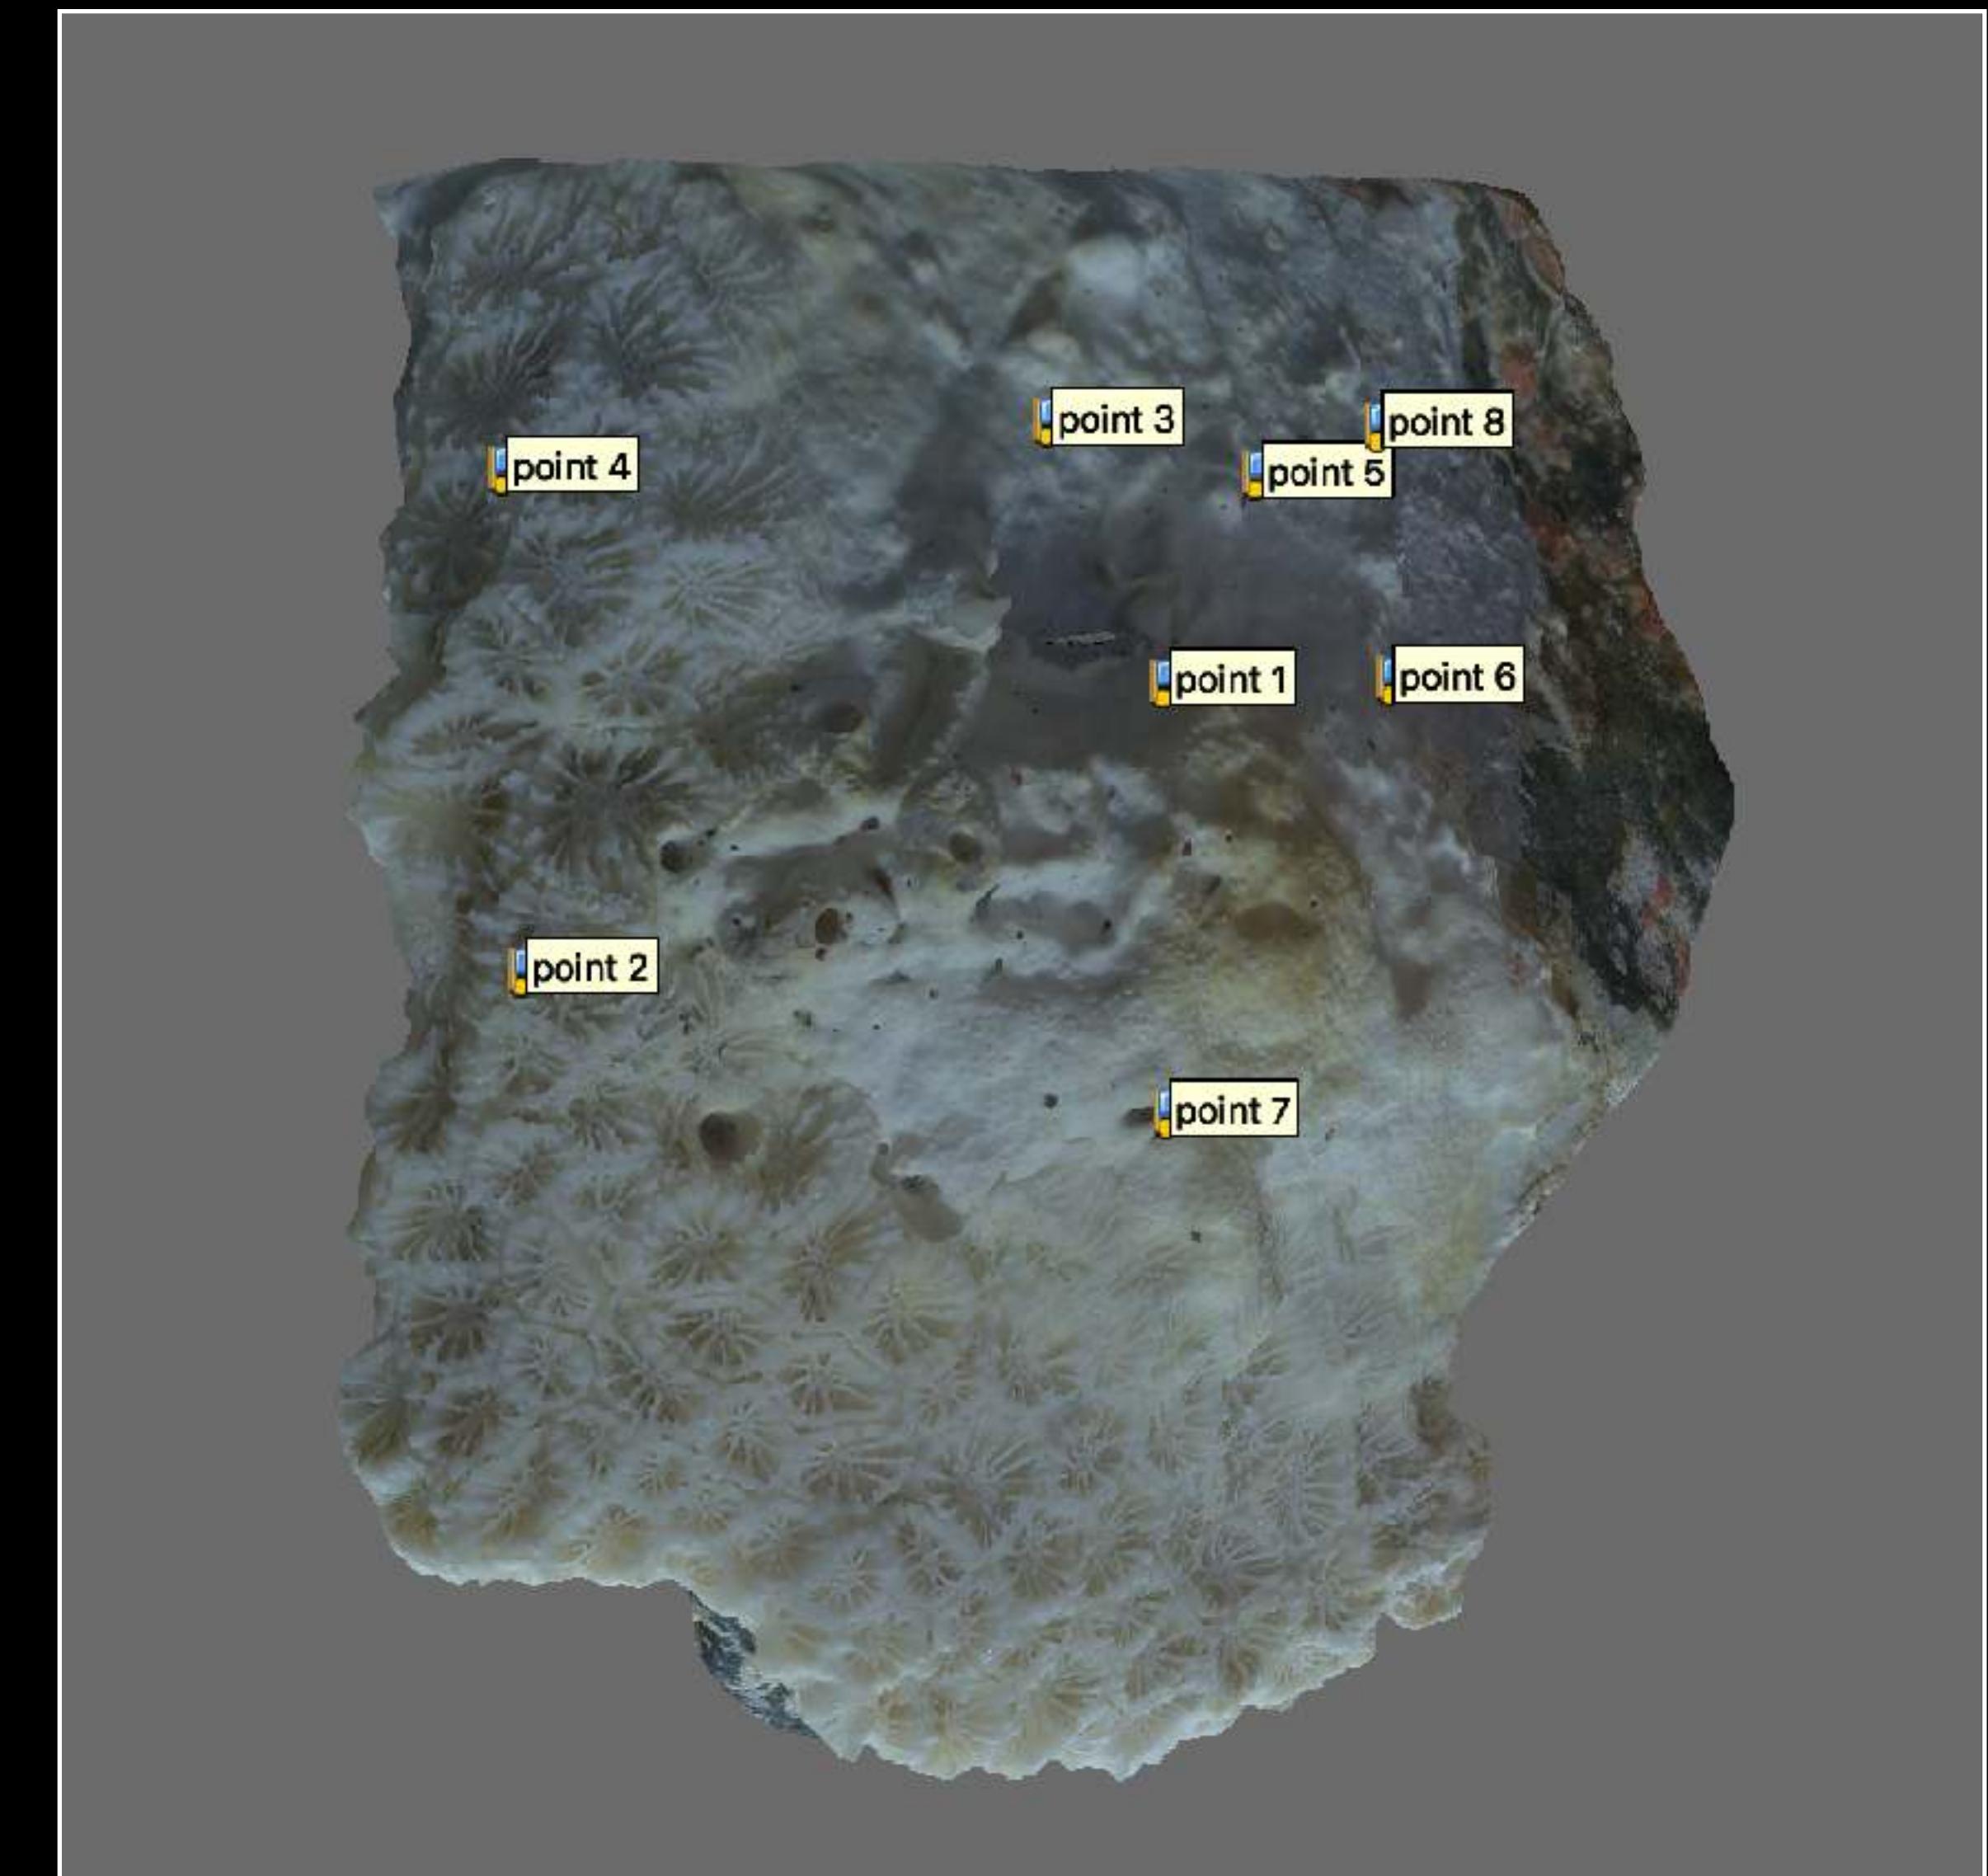
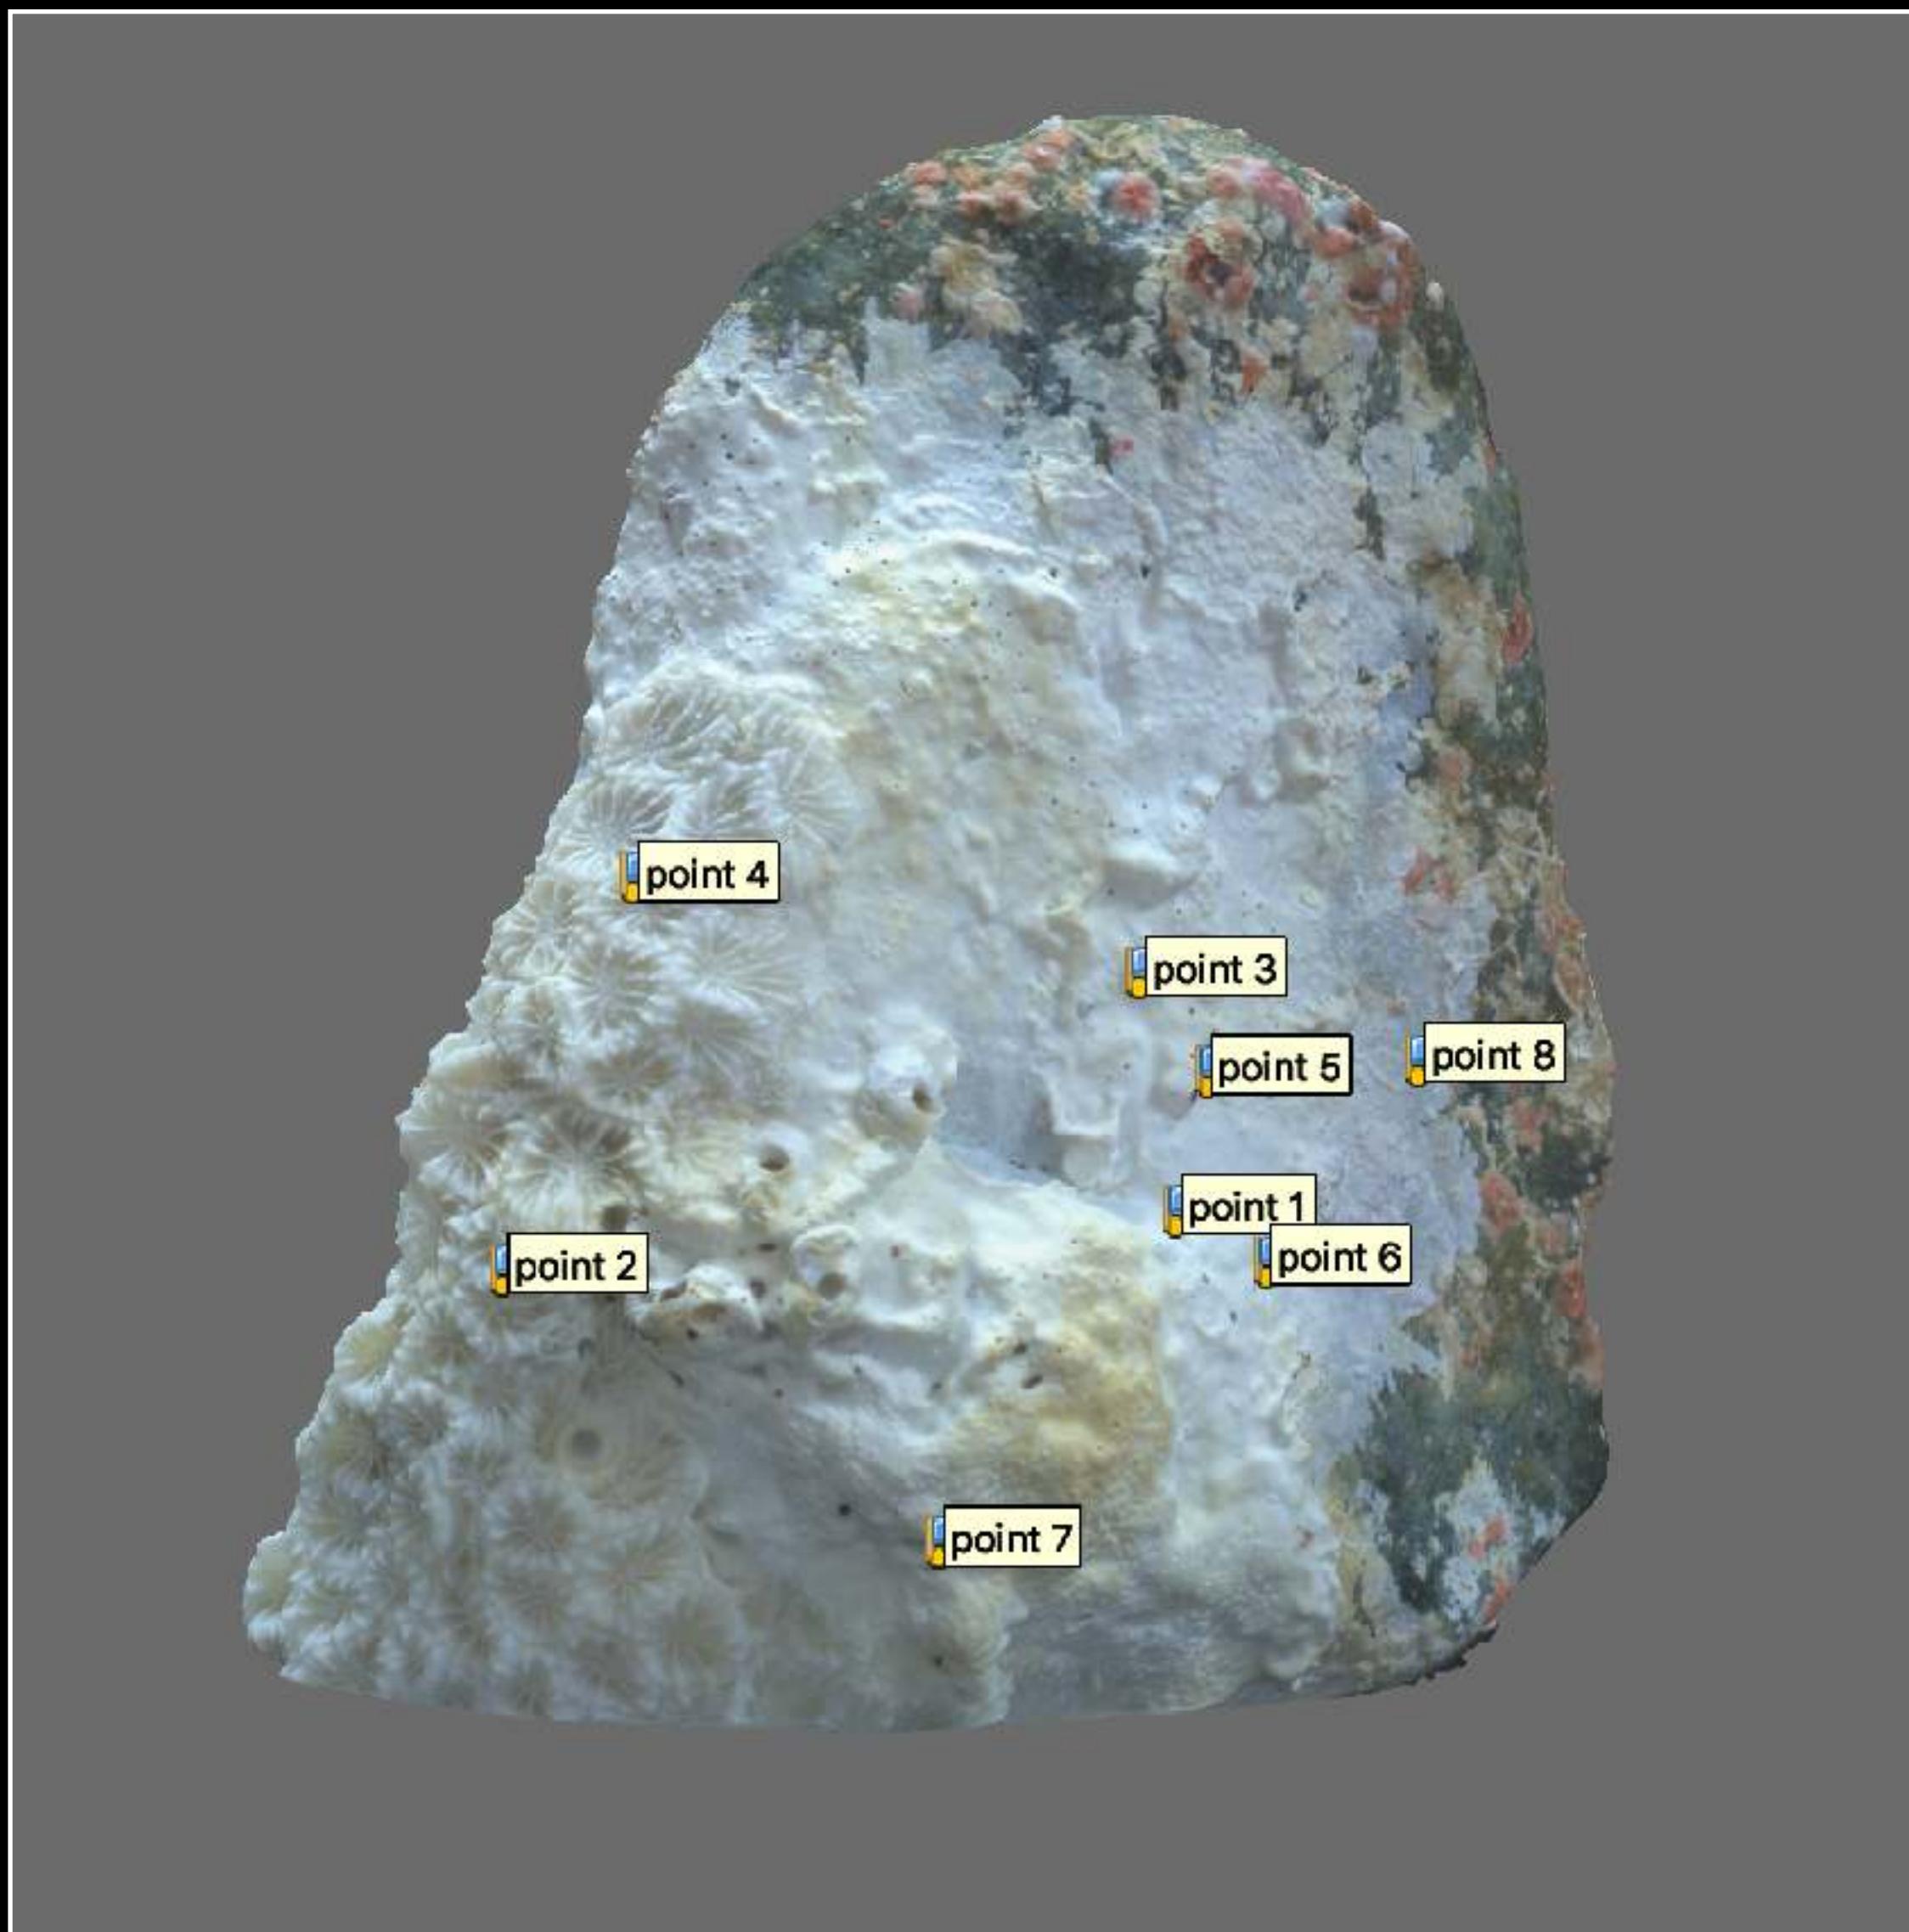
Automatically detected markers

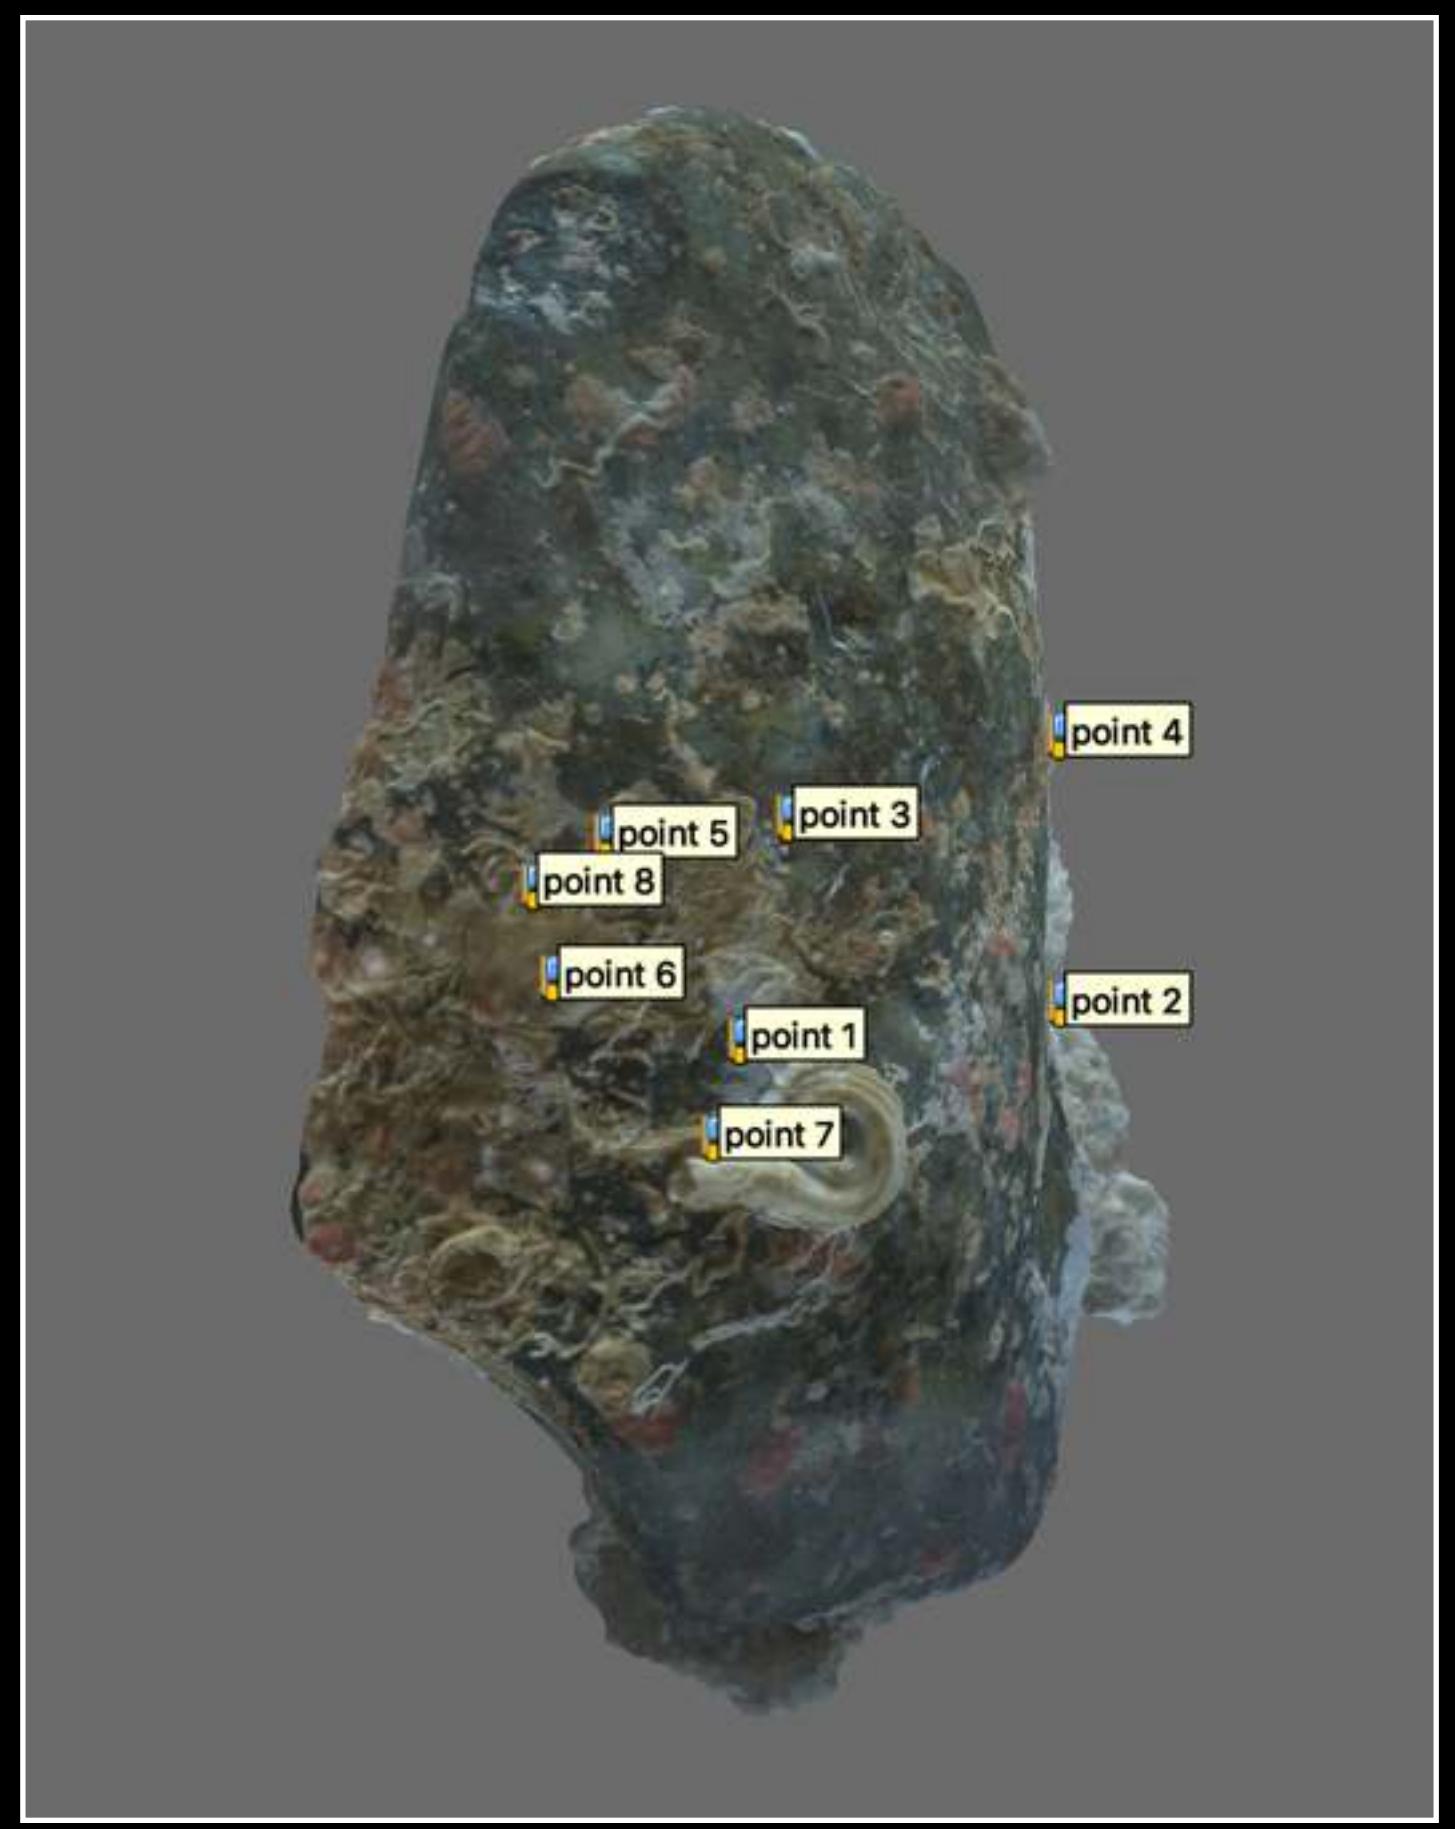
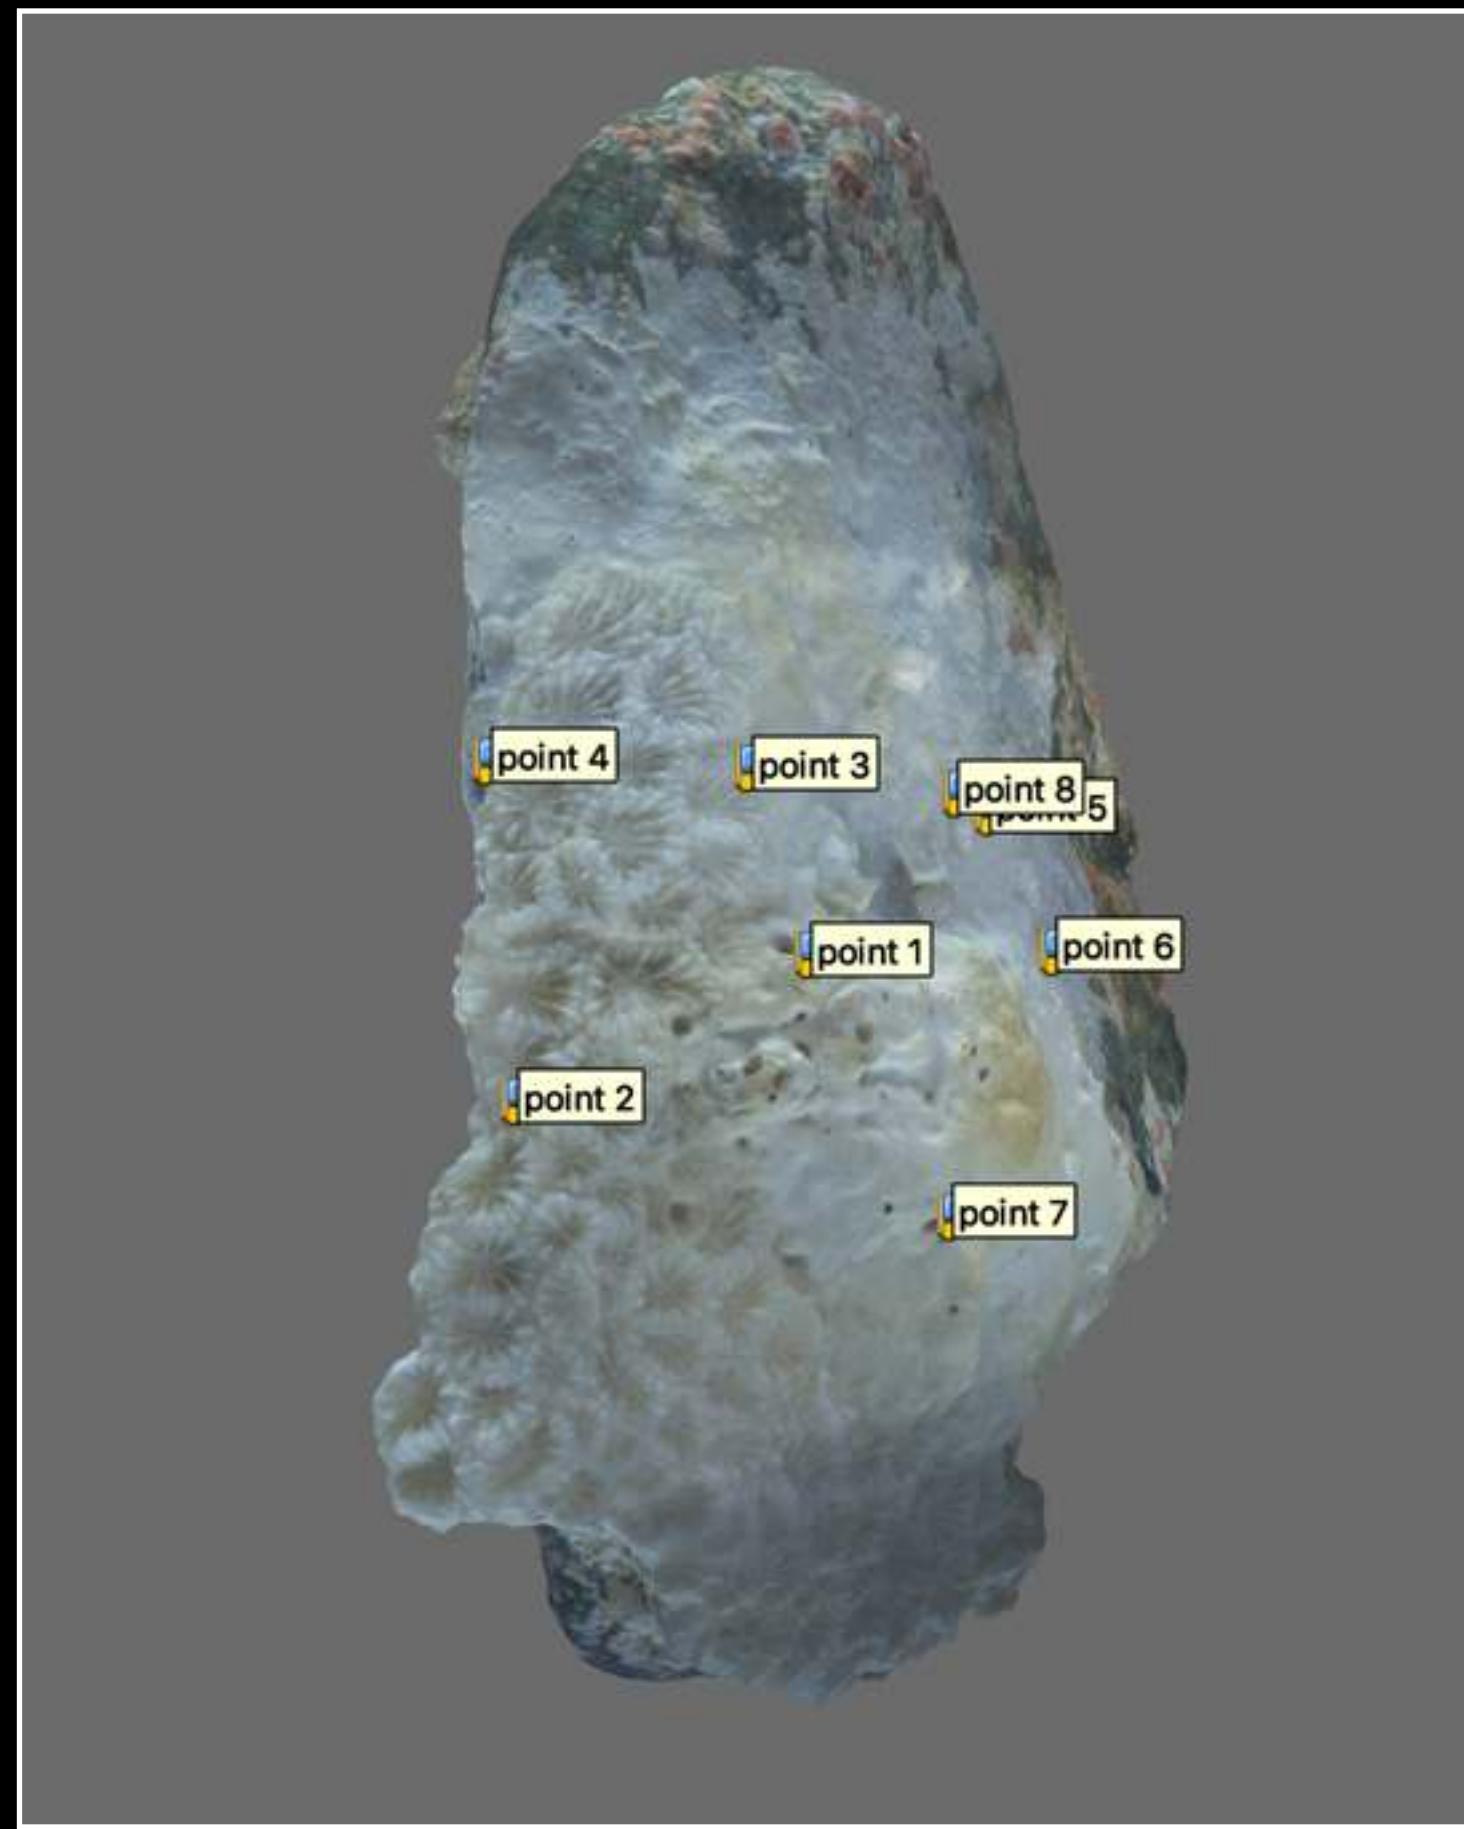
Lens calibration

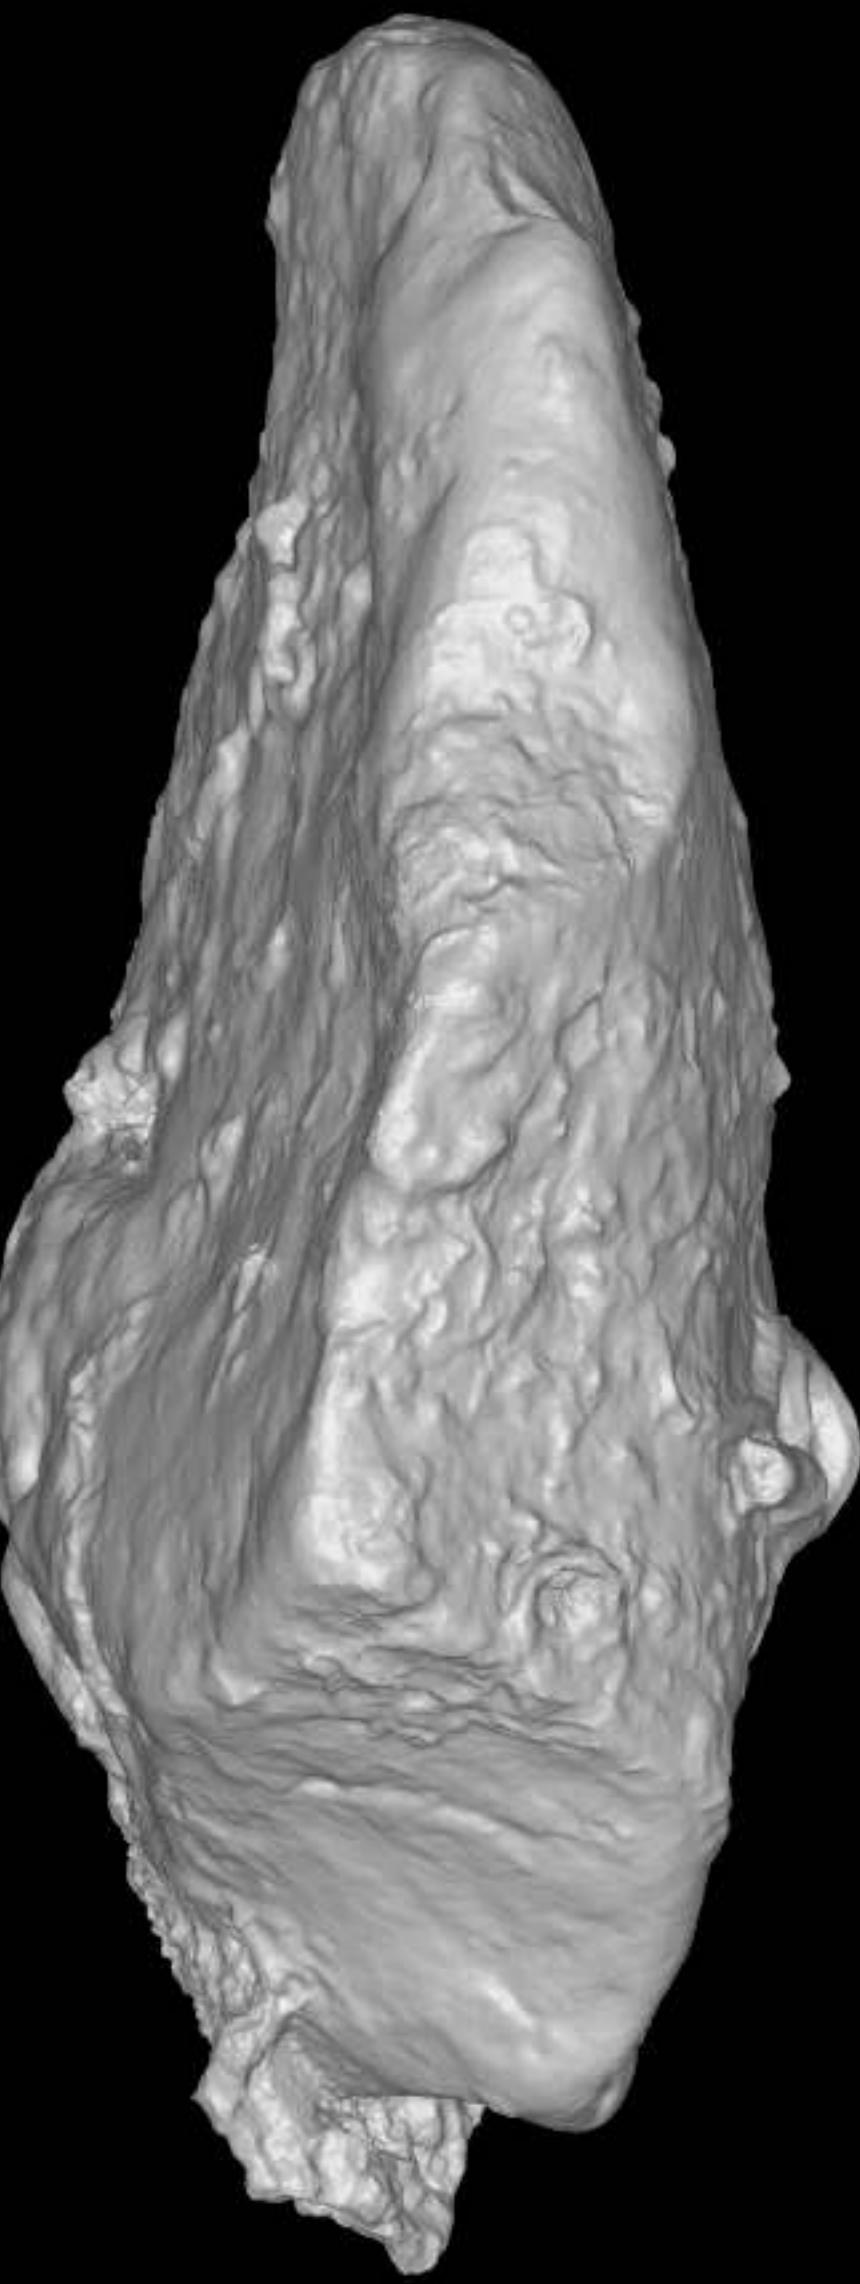


Merging multiple scans





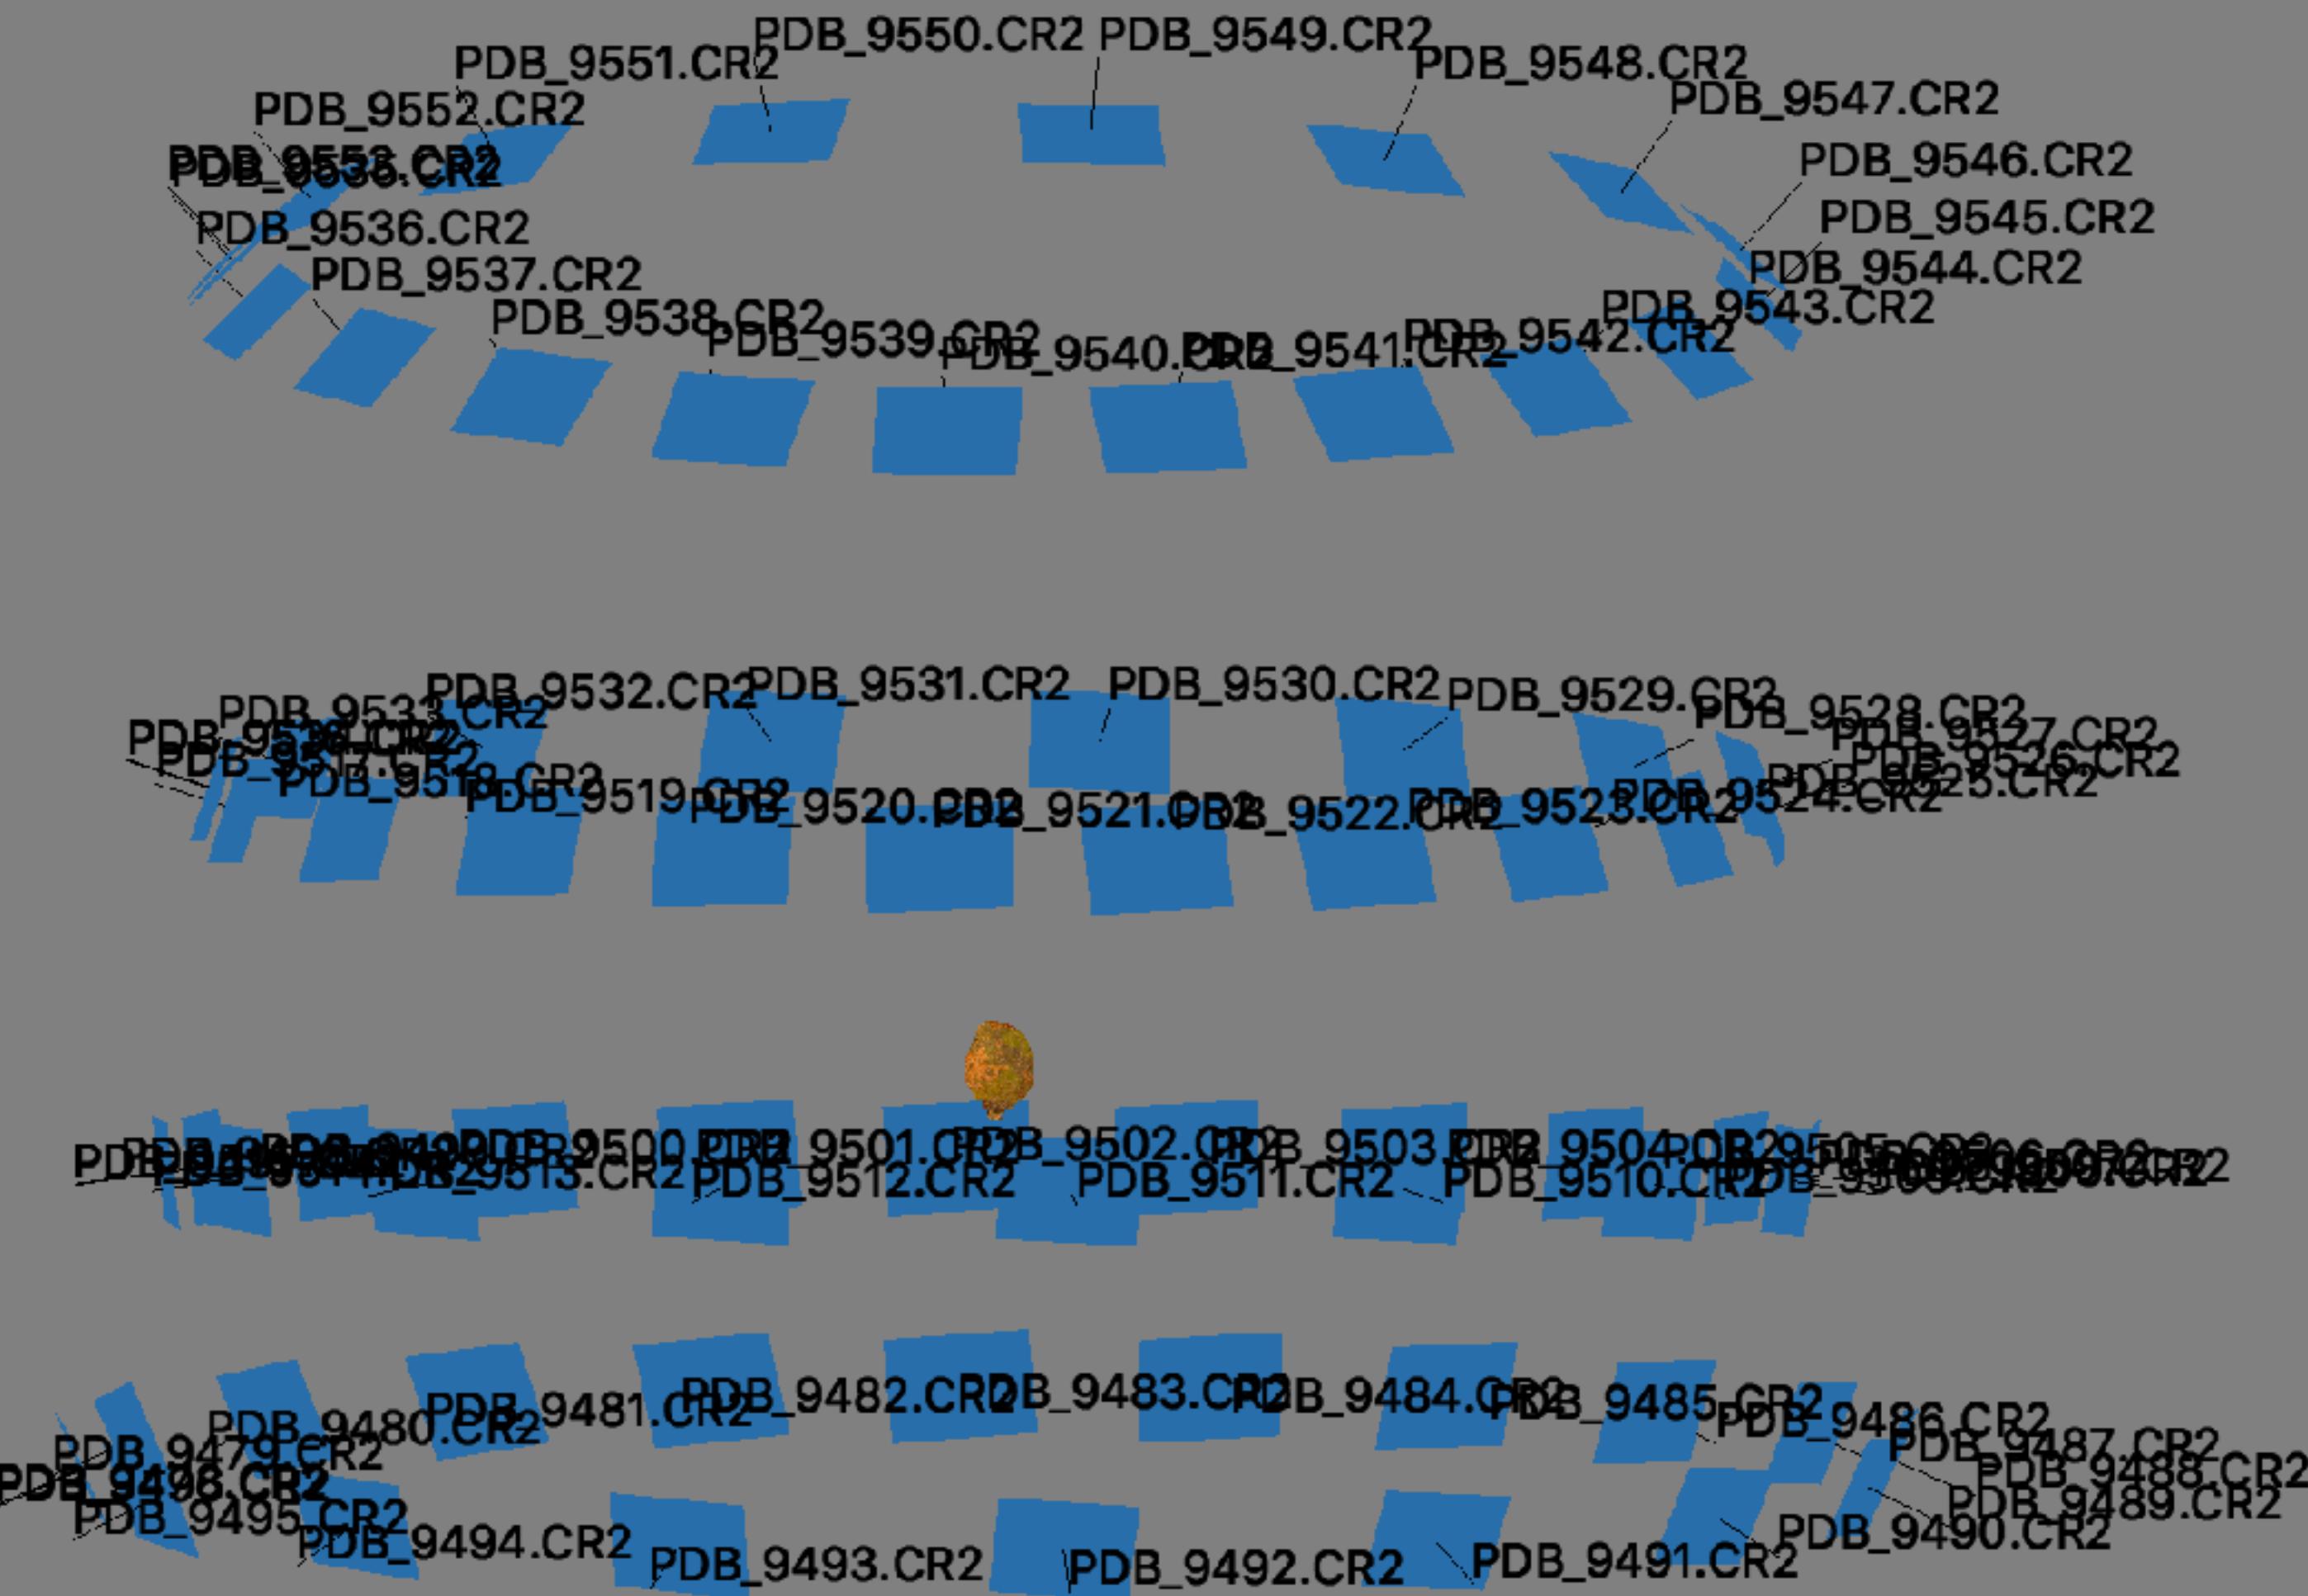


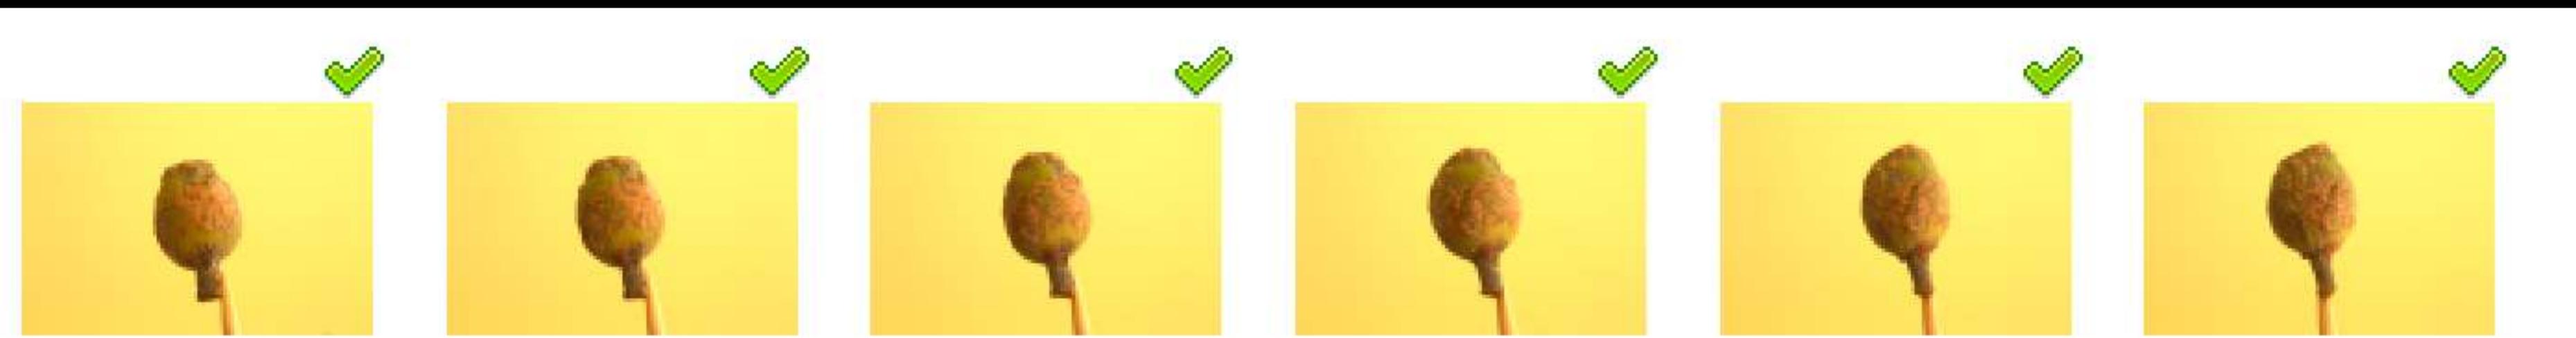


Turntables









PDB_9478.CR2

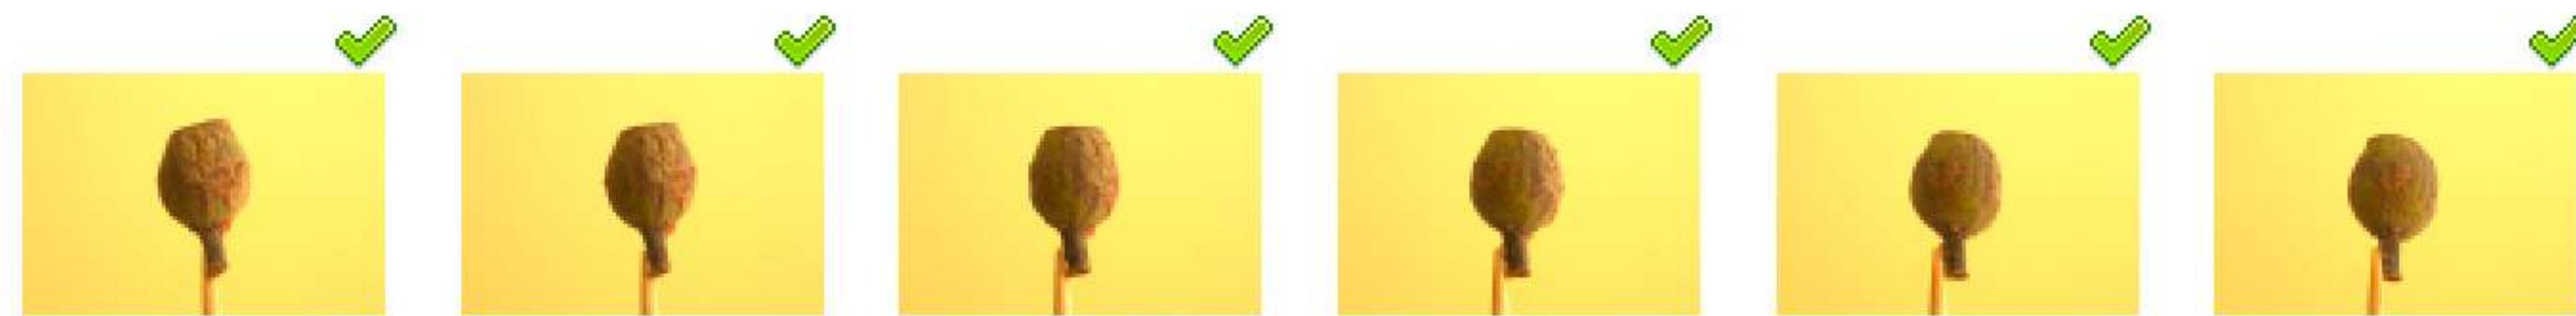
PDB_9479.CR2

PDB_9480.CR2

PDB_9481.CR2

PDB_9482.CR2

PDB_9483.CR2



PDB_9484.CR2

PDB_9485.CR2

PDB_9486.CR2

PDB_9487.CR2

PDB_9488.CR2

PDB_9489.CR2



PDB_9478.CR2

PDB_9479.CR2

PDB_9480.CR2

PDB_9481.CR2

PDB_9482.CR2

PDB_9483.CR2



PDB_9484.CR2

PDB_9485.CR2

PDB_9486.CR2

PDB_9487.CR2

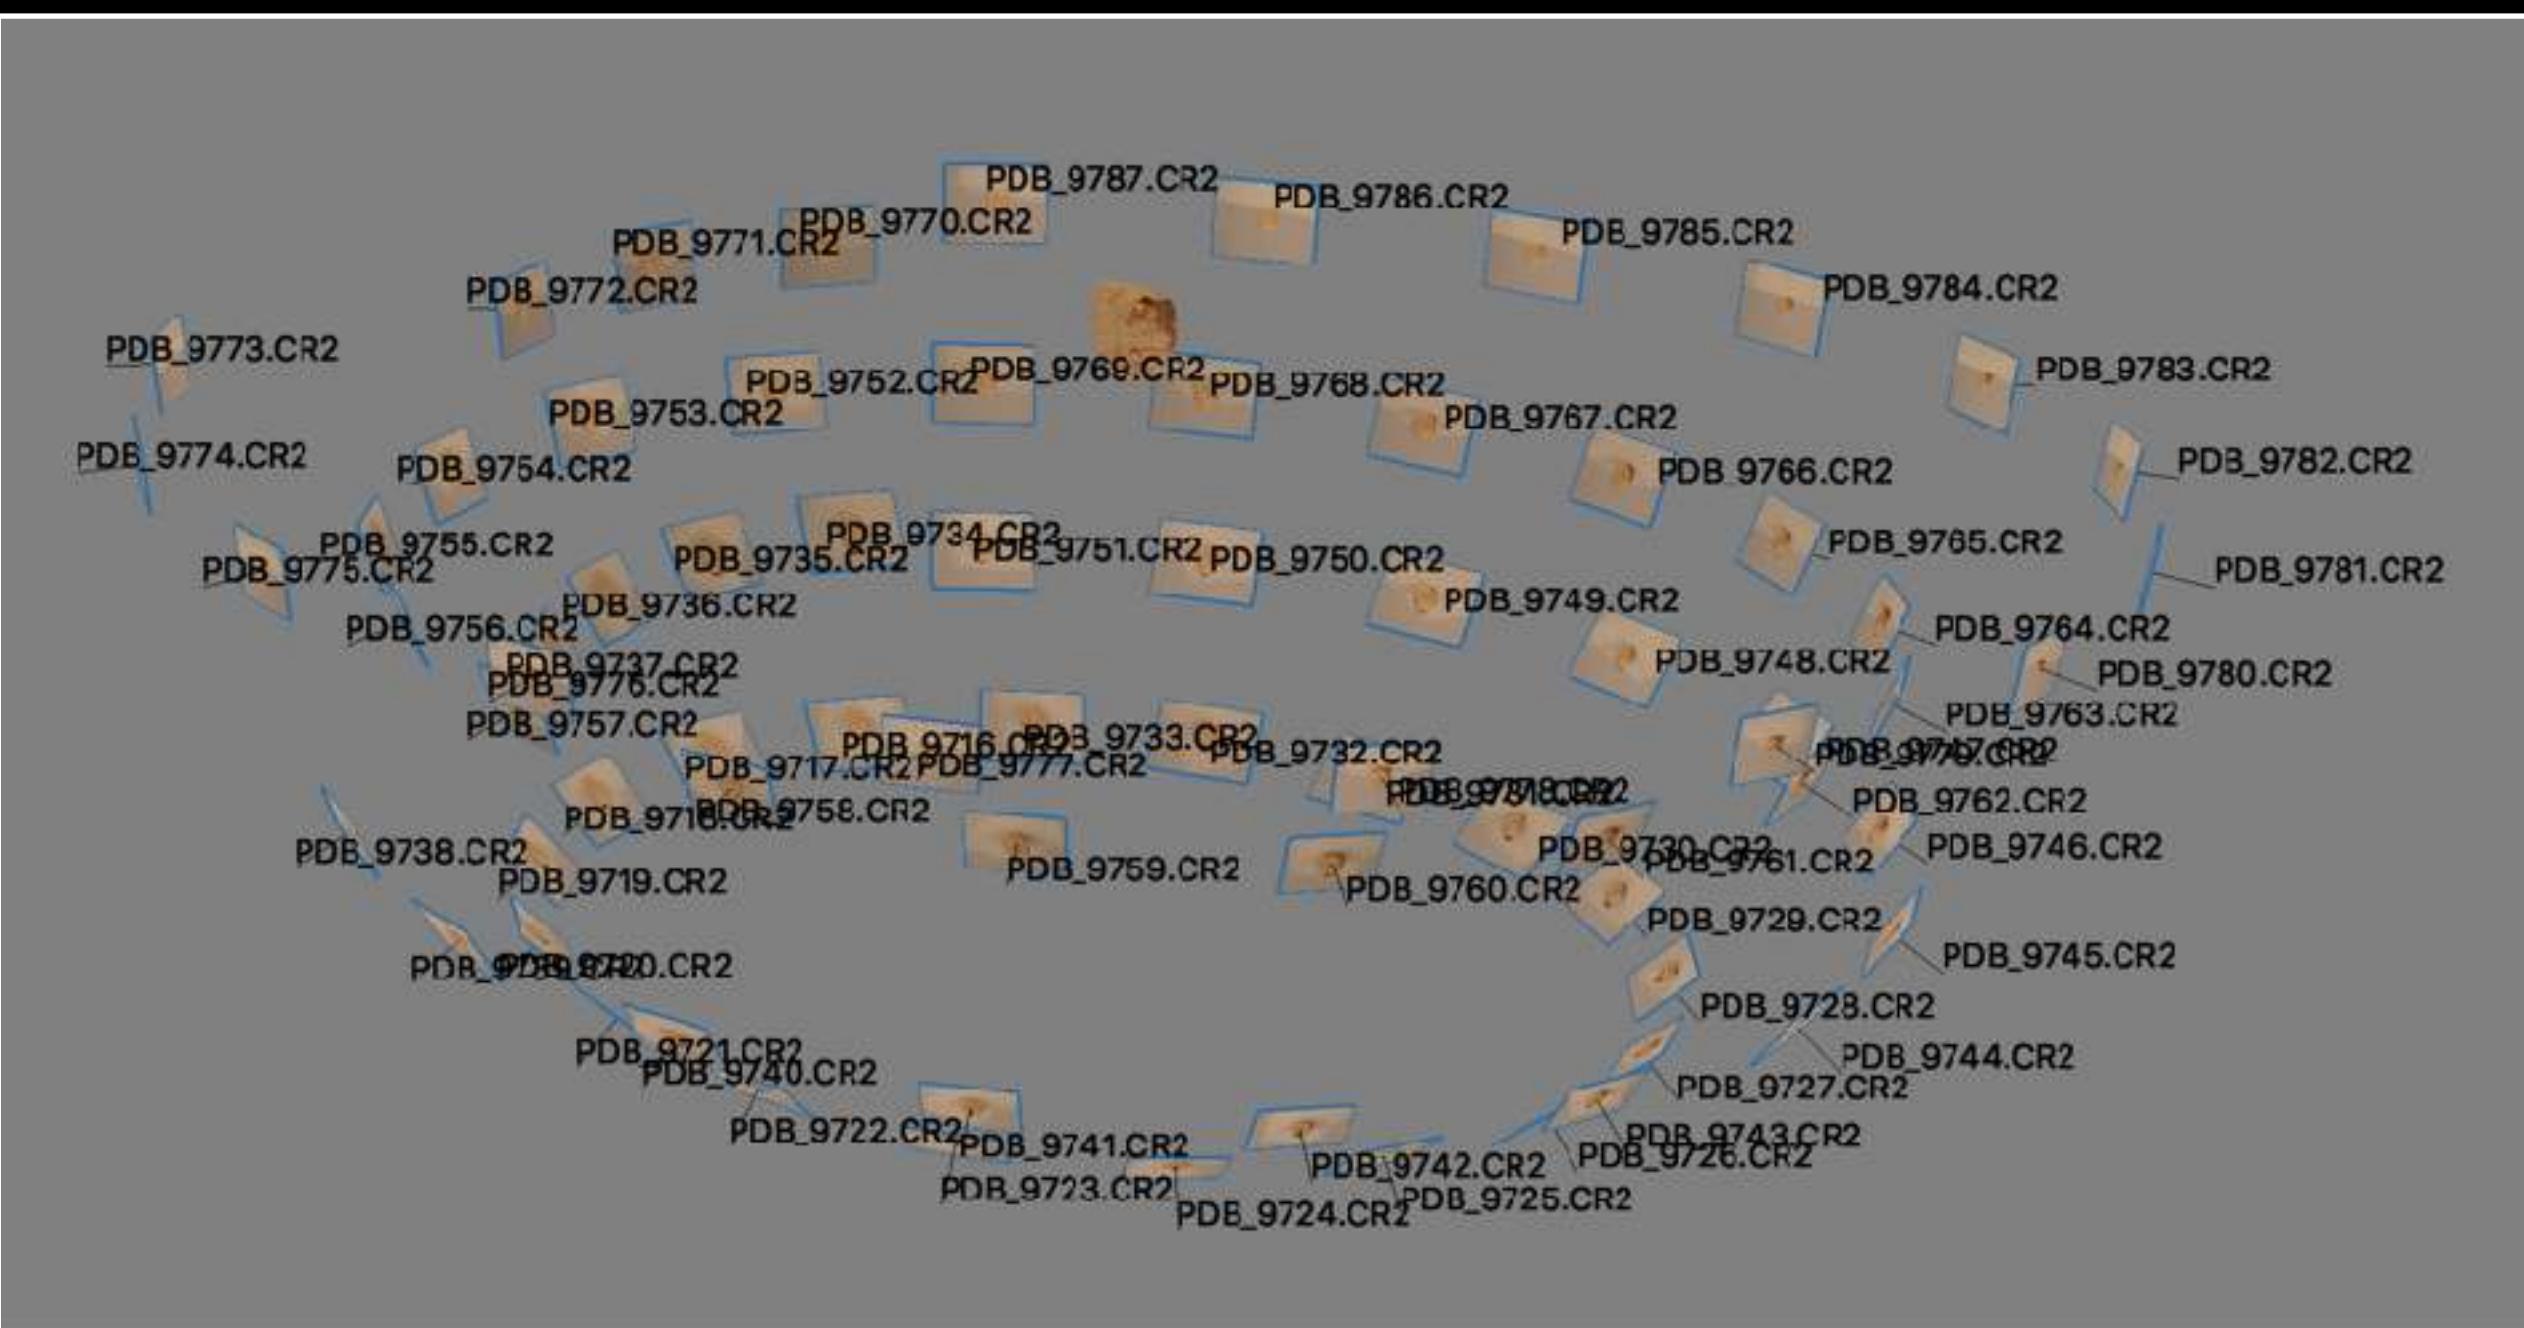
PDB_9488.CR2

PDB_9489.CR2

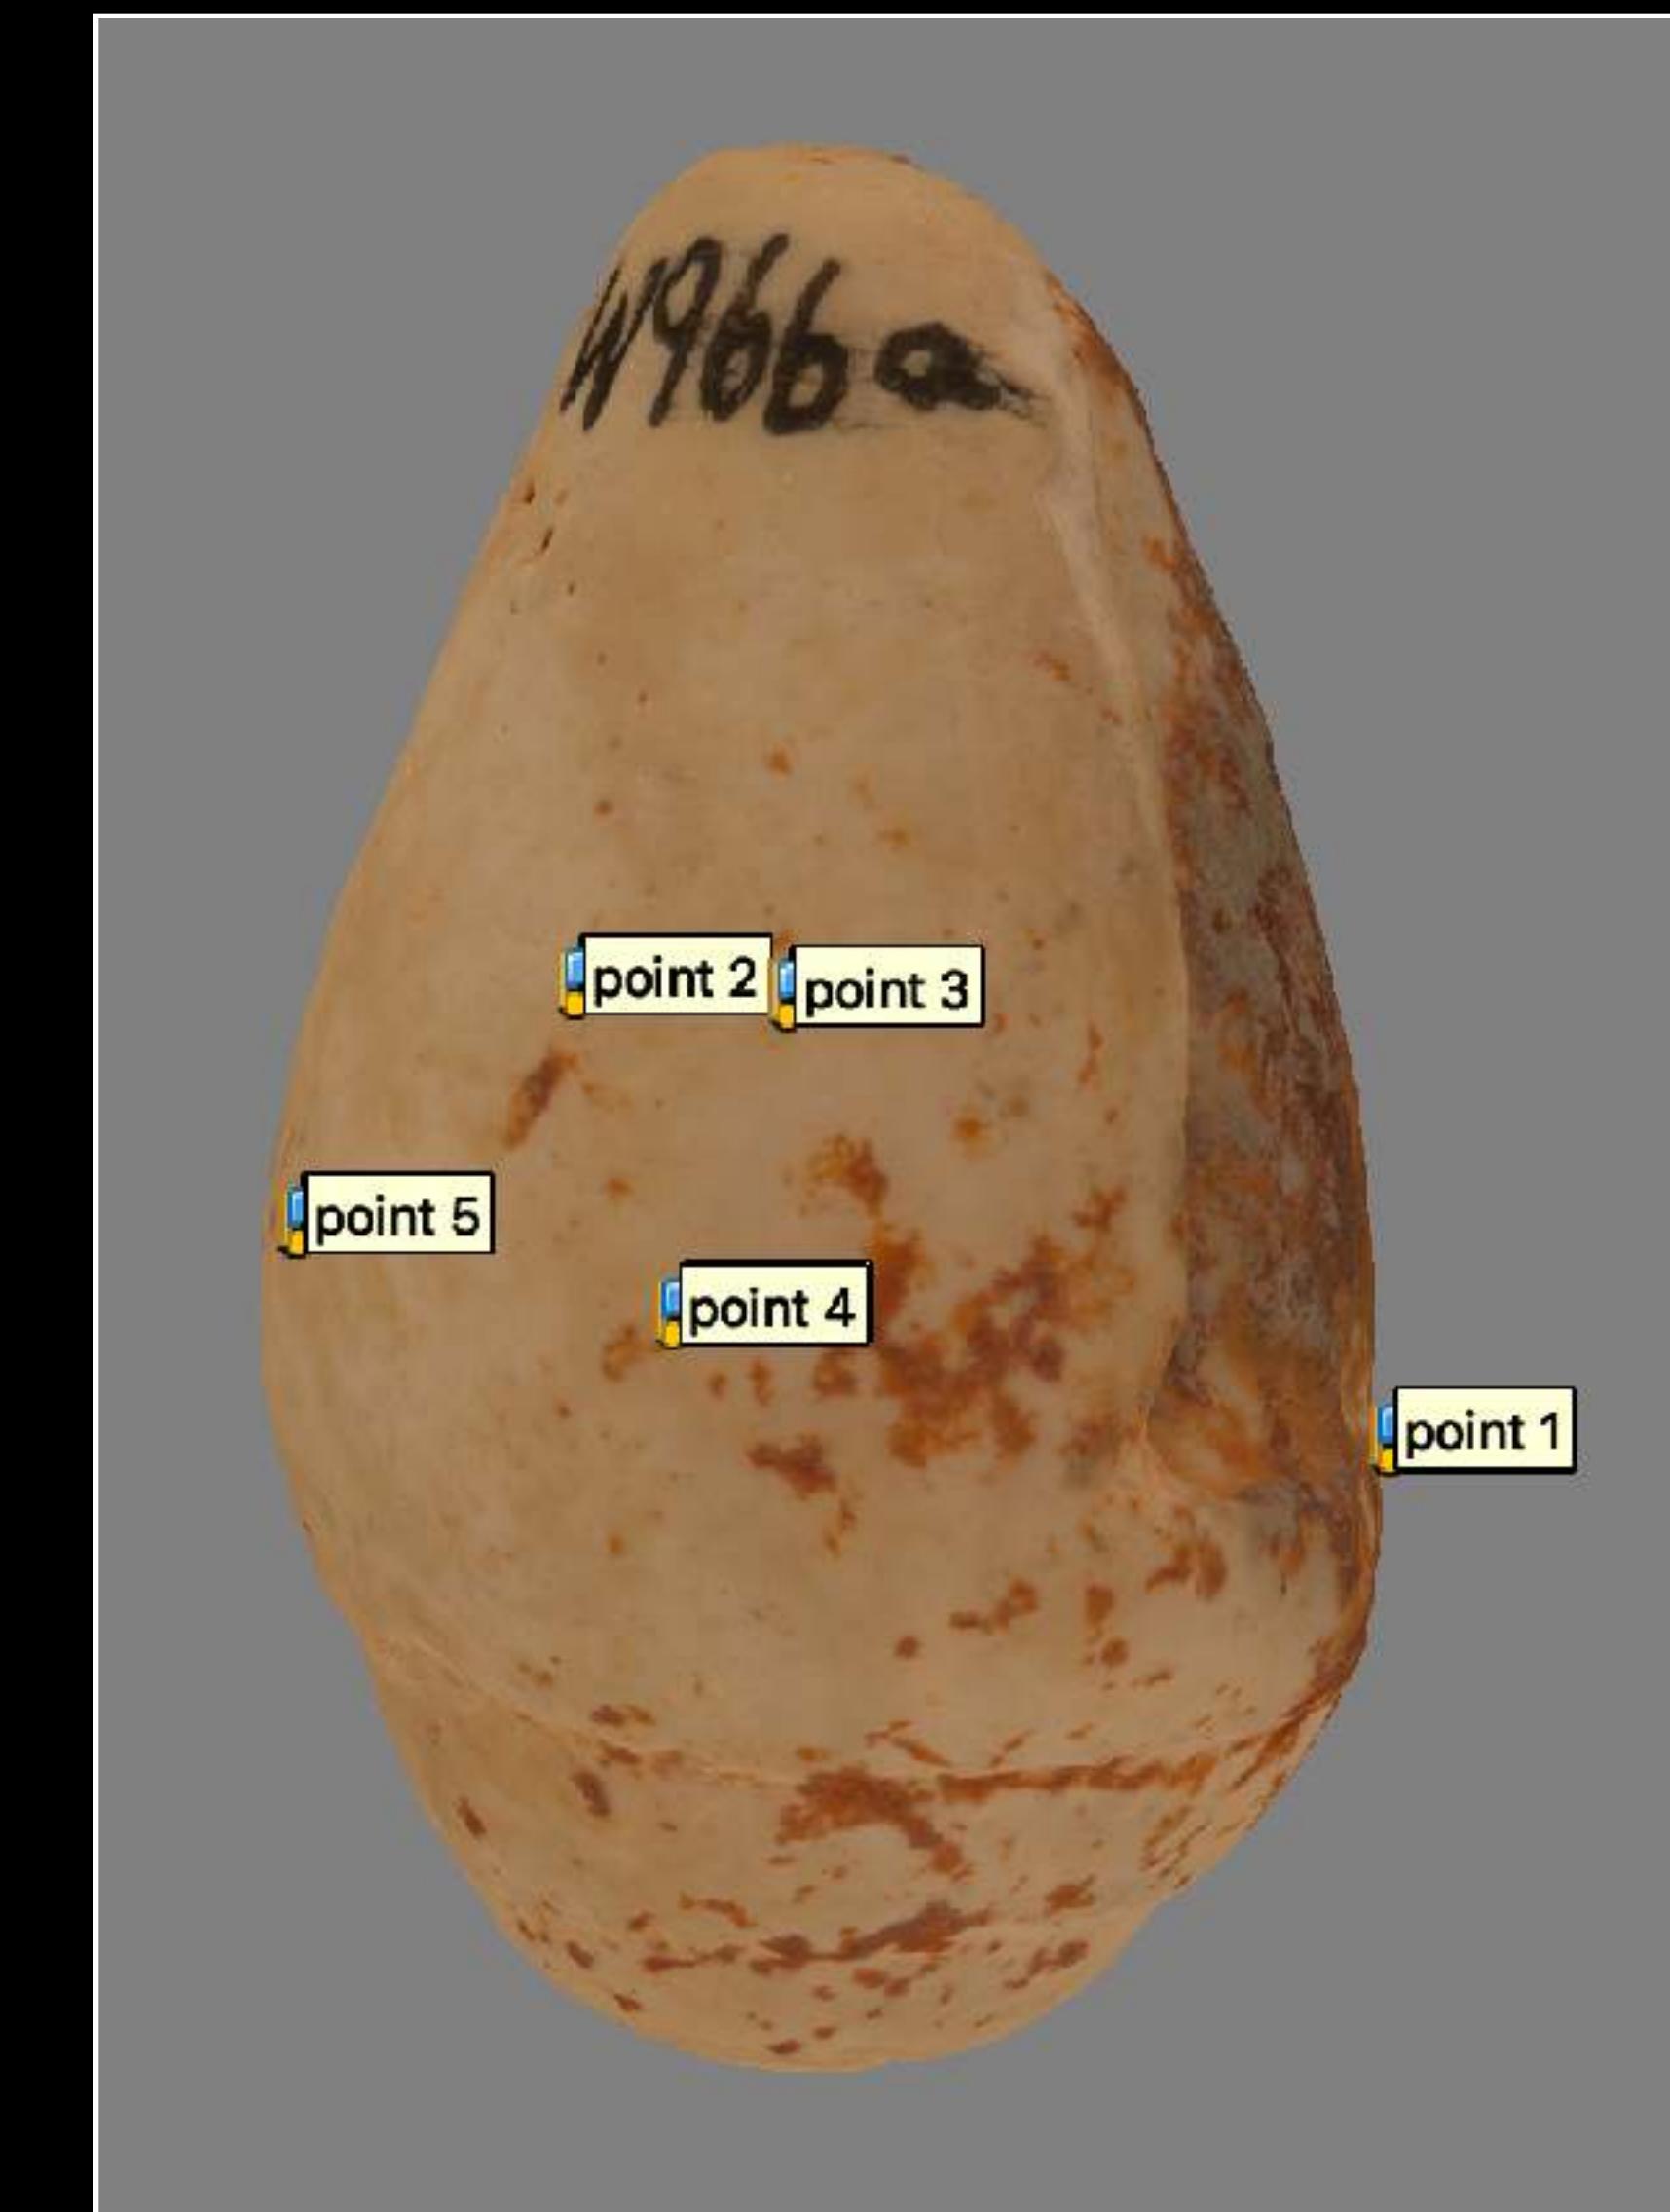
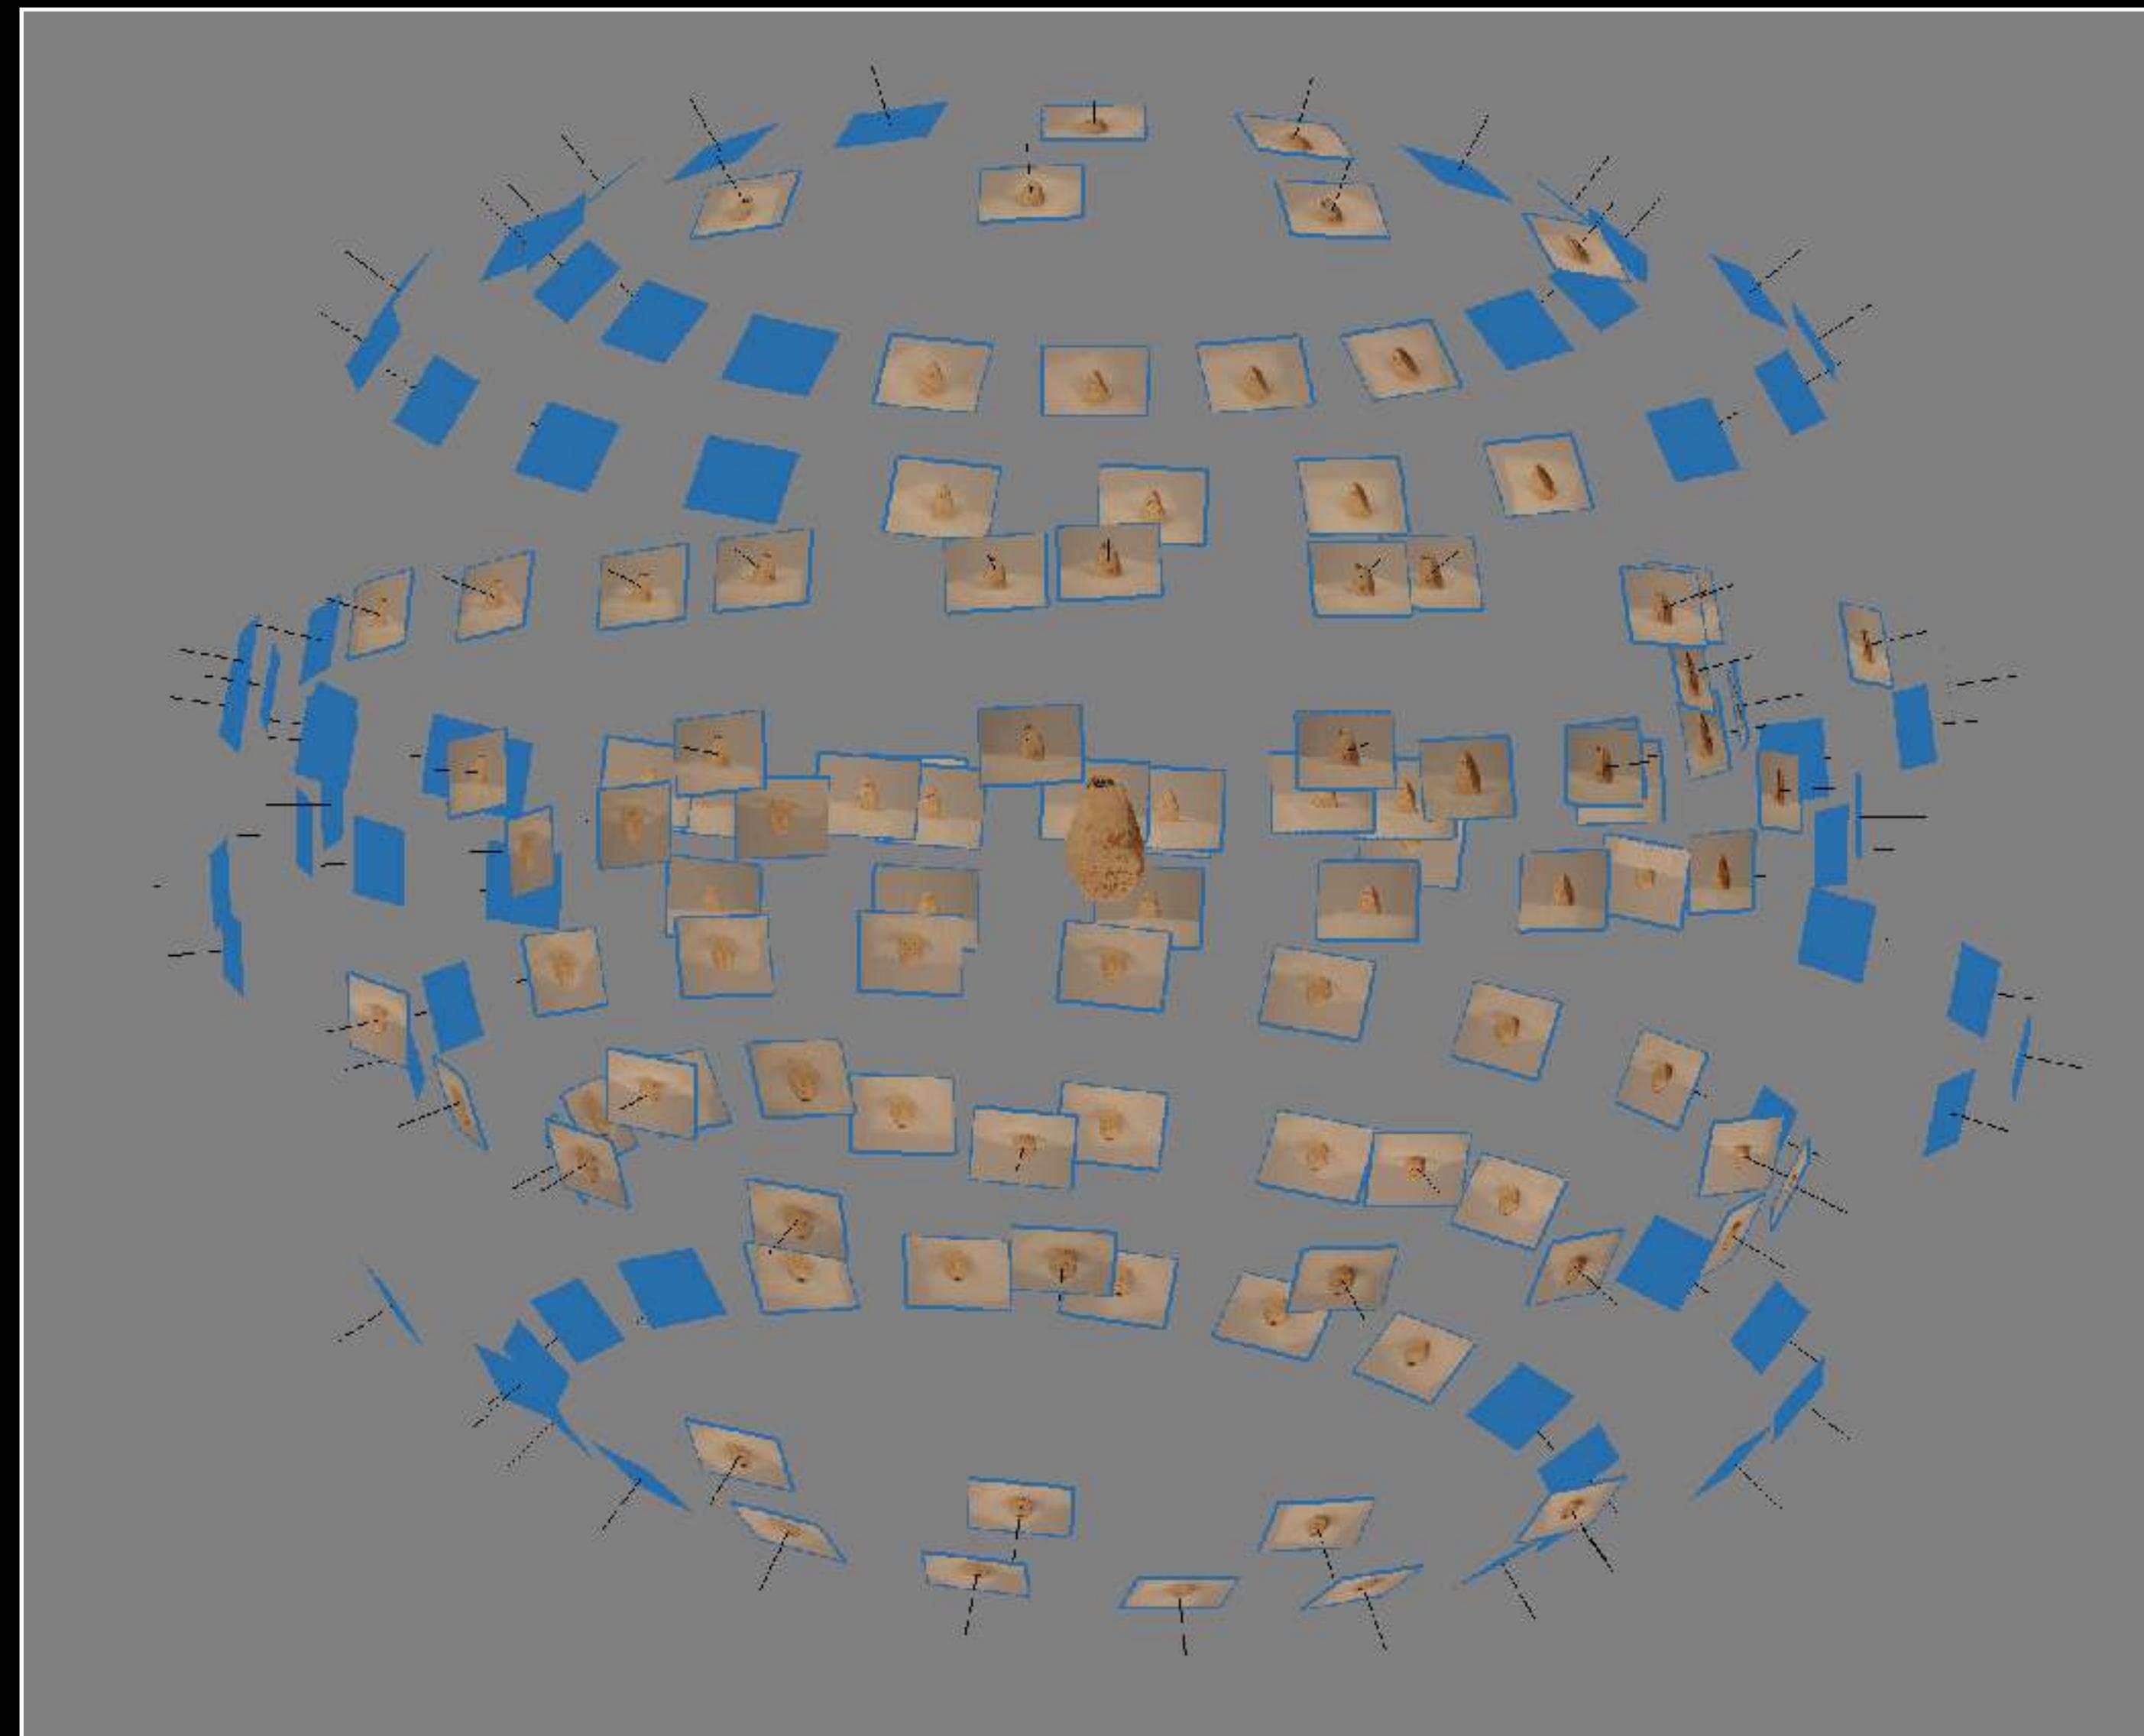


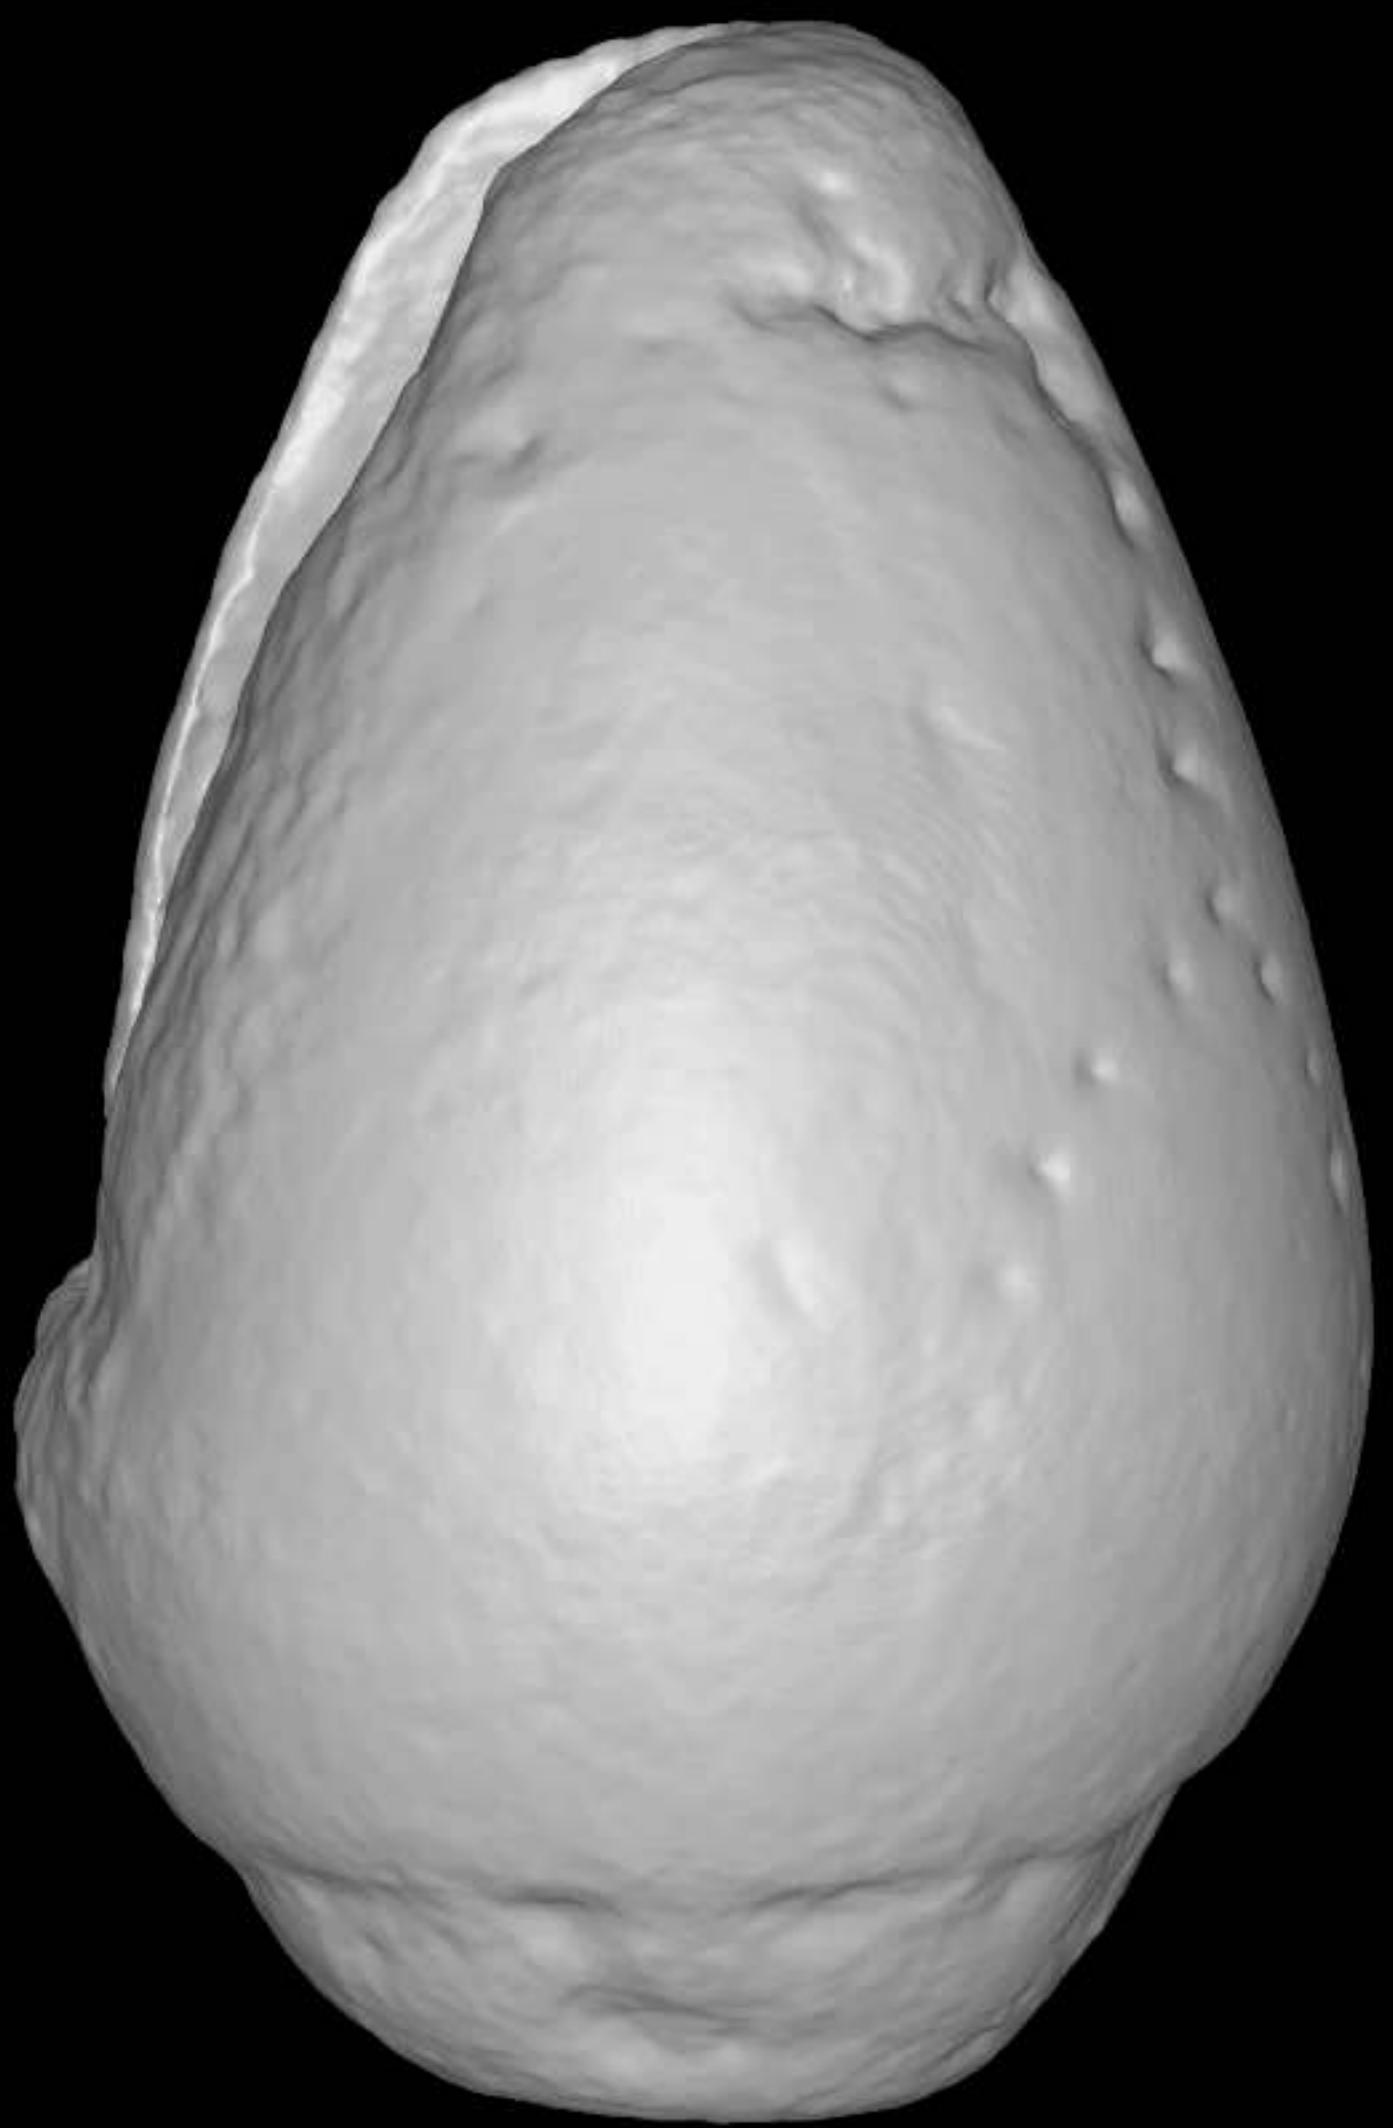


Top half + overlap



Bottom half + overlap

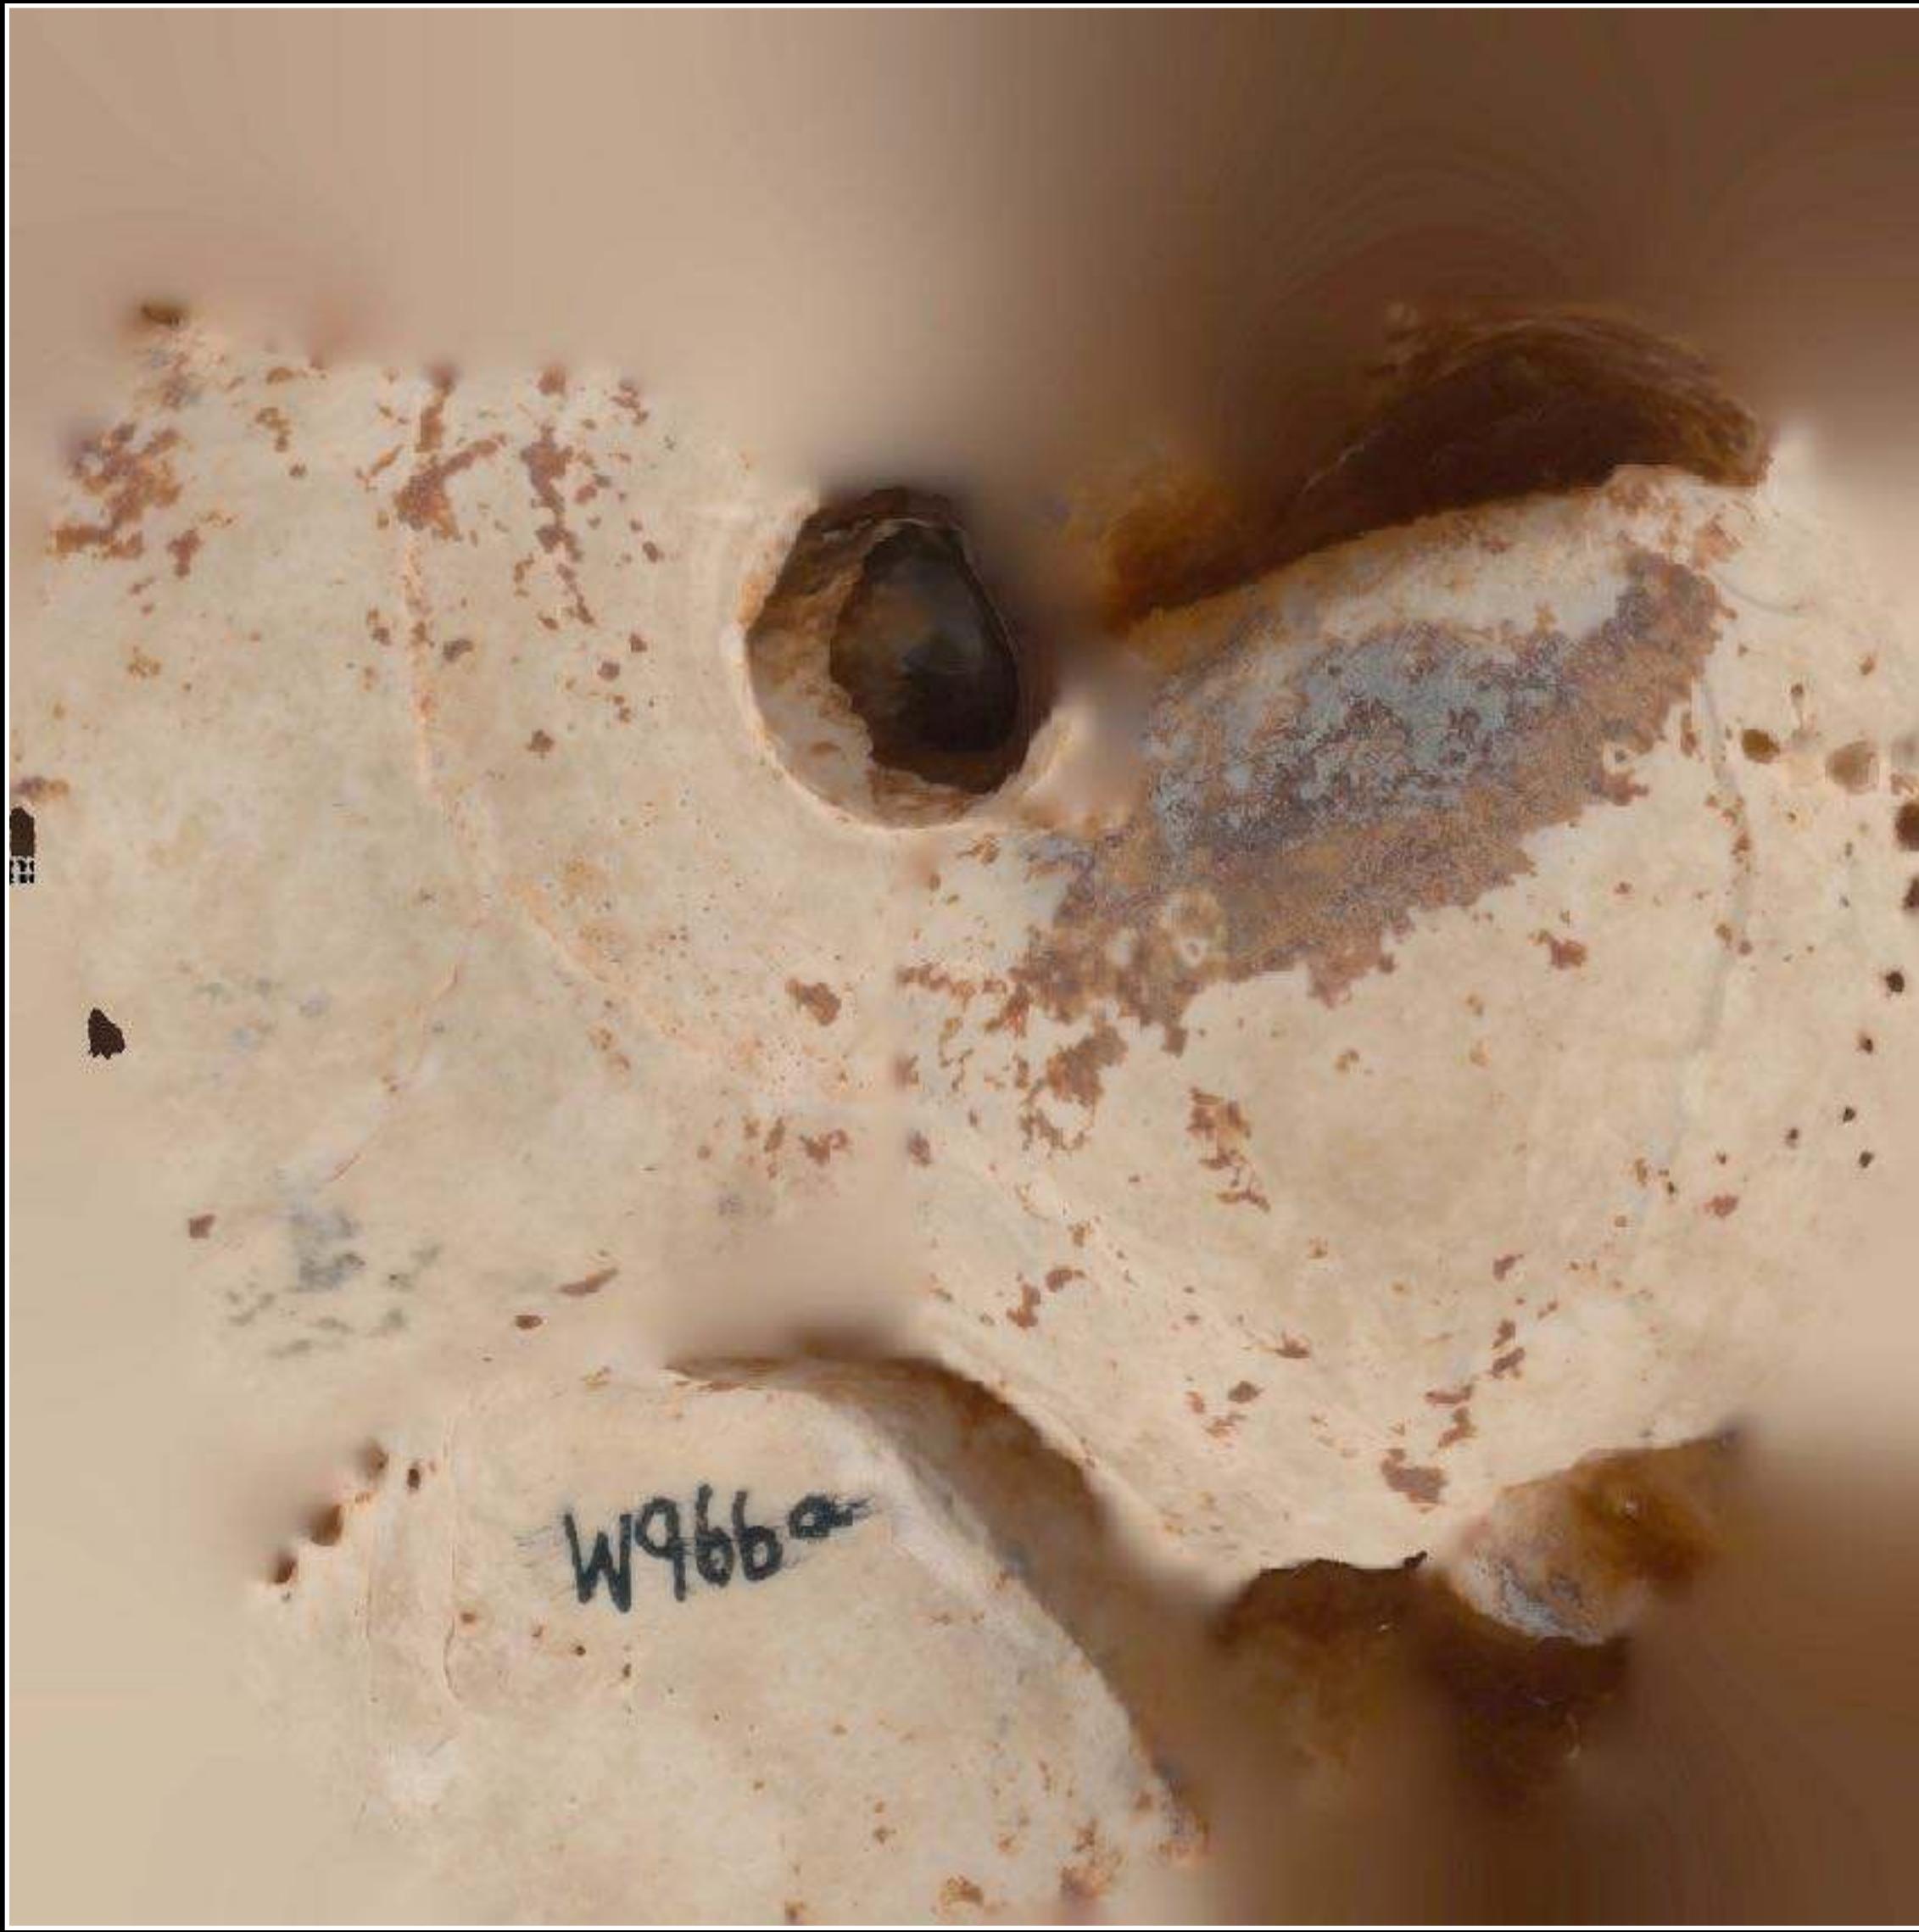




Texture editing

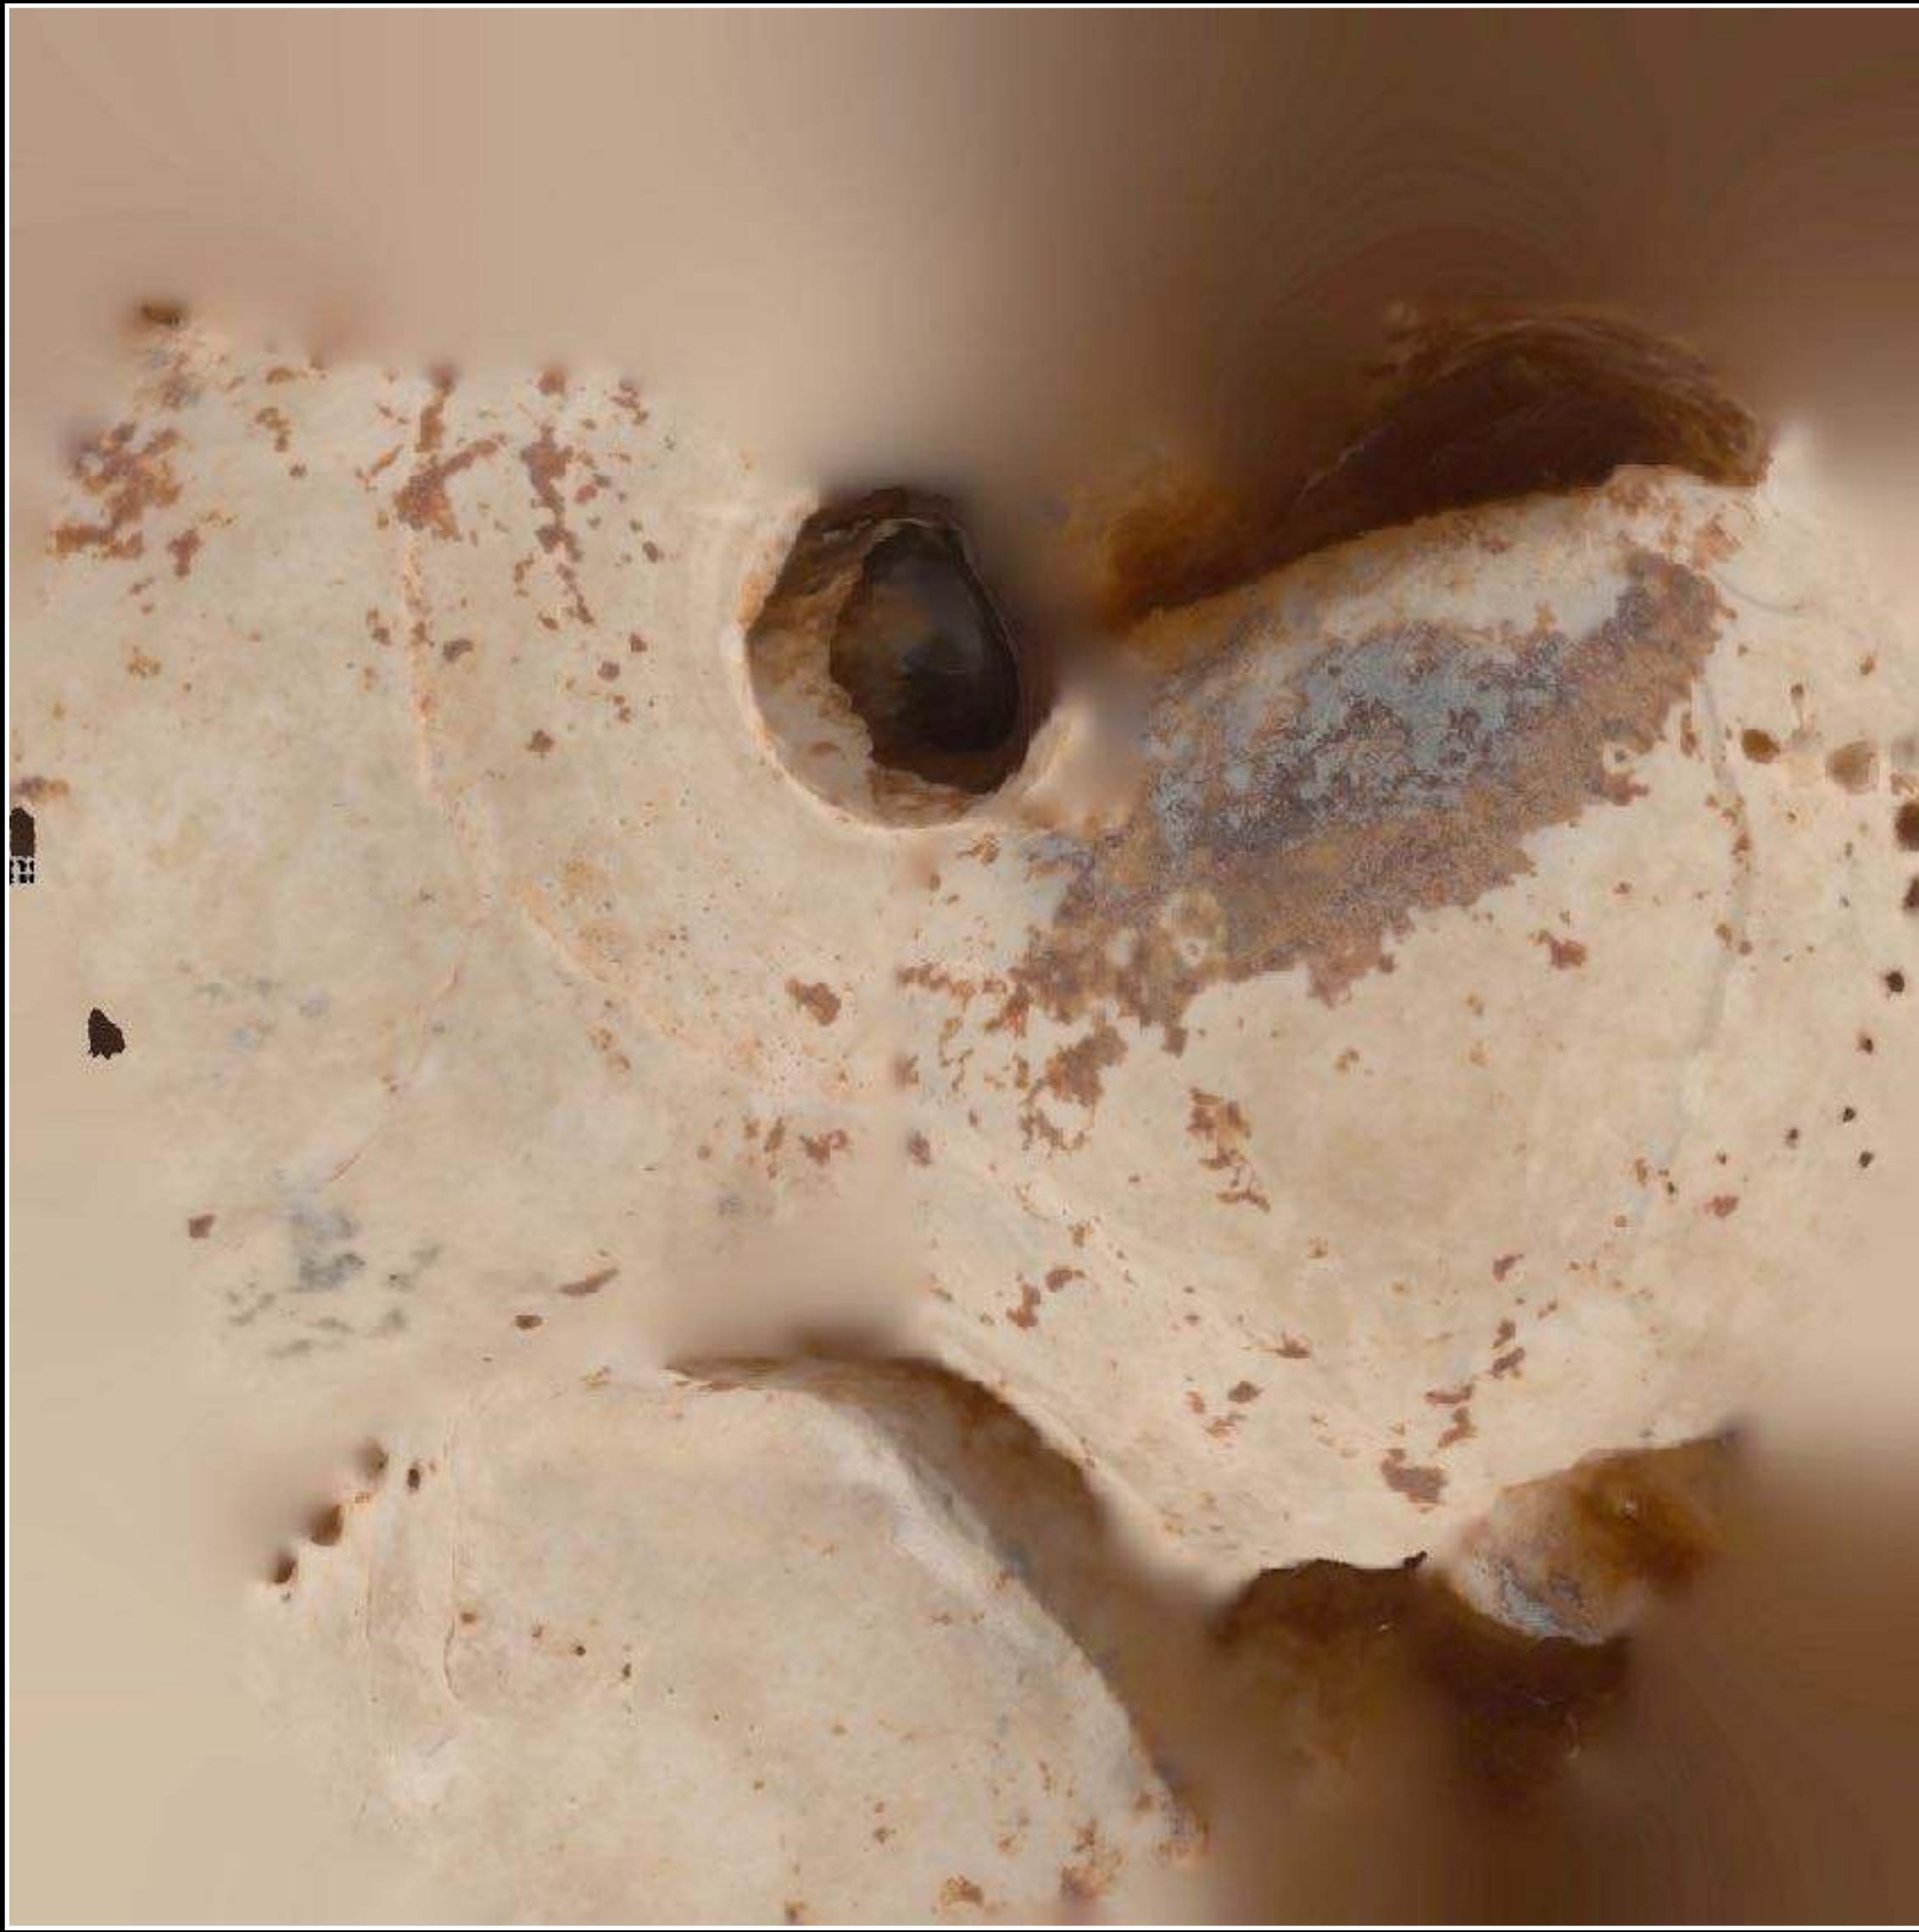
- Lots of applications to editing the texture file created.
 - Colour matching to a colour standard. Swatches in the reconstructed scene.
 - Applying filters for effect, eg: sharpening, increasing vibrance and so on.
 - Annotating, eg: labels, contour lines, segmentation.
 - Exposing shadow detail.
- Sometimes the baked in shadows can be removed, or differences in shading reduced.
- Holes in the model will generally not have any texture so may be filled in.
- Textures are not always clean and may need manual correction.





W966a

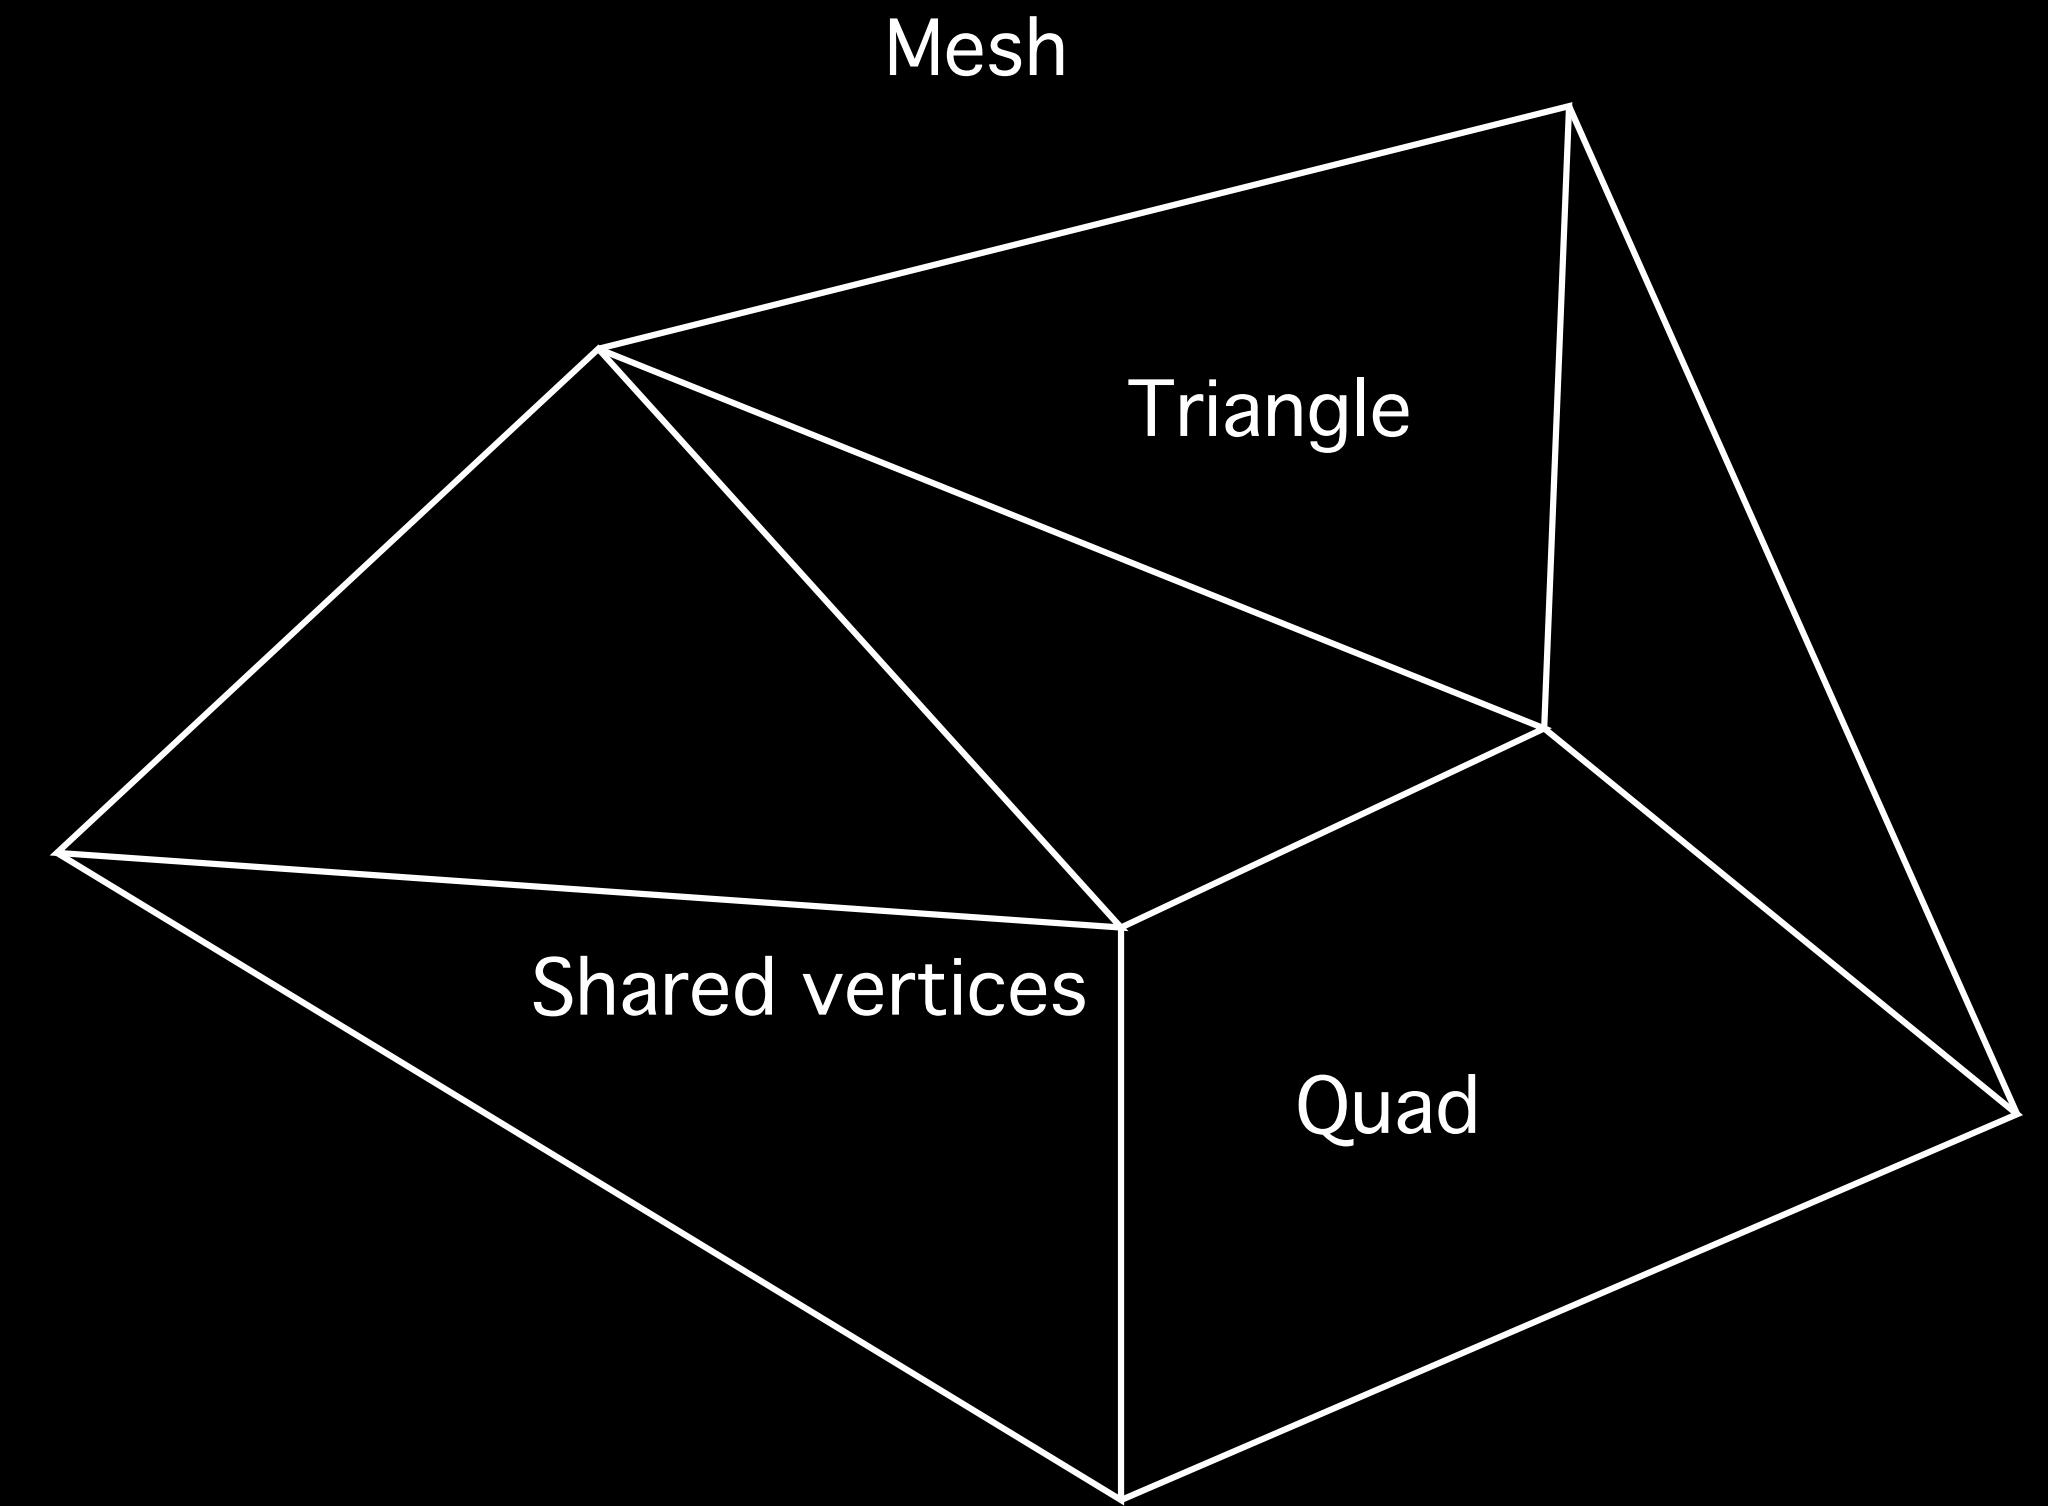






Data formats

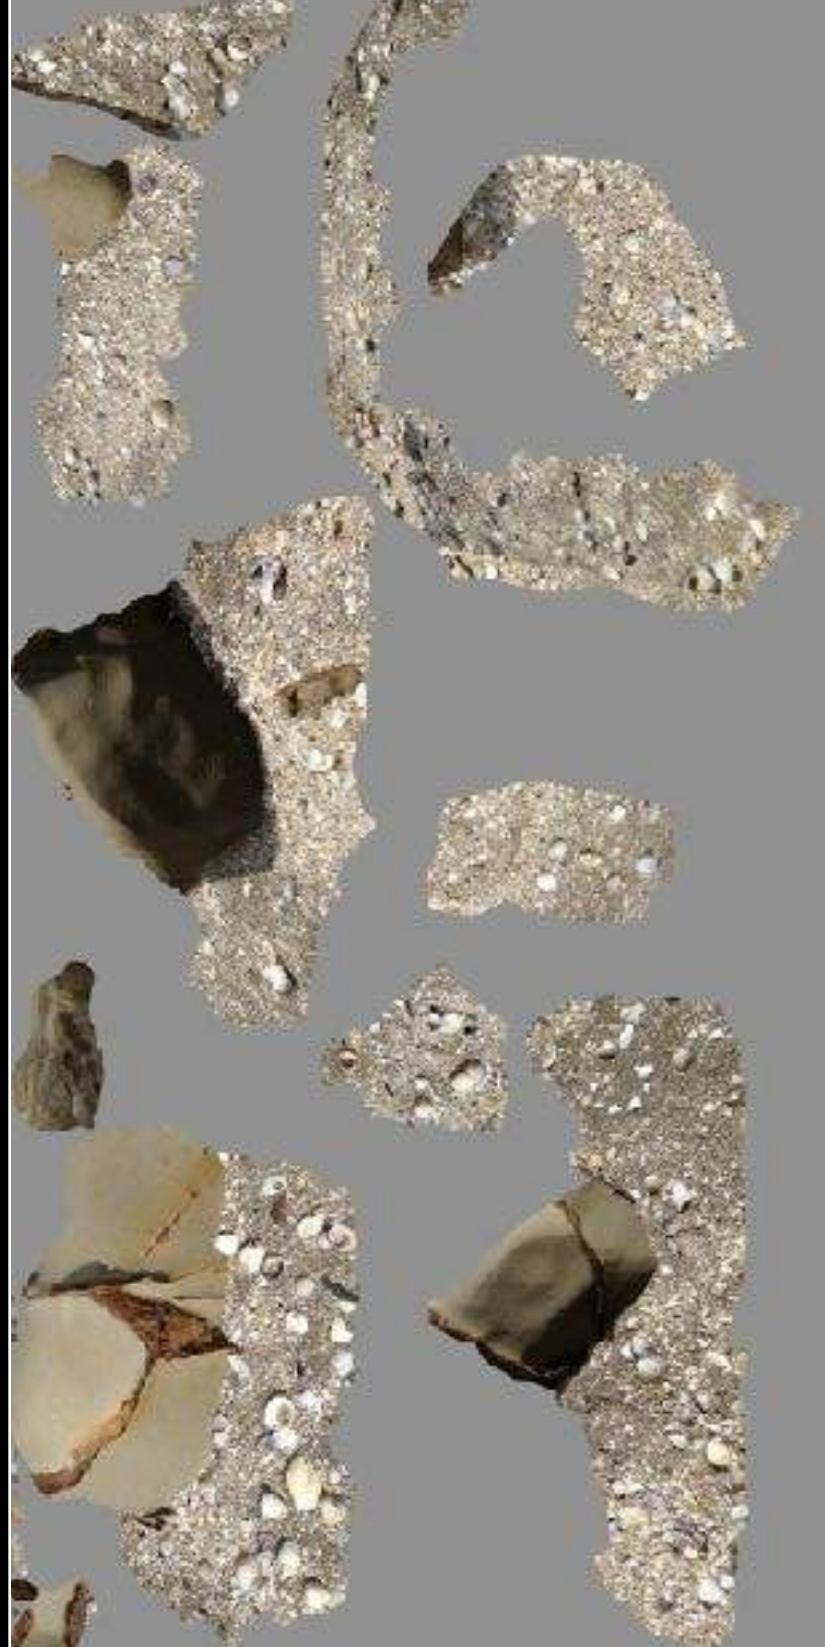
- Requirements: unstructured triangular mesh.
 - mesh (vertices - edges - triangles - polygons)
 - texture coordinates
 - image based textures
- Common options
 - 3ds (3DStudioMax)
 - vrml, x3d
 - obj (Wavefront)
 - dae (collada)
- Pretty much standardised on obj, desirable characteristics.
 - text only so human readable
 - relatively easy to parse by software for post processing or custom utilities
 - well supported by commercial 3D applications (import/export)
 - shared vertices so no chance of numerical holes
 - supports multiple texture materials and images



Anatomy of an OBJ description

Consists of 3 parts

- obj file: vertex, face, normals, texture coordinates
- materials file: colours and/or texture filename
- texture file(s): image



texture file

material file

material name

```
newmtl rockmaterial
Ka 0.2 0.2 0.2
Kd 0.752941 0.752941 0.752941
Ks 1.000000 1.000000 1.000000
Tr 1.000000
illum 2
Ns 0.000000
map_Kd stone_tex_0.jpg
```

obj file

```
mtllib stone.obj.mtl
```

```
v 7.980470 5.627900 3.764240
v 8.476580 2.132000 3.392570
v 8.514860 2.182000 3.396990
```

```
:
```

```
:
```

```
vn -0.502475 -1.595313 -2.429116
```

```
vn 1.770880 -2.076491 -5.336680
```

```
vn -0.718451 -4.758880 -3.222428
```

```
:
```

```
:
```

```
vt 0.214445 0.283779
```

```
vt 0.213670 0.287044
```

```
vt 0.211291 0.287318
```

```
:
```

```
:
```

```
usemtl rockmaterial
```

```
f 5439/4403/5439 5416/4380/5416 7144/6002/7144
```

```
f 5048/4013/5048 6581/5437/6581 5436/4400/5436
```

```
f 5435/4399/5435 5049/4014/5049 5436/4400/5436
```

```
:
```

```
:
```

```
:
```

vertex

index

normal

index

texture

coordinate

index

material filename

vertices

normals

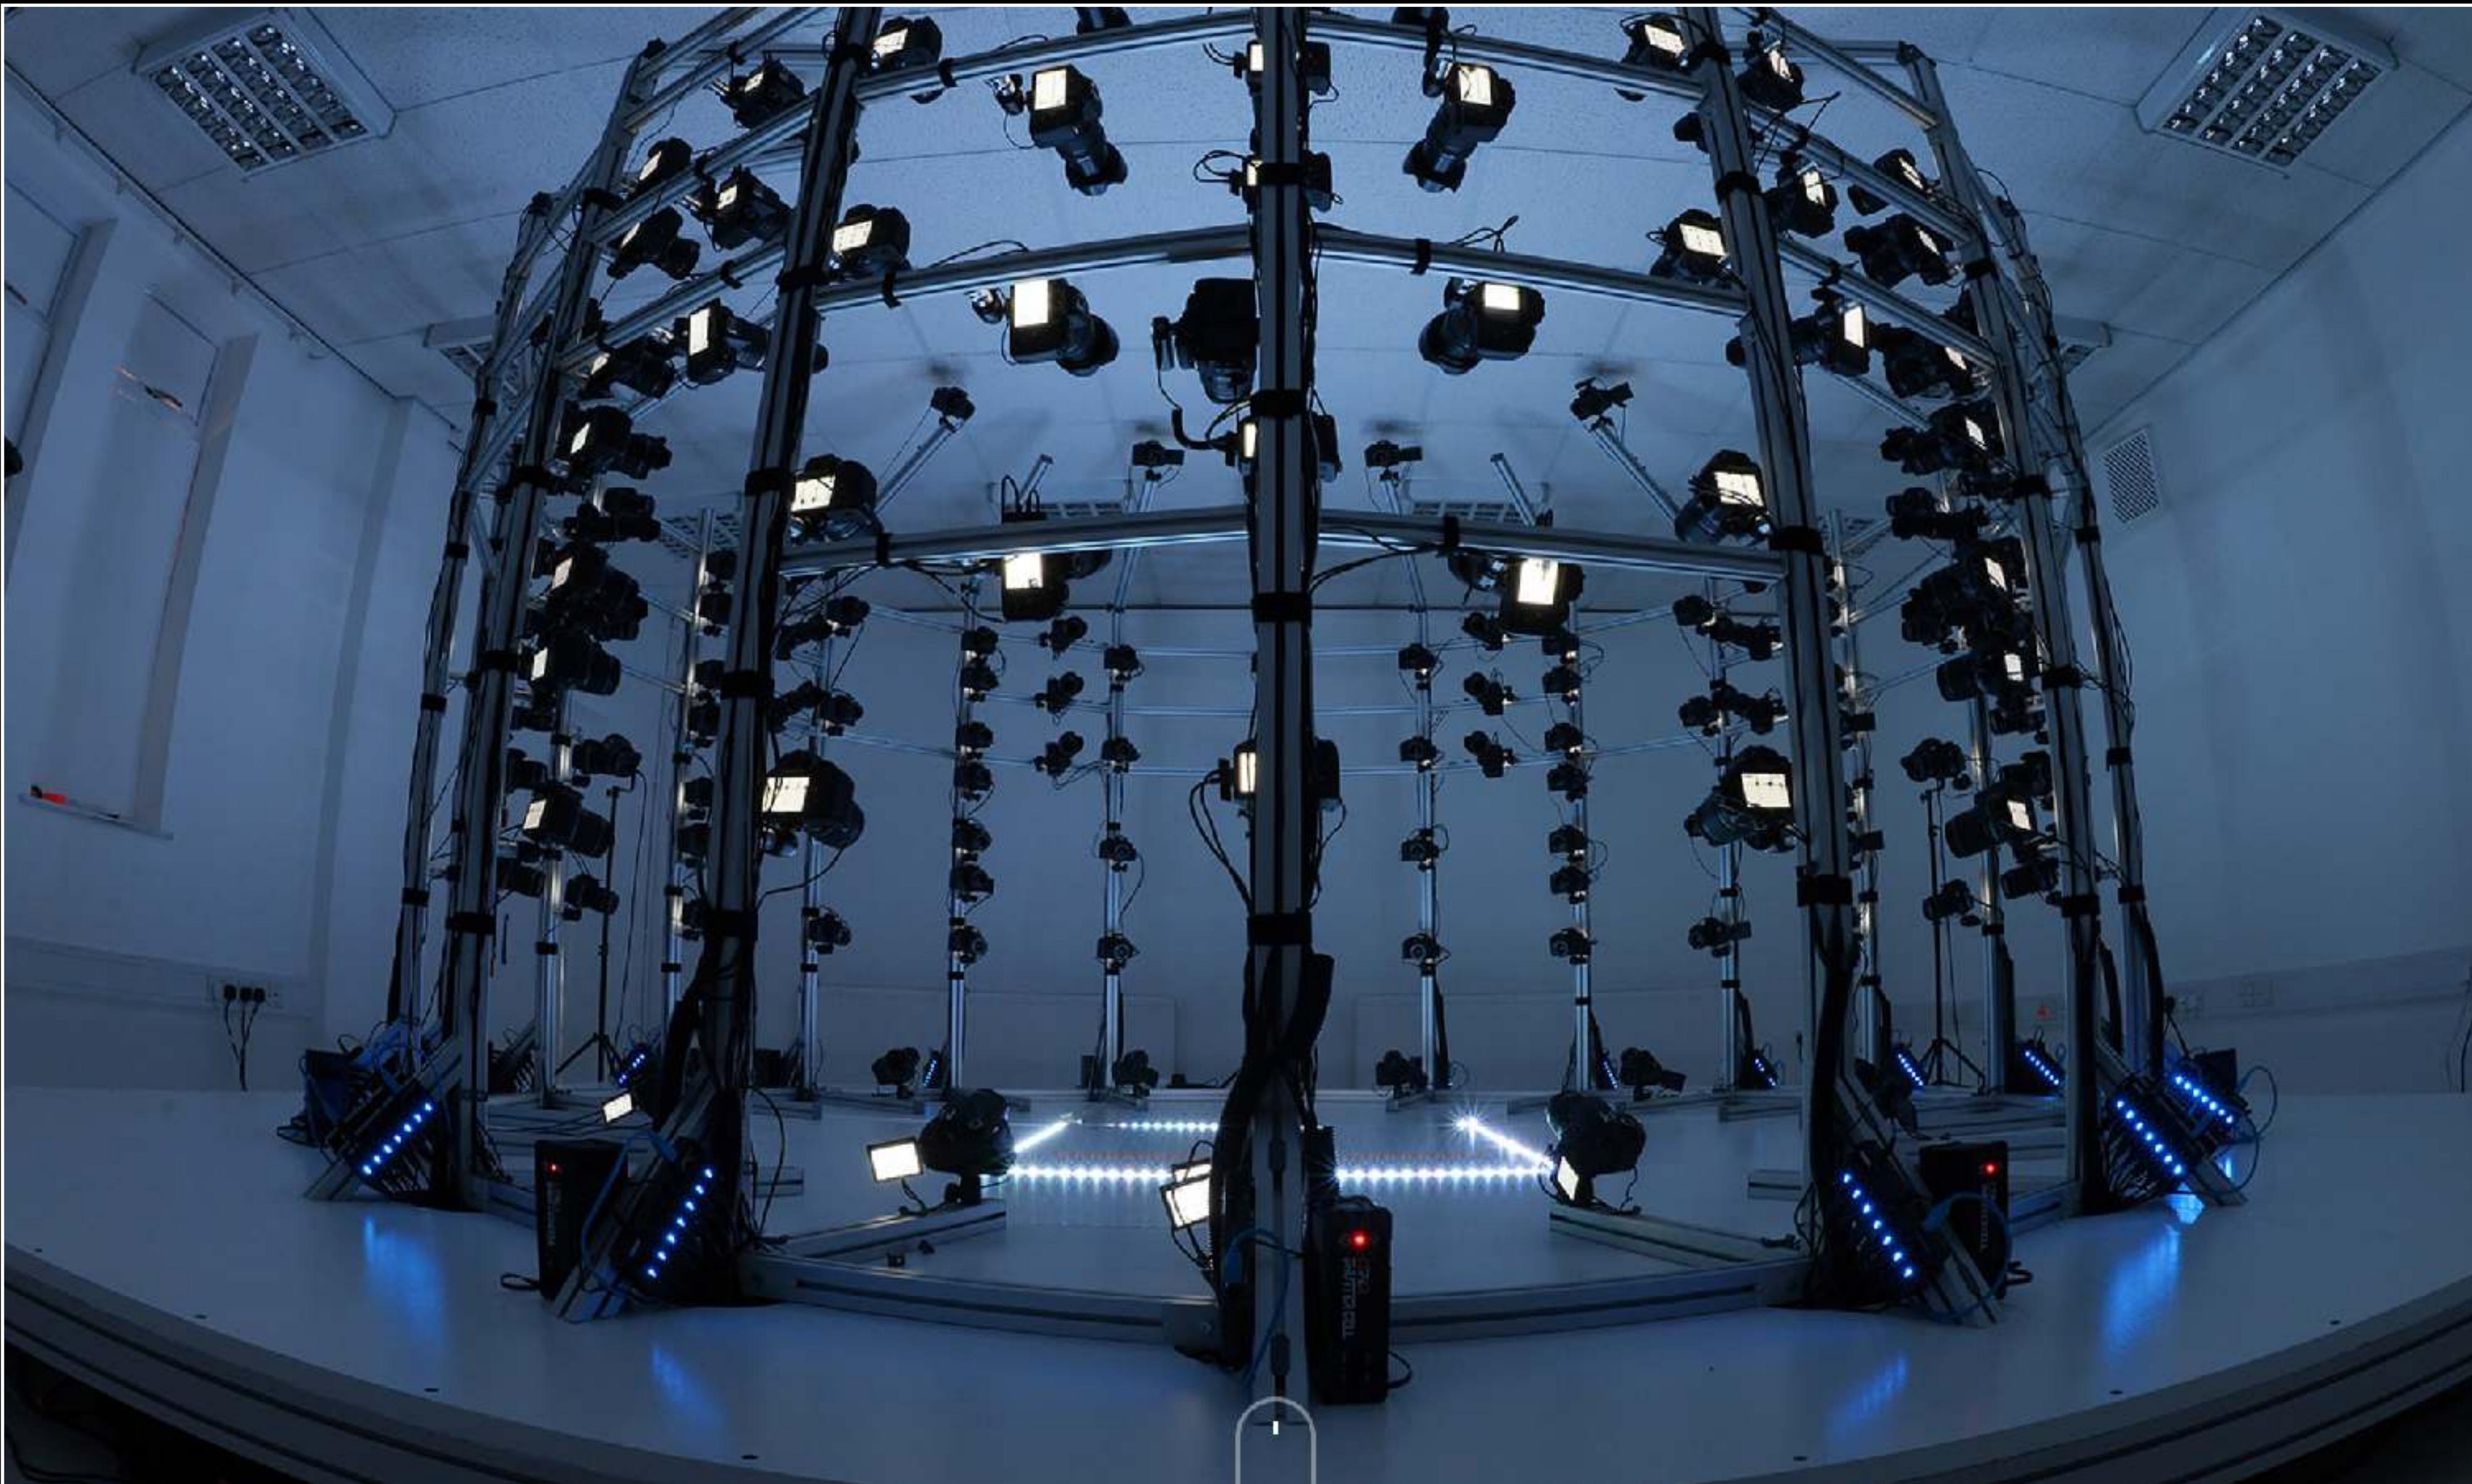
texture coordinates

triangles

Limitations

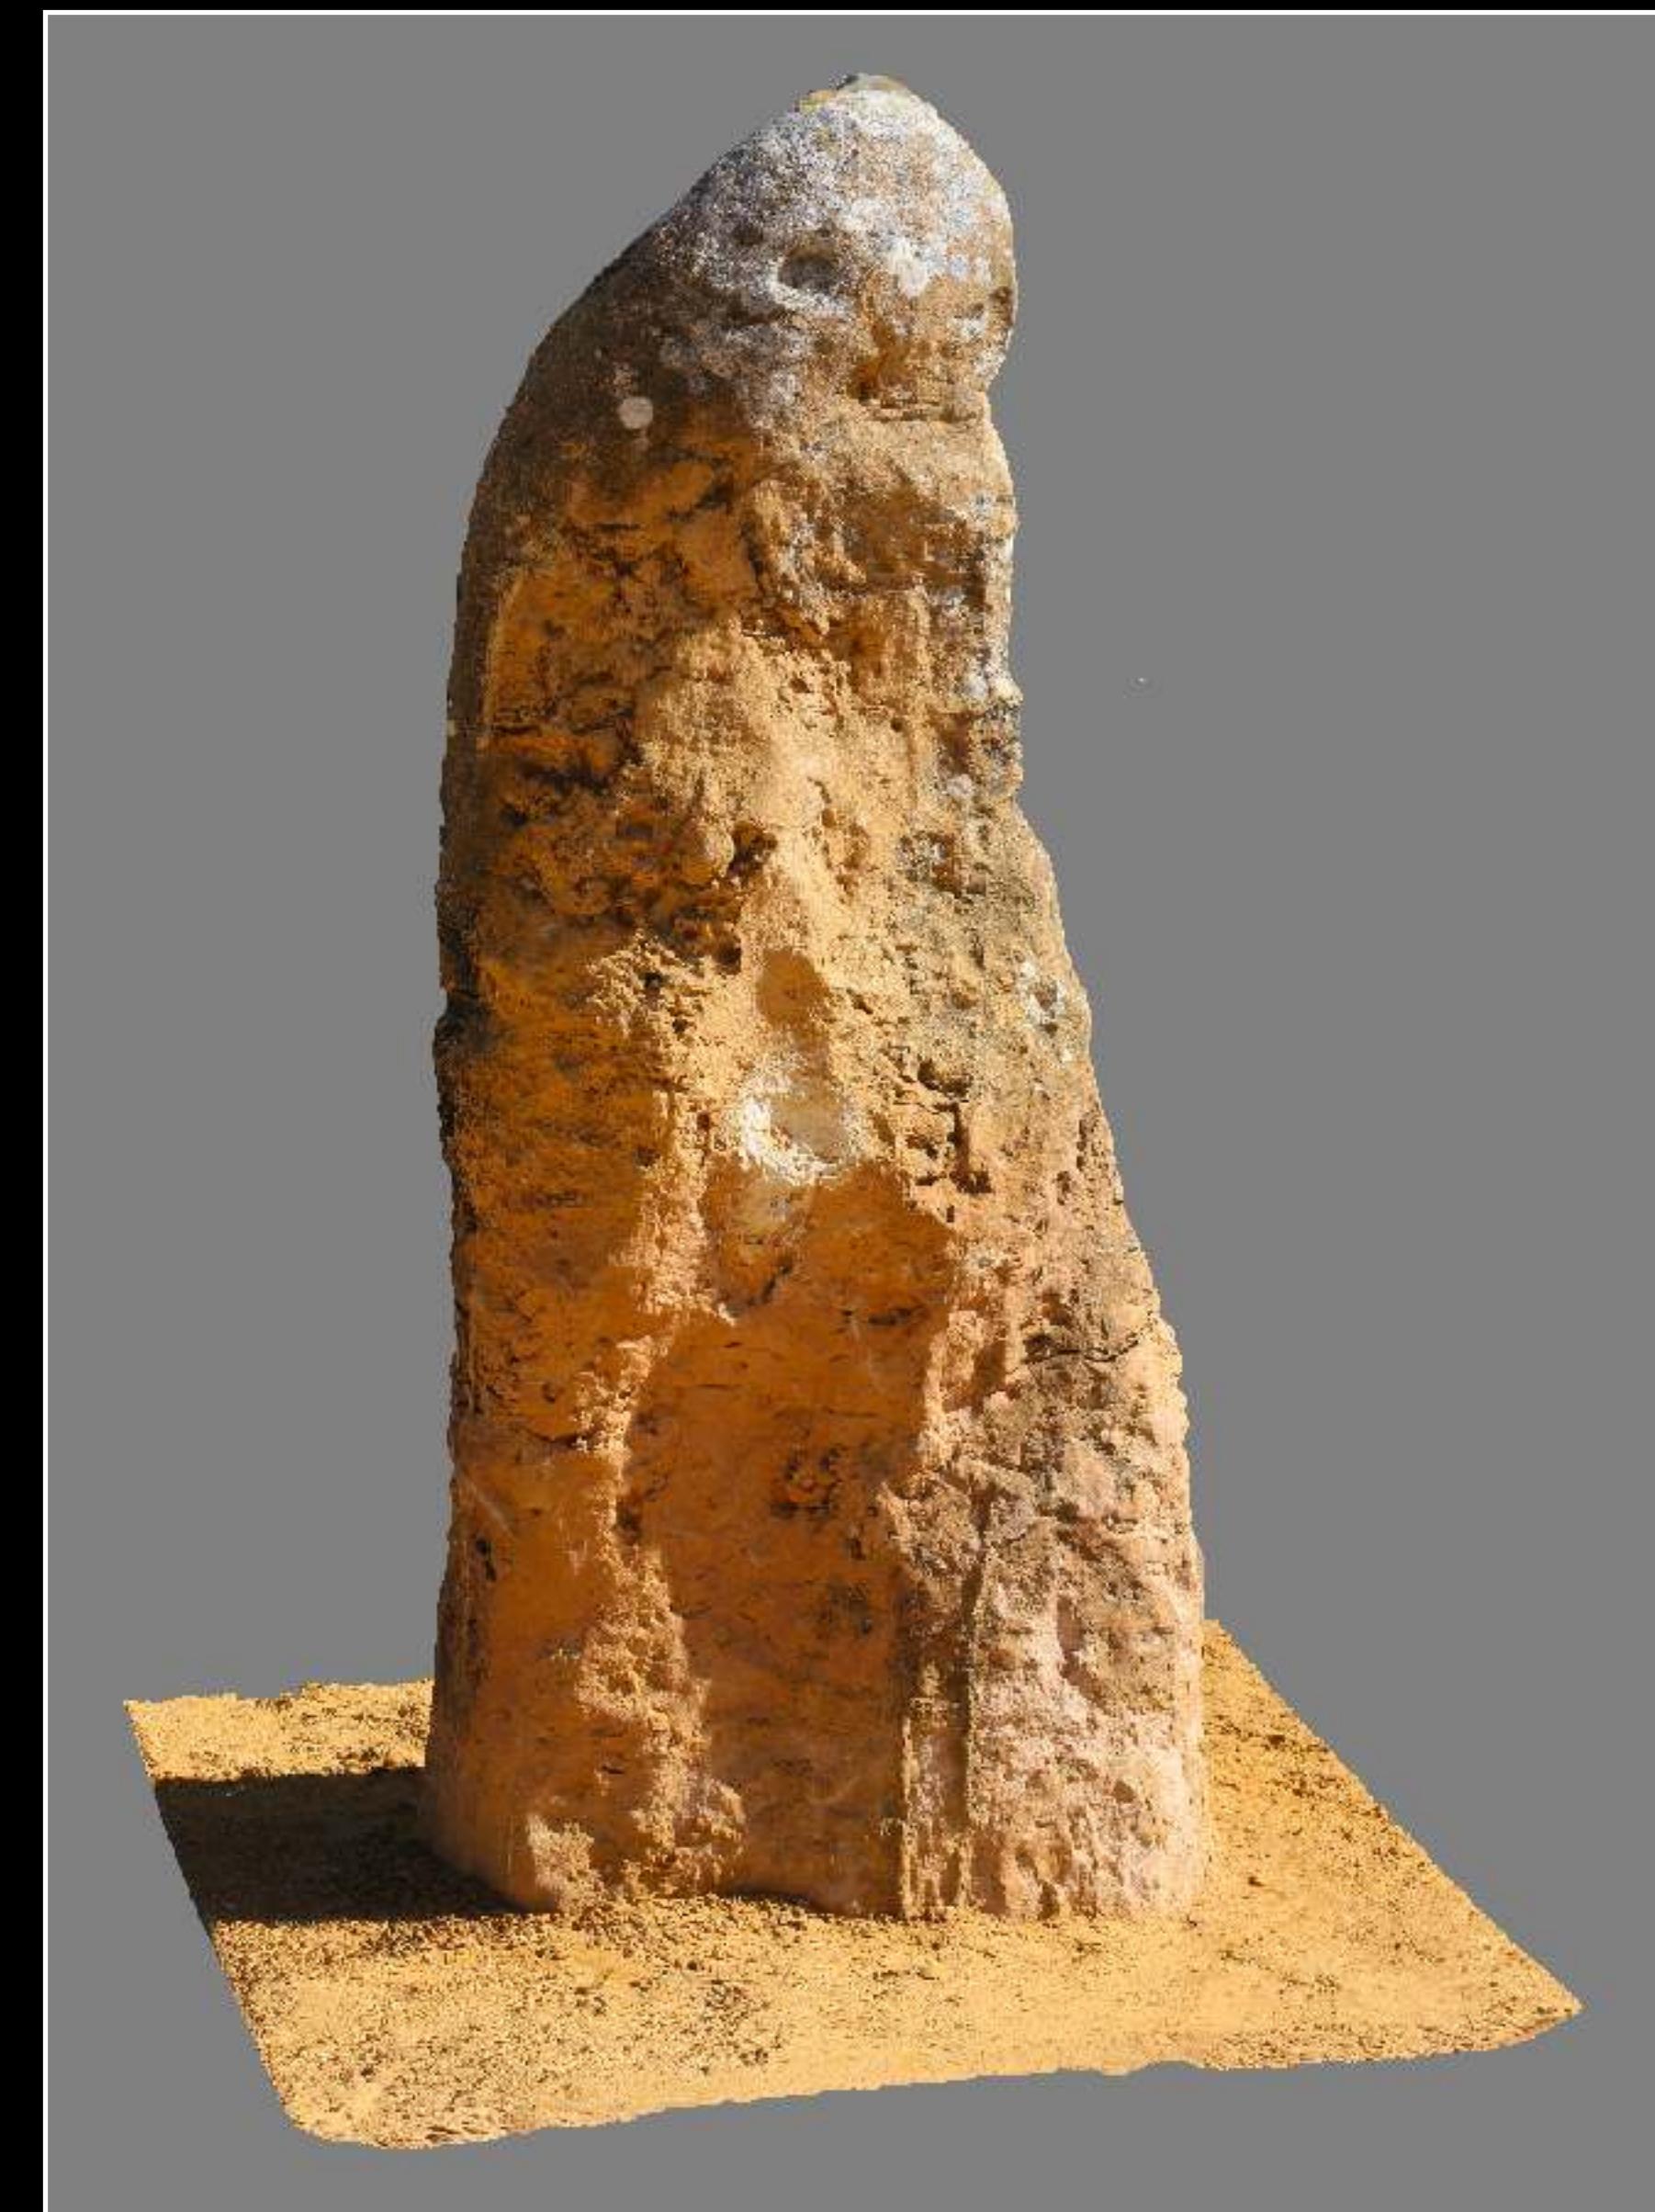
- Require static objects otherwise a multiple camera rig required.
- Shadows and shading are baked into the textures.
Shadows of the photographer.
Diffuse lighting ideal = cloudy skies
- In the field, one cannot always photograph from desired locations.
Poles - Drones - Cranes - Ladders
- Cannot capture structures that are below a few pixels wide in photographs.
- Reflective, specular surfaces problematic.
Dusting - Painting - Polarisation technique
- Smooth featureless surfaces problematic.
Dusting - Painting

Camera rigs for human scanning



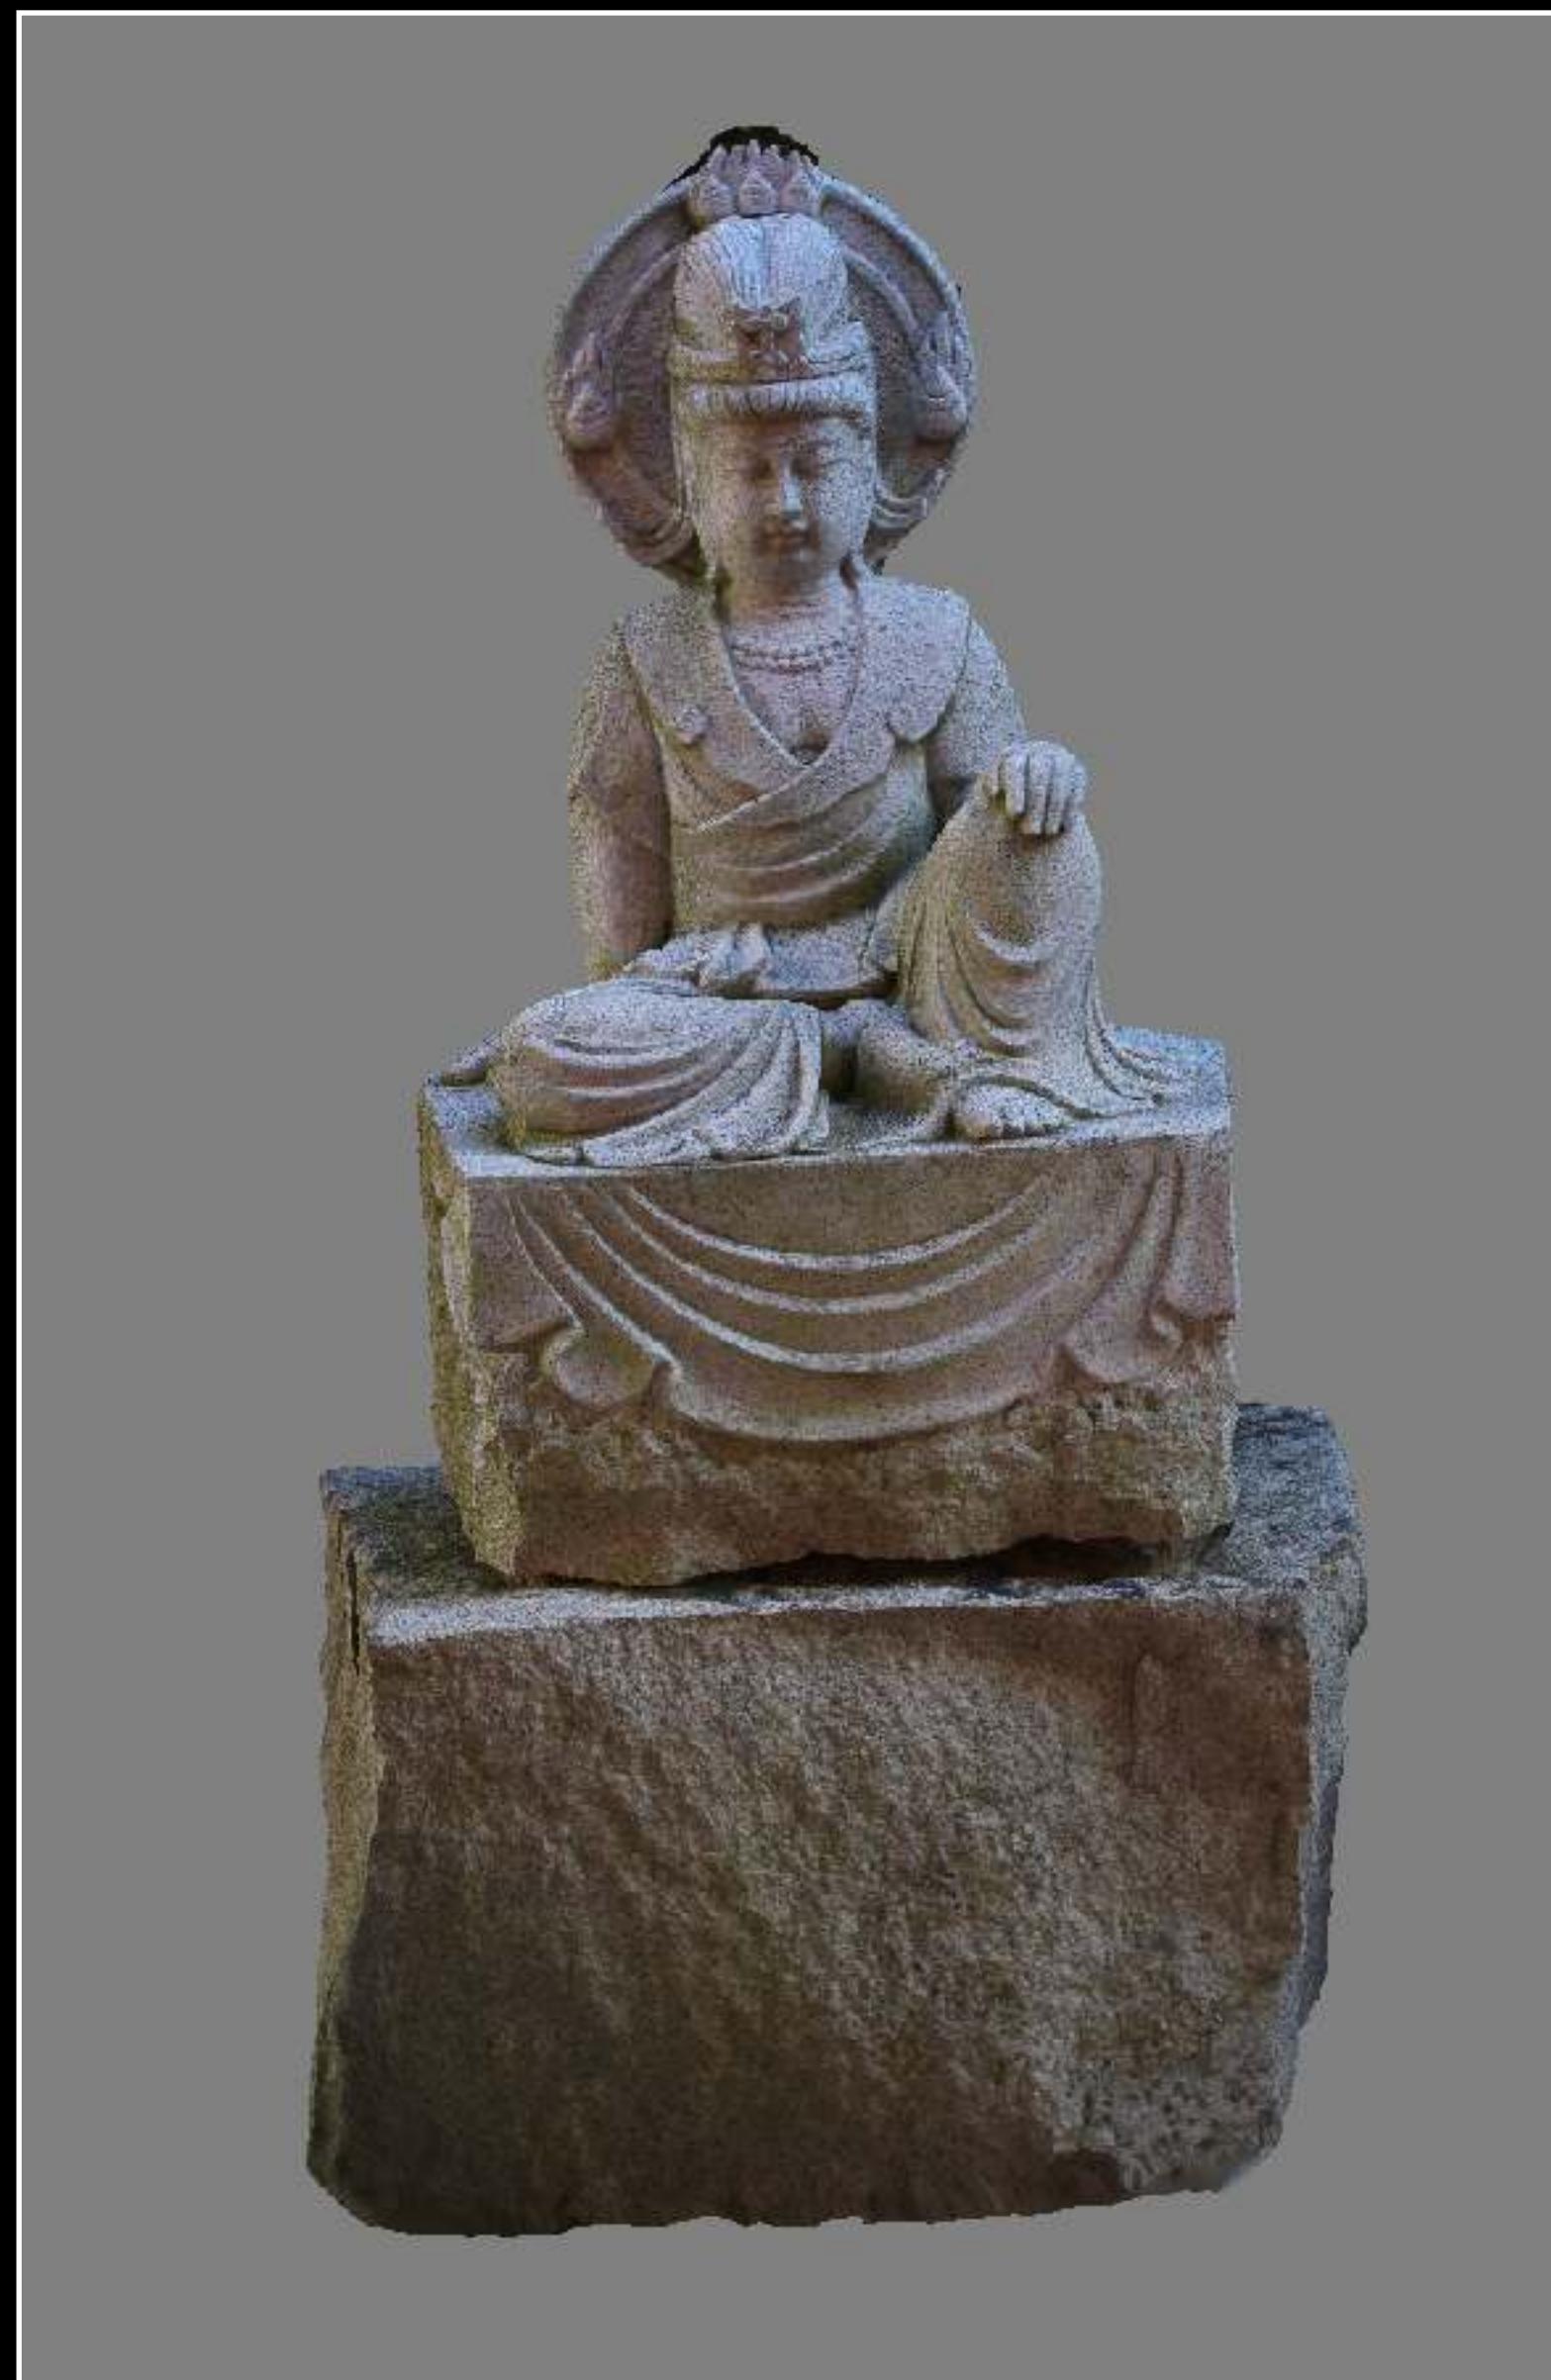


Shadows and Shading



Access to camera positions





Post Production

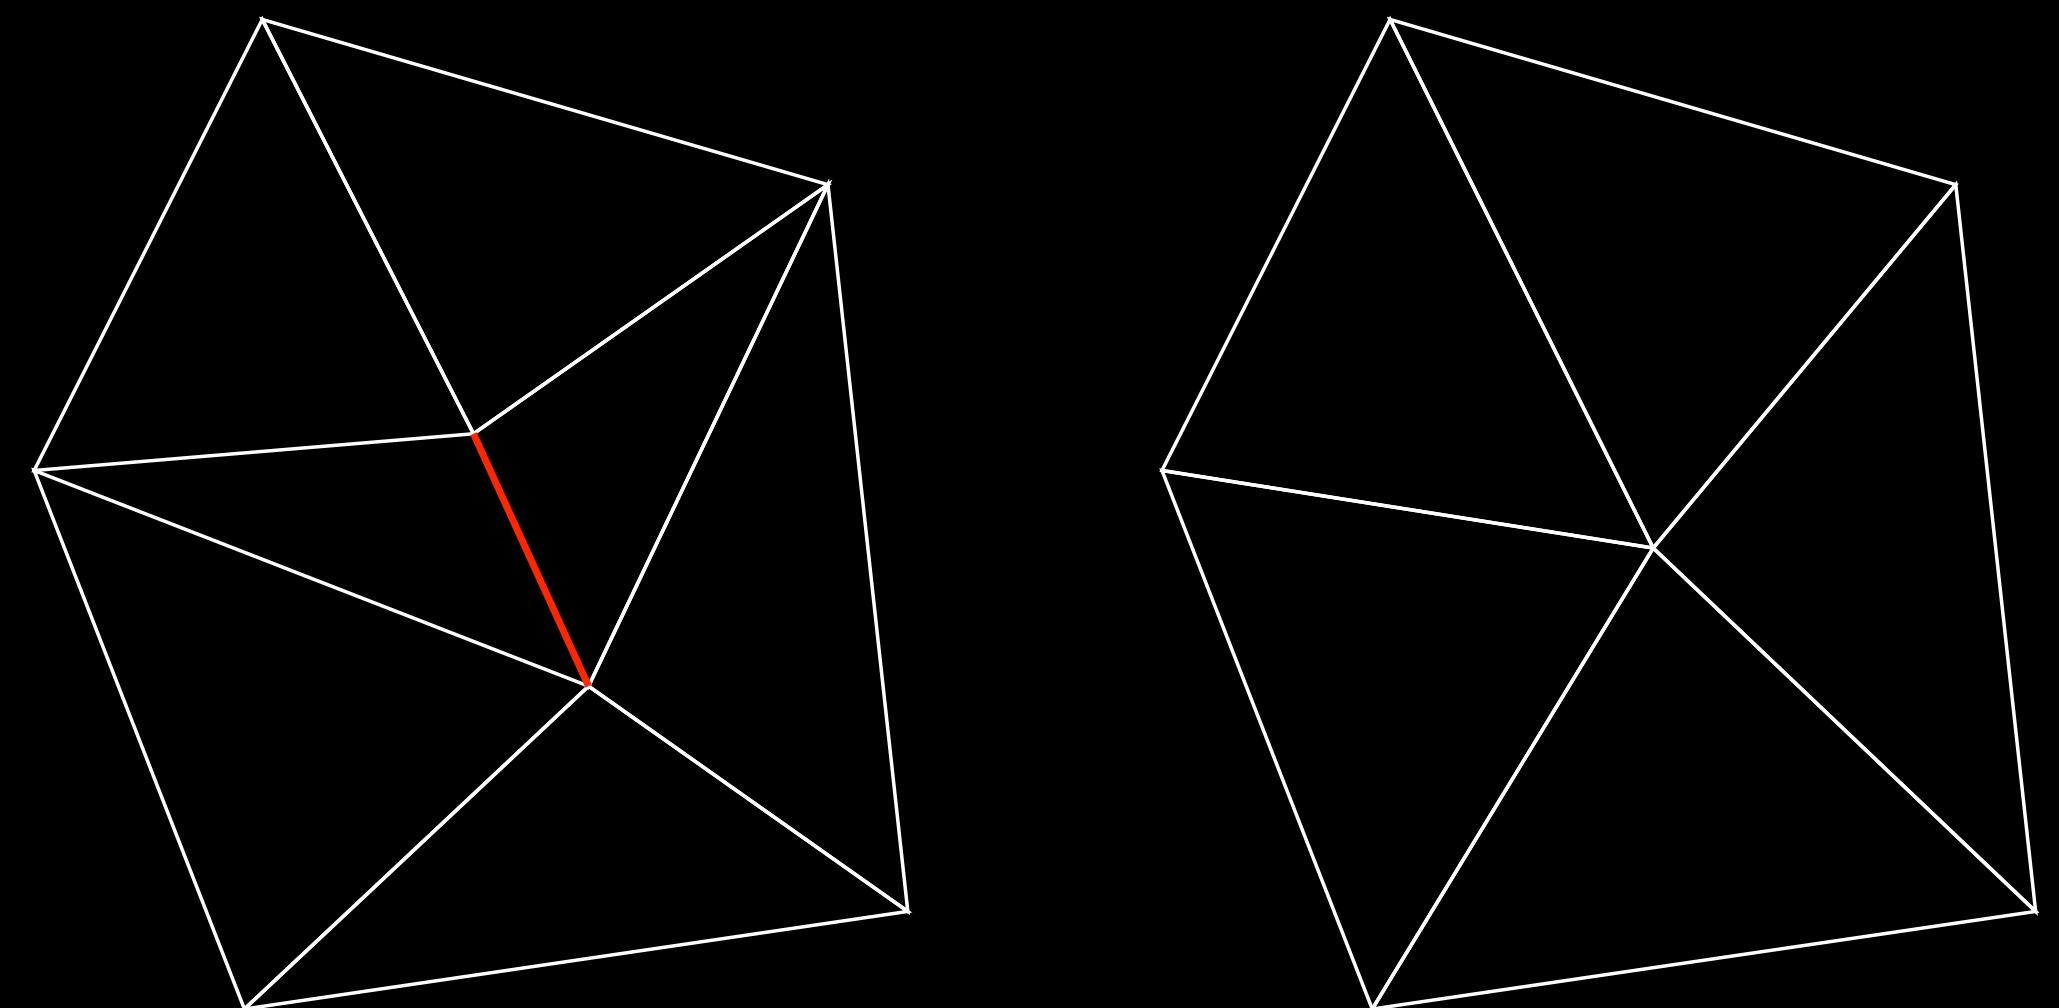
- Mesh simplification
- Thickening (3D printing)
- Relighting
- Hole closing ("watertight" for 3D printing)
- Removing shrapnel
- Per vertex editing
- Smoothing

Mesh simplification

- Meshes directly from the reconstruction (generated from the dense point cloud) are generally inefficient. Often need to reduce them for realtime applications and/or web based delivery.
- Also used to create multiple levels of details (LOD) for gaming and other realtime applications.
- The goal is easy to understand: remove mesh density where it will make minimal impact on the mesh appearance. For example, don't need high mesh density in regions of low curvature.
- Most common class of algorithm is referred to as "edge collapse", replace an edge with a vertex.
- A texture and geometry approximation ... need to estimate new texture coordinate at new vertices.
- Need to preserve the boundary.

Edge collapse

- Most edge collapse algorithms involve replacing an edge with a vertex
 - How to choose the edges to remove is the “trick”.
 - Where to locate the new vertex so as to minimise the effect on the surface.
 - How to estimate the new texture coordinate.
- Number of triangles reduces by 2 on each iteration.
- Can calculate the deviation of the surface for any particular edge collapse.
Choose edges that result in the smallest deviation.
For example: remove edges on flat regions, retain edge in regions of high curvature.



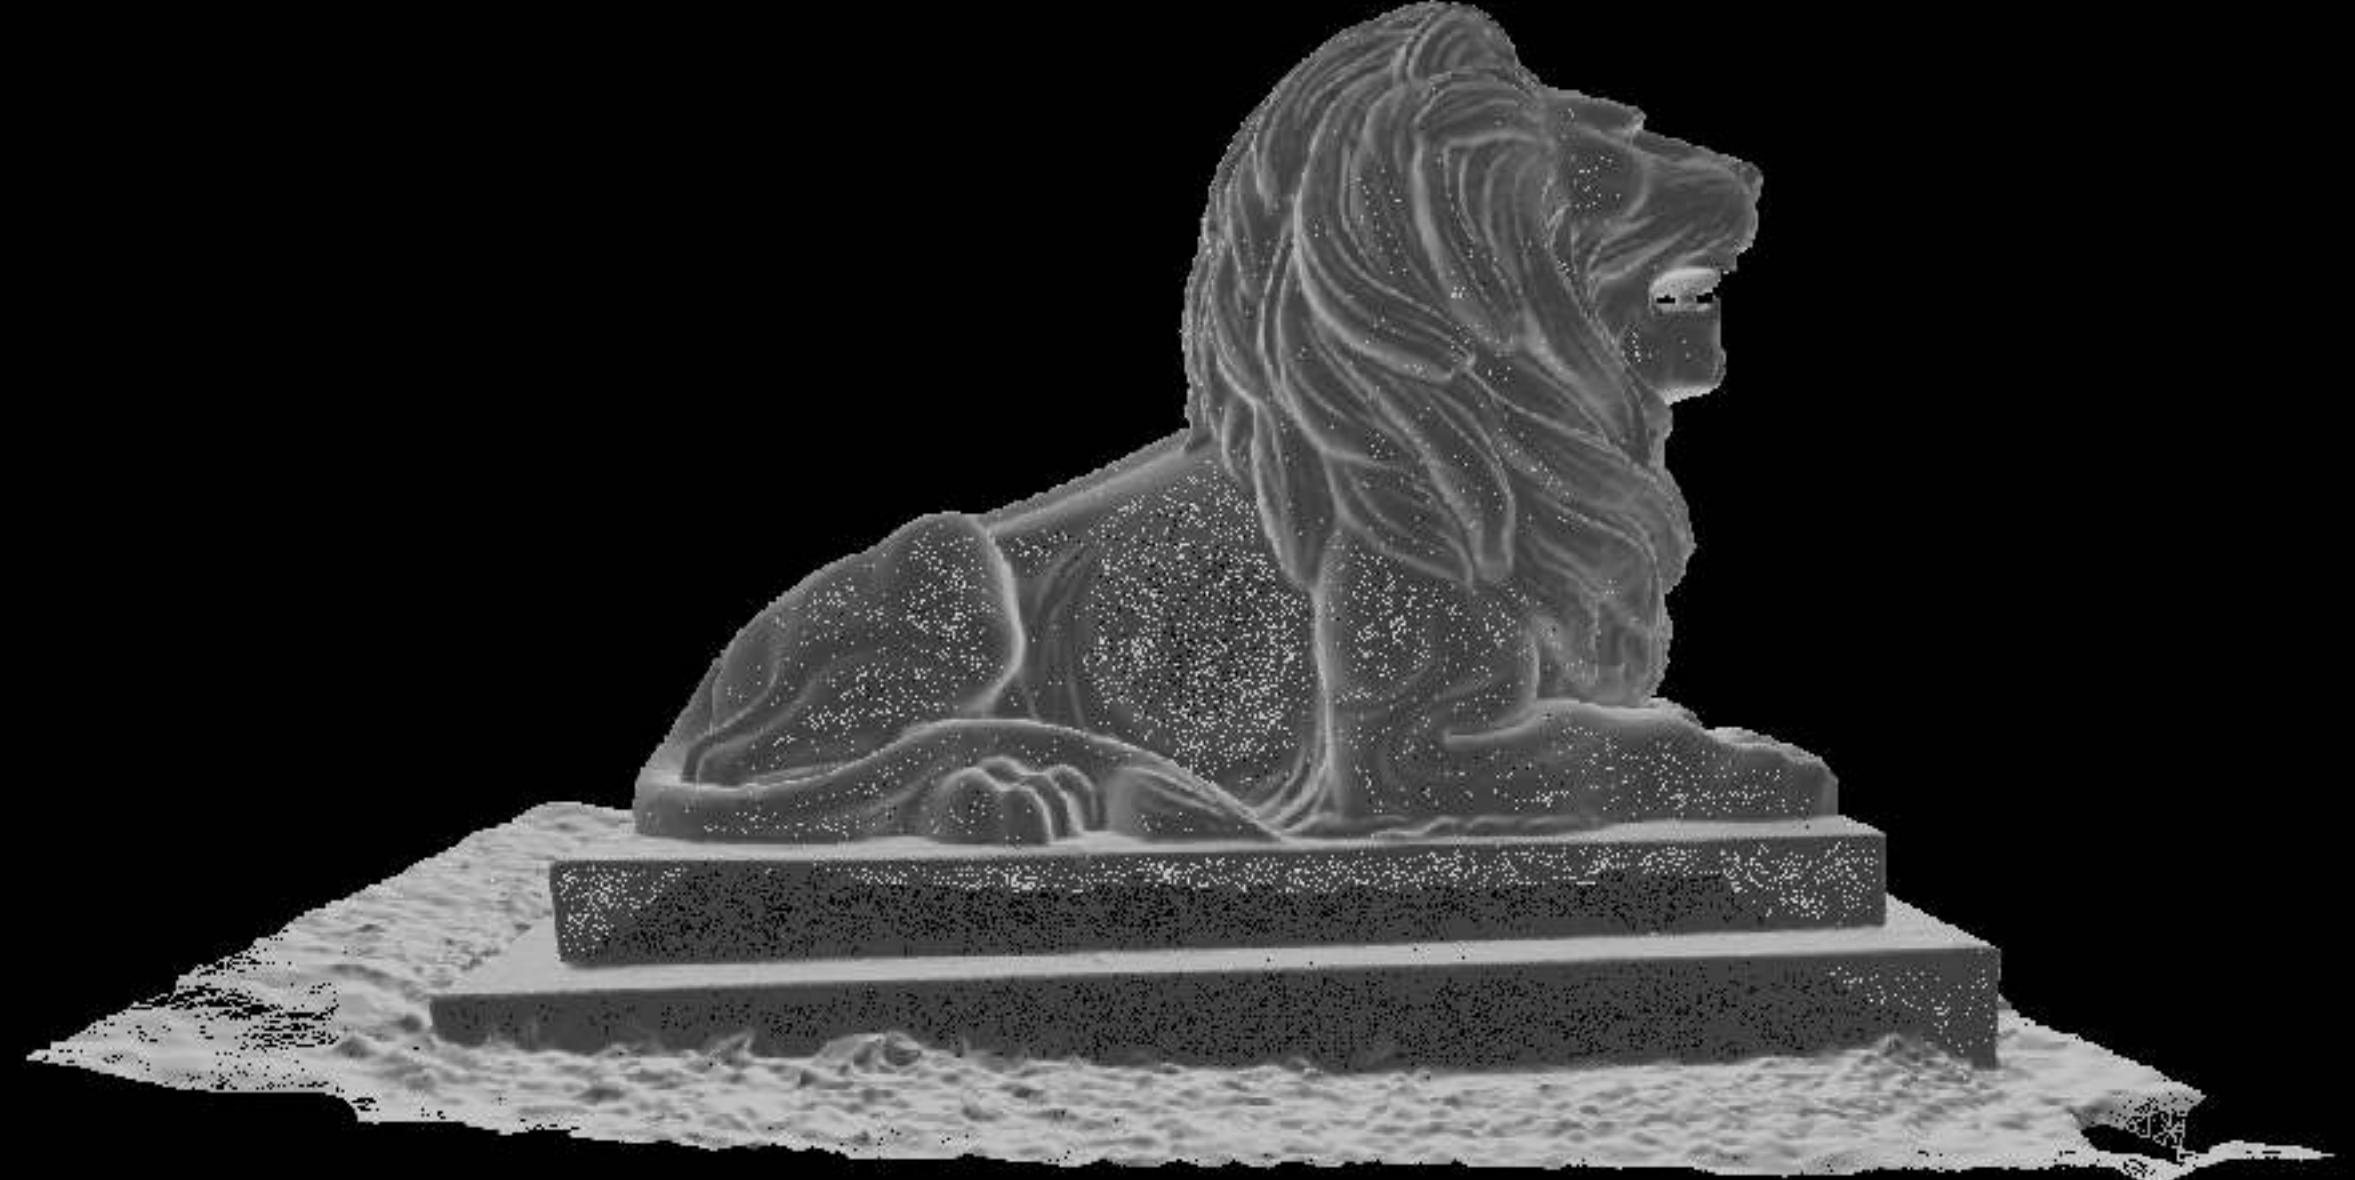
**Red edge removed,
results in two fewer
triangles**



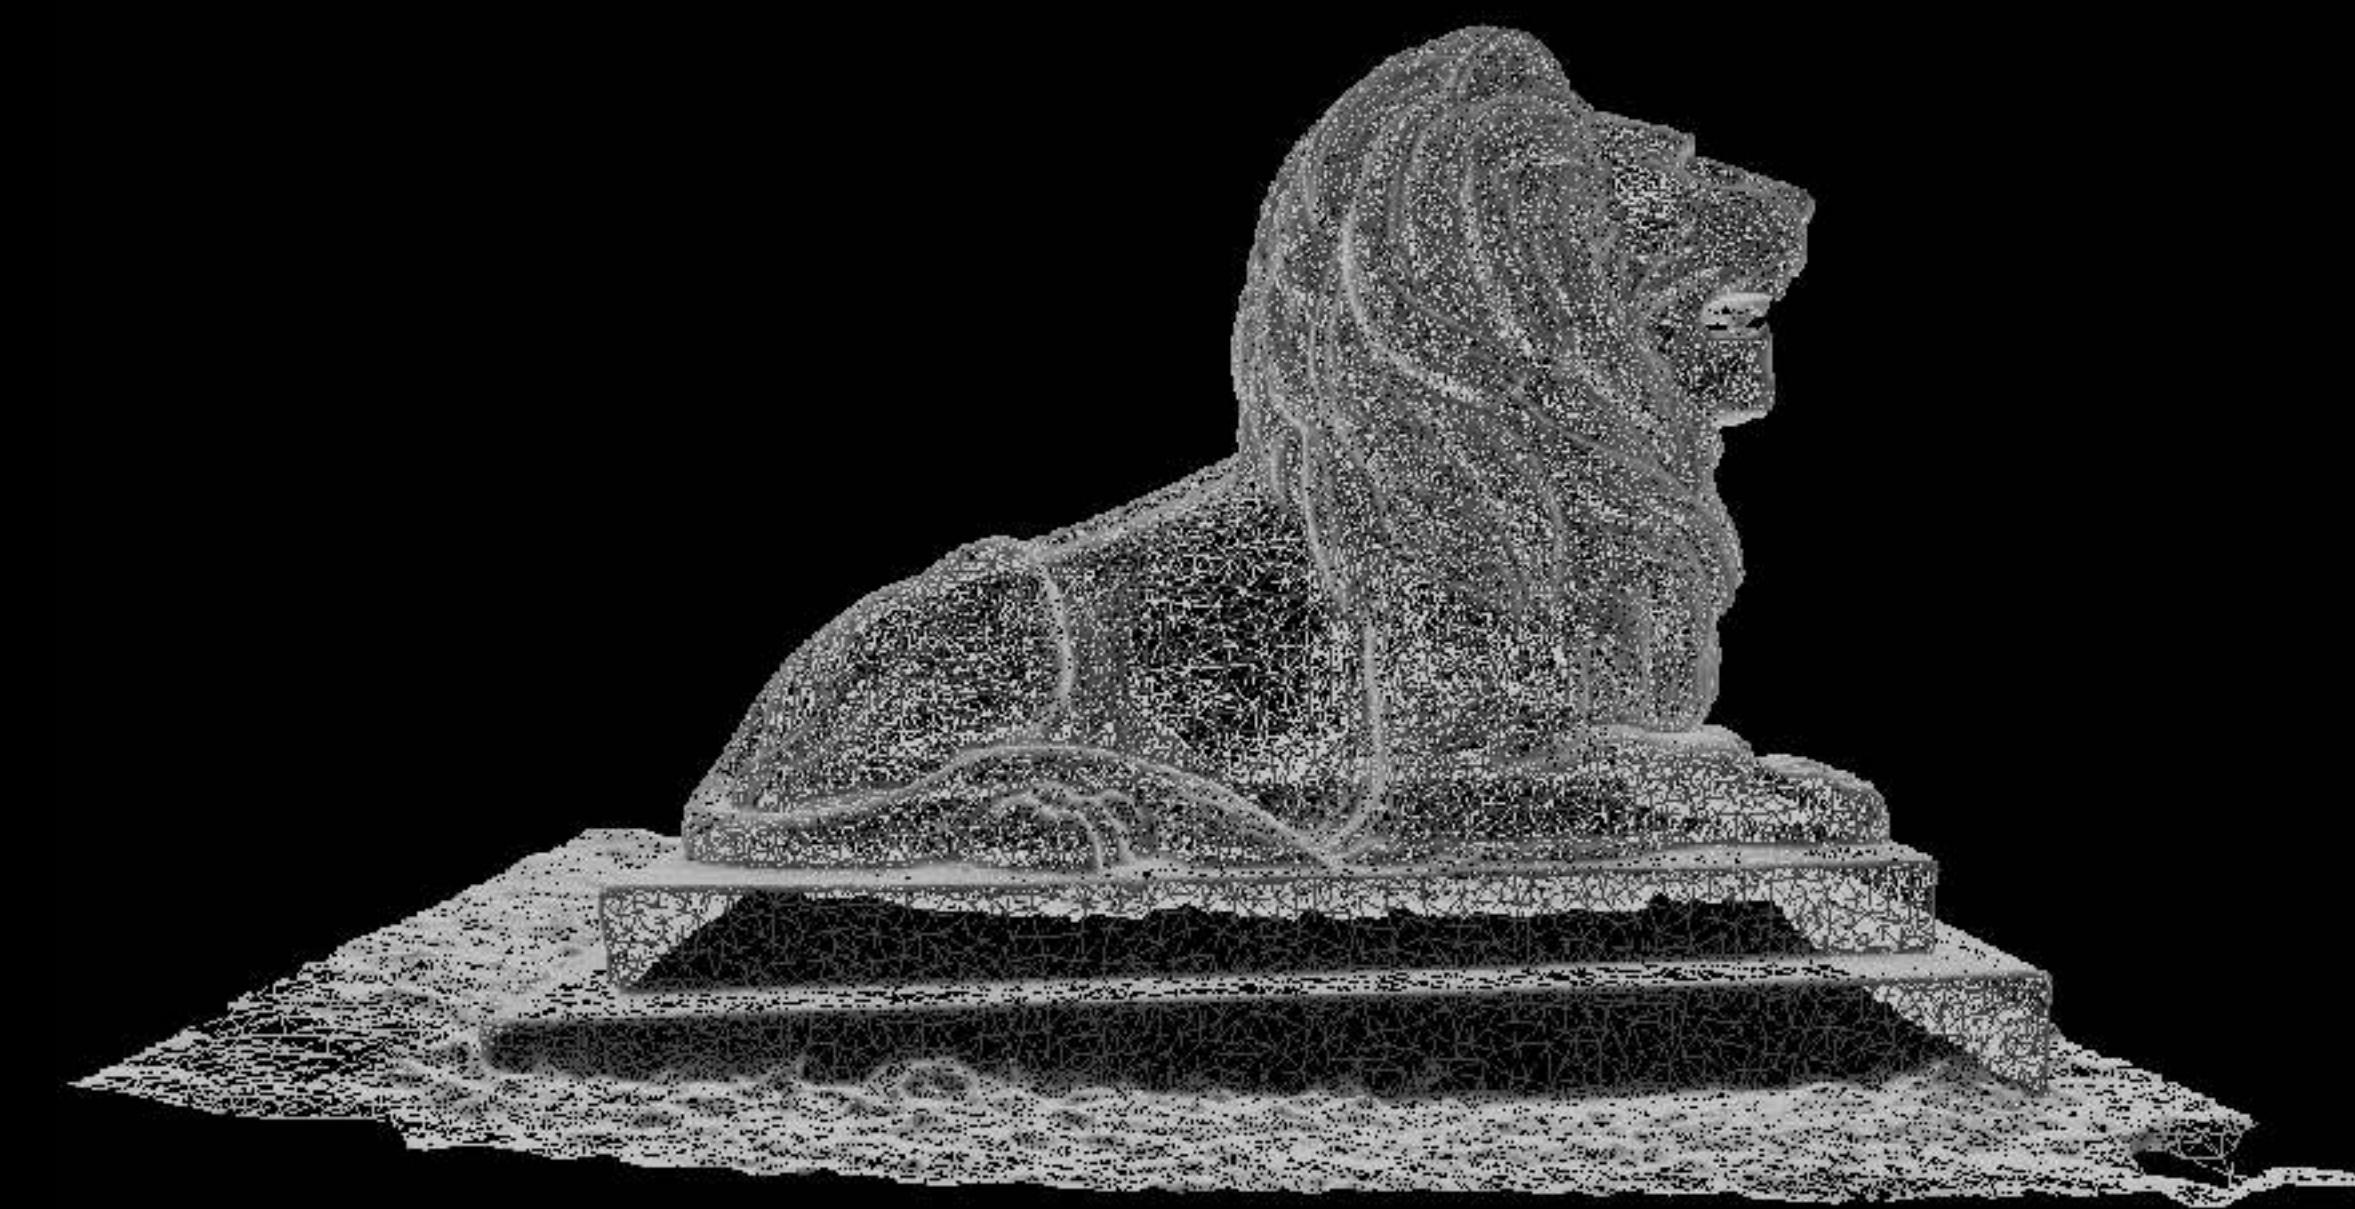
1,000,000 triangles



100,000 triangles



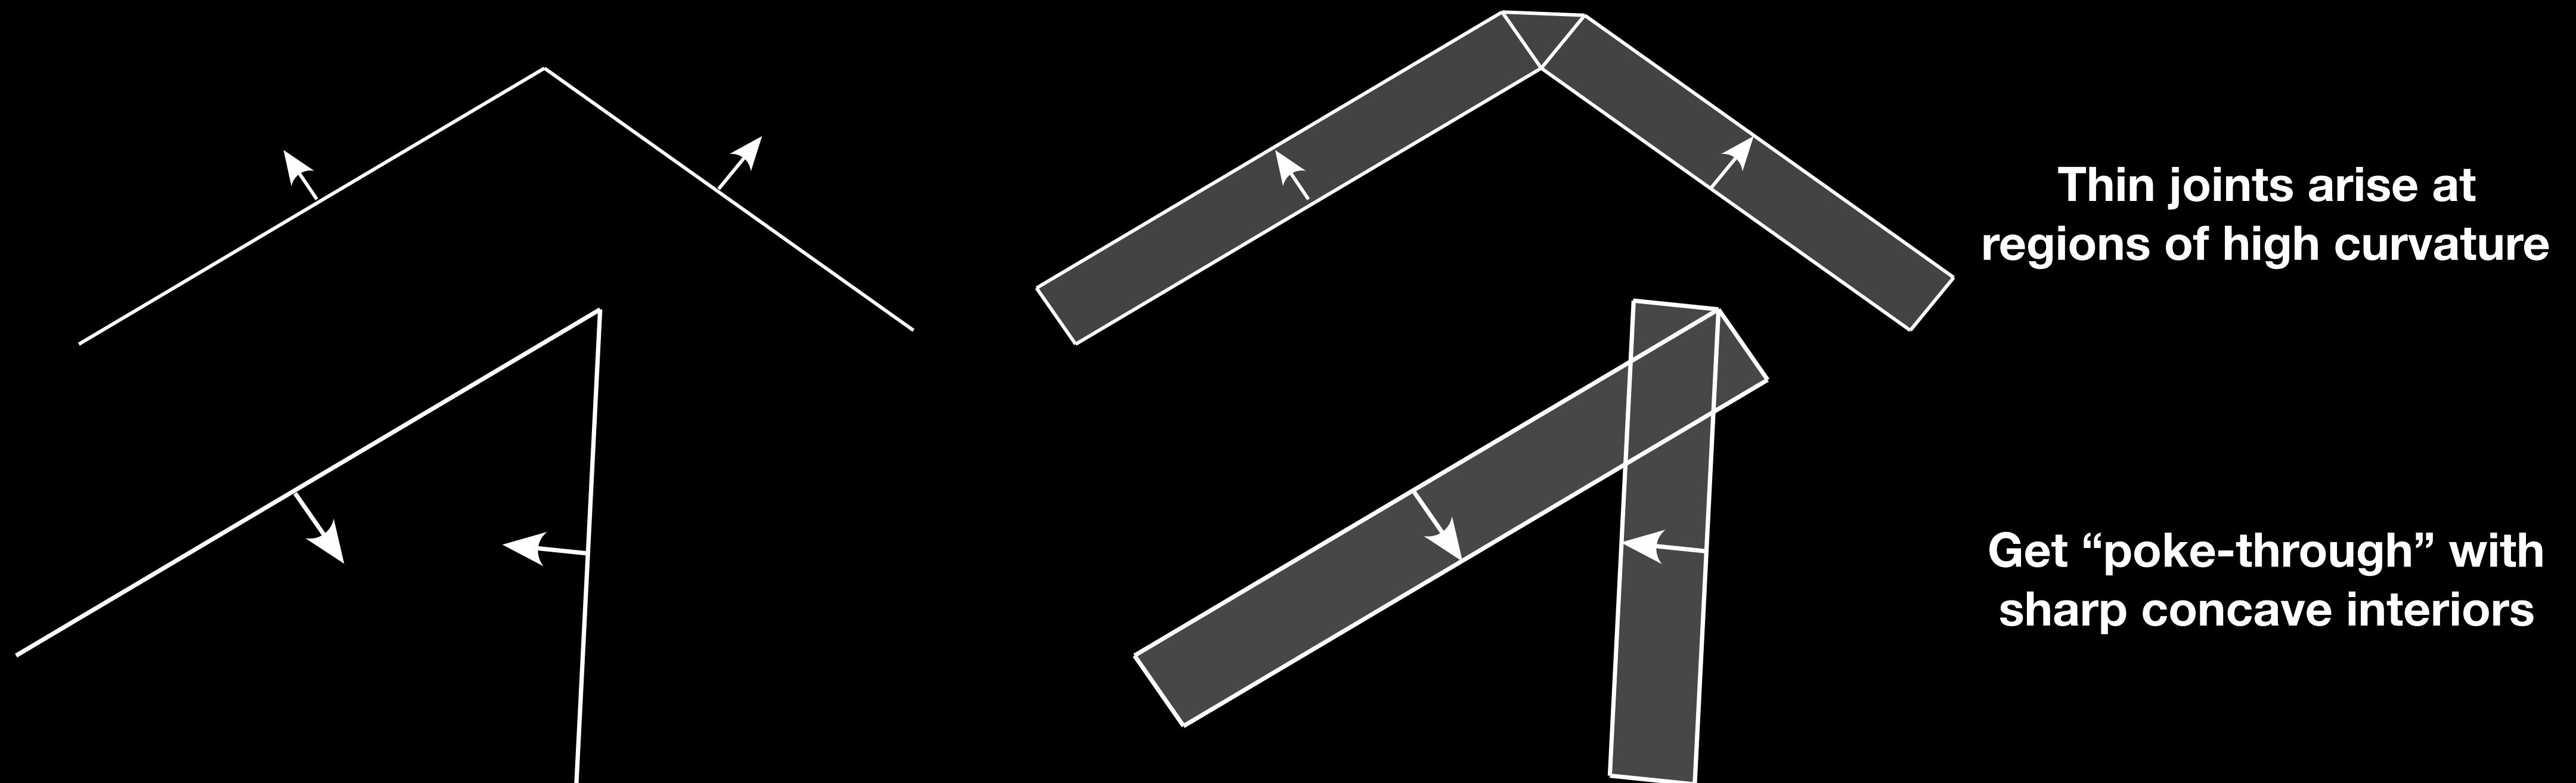
1,000,000 triangles



100,000 triangles

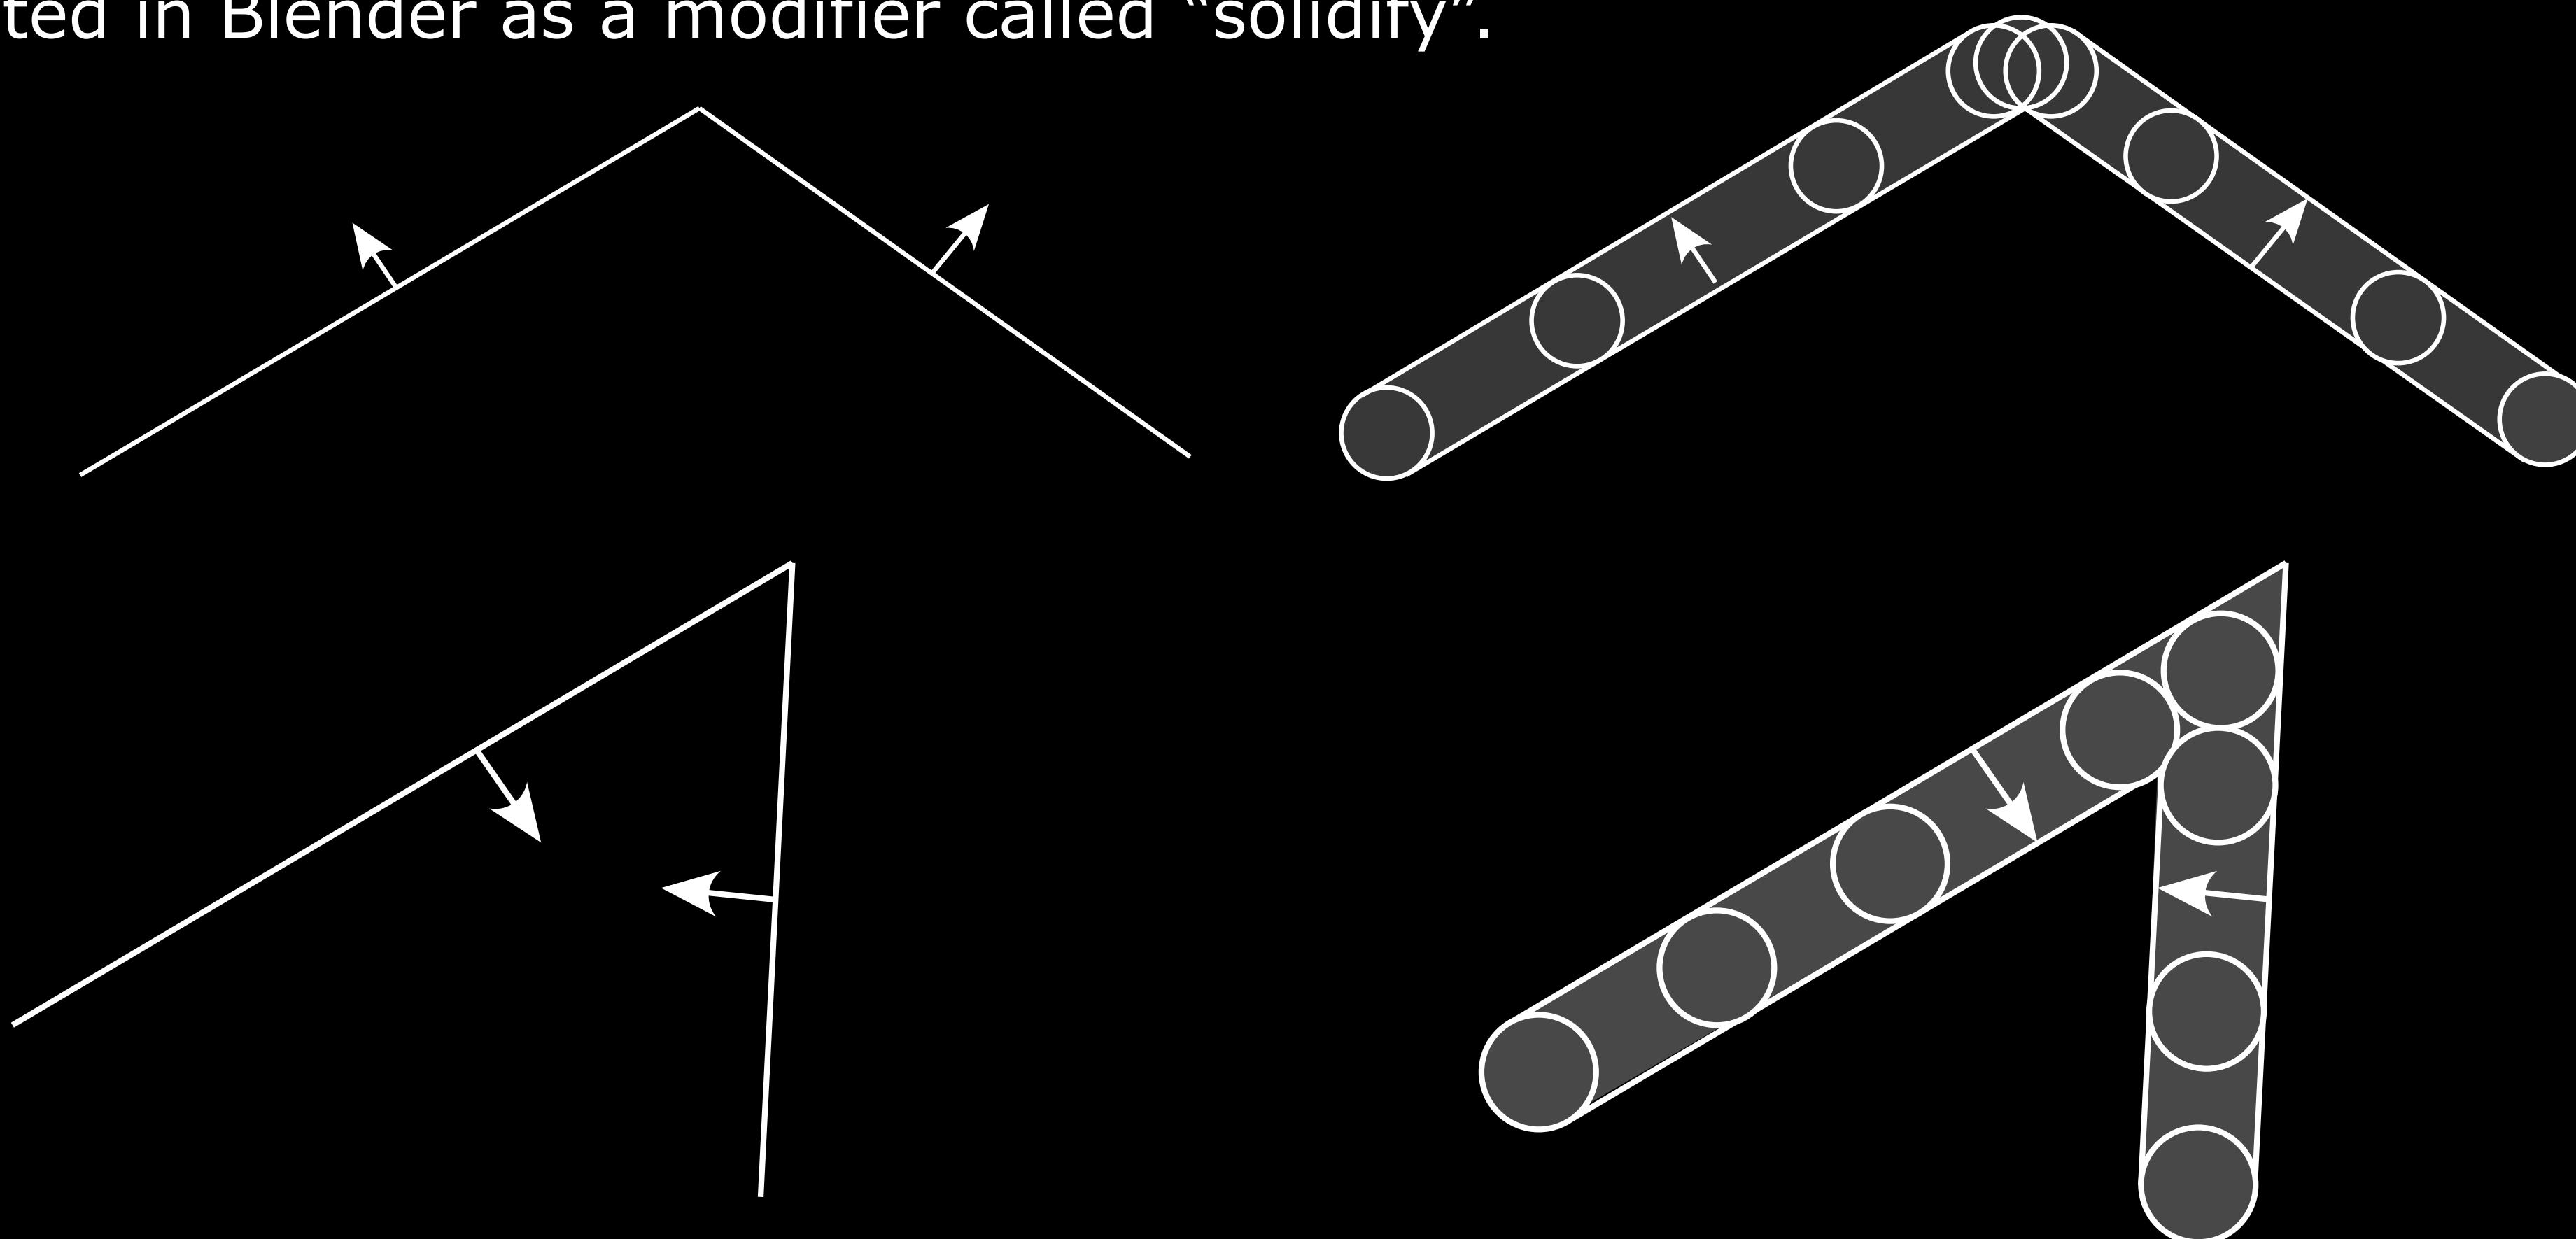
Mesh thickening

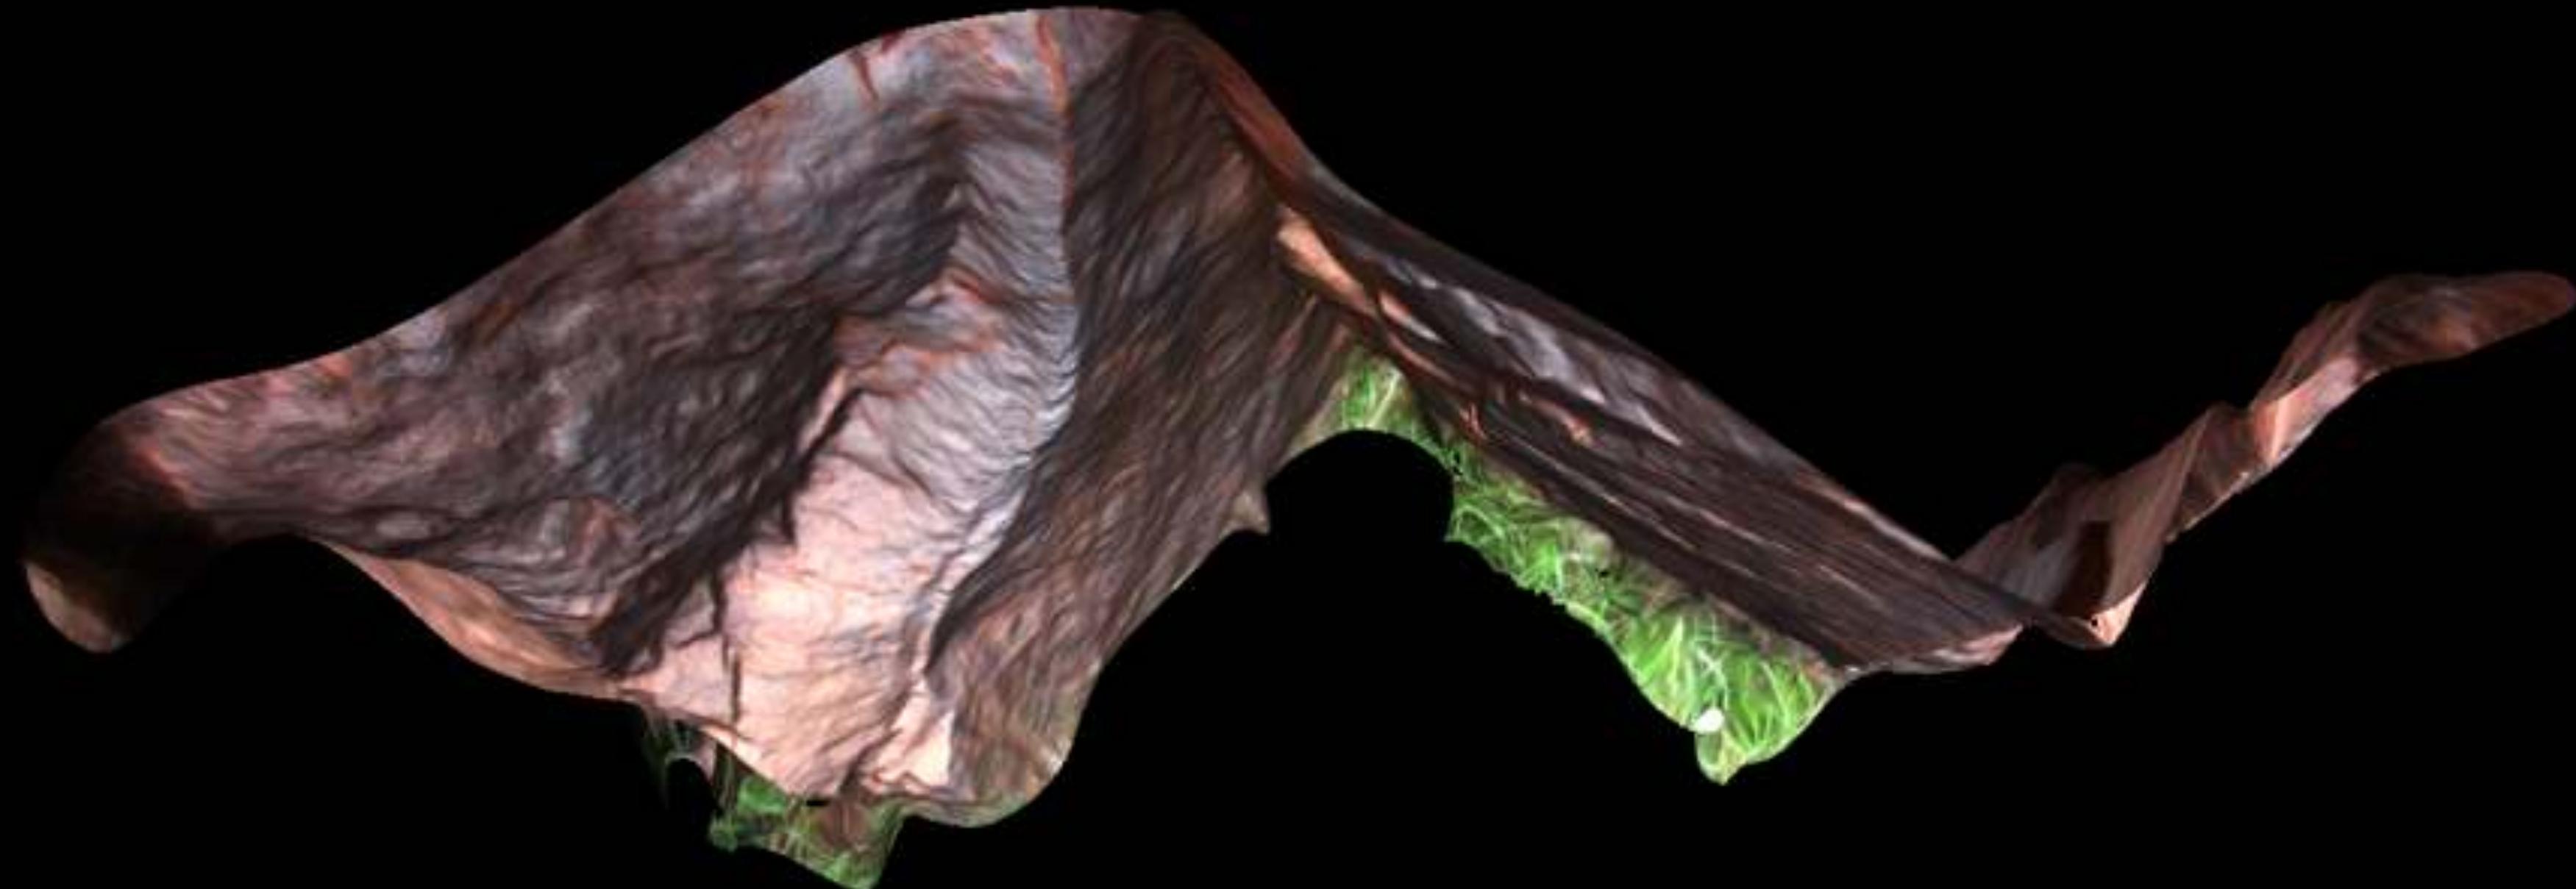
- Cases exist where one does not want idealised “infinitely thin” surfaces.
- Double sided rendering in realtime APIs is not quite the same visual effect as physical thickness.
- Required to create physical models, see rapid prototyping later.
- Perhaps the most common algorithm is known as “rolling ball”.



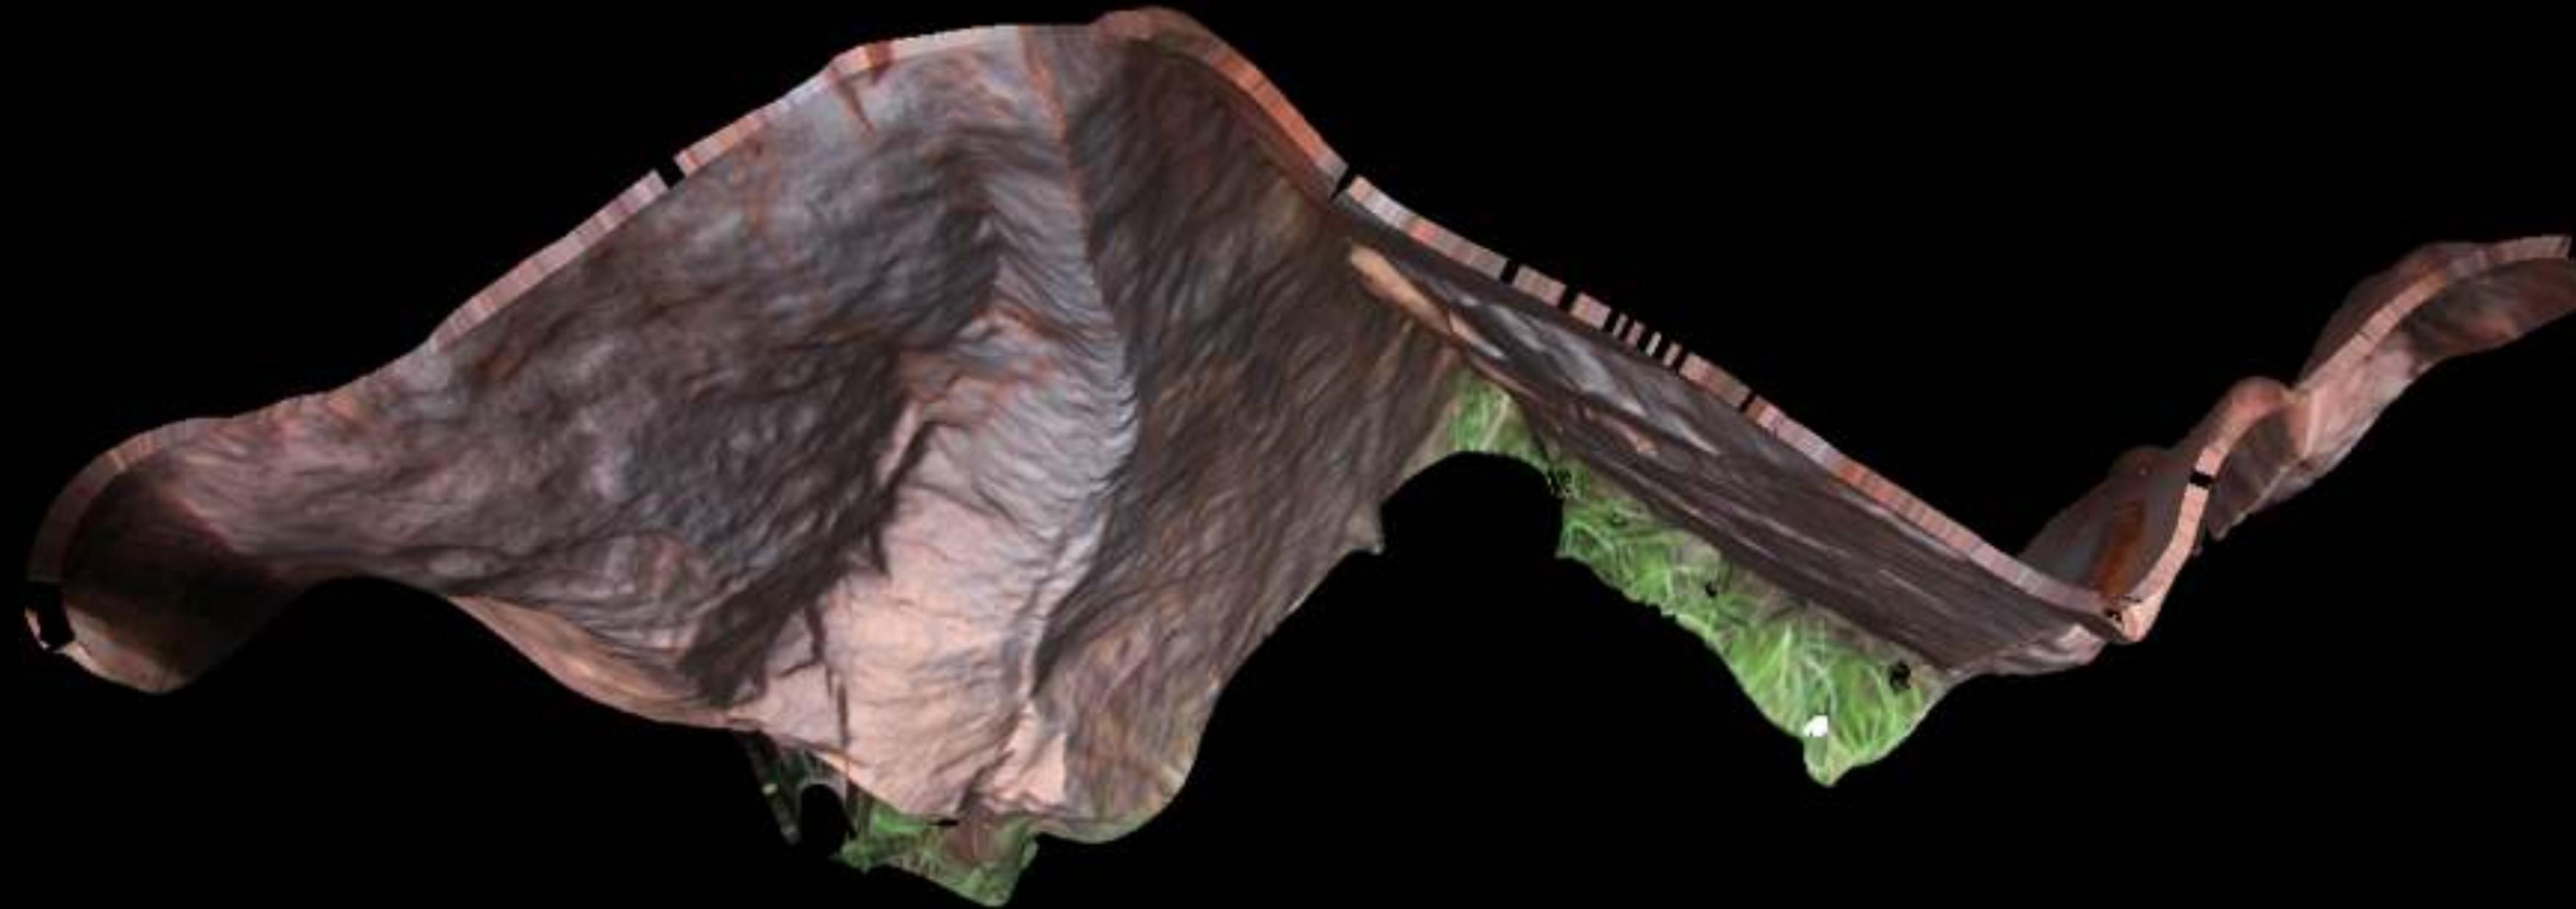
Ball rolling algorithm

- Solution is called “rolling ball” thickening.
- Imagine a ball rolling across the surface, form an external mesh along the ball path.
- Implemented in Blender as a modifier called “solidify”.

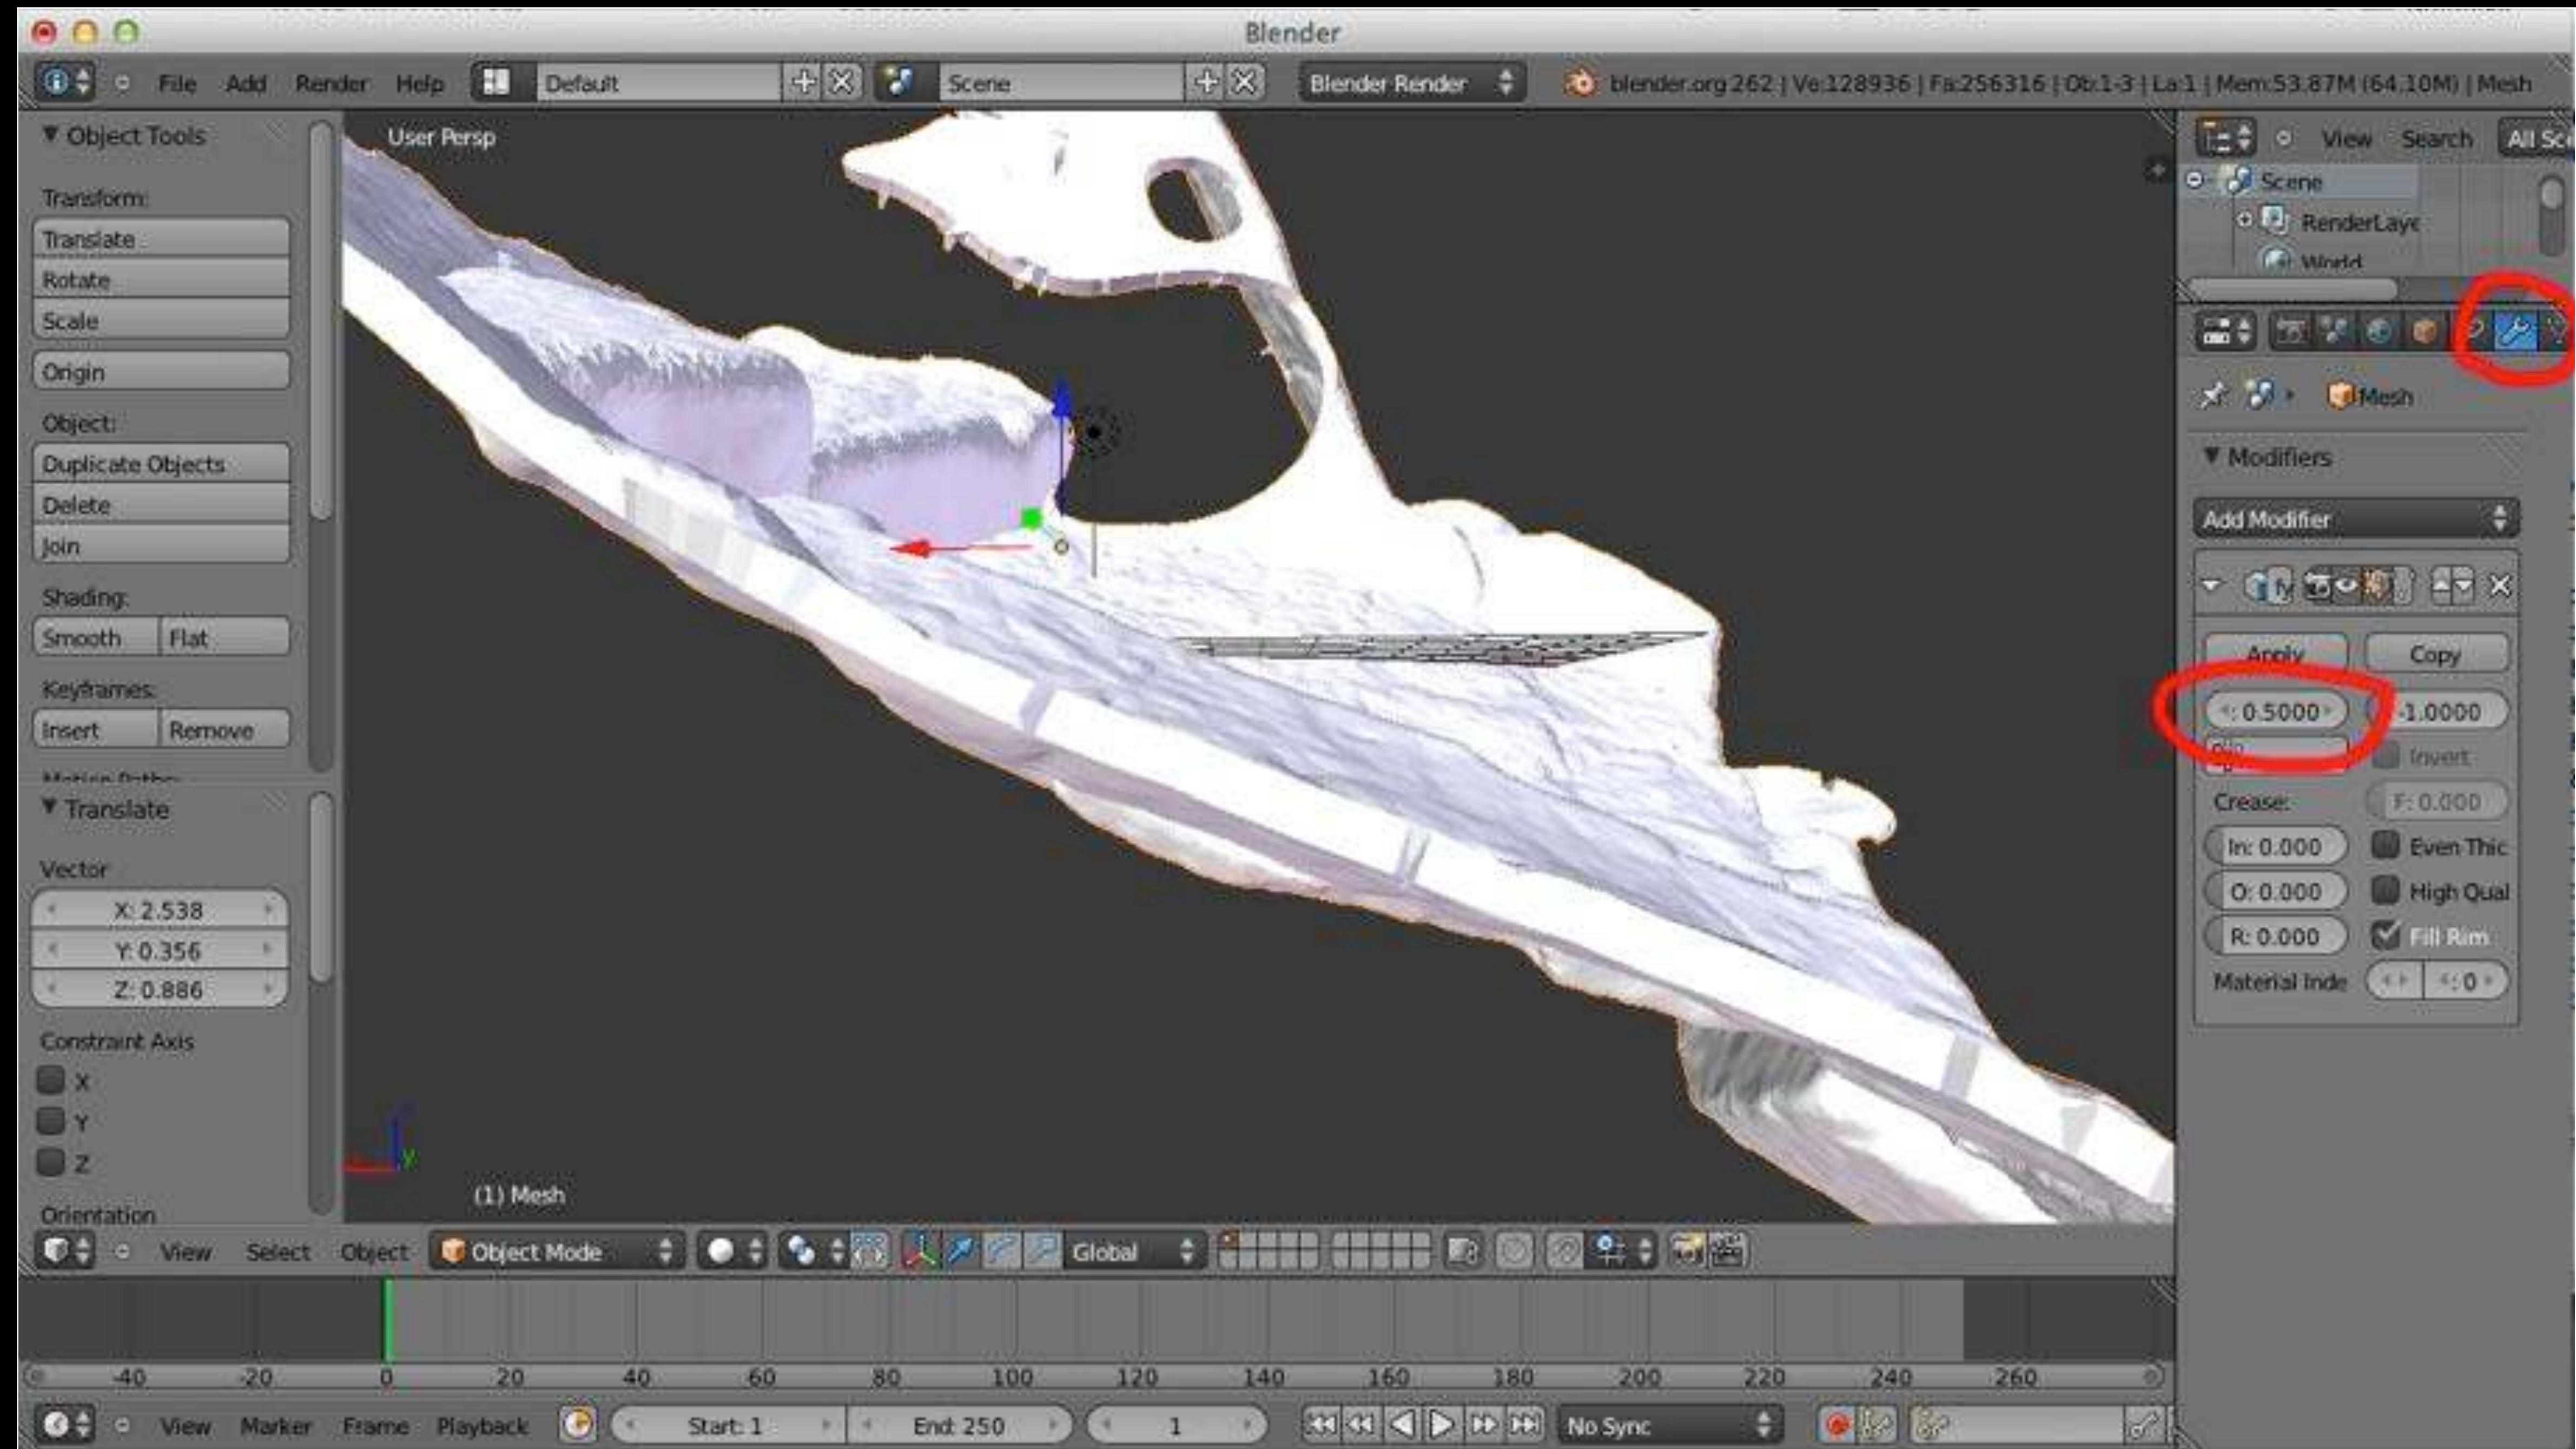




Infinitely thin surface
Unrealisable



Thickened “realisable” surface





Relighting





Post Production tools

- Smoothing.

Z-brush

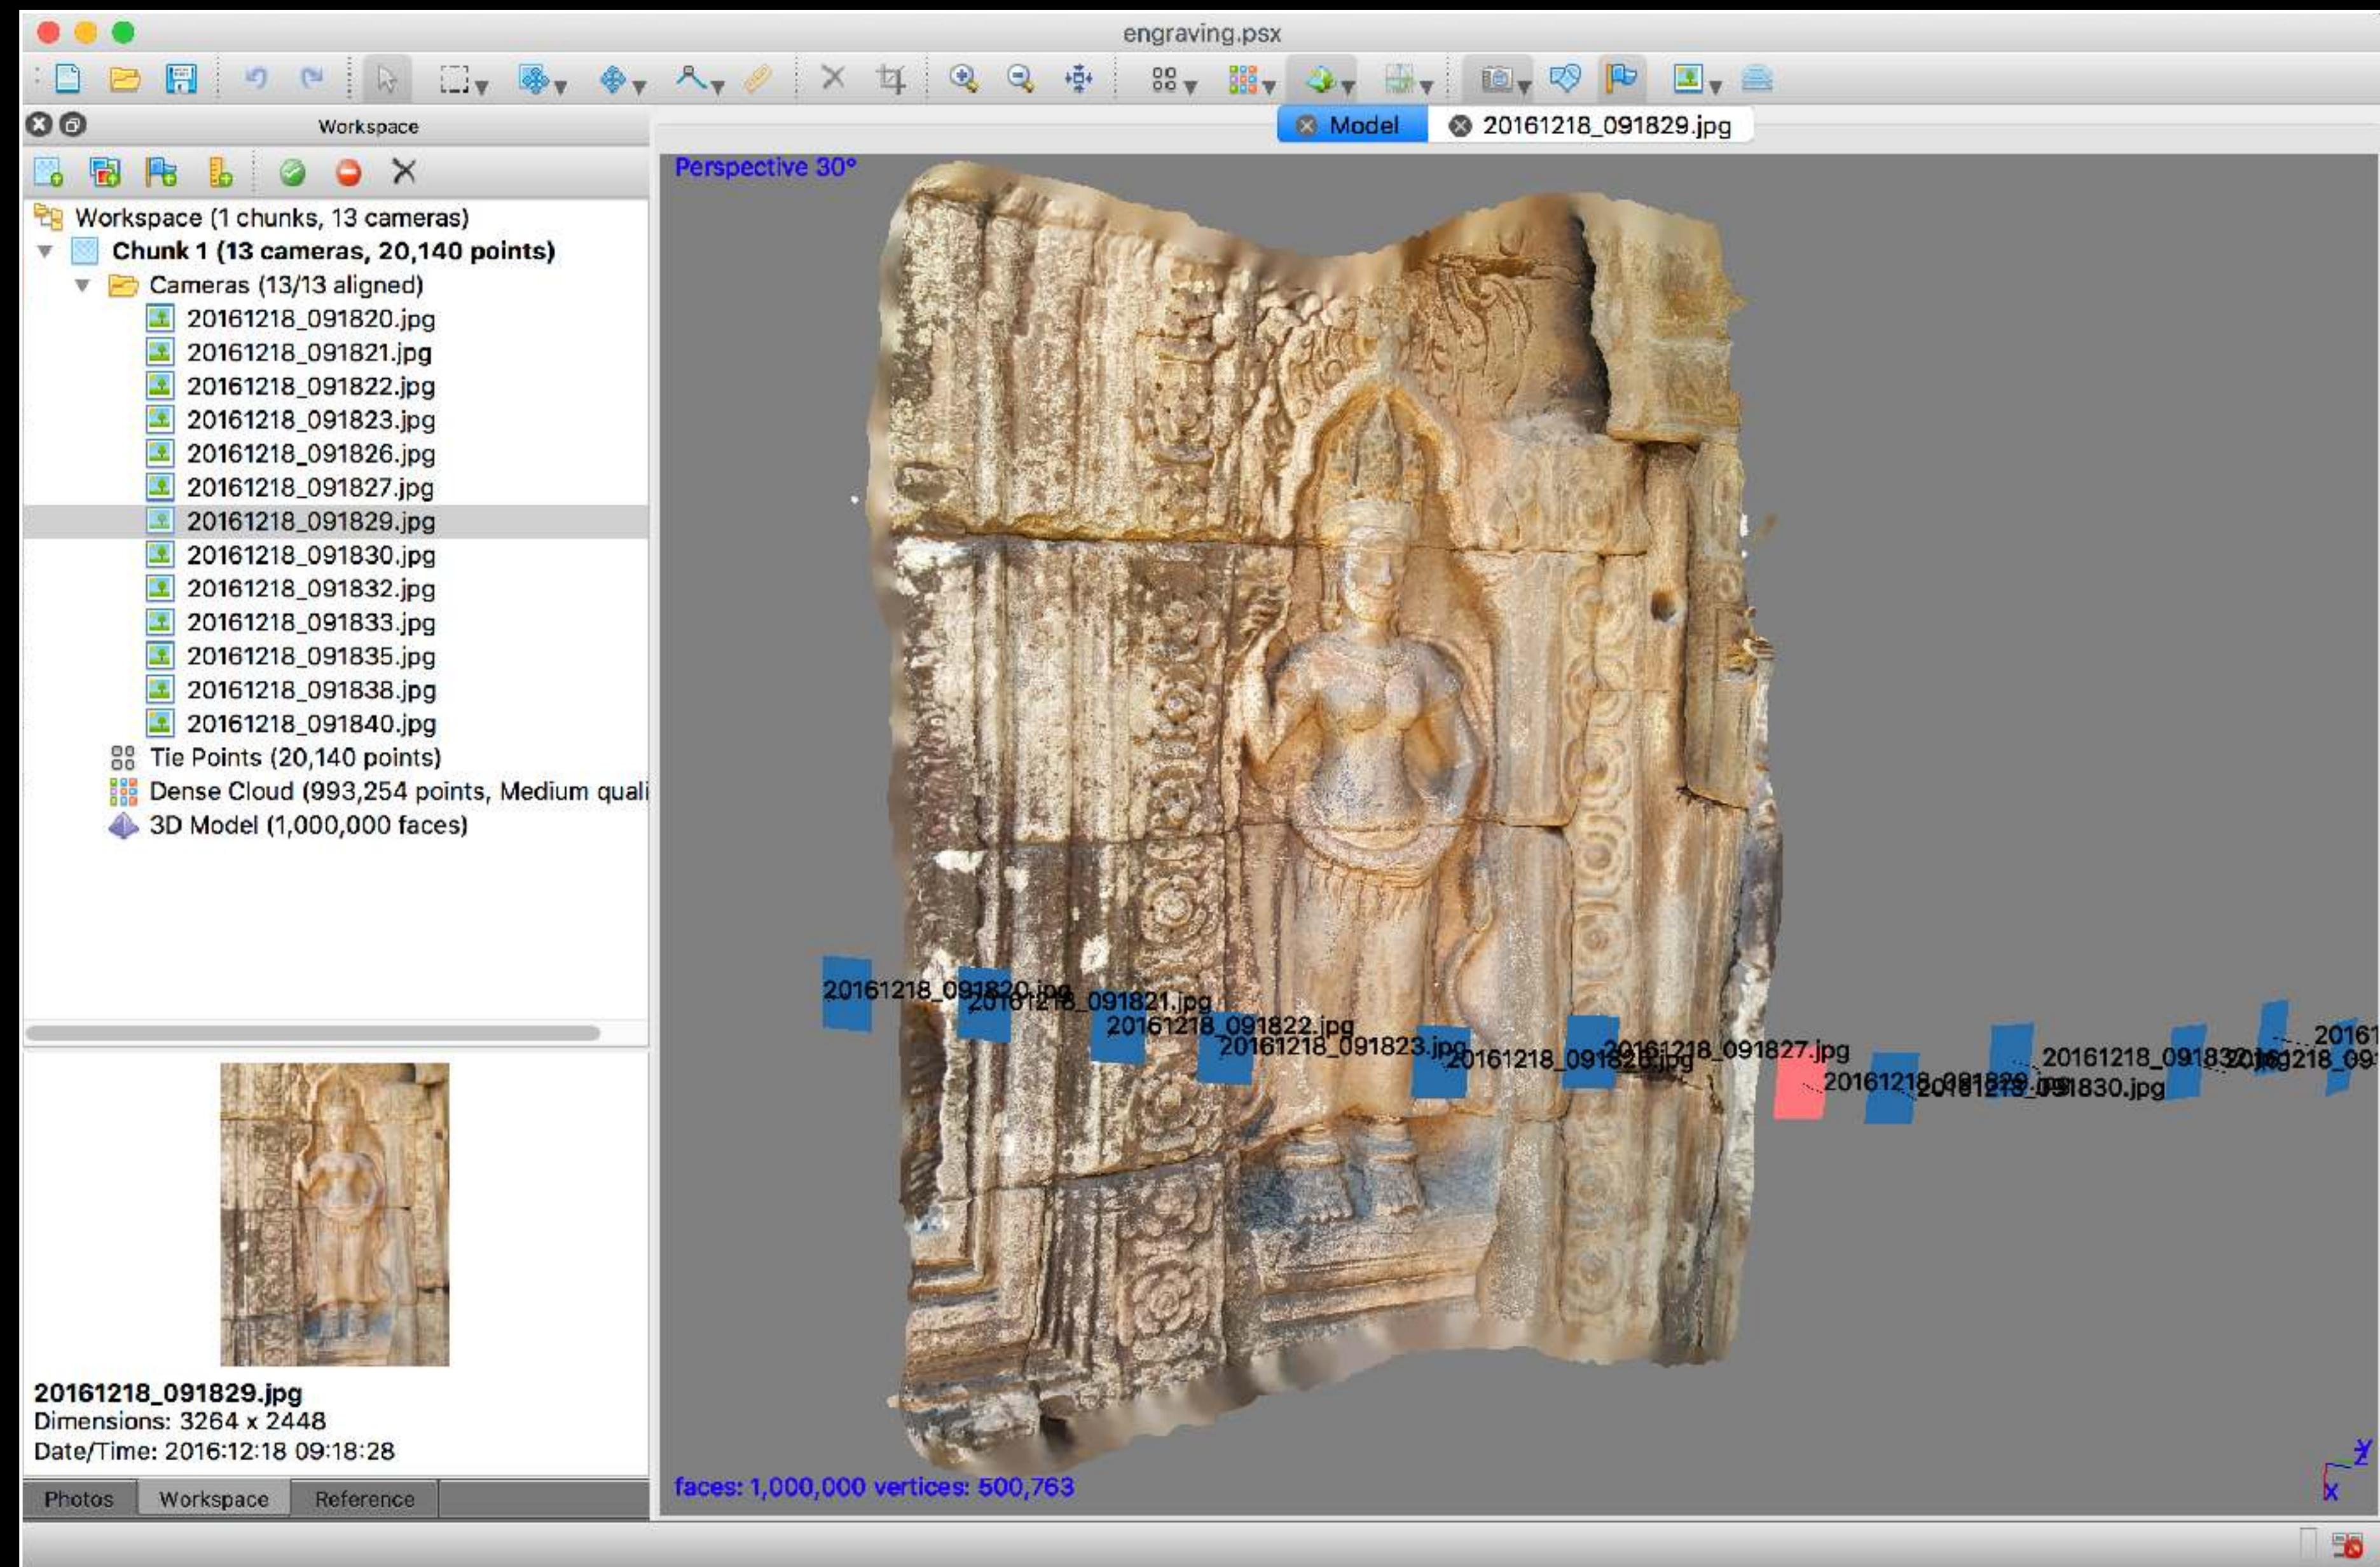
- Algorithmic. eg: mesh density reduction

MeshLab

- Almost any 3D editing package you may be familiar with that can handle unstructured meshes.

3DStudioMax - Maya - Cinema4D - Blender

Worked example



Your turn

- Free 30 day trial
<https://www.agisoft.com/downloads/installer/>
- MeshLab (free)
<http://www.meshlab.net/>
- This is about experimenting
Start off with a modest number of photographs
so compute time will be modest.